

BDL Design and Formalization Monograph

A Behavior Design Language for Interactive Physical Products:
Why Behavior Needs a Design Medium, the Architecture,
One Product Through the Lifecycle, the Derived Formal Model,
the Toolchain, Studio, and the Evidence

ZHU ZHEHAO

Living technical record — revision of 2026-09-20

Formal development KCN-judu/BDL_FV — this working tree, through Phase 15 (previous commit dce5ac4)

Production KCN-judu/BDL at 081296df606d577eece7e269ed250b255547d497 (2026-09-20, protocol 0.24)

This document is the authoritative narrative design record of BDL — a behavior-design language for interactive physical products — written as one argument, from the design problem through the architecture, one product's lifecycle, the derived formal model, the toolchain and the authoring environment to the evidence, and a bridge between its formal model and its production system. Its canonical source is `paper/` in KCN-judu/BDL_FV; the production repository carries a reference mirror that does not evolve on its own. It is not a paper, not a submission and not written to a page limit; it records what the language is, why each construct exists or was rejected, which theorem, counterexample or executed example supports each claim, how production implements it, where production deviates from the formal model, and what remains open — including the alternatives that failed. Every claim carries a strength label (formally proved; formally characterized under restricted hypotheses; mechanically executed example; counterexample; informed by FV; production implemented and tested; production architecture decision; proposal / not implemented; design recommendation; open empirical question). No usability claim in it has been tested with users.

Contents

Preface	6
How to read this monograph	6
Part I — Why Behavior Needs a Design Medium	8
The asymmetry of design media	8
The question	8
The working hypothesis: the typed relationship first	8
The Representation Problem	10
The target user is not a programmer with fewer syntax skills	10
A flow graph is a dependency view, not a program counter	10
Cognitive budget	10
Progressive disclosure is a semantic property	11
What BDL is not	11
Who BDL is for	11
Scope	11
Part II — The BDL Architecture	13
The design path	13
The tooling path	14
Architectural principles	14
Three layers: surface, kernel, validation	15
The state of the record	16
Part III — Designing an Interactive Physical Product with BDL	18
III.1 Naming what the product is about: concepts	18
III.2 Stating relationships before defining them	18
III.3 Giving a relationship a formula	20
III.4 Where the environment enters: the Source	20
III.5 Where the design leaves: the logical Output	21
III.6 Time, only where the product needs it	21
III.7 Organizing a larger design: behaviors and components	22
III.8 The Standard Library, and computation supplied from outside	23
III.9 Simulating before any hardware exists	23
III.10 Choosing a target and realizing the boundary	23
III.11 What the workspace says at each step	24
III.12 Looking underneath	24
III.13 The lifecycle, and where the evidence stops	25
Part IV — Deriving the Formal Model	26
Method	26
Scale and trust base	26
IV.1 Persistent declarations and refinement	26
IV.2 Concept identity and physical quantities	28
Nominal concepts	29
Representation binding and the grant	29
Physical dimensions	30
The concept ladder (Phases 19–22)	31
IV.3 Data and equations	31
The question and the hypothesis	31
What entered the kernel, and why each item could not be derived	31
Polymorphism: rank-1, by families, no kernel type variable	32
Capabilities: the audit that reverted a decision	33
The definitional library	34
Enumerations, records, sets: what was encoded and what was deferred	34

The expressiveness ceiling	35
The natural expression surface	35
IV.4 Units, coordinates and charts	36
Three notions kept apart	36
Linear units: no kernel construct	36
The numeric domain, in the open	36
Presentation, and the realistic formula	37
Unit ownership in authoring	37
Affine units: the failure of scale-only units, and the revision	37
Charts: coordinate erasure and conversion functoriality	37
The production chart model	38
The Formula Composer's formal basis	39
IV.5 Time: memory, causality and clock domains	39
A Minimal Reactive Semantics	39
Clock Domains	41
Cross-domain occurrences and lossless buffering	42
IV.6 Behavior composition: components and groups	44
Behavior components	44
Behavior groups	46
Hardware interaction	47
IV.7 The physical boundary: Sources, logical Outputs, provision and realization	47
The canonical type of a relationship that reads nothing	47
Physical outputs without arbitration	50
Provision: how a Source gets its value on a product	51
Realization: how an output reaches the world	53
The adapter boundary and the device clock	56
Communication as state, catalogue profiles and assign	57
The provider's contract and the occurrence-preserving window	59
The Source side completed: provider state, the device clock, initialization, commitments, readings	60
The physical boundary as one whole	61
Two remarks across §IV.1–IV.7	62
IV.8 Validation and deployment	63
Three verdicts that must not be confused	63
Target-specific hardware validation	63
Deployment capacity	65
Provision and realization as deployment architecture	65
The generated-code boundary	65
Open hardware questions	66
Part V — From Semantics to an Executable Toolchain	67
Four trust layers	67
V.1 Authoring surfaces and the semantic project	68
V.2 Elaboration and checking	68
V.3 The reference evaluator	70
V.4 Lowering and the generated core	70
The generated core	70
Embedded execution model	71
V.5 Platform adapters: one command plan, several host mechanisms	72
V.6 The IDE service, the daemon and the protocol	73
V.7 The implementation map	74
Part VI — Designing Studio and the IDE	76
VI.1 The user, the workflow and the constraints	76
VI.2 Principles	76
VI.3 Professional software references: what is borrowed, what is refused	77

VI.4 Semantic projection: one service, two surfaces	78
VI.5 The canvas as a diagram of product relationships	79
VI.6 Roles, gestures and context menus	80
VI.7 Diagnostics, fixes and the three information levels	81
VI.8 The Code view: the same project as text	81
VI.9 The definition editor and the Formula Composer	81
VI.10 The Library and the Source sheet	83
VI.11 Simulate, Deploy and the boundary of Design	83
VI.12 What Studio does not decide, and what is not claimed	84
Part VII – Evidence, Rejected Designs, and the Open Agenda	85
VII.1 Formal $\leftrightarrow$ production correspondence	85
Construct by construct	85
Deviations, stated once with their reasons	91
The snapshot	92
VII.2 Rejected designs, as arguments	92
Kernel constructs retained	94
Rejected constructs	94
Decisions that changed	99
VII.3 Minimality, bounded	100
VII.4 The open agenda	101
Formal limits, stated once	101
Open formal questions	101
Production engineering work	103
Empirical questions, explicitly	103
Risks of the design	105
Intellectual Context and Related Work	106
Industrial design tools and physical prototyping	106
Visual programming and end-user programming	106
Synchronous languages and functional reactive programming	106
Model-based design, statecharts and systems engineering	106
Typed holes, live and structured editing, projectional editing	107
Dimensional typing and units	107
Effects	107
Resource allocation	107
Formal verification, verified compilation and differential testing	107
Embedded DSLs and toolchains	108
Closing	109
Appendix A – Formal notation	110
Appendix B – Theorem and result index, by concept	112
Refinement and client stability (Core/Interface, Core/Decl, Core/Satisfaction, Core/Env)	112
Typing, unfolding and dependency (Core/Typing, Core/Dependency)	112
Concept identity, representation and the grant (Core/Base, Experiments/SemanticTypeAlternatives, Experiments/ RepresentationBindingAlternatives)	112
Quantities and dimensions (Core/Base, Experiments/DimensionAlternatives)	113
Reactive semantics (Core/Reactive, Experiments/ReactiveAlternatives)	113
Clock domains and transport (Core/Clock, Experiments/ClockAlternatives)	113
Lists, the buffer and capacity (Core/ListData, Surface/Buffer, Validation/Capacity, Experiments/ BufferAlternatives)	114
Products, the recursor, equality and order (Core/Base, Surface/Poly, Experiments/PolyAlternatives)	114
The equation library and rank-1 polymorphism (Surface/Poly, Surface/Stdlib, Surface/Generic, Experiments/ EquationExamples)	115
Natural surface (Surface/Natural, Experiments/NaturalExamples)	115

Units, coordinates and charts (Surface/Units, Surface/Composer, Surface/Affine, Surface/Rational, Surface/Charts)	116
Physical outputs (Core/Output, Experiments/OutputAlternatives)	117
Unit domain and the source role (Surface/UnitDomain, Experiments/UnitDomainExamples)	117
Source provision (Surface/Provision, Experiments/ProvisionExamples)	117
Output realization (Surface/OutputRealization, Experiments/OutputRealizationExamples)	118
The adapter boundary and the device clock (Surface/Adapter, Surface/DeviceClock, Experiments/AdapterExamples)	119
Communication as state, catalogue profiles and assign (Surface/Assignment, Experiments/CommunicationExamples)	120
The provider's contract and the output window (Surface/Provider, Surface/OutputWindow, Experiments/ProviderExamples, Experiments/OutputWindowExamples)	120
The Source-side boundary (Surface/SourceBoundary, Experiments/SourceBoundaryExamples)	121
Hardware validation (Validation/Hardware, Experiments/HardwareAlternatives)	122
Behavior systems and groups (Behavior/*, Experiments/BehaviorAlternatives, Experiments/GroupAlternatives)	123
Appendix C — Decision index	124
Appendix D — Evidence-strength ledger	142
Appendix E — Migration map from the old decision numbers	146
Open items	152
Appendix F — Production snapshot	155
Milestones since the snapshot de8154f, newest first	155
Protocol	156
Implementation status, by area (production's own words, abridged)	156
Facts this document relies on that could move	157
Appendix G — Development chronology	158
Appendix H — Revision log	161
Appendix I — Record conventions: claim strength and identifiers	163
Claim strength	163
Identifiers	163
Snapshots	163
Canonical source and the production mirror	163
Appendix J — Section migration map from the previous revision	165
References	167

Preface

This is the *BDL Design and Formalization Monograph*: the living technical record of the Behavior Design Language — a language in which the behavior of an interactive physical product can be designed before it is implemented. It is written as an argument that a reader can follow from beginning to end, and it is also the authoritative record of what BDL is, why each of its constructs exists, which alternatives were tried and failed, what its formal model proves, what its production implementation does, where the two deliberately differ, and what remains open. The two purposes are not in tension: the argument is stronger for carrying its evidence, and the evidence is easier to trust when it appears where the question it answers is asked.

The document is not a conference paper and is not written to a page limit. It keeps its negative results, its reversed decisions and its unproved boundaries, and it says at every substantive claim how strong the claim is. It lives in `paper/monograph/` beside the core-calculus paper (`paper/core_calculus/`); reader-facing text says *monograph* or *record*.

Two repositories are described. The formal model is `KCN-judu/BDL_FV`, a Lean 4 development with no external libraries; it is described as of the working tree of this revision, through Phase 18. The production system is `KCN-judu/BDL`, a Rust toolchain and a Flutter authoring environment; it is described as of one audited commit, `aa6e7f4ee249556ca7561c854b43607f45f229fe` (2026-09-20, protocol 0.24), and every sentence about production is a sentence about that commit. The volatile details of that snapshot are in Appendix F so that the argument does not go stale with the next milestone.

How to read this monograph

The argument runs in seven Parts, and a first reading should take them in order.

Part I asks why product behavior needs a design medium at all: what industrial designers have for form and lack for behavior, what the gap costs, and what BDL is *not*. **Part II** gives the whole architecture in one map — from a designer’s intent, through concepts, typed relationships and behavior, to the logical boundary where a product meets its environment, and on through deployment, provision and realization to a raw command, a platform adapter and the physical world — together with the design principles the rest of the document justifies. **Part III** walks one interactive physical product, a tilt-dimmed lamp with a warming base, through the entire lifecycle: concepts, relationships left incomplete, formulas, a Source, a logical Output, timing, a target, realization, validation, generated code and a board. A reader who stops after Part III knows what the system does and why each layer is there.

Part IV derives the formal model from the design requirements the example exposed, in the order the constructs depend on one another rather than the order they were built: persistent declarations and refinement; concept identity and physical quantities; data and equations; units; time; behavior composition; the physical boundary on both sides; validation and deployment. **Part V** shows how that model becomes a toolchain — parser, elaborator, checker, reference evaluator, lowering, generated core, platform adapters, daemon — organized by what each layer owns and must never decide. **Part VI** is the design argument for Studio and the IDE: who the tool is for, what professional software it borrows from and what it refuses, and the one architectural idea that organizes it — the compiler’s semantics projected into a canvas and a text editor alike. **Part VII** gathers the evidence: the formal $\leftrightarrow$ production correspondence construct by construct, the rejected designs as arguments, the bounded minimality claims, and the open agenda in three separated classes — formal, engineering, empirical. The related-work chapter then places BDL against the traditions it borrows from, after the reader knows what is being compared.

The appendices are the reference apparatus: notation, the theorem index by concept, the decision index, the evidence ledger, the map from the retired decision numbers, the production snapshot, the development chronology, the revision log, the record conventions, and the map from the previous revision’s sections to this one.

Provenance. The formal development was done in phases (Phase 0 ... Phase 18), and the phase numbers appear throughout as provenance — the place to find the experiment behind a claim — never as the structure of the exposition. The order in which the constructs are explained here is their conceptual dependency; it is not the order in which they were discovered, and the chronology (Appendix G) says which is which.

Evidence. Every substantive claim carries one of a fixed set of strength labels. The short legend is below; the full vocabulary, with the words production uses in its own records, is Appendix I.

label	means
formally proved	a named Lean theorem proves the stated property of the formal model — never of the Rust or Dart code
formally characterized under restricted hypotheses	a named theorem proves the property for a stated fragment; the restriction is part of the claim
mechanically executed example · counterexample	a concrete case run inside the proof checker; evidence for that case, or a candidate refuted by it
informed by FV	a formal result bounded an engineering choice; the choice itself is production's
production implemented and tested	present at the pinned snapshot and exercised by a named test
proposal / not implemented · design recommendation	a construction or guidance with no implementation, or no theorem, behind it
open empirical question	a claim about designers or usability; no study has been run, and nothing here is evidence for it

“Minimal” is never claimed globally: where the model says minimal it means *minimal among the tested candidates*, and the text says which.

Identifiers. Formal decisions are FVD-NNNN, formal open items FVI-NNNN; production's records are ADR, PRP and ISS-NNNN; a formal result and the production record it concerns are cited together, ADR-0032 (FVD-0118). Theorem names are Lean identifiers and never change to track a document. The rules, and the map from the older D-NN numbers, are Appendices I and E.

Part I — Why Behavior Needs a Design Medium

The asymmetry of design media

Industrial design is increasingly concerned with products whose behavior is determined not only by geometry, materials, and mechanisms, but also by sensing, logic, timing, software, and networked control. A cup may infer that it has been lifted, a lamp may adapt to ambient conditions, a medical device may gate an action on multiple safety conditions, and a consumer robot may continuously map sensor estimates to actuator behavior. In such products, behavior is no longer a late implementation detail. It is part of the product concept.

Yet the dominant design media remain asymmetrical. Industrial designers have mature media for shape, layout, appearance, and physical assembly, while product behavior is usually externalized through prose, storyboards, flowcharts, state diagrams, interactive mock-ups, or ad hoc embedded code. The first four are easy to sketch but weak as executable specifications; the last is executable but forces the designer to adopt implementation-oriented concepts such as mutable variables, callbacks, polling loops, and device APIs. Existing physical prototyping systems such as Phidgets and d.tools substantially lowered the cost of building interactive prototypes [1], [2], but their behavioral representations still inherit important assumptions from programming and state-machine formalisms.

A concrete instance of the resulting cost appeared in a two-day introductory hardware workshop that the author designed and taught for ten participants with no prior embedded experience. Two of the available sensors behaved in opposite ways: the ambient-light sensor reports larger values as illumination increases, while the distance sensor reports smaller values as the target recedes. Participants lost track of which was which, repeatedly and across the whole group. This pattern suggests a representational problem rather than merely a syntactic one. Whether a reading rises or falls with the quantity it measures is a stable fact about a device, and in the code they were writing there was nowhere to record it. It survived only as a sign buried inside an expression and had to be reconstructed each time it was needed. The obstacle was not syntax. A piece of semantic information had no place to live.

The question

BDL draws a different boundary. The goal is not to make engineering representations merely easier for designers to use, and the question is not how industrial designers might write ordinary software with less friction. The question is closer to this: *can the behavior of an interactive physical product itself become a design material — something sketched, compared, checked and simulated in its own terms — before implementation details dominate the representation?* The goal is therefore to define a *native representation of product behavior for design itself*, while retaining enough formal structure for static checking, simulation, and eventual realization on a board. This is a design and research position, stated as one; whether designers work this way when the position is available is an empirical question this document does not claim to have answered (§VII.4).

The behavior meant here is a particular kind. It is the behavior of products whose value lies in the crossing of four things — a person’s action, an environment’s state, a computation, and a physical response — where a tilt becomes a brightness, a temperature becomes a cut-out, a button becomes a mode, and every one of those crossings has a sensing side, a timing, and an actuator on the other end. Sensor-driven, actuator-driven and embedded interactive systems are the home ground; the primary objects of the language — concepts with physical value forms, relationships that may be declared before they are defined, timing domains, and outputs that reach the world through one explicit driver — are shaped by that ground. Software interaction with no physical boundary can be written in BDL where it helps, and nothing in the kernel forbids it, but it is not the design problem the language is built around, and this document does not frame BDL as a general-purpose language for user-interface or software behavior.

The working hypothesis: the typed relationship first

The system is **BDL**, a Behavior Design Language. BDL is organized around a working hypothesis: when a behavioral relationship is first specified, *what kind of relationship should exist* is frequently settled before its exact implementation is. A designer may know that **Tilt influences Brightness** before deciding the transfer function; that **CupPickedUp** should activate a behavior before deciding which sensor and threshold detect pickup; or that a safety condition should suppress an actuator before choosing the device driver. Therefore the primary design artifact should be the *typed relationship*, not the procedure that computes it.

The central example is a relationship. Instead of decomposing a design into procedural steps such as “read tilt,” “calculate brightness,” and “set the LED,” BDL represents one mapping:

$$?f : \text{Tilt} \rightarrow \text{Brightness} .$$

The mapping may exist before its body. A formula, curve, examples, or a fitted function can later be attached as a definition of the same block. Once a formula is supplied, for example

$$f(\theta) = \text{clamp} \left(0.2 + 0.8 \frac{\theta}{60^\circ}, 0, 1 \right),$$

it inhabits the previously declared signature. The claim attached to this **signature-first** model is deliberately narrow. It is not that designers universally think in signatures before bodies, nor that they should be trained to. It is that an unresolved typed relationship is a legal, statically meaningful state of the design, so that stopping between declaring a relationship and realizing it costs nothing. Whether designers make use of that position, and at what level of task complexity, is an empirical question that remains open; §VII.4 states it as such and lists the studies that would answer it.

Signature first; formula attached later




Figure 1: A signature-first relationship. The flow graph contains one semantic relationship, $\text{Tilt} \rightarrow \text{Brightness}$; the formula is attached to the block rather than represented as an additional execution step.

BDL separates a designer-facing surface language from a small formal kernel. The kernel presented here was not designed on paper and then implemented. It was derived by a mechanized design-space exploration in the Lean 4 proof assistant [3]: each candidate construct from an earlier draft of the language was formalized, attacked with counterexamples, reduced to other constructs where possible, and kept only when the tested alternatives failed for a stated reason. The result is smaller than the draft it replaced. Clock-indexed signal types, a separate event type, effect rows, action requests, and runtime actuator arbitration were all removed, each for a reason recorded in the accompanying development. What remains is an environment of named declarations with frozen expected types, monotone public commitments, and write-once realizations; nominal concept types whose values can be constructed only inside a declaration whose own signature announces the concept; physical dimensions carried in the types of primitive operators; one temporal primitive that reads a clock domain at its previous activation; nominal clock domains with a separate domain judgment; and nominal physical outputs with a single explicit driver each. A validation layer outside the kernel decides whether a design can be placed on a declared target board. Later phases added, each for a reason recorded in §IV.2–IV.5, list and product data with one recursor, behavior components and groups, an exact model of units and charts, and an interface-level account of the canonical type $() \rightarrow B$ that leaves the kernel without a unit; nothing that was removed has returned.

The rest of this document is the working out of that hypothesis: the architecture it implies (Part II), the architecture at work on one product (Part III), its derivation from the requirements the product exposes (Part IV), its implementation (Part V), its exposure to the designer (Part VI) and the evidence for all of it (Part VII). What has *not* changed since the first draft of the language is that no user study has been run: every statement here about what designers find natural is a hypothesis, and the last chapter of Part VII lists the studies that would test it.

The Representation Problem

The target user is not a programmer with fewer syntax skills

Many low-code and visual programming systems reduce textual syntax while preserving the underlying computational ontology: variables, assignments, loops, callbacks, functions, transitions, and scheduling. For software developers this can be convenient. For industrial designers, however, the dominant difficulty is often not syntax but *semantic translation*. The designer begins with a product statement such as “while the cup is held, brightness follows tilt” and must translate it into implementation machinery.

BDL therefore follows a stronger criterion: a surface primitive should be exposed only when it corresponds to a concept that is independently meaningful in the design task. A mutable accumulator used to count samples is generally not such a concept; “three pickup events within ten minutes” is. A polling loop is generally not; “while the product is held” is. A callback is not; “when the button is pressed” is.

The language should directly expose concepts, time-varying quantities, discrete occurrences, typed mappings, conditions, behavioral contexts, temporal relations, physical outputs, and safety constraints. By default it should hide program counters, threads, callbacks, continuations, clock variables, and bus transactions. These may remain inspectable in an expert or debugging view, but they are not the primary design medium.

A flow graph is a dependency view, not a program counter

BDL uses a flow-like canvas because causal and functional paths are useful visual structures. The arrows do **not** mean “execute the left node and then the right node.” They mean that the target relationship depends on the source relationship or value.

Consider the procedural decomposition:

```
[Cup picked up]
  |
[Read tilt]
  |
[Compute L = f(theta)]
  |
[Set LED brightness]
```

This decomposition is rejected as a primary design representation because the “read,” “compute,” and “set” nodes are artifacts of an execution model. The corresponding BDL view is a `Held` context containing a mapping from `Tilt` to `Brightness`. The formula is a property of that mapping, and the realization layer later binds the mapping’s value to a physical light output.

The distinction matters because it changes what the designer edits. In a procedure-centric editor, changing an implementation may require rewriting steps and intermediate state. In BDL, changing a transfer function edits the definition of one relationship while preserving the surrounding product logic. This is not only a matter of presentation. Realizing or refining one declaration leaves every other declaration’s typing untouched, and leaves every commitment that was already discharged in place, provided the evidence for those commitments does not depend on what was still unknown.

Cognitive budget

The surface language must remain intentionally small. Each additional primitive has a cost in learnability, visibility, consistency, and error-proneness, all familiar concerns in the Cognitive Dimensions tradition [4]. BDL therefore uses a *cognitive budget*: a new visible construct is justified only when it captures a distinct design concept that cannot be expressed cleanly as a property of an existing construct.

The same budget was applied to the kernel, under a stricter test. A kernel construct earns its place only if removing it makes some design either unrepresentable or ambiguous, and the argument for each is a theorem or a counterexample rather than a preference. Several constructs that the surface exposes as distinct concepts turned out, under that test, to be derived forms. Every temporal modifier reduces to one delay primitive. Every cross-domain policy reduces to one transport primitive plus ordinary data. Every output-selection policy reduces to ordinary computation upstream of a single drive edge.

Progressive disclosure is a semantic property

A design tool can be visually simple and still force premature decisions. BDL instead treats progressive disclosure as part of the language semantics. A design may contain a declaration whose realization is intentionally absent. Later steps may refine the same declaration with mathematical properties, a body, a clock domain, an output binding, and a target board. Each step is either a *refinement*, which preserves everything previously established, or an *edit*, which is permitted but reopens the validation of dependents. That distinction is made precise in the section on declarations, and it is the organizing principle of the interaction model.

What BDL is not

The shape of the problem is easier to see against the things BDL is sometimes mistaken for, so they are set aside here rather than at the end.

BDL is not a *low-code C*: it does not make the operational vocabulary of embedded programming — variables, loops, callbacks, device registers — easier to type, because that vocabulary is the cost, not the syntax around it. It is not a *generic visual programming language*: its canvas is a diagram of product relationships, not a program counter or a dataflow execution graph, and no node exists per arithmetic operator (Part VI). It is not a *GUI or software-interaction language*: a screen-only interaction with no sensing, no timing domain and no physical output is representable only as a degenerate case, and nothing in the language is shaped for it. It is not a *PLC replacement*: BDL has no ladder logic, no scan cycle as a designer-facing concept, and no ambition to run plant control; its clock domains and single-driver outputs are a semantics for products, not for industrial automation. And it is not a *programming tutorial for designers*: the claim is not that designers should be taught a smaller language, but that product behavior can be authored in the terms of the product, with the implementation machinery elaborated underneath.

Software behavior may be written in BDL where it helps — a mode, a menu state, a timeout — and the kernel does not forbid it. The core problem the language is built around is the crossing of a person, an environment, a computation and a physical response in one artifact: sensing, semantic mapping, timing, and physical output.

Who BDL is for

The target practice is industrial and product design: designers who own a product's form and must now own its behavior; interaction designers working with physical products rather than screens; embedded prototyping teams who want a checked behavior artifact before firmware; and, increasingly, individual designers who work across sensing, logic and actuation on small boards. What these practitioners share is a mature medium for everything except behavior, and a workflow in which the behavior is decided late, in someone else's vocabulary, and revised at the cost of a rebuild.

Two disciplines govern every claim about these people in this document. First, nothing here is evidence that designers *prefer* BDL, *understand* it faster, make *fewer mistakes* with it, or *think* in its abstractions; every such statement is an open empirical question, listed as such in Part VII, and the studies that would answer them have not been run. Second, the one anecdote the record contains — a two-day workshop, and one author's own practice on one hardware project — is reported as what it is. The design position is stated as a position; the language and its formal model are the part of the position that can be checked.

Scope

In scope. A designer-facing language for the behavior of physical products — concepts, relationships, timing, outputs, components — with a formal kernel small enough to reason about and a production toolchain that follows it: authoring in a graph and in text, checking while incomplete, simulation, target-relative deployment validation, and code generation for a deterministic embedded core. The formal development's job is to say what the kernel must contain and to reject what it need not; the production system's job is to build everything the kernel deliberately omits — the elaborator, the tooling, the execution path — without redefining the language.

Out of scope, by decision. In the technical sense BDL is not a general-purpose programming language either: there is no general recursion, no user-defined type abstraction, no effect system, and the equation language is total and first-order in its data (§IV.3). It is not a continuous-time or hybrid modelling language: time is a global tick with named domains, and continuous dynamics are outside the model (§IV.5). It is not a systems-engineering framework: requirements, verification of physical properties beyond typing and dimension, and electrical or thermal budgets

are not modelled (§IV.8). It is not a runtime with scheduling policy: the strictly-before rule removes the scheduler from the semantics, and the generated core has no tasks (Part V). It does not verify its own compiler: the production code is held to the model by tests, and the refinement proof is an open item (§VII.4).

Not claimed. That designers think signature-first; that Studio is usable; that the kernel is minimal in any absolute sense; that any theorem says anything about Rust or Dart. Each of these is either an open empirical question or a restriction on a claim, and the text says which wherever it arises.

Part II — The BDL Architecture

Before any syntax and before any formalism, this Part gives the whole architecture in one map, so that every later chapter can be read as the justification of a piece the reader has already seen. The map has two directions. The first is the *design path*: how a designer’s intent becomes a physical effect on a board, and where along that path the behavior stops being the designer’s and starts being deployment’s. The second is the *tooling path*: how an authoring environment, a compiler and a runtime divide the work so that exactly one of them owns the meaning of a design. The Part ends with the architectural principles the two paths embody; Parts IV–VI justify them.

The design path

A BDL design begins as **design intent** — *brightness follows tilt; the heater cuts out above a critical temperature* — and is written down in five kinds of object, all of which exist before any formula does.

Concepts are the things the product is about: Tilt, Brightness, Temperature. A concept has an identity that the language tracks, and a *value form* — a physical quantity with a dimension, an on/off state, a count, a collection — that says what kind of value carries it. Two concepts of the same value form are never interchangeable; that is the first architectural fact, and Part IV derives it from a counterexample.

Typed relationships are the design’s statements: `dimByTilt : Tilt -> Brightness` says that brightness depends on tilt and nothing about how. A relationship may be *declared* and left undefined; it may later be given a formula; and it is one of three *roles* by two authored facts alone — a **Rule** if it reads something, a **Value** if it reads nothing and has a definition, a **Source** if it reads nothing and has none. A Source is the point where the environment enters the design.

Behavior structure gathers relationships: timing domains that say which values update together, memory and transport across domains with explicit first values, behaviors (groups) that organize a large design, and components with promised interfaces that can be instantiated and bound.

The **logical boundary** is where the design meets the world without saying how. On the way in, a Source `tilt : () -> Tilt` is a value the environment provides at the concept’s type. On the way out, a **logical Output** `light : Brightness` is the design’s intent that the product carry a brightness, driven by exactly one relationship — *drive light by brightness* — and saying nothing about PWM, GPIO or I²C. Everything above this boundary is *behavior semantics*: it is what the simulator runs, what the theorems are about, and what a designer edits.

Below the boundary is **deployment**. A target board is chosen, and the design’s needs — a PWM line here, an I²C pair there — are derived from device kinds and placed on the board’s pins by a solver, or refused with a reason. **Source provision** realizes a Source by a raw device reading and a pure *transducer* that turns the raw value into the concept’s representation; **Output realization** realizes a logical Output by a pure *encoder* that turns the concept’s representation into a **raw command** — an 8-bit duty, a level — and a machine sink that carries it. Both are constructions the behavior cannot observe: the design’s traces are unchanged.

Below the raw command, the **platform adapter** turns each tick’s commands into **adapter operations** on peripherals — set, refused, held — under an explicit numeric policy, on a **target platform** (a Raspberry Pi Pico over Embassy, an Arduino Nano over `avr-hal`), and the **physical world** receives light or motion. The behavior semantics ends at the raw command; the operation is modelled and proved one step further; the register write and the photons are not modelled at all.




Figure 2: The BDL architecture. Left: the design path from intent to physical effect; the horizontal rule after the logical boundary marks where behavior semantics ends and deployment begins. Right: the tooling path; the compiler and its IDE service own every semantic verdict, and Studio’s canvas and Code view are two projections of them.

Which arrows are proved, which are tested and which are open is Part IV’s boundary chapter and Part VII.

The tooling path

The same design has one semantic home and several surfaces. **Studio** shows a project as a graph on a canvas, as its source files in a Code view, or both side by side, with an inspector for the selected object and pages for simulation, deployment and the library. Studio decides no semantic fact: it never types a formula, never infers a dimension, never decides who drives an output. Every verdict it shows comes from the **IDE service** behind the **daemon** `bdld` — one process per open project, one strictly monotone revision stream, a typed protocol — which answers hover, completion, navigation, references, semantic actions and highlighting for the canvas, the Code view and an external language server alike. The service is a thin set of queries over the **compiler’s** analysis of the semantic model: elaboration of concepts, signatures and formulas into the kernel’s terms; type, grant and dimension checking; dependency, causality and clock analysis; output and deployment analysis; and a **reference evaluator** that is the executable definition of runtime behavior. Below the compiler, **lowering** turns a checked design into an executable IR and a generated `no_std Rust` core whose every tick yields values, outputs and raw commands, and a **target** entry turns the solved deployment into adapter glue and firmware for a board family. The rule the whole path is built to keep obvious is that *BDL semantics flows downward and implementation mechanisms never flow upward to redefine the language*.

Architectural principles

The following principles are the shape of the map. Each is stated here as a rule; the chapter named after it is where the rule is derived from a design pressure and, where it can be, proved.

Behavior is first-class. The primary artifact is a typed relationship between product concepts, not a procedure that computes it. A designer edits what a relationship *is* — its signature, its definition, its timing — and the operational machinery is elaborated underneath (Part IV, the core declaration model).

Incomplete relationships are legal design states. A relationship may be declared and left undefined; the design around it is checked as far as it can be, other relationships may already depend on it, and later giving it a definition is a *refinement* that disturbs nothing already established, where changing its type or removing its definition is an *edit* that reopens its dependents (Part IV, refinement versus edit).

Identity is not display name. Every concept, relationship, clock domain and output has a stable identity that survives renaming, so that a rename is a refactoring and a reference never dangles (Part IV; Part V, the identity sidecar).

A semantic concept is not its representation. `Tilt` and `MotorAngle` may both be angles and are never interchangeable; a concept's value can be observed as its representation anywhere but constructed only where a signature announces the concept (Part IV, concept identity).

Behavior semantics is separate from deployment. Everything above the logical boundary — values, timing, outputs — has a meaning that does not depend on a board, a device or a pin, and a board can be chosen last (Part IV, the physical boundary; Part IV, validation).

A logical Output is not the physical device. An output says what concept the product carries and in which timing domain; how a duty cycle, a level or a bus write realizes it is deployment data chosen on the Deploy page, never on the output (Part IV, realization).

Device adaptation is downstream of behavior. Provision on the way in and realization on the way out are pure constructions added *below* the design; the design's traces are unchanged by them, and the platform adapter that follows interprets nothing (Part IV, the physical boundary).

BDL models the evolution of product meaning, not the transport of bits. A message, a frame, a packet or a byte has no place in a design; what a communication *does to the product* — a job queued, a position reported, a fault latched, an acknowledgement received — is state, written with the ordinary constructs, and the carrier stays below the physical boundary, where replacing it under the same semantic trace changes nothing (Part IV, communication as state).

The compiler owns semantic truth. One semantic model, one analysis, one evaluator; every other component consumes their verdicts (Part V).

The UI is a projection of semantic truth. Studio's canvas, inspector, Code view and Formula Composer render what the compiler says, in the designer's vocabulary, and hold no second definition of the language (Part VI).

Hidden coercion and hidden clock crossings are refused. No implicit conversion between concepts, no implicit resampling between timing domains, no implicit arbitration between two drivers of one output: every crossing is a visible object with a stated first value (Part IV, time and outputs).

Three layers: surface, kernel, validation

The design path above is a layering of *meaning*; the formal model draws a second, orthogonal layering of *what may be consulted to decide what*. The **surface** is what the designer authors, and every surface form elaborates one way into kernel objects. The **kernel** is the small set of judgments that give a design its meaning — typing that reads only the interfaces of declarations, a tick-indexed evaluation relation, a clock-domain judgment, and a few global conditions. The **validation** layer is everything that may depend on more than types: the discharge of a declaration's commitments, the fit of a design to a board, the capacity of its collections, and the admissibility of a chosen realization. The boundary between the three is the main architectural result of the formal development, and it is what keeps a design's validity separate from its deployability.




Figure 3: The three layers. Typing consults only the type view of declarations and the representation view of concepts; validation may consult commitments, evidence, and the target; the surface is derived forms over the kernel.

The **surface** is what the designer authors: concepts, relationships, temporal modifiers, contexts, device bindings, units, generic equations, groups and components, and display names. Everything in it elaborates to kernel objects, and the elaboration is one-directional: the kernel never needs to recover surface structure.

The **kernel** is the formal object of Part IV. It consists of an environment of declarations, a typing judgment, a tick-indexed evaluation relation over one or several clock domains, a domain judgment, and a small number of global well-formedness conditions: every realization satisfies its interface, the instantaneous dependency graph is acyclic, every reference respects domains, and every physical output has at most one driver. The typing judgment reads only the *type view* of declarations — their expected types — and the *representation view* of concepts. It does not read realizations, commitments, evidence, clocks, or bindings.

The **validation layer** is everything that may depend on more than types. It discharges the commitments a declaration makes, and it decides whether a design fits a target board. Two kinds of evidence live here and are kept apart. Evidence that is meant to survive refinement — a monotonicity commitment discharged compositionally through the commitments of other declarations — must be stable under monotone extension of the environment, and the kernel imposes that condition. Evidence that is not meant to survive refinement — the existence of a pin assignment on a particular board — is re-established after every change and is never merged with the first kind.

Figure 3 shows the layers as they stand at Phase 18; the kernel band also holds list and product data with one recursor (§IV.3) and the behavior-component constructs (§IV.6), the surface band holds the deployment construction of §IV.7, and the validation band holds deployment capacity (§IV.8). What is notable about the arrangement is how much of the earlier draft of BDL is absent from the kernel band. Reactive types, event types, effect rows, action requests, policy transformations, and a five-phase tick with a resolve step were all part of the draft kernel. Each was removed because it either added no rejection the smaller kernel lacked, or made a design decision on the designer’s behalf that should have been visible in the design. The states of Figure 4 are the designer-facing face of the same structure: *declared* and *type-valid* are the typing judgment and satisfaction; *temporally valid* is causality; *clock-consistent* is the domain judgment; *output-complete* is the single-driver and completeness conditions; *hardware-feasible* is the validation layer’s solver. Each is decidable for finite designs.

The state of the record

The map above is the architecture as it stands; the following three paragraphs are the state of the two repositories that realize it, kept here so that the reader of Part III knows what exists and what is proposed.

Current result. A small kernel, derived by a mechanized design-space exploration in Lean 4 and stated in Part IV: an environment of named declarations with frozen types, monotone public commitments and write-once realizations; nominal concepts constructed only where a signature announces them; dimensions in the types of primitive operators; one temporal primitive that reads a clock domain at its previous activation; nominal clock domains and

nominal logical outputs with one explicit driver each. Above it, proved constructions rather than kernel constructs: list and product data with one recursor, a definitional equation library with rank-1 instantiation by matching, units as coordinates with exact affine charts, natural binder syntax as conservative desugaring, behavior components and groups that flatten into the same kernel, the canonical interface type $() \rightarrow B$ whose kernel value is B , the provision of a Source at deployment by a raw reading and a pure transducer, proved transparent to the design, and — its output-side counterpart — the realization of a logical output by a pure encoder $\text{Rep}(C) \rightarrow \text{Raw}$ and a machine sink, proved to add downstream structure only: the logical output stays part of the behavior semantics, realization is deployment structure below it, the encoder is typed $\text{rep} \rightarrow \text{raw}$ in the empty design under no grant and is stateless, and the behavior’s environment and every trace it produces are literally unchanged under the proved hypotheses; two realizations of one output evaluate every behavior term alike; quantizing encoders are admitted; the machine boundary is a relation on raw commands, not a term. A design is *admissible* for a device only when three separate judgments hold — the encoder’s typing, its fit to the concept’s representation (EFits), and a solvable board — and the formal record shows that the first cannot be dropped: an encoder that fits and allocates but is ill-typed is not admissible (Phase 14 and its hardening, FVD-0139 superseding FVD-0137). What is *not* proved is the step from a raw command to a physical effect (FVI-0022). A separate validation layer decides hardware feasibility and collection capacity. Production implements the language as a Rust toolchain — model, elaborator, checker, reference evaluator, `no_std` code generator with output realization and generated raw commands, deployment analysis, daemon, IDE service — embedded platform adapters for the Raspberry Pi Pico and the Arduino Nano that consume those commands (production-tested, never formally proved), and a Flutter authoring environment, Studio, whose every semantic verdict is a projection from the compiler.

Formal scope. Phases 0 through 15 — twenty-one phase reports, counting the sub-phases and the post-Phase-1 migration — in 60 Lean modules (11 Core, 12 Behavior, 15 Surface, 2 Validation, 20 Experiments), about 1 300 theorem and lemma declarations, 142 decision records (one superseded) and 28 open items (8 open, 9 deferred, 11 resolved after the audit of 2026-09-20); no sorry, propositional extensionality and quotient soundness as the only axioms, no classical choice. Every theorem is about the *model*; none is about the Rust or Dart code, and none is about a board.

Production scope. Production is described as of commit `aa6e7f4ee249556ca7561c854b43607f45f229fe` of KCN-judu/**BDL**, 2026-09-20 — the head of `main` after the platform adapters for two target families (ADR-0037, amended: the RP2040 over Embassy and the Arduino Nano over `avr-hal`), output realization (ADR-0036), the Source sheet, the Code view as an IDE surface and the `drive ... by ...` spelling; protocol 0.24. The formal development is described as of the working tree that contains this revision of the document; the last commit before it is `f2b81c0` (the Phase 18 records), and the canonical copy of the production hash in the formal repository is the snapshot field of `docs/project/production-correspondence.md`. Every sentence about production is a sentence about that commit; volatile details are gathered in the production snapshot (§VII.1 and Appendix F) so that the conceptual chapters do not go stale with the next milestone. Where a statement was only checked at an earlier snapshot, the text says so.

Part III — Designing an Interactive Physical Product with BDL

This Part walks one product through the whole of the architecture of Part II, from the first concept to a board, in the order a designer meets the steps. The product is small and chosen because it exercises everything: a table lamp that dims when tilted and holds its brightness when set down, with a warming base whose heater must cut out above a critical temperature whatever else the lamp is doing, two quantities that update at very different rates, two physical outputs of different kinds, and a board chosen last. At every step the text says what the designer does, what the tool checks, and — briefly — what failure the layer exists to avoid. Where a step rests on a theorem the text says *later proved* and moves on; the derivations are Part IV, and nothing here stops the walkthrough to prove anything. Where the product’s vocabulary goes beyond what the tools offer at the pinned snapshot, the text says so in place, so that the lifecycle a reader sees is the one that exists.

The language forms below are those of the production snapshot: the .bdl text and the Studio surface are two views of one project (Part VI), and the same design can be authored in either.

III.1 Naming what the product is about: concepts

The designer begins by naming the things the lamp is about — Tilt, Brightness, Temperature, Held — before naming any sensor or any number. None of these is a device or a scalar.

The first-class object is a **concept**, not a raw scalar: Tilt, Brightness, Temperature, CupContact, MotorAngle. A concept has a **value form** — a quantity with a physical kind and a unit (Tilt : Angle), on/off, a count, a collection of another concept’s values, a grouped value, an optional value — or none yet (*decide later*); it may carry a description, a preferred display unit and, for a quantity, a declaration that it is **ordered**. Its identity is what the language tracks: Tilt and MotorAngle may both be angles and are never interchangeable, and a link between sockets of different concepts is refused before any formula exists.

Designer problem: the workshop cost of Part I — a stable fact about a kind of value with no place to live. The concept is the *type* of that value, a template; each place the product carries such a value is a **Sem block**, an instance of the concept holding one value per tick, produced by one **mapping block** or by the environment (§IV.2, the concept ladder). *Language concept:* the named concept with its value form. *Formal mechanism:* the nominal type `sem s` and the write-once representation binding Θ (§IV.2). *Owner:* `bdl-model` (identity, value form), `bdl-elab` (Θ), `bdl-check`. *Hidden:* the type constructor, the grant, the representation; the canvas shows a hue and a socket shape.

In text:

```
concept Tilt : Angle
concept Brightness : Scalar
ordered concept Severity : Scalar
concept Readings : List<Tilt>
```

Failure avoided. A design in which Tilt and a motor’s MotorAngle are both “a number in degrees” cannot say that wiring one into the other is a mistake; the workshop of Part I is the cost of that silence. Making the concept the object, with its identity tracked independently of its representation, is the first architectural decision, and Part IV derives it from the counterexample that decided it (§IV.2).

III.2 Stating relationships before defining them

The first relationship is drawn as an arrow from Tilt to Brightness, or written as one line of text, and the tool asks for nothing else.

A **relationship** is created from a **signature**: what it reads and what it produces.

```
mapping dimByTilt : Tilt -> Brightness    // reads Tilt, produces Brightness
mapping brightness : () -> Brightness      // reads nothing, produces Brightness
mapping tilt : () -> Tilt                  // reads nothing, no formula: a Source
```

The keyword is mapping; the product word is *relationship*. A relationship has a stable identity, a display name, a signature, an optional **formula** (its realization) and, when it lives in a timing domain, a domain (`@interaction`). Every relationship has one canonical type `domain(inputs) -> B`, with the empty product `()` as the domain of a relationship

that reads nothing; mapping $f : B$ is accepted as legacy shorthand with a hint and a fix, and f , $f()$ and $f(())$ are one reference (ADR-0029, §IV.7).

What a relationship *is* is one of three **roles**, derived from two authored facts — whether it reads anything, and whether it has a realization — never persisted, never chosen, never an identity (ADR-0032 (FVD-0118), production’s `bdl_model::RelationshipRole`):

role	when	has a value at a tick?	how a formula uses it	may drive an output?
Rule	it reads something (its type is an arrow)	no — it is a function	applied: <code>dimByTilt(tilt)</code>	no — its type is an arrow
Value	it reads nothing and has a formula	yes	by reference: <code>brightness</code>	yes, when its concept and domain are the output’s
Source	it reads nothing and has no formula	yes, once the environment supplies it	by reference: <code>tilt</code>	yes, on the same condition

The rule has one home, the model, and is stated on every projection (protocol 0.20); Studio maps the enum and re-derives nothing. *States* — declared, open, invalid, valid, not applied, driven, bound, clocked — vary within a role and never move it. A Source is complete, not unfinished: nothing is missing from it, the environment provides its value once per activation, and it is drawn with a boundary bar on the environment side and the word *Source*, never dashed. The word *sensor* is not used: a button state and an external configuration value are Sources too, and what realizes a Source is deployment’s business (§IV.7).

Designer problem: stating that brightness depends on tilt before deciding how; distinguishing a value the product computes from one it observes. *Language concept*: signature, relationship, the three roles. *Formal mechanism*: a declaration with an interface and an optional realization (§IV.1); the unit-domain normalization and the source role as a realization state (§IV.7). *Owner*: `bdl-model` (`MappingBlock::role`), `bdl-compiler` (`MappingAnalysis.role`), the daemon (`MappingView.role`). *Hidden*: currying, unit elimination, the kernel’s single input stream.

Stopping between declaring a relationship and defining it costs nothing, and the product has several such stopping states, each drawn and explained rather than flagged as an error:

- a **declared** rule — a relationship that reads something and has no formula — drawn dashed with the one word *declared*; the design around it is checked as far as it can be, and other relationships may already apply it;
- a Source with **no value yet** in the simulator — the input control is drawn empty with a dashed outline; nothing is defaulted to false or 0, and the run does not step until every Source has a value;
- an **open** formula — one that waits on a concept whose value form is not yet chosen — a solid node with a hollow socket and the sentence *checked once Temperature’s value is decided*;
- a rule **nothing applies** — `reactive.rule_unapplied`, an informational finding, never an error — drawn with a hollow output socket and the words *not applied*, with the fix `rule.apply` offered: *Add a value that applies `dimByTilt`*, which creates `brightness : () -> Brightness = dimByTilt(tilt)` when each read concept has exactly one producing value, offers a choice when one has several, and is blocked with the reason when one has none;
- a formula with a **slot?** — an expression not yet written, ordinary draft text that is saved with the file and refused by elaboration until filled (Part VI);
- a file that **does not build yet** — the text is kept as typed, the graph shows the last version that did build, and the reasons are listed (ADR-0023, ADR-0030).

Designer problem: conventional tools have no legal position for an undefined relationship, so a signature-first strategy cannot be adopted even when it would help. *Formal mechanism*: an unresolved declaration is a declaration whose realization is none and clients are typed against its interface (§IV.1). Whether designers use these positions is an **open empirical question** (§VII.4).

In the lamp, `dimByTilt : Tilt -> Brightness` now exists as a *declared* rule — dashed on the canvas, not an error — and the design already says that brightness depends on tilt. A property such as *increasing in tilt* cannot yet be recorded: commitments are a slot in the kernel with no production authoring, and Part IV says exactly what that slot is.

Failure avoided. In a node editor the arrow cannot exist without something to compute; in a statechart the relationship is not a first-class object at all; in code it is a comment. Here it is the primary object, and its incompleteness

is a state the tool understands — one that later giving it a definition will *refine* without disturbing anything already established (later proved, §IV.1).

III.3 Giving a relationship a formula

Opening `dimByTilt` shows the Formula Composer.

A formula is written in a small expression language: arithmetic and comparison, `&&` / `||` / `!` (drawn as *and*, *or*, *not* in the Composer), `if ... then ... else`, match over on/off, optional and numeric values, `let`, calls of the **equation library** (`clamp`, `min`, `max`, `inRange`, `any`, `all`, `map`, `filter`, `sum`, `zip`, `contains`, `getOrDefault`, ...), rules `x => ...` as arguments, membership `x in [a, b, c]`, list and pair literals, quantity literals with units (`90 deg`), memory delay(`init`, `e`), carrying across domains `sync(domain, init, e)`, and the natural forms:

```
all reading in readings: reading in 10 deg .. 45 deg
map r in readings: r / 90 deg
tilt ?? 0 deg
```

The formula's inputs are the concepts of its signature by their names; a formula attached to a relationship that produces `Brightness` is written as a scalar expression and the elaborator supplies the wrapping — there is no constructor to type, and no other concept can be manufactured inside it. Units belong to literals: `90 deg` is the designer's, `tilt` is a `Tilt` and carries no unit of its own; changing the unit of a literal with the quantity kept is a spelling change, editing its coordinate is a change of meaning (§IV.4). The natural forms are one-way spellings of library equations: a binder local is the equation's lambda parameter, a range is `inRange`, `??` is `getOrDefault`, and nothing new is evaluated (§IV.3).

In the Code view the same compiler service completes at the caret, shows a card for the name under the pointer, jumps to a declaration and lists its references across files, and lays a file out canonically as one edit; the graph and the text are two views of one project and one service (Part VI). The **Formula Composer** shows the same formula as the compiler reads it — reference chips with socket glyphs, literals as a coordinate and a unit pop-up, dashed slots, operators, calls, binders over an indented body, a choice as `if` with its condition over `then` and `else` — with, for the selected component, the compiler's expectation and its reason (*Expected: an angle, because an angle ÷ an angle = a dimensionless quantity*) and the candidates of that kind. Every structured action is a text edit of the same formula. Richer forms — `let`, `match`, blocks, rules as arguments, collection literals — are drawn as text inside the Composer and edited as text (Part VI).

Refining locally. Opening `dimByTilt` shows the Formula Composer. The designer chooses *Function* → *clamp*, fills the slots — `clamp(tilt ÷ 90 deg, 0, 1)` — with the expectation and the units offered at each slot, and the rule checks. The canvas still shows one arrow. The rule is *not applied* until something applies it; the offered fix creates `brightness : () -> Brightness = dimByTilt(tilt)`, a `Value`, which is what the simulator shows and what an output can read.

Failure avoided. The formula is checked against what the relationship *promises*: an expression that divides a tilt by a time is refused in the relationship's own terms — *this produces an angular rate, but this relationship promises brightness* — and no formula can manufacture a value of a concept its signature does not announce (later proved, §IV.2). The unit `90 deg` belongs to the literal, and switching it to radians rewrites the number, never the meaning (§IV.4).

III.4 Where the environment enters: the Source

The lamp needs a tilt and a temperature from outside. The designer opens the Source sheet, chooses the concept the Source provides — `Tilt`, an existing concept, or a new one created in the same step — and names it: `tilt : () -> Tilt`, a relationship that reads nothing and has no formula. By the derived rule of §III.2 that makes it a **Source**: complete, not unfinished; the environment provides its value once per activation; drawn with a boundary bar on the environment side and the word *Source*, never dashed. Nothing about a sensor is written, and nothing about a sensor *can* be written here: a Source is an input for a concept the designer chose, and which device provides it is a deployment decision (§III.10).

Failure avoided. A raw ADC count or a register image is never part of the design's semantics. The design says `Tilt`; how a value of `Tilt` comes to exist on a product is a construction added below the design at deployment — provision by a raw reading and a pure transducer — that the design cannot observe (later proved, §IV.7).

III.5 Where the design leaves: the logical Output

A design computes values; it does not move hardware. Physical effect happens only through an **output** — output light : Brightness @interaction — and one **drive** connecting it to exactly one relationship that produces its concept in its domain — drive light by brightness: a Value, or a Source (a value the environment supplies may be passed straight through to a light); never a Rule, whose type is an arrow. The spelling is relational on purpose, *the light is driven by the brightness*, because that is what the designer means; it is not natural language, and drive light = brightness is accepted as legacy syntax with a hint. Where several behaviors would influence one output — a safety override and an interaction — they are combined by an ordinary relationship that becomes the one driver, and the combination rule (priority, blend, maximum, clamp) is written in the design where it can be read. A second drive on the same output is refused as *contested*, with the message on the output, in the product’s words: *this output already has a driver; combine the two brightness values before connecting it*. An output with no driver is reported until the design is executable.

Below the output is the **device** that realizes it on a board — device pwmLight : pwm_channel for light { realization pwm_duty8 } : a device kind, whose requirements the board must carry, and a *realization profile*, whose encoder turns the output’s value into the raw command that kind takes — and below that the board’s pins. None of it enters the behavior model: the output says Brightness, and whether the product dims a lamp by an 8-bit duty, a 4-bit duty or a relay level is a Deploy-page choice that changes no formula, no value and no simulation (§IV.7–IV.8).

Designer problem: two behaviors reaching for one light, and a runtime rule chosen silently; a design that would have to know its PWM resolution before it is a design. *Formal mechanism*: nominal output identity, the write-once drive edge, DriveWF, SingleDriver, CompleteOutputs; realization as a lowering that adds a pure encoder and a machine sink below the output (§IV.7). *Owner*: bdl-output, bdl-output::realization, the Deploy page. *Hidden*: the encoder and the machine sink — the arbitration itself is on the canvas.

One output, one driver. output light : Brightness @interaction is driven by brightness. The warning pulse, warmPulse, cannot be connected to light as well; the repaired design has one more relationship, lampTarget, whose formula says whether the pulse takes priority while warm, or the maximum is shown, and lampTarget drives light. The heater’s cutoff is the same shape: heaterTarget forces the demand to zero above the critical temperature, visibly, upstream of the output.

Failure avoided. Two behaviors reaching for one light with a runtime rule chosen silently — last wins, highest wins, the one drawn first — is the hidden arbitration the single-driver rule refuses (later proved observable, §IV.7); and an output that named its PWM resolution would bake the mechanism into the behavior, which is what the logical Output exists to prevent (§IV.7).

III.6 Time, only where the product needs it

The designer sees two kinds of temporal thing: a **quantity** that has a value whenever its domain is active, and an **occurrence** that may or may not be present at an activation. Underneath, the distinction is not one of type: every relationship is a stream under the tick semantics, and an occurrence is an optional value (§IV.5).

Memory. delay(init, e) is last activation’s value of e, and init before there was one. Every remembered value has an explicit first value, and the tool asks for it — what brightness does the lamp show before it has ever been tilted? — because it is a product decision, not an implementation default. A relationship that depends instantaneously on itself is reported at the loop, in the names of the relationships, with the suggestion to remember one of them.

Timing domains. “Contact and orientation move with the interaction; temperature moves with the environment” is a statement about which values update together, not about how often. It is recorded by naming a **timing domain** on a relationship (@interaction), with no rate attached. The first time a relationship reads across domains the tool stops at that reference: the two values update in different rhythms, and it asks how the reader should see the source — as the last value the source produced, with a stated value to use before the source has ever reported (sync(environment, init, e), chosen in the *Starts at* sheet), or by moving the reader into the source’s domain. Both are legitimate designs; the initial value of a safety interlock is exactly the decision worth writing down. Rates belong to deployment.

Temporal modifiers and contexts — designed, not offered. The earlier draft of BDL (archived in paper/monograph/archive/) described temporal phrases as qualifiers on relationships — p for 300 ms, after e by 2 s, while

p, until e, since e, once e, every 1 s, rise p, previous x, count e, hold x e — and **contexts** (StateHandlers): named product situations, activated by a condition or entered and left by occurrences, containing a local flow graph and nested contexts, so that “while the cup is held, brightness follows tilt” is a region rather than a transition table. The formal development elaborated and executed each of these over the one memory primitive (§IV.5): previous, hold, count, since, once, every and rise as self-delayed declaration shapes; a context as an activation declaration, an entry edge, gated and reset local state, an inactive default, and a conditional in the one relationship that drives an output. The reduction is **formally characterized for the tested cases** — condition-scoped activation, entry, reset-on-entry state, inactive default, event-latched activation with exit-wins, state-local output selection, nested selection with an output — and contexts with their own timing domain and independently clocked nesting were not examined (FVI-0009). Production offers neither the phrases nor contexts at the snapshot: a designer writes `delay` and `sync` directly, and the shapes above are written out by hand (ISS-0010). Their surface is a **design recommendation** carried here so that when they are built they are built as elaborations and not as kernel constructs.

Designer problem: history without explicit state; different rhythms without a scheduler. *Language concept:* memory with a first value; named timing domains; carrying across domains with a first value. *Formal mechanism:* `delay`, `sync`, strictly-before, the domain judgment (§IV.5). *Owner:* `bdl-reactive`, `bdl-lower`. *Hidden:* the tick, the schedule, the two-phase step.

In the lamp the phrase the designer would like to write — `hold while not Held` — is the temporal-modifier surface, not offered at the snapshot; the designer writes the shape the phrase would elaborate to, `brightness() = if held then dimByTilt(tilt) else delay(0, brightness)`, and is asked for the first value. A `Held` region containing `dimByTilt` is likewise the context surface, not offered; the activation condition is an ordinary on/off relationship and the gating is written in the driver.

Two rhythms. temperature updates with the environment, `tilt` and `held` with the interaction: two timing domains, `@environment` and `@interaction`, no rates. `heaterTarget @interaction` reads the critical condition from `@environment`; the tool stops at the reference and the designer chooses *carry across, starting at released* — `sync(environment, false, critical)` — writing down what the interlock does before the first temperature reading.

Failure avoided. A reader of a value from another rhythm that silently took whatever was latest would make the scheduler’s order observable — two orderings, two products (later proved, §IV.5). The crossing is a visible object with a stated first value, and the first value of a safety interlock is exactly the decision worth writing down.

III.7 Organizing a larger design: behaviors and components

Relationships can be gathered into a **behavior** — a named group, drawn as a box, collapsible to its boundary sockets, moved, merged and split with no change to the design and no re-check (ADR-0019 (FVD-0074 ... FVD-0082)). *Package as Component* turns a behavior into a **component**: a reusable design with a **promise** — what it requires, provides and takes as a parameter, which timing domains it is written against, which concepts it shares with the system — inferred from the boundary and widened by the designer, never narrowed below what the members read. An **instance** is one use of a component with fresh identity; a **binding** connects a provided port or a value to a required port or an open relationship, directly in one domain or carried across domains with a first value; a **version** is a copy with the same promise. Every project is a **system** — a design with no components is the degenerate one — and the flat design every check runs on is derived, never persisted (ADR-0021, ADR-0022 (FVD-0064 ... FVD-0073)).

```
component AdaptiveLamp {
  use concept Tilt
  use concept Brightness
  param clock main
  requires tiltValue : Tilt @main
  mapping dimByTilt : Tilt -> Brightness
  dimByTilt(t) = t / (90 deg)
  provides brightness : Brightness @main
  brightness() = dimByTilt(tiltValue)
}
instance lamp : AdaptiveLamp { main = interaction }
bind lamp.tiltValue = tilt
```

Inside a component, the relationship behind a required port is a Source of the body — provided through the port — and Studio wears the port’s word for it; an unbound required port of an instance is a Source of the system, a

simulation input. The role is a fact of one design, and the same declaration can have a different role in the body, in the flat design and at the top level (production’s relationship-roles matrix; Theorem H of §IV.6).

Designer problem: reuse without copying; organizing a large design before deciding its interfaces. *Formal mechanism:* groups as authoring metadata proved transparent; components as templates instantiated with fresh identity and bound by ordinary realization steps; flattening (§IV.6). *Owner:* .bdl/authoring.json (groups), bdl-system. *Hidden:* renaming, the origin map, the flat design.

Failure avoided. Reuse by copying a group of relationships duplicates identities and drifts; a component instantiated with fresh identity and bound by ordinary realization steps keeps one design in one semantics, and grouping — the authoring gesture before packaging — changes nothing the kernel checks (later proved, §IV.6).

III.8 The Standard Library, and computation supplied from outside

The Library is an **authoring catalogue**: items that create ordinary objects in the design in one transaction. A *Concept* item creates a concept with its value form. A *Source* item is a **preset** for the Source sheet: a Source is never created without a concrete concept, so the sheet’s one decision is the concept the Source provides — an existing concept of the design, chosen by identity, or a new one created in the same transaction — and the preset (*Temperature Input*, *Tilt Input*, *Button Input*, ...) suggests a concept name, a value form, a unit and a Source name without inferring anything from them. What is created is a () -> concept relationship with no formula — a Source by the derived rule, exactly as one made by hand. The library carries no role and no flag; instantiating an item is the same edit sequence a designer could perform, undone as one step. Item names and descriptions are localized by Studio; identifiers never are. The Library is not a device catalogue: what realizes a Source on a board is a deployment concern (§IV.7), and the device catalogue that concern needs does not exist at the snapshot.

Supplied computation — designed, not built. Some behavior is easier to state as code than as a diagram: recursive filters, estimators, self-tuning controllers. The design admits computation blocks written in the host language behind a signature the designer places first — `smooth : Distance -> Distance` as the specification an engineer works against — with the properties the kernel cannot derive for supplied code (determinism, totality, range, state size, the rate it was designed for) declared as validation obligations, and with one rule the kernel does enforce: supplied code may not drive an output (ADR-0005 (FVD-0050 ... FVD-0056)). The trust boundary is designed (docs/architecture/component-boundary.md in production); no implementation exists.

III.9 Simulating before any hardware exists

Simulating. The Simulate page lists the Sources — *tilt*, *temperature*, *held* — as the inputs, refuses to step until each has a value, and shows every Value and every output per tick, with a probe naming for each concept what carries it and, for a rule, where it is applied.

Everything to this point is board-independent, and the simulator runs the same reference evaluator the generated code is held to (Part V). The design is *behavior-valid* when every formula checks and every reference respects its timing domain, and *executable* when, in addition, every required output has exactly one driver; both are properties of the design alone (§IV.8).

III.10 Choosing a target and realizing the boundary

Only now does the design meet a board, and it meets it on the Deploy page, never on the canvas.

Devices and placement. The outputs are given devices — the light a PWM channel, the heater a switched load — and, on selecting a board, the page derives what the design needs from those device kinds (one PWM line, one digital output), proposes a placement on the board’s pins or reports on the requirement that could not be placed and names what occupies each candidate pin. A manual pin is a constraint on placement, not a change to the design; a larger board places the same design; nothing on the canvas changes. *Failure avoided:* a design that had to know its board’s pinout before it was a design; and, the other way, a board’s capacity smuggled into typing — feasibility is a separate, non-monotone verdict about the pair (design, board), later proved decidable by a solver that is sound and complete for this scope (§IV.8).

Realizing the outputs. For each device the page lists the realization profiles whose encoder fits the output’s concept — for the light an 8-bit or a 4-bit PWM duty, for the heater a level — with the three judgments a profile must pass: its

encoder is well typed, it fits Brightness, and the board carries the kind’s requirements. Choosing a profile changes the raw command the machine will receive and nothing the behavior can observe; the canvas, the inspector and the simulation are untouched (**production implemented and tested**; later proved for the model, §IV.7). *Failure avoided*: a transducer written as unchecked host code in the firmware, and an “admissible” device whose encoder fits and allocates but is not even typed — the hardening result of §IV.7.

Realizing the Sources. The same page is where the tilt and the temperature would be bound to a device profile — a raw reading and a pure transducer per Source — and it is the one step of the lifecycle that the tools do not offer at the snapshot. The construction is proved (§IV.7) and not built (production’s PRP-0001, ISS-0016); on a board today, a design with a Source is *refused* by the adapter rather than given an invented value, and the lamp as written stops here.

Generating and running. With the board chosen and every output placed and realized, `bdld compile --target rp2040_pico` writes, beside the target-independent core, the firmware that ticks the core at the base tick, activates the timing domains on the compiled schedule and applies each tick’s raw commands to the assigned pads on a Raspberry Pi Pico; `--target arduino_nano` does the same for an Arduino Nano from the same plan, with a different board table, a blocking tick and no collection arena. The adapter applies a duty of 102 to the light’s PWM slice, refuses a duty of 306 and holds the line, and leaves a sink untouched on a tick in which its domain is not due. Every one of those sentences is **production implemented and tested** through recording sinks and a cross-build; the correspondence from the raw command to the operation is modelled and proved one step below the command (§IV.7), and from the register to the light it is not modelled at all.

III.11 What the workspace says at each step

Figure 4 lists the states a design passes through, as the product reports them; each answers what is unresolved, what is settled and what can be done next.

State	What is settled	Where
declared	a rule is named and typed; others may apply it; no formula yet	<i>canvas</i>
open	a formula waits on a concept whose value form is not yet chosen	<i>inspector</i>
invalid / valid	the formula does not check / produces what the signature promises	<i>inspector</i>
not applied	a rule no value applies; the value that would is offered	<i>canvas, Simulate</i>
temporally valid	every remembered value has a first value; no instantaneous loop	<i>canvas</i>
clock-consistent	every read across timing domains has a transport and a first value	<i>canvas</i>
executable	every needed relationship valid; every required output driven exactly once	<i>outputs</i>
feasible	the chosen board can carry every derived requirement, and every collection fits	<i>Deploy</i>

Figure 4: States a design passes through, as the workspace reports them. All but the last are properties of the design alone; the last is a property of the design together with a board, and is re-established whenever either changes.

A relationship that is declared but not defined is shown as open, and the design around it is checked as far as it can be. A definition that does not check is marked at the definition, in the vocabulary of the relationship. A loop or a missing first value is marked at the wire. A missing driver is marked at the output. A conflict on the board is marked on the Deploy page, at the requirement that could not be placed. *Feasible* is separate from everything above it because it is the only state that depends on something other than the design: six PWM channels fit a Nano, a seventh may not, and nothing about the first six changes.

III.12 Looking underneath

Most of the time the compact surface is all a designer sees. **Explain**, a collapsed disclosure at the end of the inspector, opens the selected object to what it became: the canonical type and the kernel interface, `role: Source | Rule | Value`, the Core term a formula elaborated to, the memory cell behind a `delay`, the transport and the tick at which a source was last read, the drive edge and the requirements derived from a device. It is where the kernel vocabulary is allowed to appear, for the engineer who receives the design and wants to know exactly what was generated, and it is the answer to the hidden-elaboration risk of §VII.4: the surface stays simple because elaboration does real work, and the work is inspectable rather than mysterious.

III.13 The lifecycle, and where the evidence stops

The walkthrough above is the whole design path of Part II, and each step has a chapter of Part IV that derives it and an evidence strength that bounds it.

step	the designer	the tool checks	derived in	strength
concepts	names what the product is about, with a value form	identity, not representation; links only between equal concepts	§IV.2	formally proved
relationships	declares before defining; three derived roles	typing against interfaces; the derived role	§IV.1, §IV.7	formally proved
formulas	writes in the relationship's own terms; units on literals	typing under the signature's grant; dimensions; exact unit laws	§IV.2, §IV.3, §IV.4	formally proved; production approximates units in f64
Sources	chooses the concept an input provides	a Source is a realization state, not a kind	§IV.7	formally proved
logical Outputs	connects one driver: drive light by brightness	one driver per output; type and domain equality	§IV.7	formally proved
time	names domains; carries across them with a first value	causality; the domain judgment; strictly-before	§IV.5	formally proved
behaviors, components	groups; packages with a promise; instantiates and binds	flattening preserves every judgment	§IV.6	proved; modular semantics restricted
simulation	supplies Sources per tick	the reference evaluator	Part V	production implemented and tested
placement	chooses a board; pins by hand	requirements → a sound and complete solver	§IV.8	formally proved for the finite scope
output realization	chooses a profile per device	three-judgment admissibility; the lowering changes no behavior	§IV.7	formally proved; production implemented and tested
Source provision	(not offered)	provision is transparent to the design	§IV.7	formally proved, not implemented
generation, running	compiles for a target	the core against the evaluator; the adapter against the command trace	Part V	production implemented and tested; the adapter's operation modelled one step (§IV.7); the register and the world not modelled

Part IV — Deriving the Formal Model

Part III showed the architecture working. This Part answers why it is shaped that way. Each chapter starts from a requirement the lamp exposed — that `Tilt` must stay distinct from every other angle, that a declared relationship must be refinable without disturbing its clients, that a value read from another rhythm must not make the scheduler observable, that a logical Output must not know its PWM resolution — states the smallest construct that meets it, records the alternatives that were formalized and failed, and names the theorem or counterexample that decided. The chapters are ordered by *dependency*, not by the order in which the work was done: persistent declarations first, because everything else is stated over them; then what a type means; then what can be computed; then time; then composition; then the two sides of the physical boundary; then what may be decided only with a board in hand. The phase in which each result was obtained is given as provenance, and Appendix G is the chronology.

Two conventions hold throughout. Every theorem named is a Lean declaration in the formal development, and Appendix B indexes them by concept; every “minimal” is minimal among the tested candidates. The notation is Appendix A.

Method

The kernel was obtained by a method, and the method is part of the record. Each phase of the development took a family of candidate constructs from the earlier draft, formalized the smallest plausible version and its alternatives in Lean 4 without external libraries, and attacked each with the same questions. What does it reject that the others accept? What does it accept that it should not? Is it a special case of another? Which operations on a design are refinements under it, and which are edits? Constructs were promoted from the experiment modules to the kernel only after surviving, and the theorems about them were re-proved in the promoted form.

The formal development labels each conclusion with one of six strengths — formally proved, formally rejected by counterexample, reduction by proof, tested formulation redundant, engineering preference, not ruled out — and this document maps them onto the five labels of the front matter: the first three are **formally proved** (a counterexample and a reduction are theorems), the fourth and fifth are **informed by FV** or **design recommendation** according to whether a mechanized example bounded the choice, and the last is stated as open. Nothing in the development is a proven impossibility, and “minimal” always means minimal among the tested designs.

Scale and trust base

The development builds with Lean 4.33.1 with no sorry. The axioms used by every theorem are propositional extensionality and quotient soundness, the latter only through function extensionality and the choice-free rational quotient of §IV.4; classical choice is absent, and each phase re-audited the whole development for it. As of Phase 18 the sources are 68 modules: 11 in Core, 12 in Behavior, 19 in Surface, 2 in Validation, and 24 experiment modules holding alternatives, counterexamples and executed examples. Every trace, assignment, unsatisfiability result and executed example reported here was obtained by running a proved-sound interpreter or solver inside the proof checker. Several theorems are recorded as trivial by definition — the typing half of client stability is a one-liner, and the well-formedness of a refinement target is unused because the invariant was moved into the definition — and they are reported as such rather than presented as content.

IV.1 Persistent declarations and refinement

The lamp’s first step was a relationship declared and left undefined, and its later steps refined it — a formula, a timing domain, a drive edge — without touching anything that already depended on it. That requires an object that persists across those changes with a stable identity, an interface its clients can be checked against while its body is absent, and a precise line between the changes that preserve what clients established and the changes that do not. This chapter is that object and that line; it is the foundation every later construct is stated over (Phases 0–1 and the post-Phase-1 migration).

The foundational kernel object is a **design declaration**,

$$\text{DesignDecl} = \langle id : \text{DeclId}, \\ interface : \text{DeclInterface}, \\ realization : \text{Option Expr} \rangle,$$

where an interface is an expected type together with a set of public commitments:

$$\text{DeclInterface} = \langle expectedType : \text{Ty}, \\ commitments : \text{PropertyId}^* \rangle.$$

A design is an environment $\Delta : \text{DeclId} \rightarrow \text{Option DesignDecl}$. An unresolved declaration is one whose realization is none; nothing else distinguishes it. Display names are not part of the kernel.

Typing through the type view

Terms refer to declarations by their constant name δ . The typing judgment $\Theta; \Delta; \Gamma \vdash_G e : \tau$ takes a concept environment Θ and a grant G , both introduced in the next section, and a declaration environment Δ which it consults through exactly one projection, the type view $\Delta \ni \delta : \tau$, the expected type of δ if declared:

$$\frac{\Delta \ni d : \tau}{\Theta; \Delta; \Gamma \vdash_G d : \tau}.$$

This is the only rule that reads Δ . A client of δ is therefore typed against δ 's interface and never against its body, whether or not a body exists. Inference is syntax-directed and decidable, and an inference function is proved sound, complete, and unique against the judgment.

Satisfaction, well-formedness, and refinement

A realization e **satisfies** an interface $\mathcal{P}$ when it has the expected type under the grant of that type and discharges every commitment:

$$\text{Satisfies } ev \ \Theta \ \Delta \ \Gamma \ e \ \mathcal{P} := \\ \Theta; \Delta; \Gamma \vdash_{\text{grant}(\mathcal{P}.\tau)} e : \mathcal{P}.\tau \\ \wedge \forall p \in \mathcal{P}.\mathcal{K}. \text{ev } \Delta \ e \ p,$$

where $\mathcal{P}.\tau$ is the expected type and $\mathcal{P}.\mathcal{K}$ the commitment list. Here $ev : \text{DeclEnv} \rightarrow \text{Expr} \rightarrow \text{PropertyId} \rightarrow \text{Prop}$ is an abstract **evidence** relation supplied by the validation layer. It takes the environment as an argument because compositional discharge needs it: “ A is monotone because B is committed to be monotone” consults B 's interface. A design is **globally well formed** when every stored declaration sits under its own identity and its realization, if any, satisfies its interface in that design.

Interfaces are ordered by monotone refinement: $\mathcal{P} \sqsubseteq \mathcal{P}'$ when the expected type is unchanged and the commitments of $\mathcal{P}$ are contained in those of $\mathcal{P}'$. A declaration takes a refinement step in one of three ways. An unresolved declaration may have its interface refined; an unresolved declaration may be realized with a satisfying body; and a realized declaration may have its interface strengthened, provided the body is re-verified against the new interface. The reflexive-transitive closure of these steps is exactly the structural order together with well-formedness of the target, where the structural order DeclLeq requires the same identity, interface refinement, and a write-once realization, and $\text{EnvRefines } \Delta_1 \ \Delta_2$ lifts it pointwise while permitting new declarations.

Client stability

Can a declaration be refined or realized without editing its clients, and without invalidating what was previously established about them? The answer has two halves with deliberately different hypotheses.

The typing half needs only the structural order. If B is declared in Δ and $\text{DeclLeq } B \ B'$, then every typing judgment in Δ holds in $\Delta[B']$. This is a one-line consequence of typing references through the type view, and it should be read as such: its content is that the decision to let clients see interfaces and never bodies is *sufficient* for client stability. It is also necessary. Change B 's expected type while keeping its identity, and every client breaks.

The commitment half needs more. If Δ is globally well formed, B takes a refinement step whose side conditions are checked in Δ , and the evidence relation is **monotone** — stable under EnvRefines — then $\Delta[B']$ is globally well formed, and the same holds for a whole lifecycle of B checked against the original environment. The monotonicity hypothesis was not part of the original design. It appeared when the theorem was attacked. An evidence relation that consults the *absence* of information — one that discharges a property because a dependency is still unresolved, say — is destroyed by a perfectly valid realization step, and the non-monotonicity of that evidence can be derived from the failure of preservation. Any discharge mechanism intended to survive refinement must therefore be positive in the environment.

Refinement versus edit

The preservation theorems cover refinement only. The table lists the operations examined and their classification; each row is witnessed by an example on a two-declaration design in which $A : \text{Nat} \rightarrow \text{Bool}$ is realized through an unresolved $B : \text{Nat} \rightarrow \text{Nat}$.

Operation on B	Kind	Effect on A
add a public commitment	refinement	typing and commitments preserved
realize	refinement	preserved; A now unfolds to a closed program
strengthen a realized interface	refinement, with re-verification	preserved
change the expected type, keep identity	edit	typing broken
drop a commitment	edit	typing silent; A 's commitment broken
replace B by a new identity	edit	dangling reference
add a new B' beside B	edit	A still depends on the stale B ; never executable
detach or replace the realization	edit	evidence that consulted the body is void
assign or change a clock domain	edit	domain judgment on clients broken
retarget an output binding	edit	completeness or single-driver may break

Two rows deserve comment. Dropping a commitment changes no type, so the type checker is silent, yet A 's own commitment was discharged through B 's and is now unsupported. Commitments are therefore part of the interface in the same load-bearing sense as the expected type. That is a stronger position than the earlier draft of the language took when it described properties as merely attaching to a name. Detaching a realization, which that draft permitted as an ordinary operation, is an edit for the same reason: clients' typing is unaffected, but any evidence that consulted the body is void. The kernel does not forbid edits. It declines to promise anything about them, and the tool must reopen the validation of transitive dependents.

What persistent identity is

The earlier draft left open whether a persistent, referable identity for an unresolved relationship is a novel abstraction. It is not. Every theorem that mentions identity uses it only to make an update land on the slot a reference resolves to, which is what a name does in any environment-based semantics. The non-trivial content lies in the environment order, and in the two facts that clients depend on it only through the type view for typing and only through monotone evidence for commitments. This is the familiar interface/implementation separation of module signatures, of a parameter later given a definition, or of a metavariable context with write-once assignment and fixed types, together with a Kripke-style stability condition on evidence. What is slightly non-standard is that the interface carries a growable commitment set whose growth is a first-class operation on a declared-but-undefined name, and that the kernel imposes a stability condition on the validation layer. Neither is a new type-theoretic mechanism, and no such claim is made.

IV.2 Concept identity and physical quantities

The lamp needs Tilt to remain distinct from any other angle-valued concept — a motor's MotorAngle , say — even though both are represented by the same physical quantity, and it needs a formula attached to a relationship that promises Brightness to be unable to manufacture anything else. Part IV.1 typed declarations without saying what a type *means*; this chapter adds the two things a product concept carries that a number does not — an identity that survives representation, and a physical dimension — and shows by the counterexamples that decided each why structural numeric identity is insufficient, why observing a concept's representation and constructing one must be treated differently, and why dimensions live in the types of operators rather than in a separate judgment (Phases 2–

3). What *equality* and *order* mean on concept values was forced later by the equation library and is settled in §IV.3; its conclusion — equality on data, order on quantities and declared-ordered concepts only — belongs to this chapter’s picture of a concept.

Nominal concepts

Suppose concepts were represented only by their representation types, so that `Tilt` and `MotorAngle` are both numbers. Then the wire `motorTarget := tiltSensor` is well typed and the design is globally well formed, because nothing in the model records the distinction. This baseline was built, the wire was accepted, and three ways of recording the distinction were then tried against it.

The one that survived is a single nominal type constructor over an internal identity:

$$\text{ConceptId}, \quad \text{Ty} \ni C.$$

Two distinct identities are distinct types regardless of representation, so the invalid wire is rejected by the ordinary rules of the simply typed calculus with no additional judgment. An explicit relationship between concepts, `tiltToMotor : Tilt -> MotorAngle`, is an ordinary declaration of arrow type — a design relationship that is itself signature-first and may remain unresolved — and it appears in the term wherever a crossing occurs. It is not a cast, coercion, or conversion; the kernel has no such mechanism.

Concept identity is independent of declaration identity, of display name, of dimension, and of hardware, and each independence was tested rather than assumed. A rename preserves identity, whereas a model in which the name *is* the identity makes renaming destructive. Treating a concept as an ordinary declaration admits two category errors: the concept becomes usable as a value, and it can be realized by a number. Keeping identity out of the type as interface metadata, checked by a direct-wire rule, is evaded by η -expansion, since $(\lambda x. x) \text{ tilt}$ has the same flow with no direct wire. A compositional role judgment strong enough to close that gap has the rule shapes of typing over types-with-sem, and the one such formulation examined duplicated nominal typing without benefit. Flow-sensitive or relational semantic analyses were not formalized and are not ruled out; nominal sem is the smallest mechanism among the designs that were tried, not the only one possible.

Erasing all concept identities is sound — a well-typed semantic term is well typed at its representation — and the baseline is exactly what erasure leaves. Generated code is thus ordinary code; the semantic layer has no runtime residue.

Representation binding and the grant

Nominal identity alone leaves Sem values opaque. Without a way to observe a representation and construct a value, no mapping can be realized by a formula; under nominal typing alone a Sem value can only originate from a declaration of concept type. That is the right state before representation is added. The question is how to add it without destroying what identity just bought.

The obvious form — global $\text{rep}_C : C \rightarrow R$ and $\text{mk}_C : R \rightarrow C$ available to every term — destroys it immediately. $\lambda x. \text{mk}_{\text{Motor}}(\text{rep}_{\text{Tilt}} x)$ is a well-typed `Tilt -> MotorAngle` in the empty environment, with no declaration and no mapping; a motor angle can be manufactured from a literal; and the crossing can hide inside a body whose signature mentions no motor. Observation alone, with no construction, is safe but cannot realize a mapping. The surviving model separates the two:

- a **concept environment** $\Theta : \text{ConceptId} \rightarrow \text{Option Ty}$ binds each concept, write-once, to a representation that mentions no concept type and contains no function type;
- $\text{rep } e$ is typed at R whenever $e : C$ and $\Theta(C) = R$, everywhere;
- $\text{mk}_C e$ is typed at C whenever $e : R$, $\Theta(C) = R$, and the grant permits C .

$$\frac{\Theta(C) = R \quad \Theta; \Delta; \Gamma \vdash_G e : C}{\Theta; \Delta; \Gamma \vdash_G \text{rep } e : R}$$

$$\frac{C \in G \quad \Theta(C) = R \quad \Theta; \Delta; \Gamma \vdash_G e : R}{\Theta; \Delta; \Gamma \vdash_G \text{mk}_C e : C}$$

The grant G is a predicate on concepts. Client code is typed under the empty grant. The realization of a declaration is typed under $\text{grant}(\tau)$, the concepts in result position of its own signature τ . A value of `MotorAngle` can therefore be constructed only inside a declaration that announces `MotorAngle` in its signature, which is exactly where a reader of the design would look for it. A well-typed term constructs C only where granted C ; the hidden crossing above is rejected under the grant of an unrelated declaration and becomes legal, and visible, once `tiltToMotor` is declared; binding an unbound concept is monotone for typing, satisfaction, and global well-formedness; and rebinding a concept to a different representation is an edit that breaks existing realizations.

Two constraints on the representation were not anticipated. Representation types must be free of concept types, because if `Tilt` may be represented *by* `MotorAngle` then `rep` itself is a hidden mapping under every policy, including observation-only. And they must be data types, a requirement that arrived later from the reactive semantics: a `Sem` value may be delayed, and a function-typed representation would carry a closure across ticks.

The grant is a known shape — a capability attached to a definition site, or equivalently the private constructor of an abstract type exported only to the module that declares it. What is specific here is where the capability comes from: the signature the designer already wrote, so no annotation is added. It is not presented as a capability calculus. One consequence should be stated plainly. After all bodies are inlined into one executable program, that program is checked under the universal grant, because each construction was authorized at its own declaration. Semantic isolation is a property of the design graph and survives inlining as provenance, not as a type property of the executable.

Physical dimensions

Physical quantities have type Q_d for a dimension d , an exponent vector over a small set of base dimensions. There is no dimension-specific typing rule; the algebra lives entirely in the types of registered primitive operators,

$$\begin{aligned} \text{add}_d &: Q_d \rightarrow Q_d \rightarrow Q_d, \\ \text{mul}_{d_1 d_2} &: Q_{d_1} \rightarrow Q_{d_2} \rightarrow Q_{(d_1+d_2)}, \\ \text{div}_{d_1 d_2} &: Q_{d_1} \rightarrow Q_{d_2} \rightarrow Q_{(d_1-d_2)}, \end{aligned}$$

and an application is checked by ordinary function application. Erasing every dimension to the zero vector is sound and accepts `length + time`, so the untyped numeric baseline is the erasure of dimensional typing in the same sense that the representation baseline is the erasure of nominal typing. Dimensions might instead have been validation metadata; the argument against is that multiplication and division *produce* dimensions, so any checker recomputes the same inference, but that family was argued against rather than excluded.

Dimension and concept identity are orthogonal. `Tilt` and `MotorAngle` both bound to `q Angle` remain distinct types. A mapping realized by the dimensioned formula $\lambda x. \text{mk bright}(\text{rep } x \cdot \text{gain})$ with $\text{gain} : q(\theta - \text{Angle})$ is typed; a dimension error inside the formula is caught by the same typing; and the formula cannot manufacture a `MotorAngle` despite the shared dimension. The association between a concept and its dimension lives in Θ , not in the identity and not in the type constructor. The earlier draft's two-index $\text{Sem}[n, d]$ becomes C together with $\Theta(C) = (Q_d)$.

Units are surface (§IV.4 in full). A literal `n u` elaborates to a dimensioned literal scaled by the unit's factor; changing the unit changes the value, never the type, and mixed-unit addition works after elaboration. Expressing a quantity in a unit — its *coordinate* — and building a quantity from a coordinate are the same arithmetic against a unit constant: `inUnit(q, u)` is q divided by the unit's scale and has dimension zero, `withUnit(x, u)` is x times the scale and has the unit's dimension, and a conversion between two units is their composition; each is elaborated, none is a kernel construct, and a unit choice never reaches a type (`1 m` and `100 cm` are equal values of one type). The unit laws — round trips, derived conversion, dimension safety — are proved over an abstract scalar domain and instantiated by a symbolic group in which π is a generator, so that a degree is exactly $\pi/180$ radian; the executable kernel truncates to naturals and production approximates in floating point. A preferred display unit is presentation: it changes the number shown and no judgment of the design. Affine scales such as degrees Celsius are not linear (0

$^{\circ}\text{C}$ is 273.15 K), yet their literals and coordinates elaborate exactly by adding an offset. Unit coordinates erase chart identity while preserving affine coordinate change: the conversion between two charts is an affine map, conversions compose and invert — compatible charts are isomorphic coordinate systems — and differences inherit the linear part of that transformation, so a difference of ten degrees Celsius is eighteen degrees Fahrenheit from any base point, while a conversion with a non-zero offset is not an additive homomorphism. The same abstraction covers sensor calibration and encoder offsets. These laws are proved over an abstract field and instantiated exactly at rationals; production floating point is held to a toleranced version of them. Whether a sum of two absolute temperatures should be permitted is a separate, optional validation question that the dimension does not decide and conversion does not need.

The concept ladder (Phases 19–22)

An audit late in the development (Phase 19) asked whether several declarations may produce values of one concept, and found that the kernel had always allowed it and that no judgment resolves a value by concept: the only ambiguity lay in a canvas that drew a concept as a node with a value. One phase (20) tried the opposite discipline — one producer per concept as a global invariant beside the single driver — and proved it preserved under refinement and composition; it was then set aside (FVD-0161 $\rightarrow$ FVD-0163) because the value already has a place: the declaration. The reading that closed the question is the **concept ladder** (FVD-0163, FVD-0164). A representation ($\Theta C = R$) is the type of a type. A **concept** — a `ConceptId`, the nominal type C — is a *type*, a template. A **Sem block** — a declaration of type C — is an *instance*: one value per tick, whose write-once realization is its **mapping block** and one producer, absent for a Source. A **value** $\text{mk}_C v$ is the instance’s state at a tick. In parallel, a rule (an arrow-typed declaration) is a template and a mapping block its instance. C reads “a Sem of C ”. Several Sem blocks of one concept are ordinary — `sensorA`, `sensorB`, `roomTemp` : `Temperature` — and every reference is by declaration identity; the theorems that state this (`producedBy_unique`, `reads_iff_dependsOn`, `new_sem_transparent`, `Surface/Sem`) are restatements of Phases 0–1, 5 and 13, which is the point. Phase 22 fixed the vocabulary in the code (`SemanticId` became `ConceptId`; theorem names keep their historical spelling, “semantic identity” in a name reading “concept identity”) and removed the proofs the misreading had produced.

IV.3 Data and equations

`clamp(tilt / 90 deg, 0, 1)` is the smallest formula the lamp writes, and a product needs more: a mode tested against a finite set of modes, every reading below a threshold, a pair of readings, a calibration mapped over a collection. The kernel of §IV.1–IV.2 computes with booleans, counts, quantities, nominal concepts and optional values and abstracts with lambdas over them; it is enough to state the theorems about identity, refinement, time and outputs, and not enough to write those equations. This chapter is the search for the smallest typed data and function basis that supports them without turning BDL into a general functional language — what entered the kernel and why each item could not be derived, what stayed a definition, which candidates were rejected, and the surface forms proved to add nothing (Phases 9b, 9c, 11).

The question and the hypothesis

The design direction was stated before the experiments: a very small core, parametric polymorphism, ordinary structured data, a rich definitional standard library, and intent-oriented abstractions at the UI. The hypothesis to test was a core of roughly `bool`, `nat/q d`, `sem s`, $A \rightarrow B$, $A \times B$, `opt A`, `list A`, plus rank-1 polymorphism — with the instruction not to assume any proposed feature belongs in the kernel.

The required cases were fixed in advance and every one is executed in `BDL/Experiments/EquationExamples.lean`: (A) `clamp` a brightness; (B) a mode in a finite set of modes; (C) every sensor below a temperature threshold; (D) any fault above a severity threshold; (E) a pair of temperature and humidity; (F) a calibration mapped over sensor values; (G) two sampled collections zipped; (H) an optional fallback; (I) dimension-preserving `min`; (J) semantic-type-preserving generic functions; (K) range membership; (L) a piecewise rule combining a range, a collection predicate and a boolean; and (M), added when production’s `enum LampMode { Off, Automatic, Manual(Brightness) }` was inspected, an enumeration with a payload.

What entered the kernel, and why each item could not be derived

Four additions were made to the kernel; each was tested against a derivation first.

Products. $\text{Ty.prod } a \ b$, Value.pair , and the operators $\text{pair} : a \rightarrow b \rightarrow a \times b$, fst , snd . A pair is data exactly when both components are (Ty.prod_data). The derivation that was tried and rejected is the Church encoding. It fails twice. First, a Church pair is an arrow, and arrows are not data: nothing of function type can be delayed or transported ($\text{arrow_not_delayable}$, **formally proved** by inversion of the delay and sync typing rules, which require $\tau.\text{Data}$). Paired *state* — a delayed reading with its timestamp — therefore needs a data product. Second, a Church pair used as a first-class value needs rank-2 types: in the toy System F of `PolyAlternatives.lean`, the type of fst on Church pairs has rank 2 (church_fst_rank), and in the prenex fragment a pair instantiated at one result type serves only one projection ($\text{church_pair_prenex_one_projection}$). Products are value composition only. They are never a component interface, an output bundle or a system boundary — Phase 8b’s Counterexample 5 shows what a tuple-returning declaration does to the dependency graph.

The list recursor. $\text{Expr.fold } f \ z \ l$, with $\text{fold } f \ z \ [x_1, \dots, x_n] = f \ x_1 \ (\dots (f \ x_n \ z))$, is a *term former*, not a registered operator. The kernel has no recursion, deliberately; a total language needs an eliminator for its inductive data, and this is the one construct in BDL that applies a function value in the course of evaluation. Registered operators still never apply closures. Its evaluation rule unrolls syntactically through the environment:

$$\frac{\Delta; I; t; \rho \vdash f \Downarrow v_f \quad z \Downarrow v_z \quad l \Downarrow []}{\text{fold } f \ z \ l \Downarrow v_z} \quad \frac{f \Downarrow v_f \quad z \Downarrow v_z \quad l \Downarrow x :: xs \quad [xs, v_z, v_f] \vdash \text{fold } \#2 \ \#1 \ \#0 \Downarrow r \quad [r, x, v_f] \vdash \#2 \ \#1 \ \#0 \Downarrow v}{\text{fold } f \ z \ l \Downarrow v}$$

The recursive premise evaluates the syntactic term $\text{fold } (\text{var } 2) \ (\text{var } 1) \ (\text{var } 0)$ in an environment holding the three values; this keeps Ev an ordinary inductive relation with no mutual recursion, so every earlier proof by induction on Ev extends by one case. Totality is a separate lemma by induction on the list (fold_total , mfold_total), and the fundamental theorem’s fold case uses it. Every collection operation — map , filter , any , all , contains , append , sum , zip , and, through toList , the option eliminators — is a definition over it (below). The alternative of one primitive per operation was rejected because a primitive cannot apply a closure and each would need its own evaluation rule; the alternative of bounded unrolling was rejected because lists (the Phase-9a buffer) are unbounded.

Equality at every data type. $\text{eq } \tau \ (h : \tau.\text{Data})$: structural equality on data values — booleans, numbers, none/some, pairs and lists componentwise, Sem values by concept and representation — with the proof of $\tau.\text{Data}$ carried *in the syntax*. This is the kernel’s only capability evidence: an equality on a function type is unwritable rather than ill-typed, which keeps the prim typing rule unconditional. Before Phase 9b, equality existed only at quantities, and production encoded boolean equality as $(a \wedge b) \vee (\neg a \wedge \neg b)$. On first-order values structural equality is equality (Value.beq_iff , **formally proved** by a mutual induction over the nested value type).

Two first-order operators. $\text{drop } \tau$ (the dual of take) and $\text{toList } \tau : \text{opt } \tau \rightarrow \text{list } \tau$. The second is the one that matters: without it an optional value has no eliminator that does not need a default — getD needs a fallback of the payload type, which a generic mapOpt cannot produce — and with it fold over a list of length at most one is the option eliminator (optElimF , mapOptF). drop lets zip be derived.

Every earlier theorem — determinism, totality in one and many domains, provenance, unfolding, the Phase-8 preservation results, the Phase-9a buffer — was re-established without change of statement. The logical relation gained the product clause; the fundamental theorem gained the fold case; the Wiring fragment gained a variant with variables for the recursor’s environment-passing sub-derivations (Ev.noCloV , Ev.foldCons_move).

Polymorphism: rank-1, by families, no kernel type variable

The five models compared, with the verdicts:

model	verdict	evidence
A — monomorphic STLC kernel	kept, unchanged	typing rules unchanged; HasType.unique
B — per-type duplication	what the kernel <i>sees</i> after elaboration	$\text{instances_are_monomorphic}$: three uses of min are three kernel terms
C — rank-1 parametric	adopted, as definitional families	every library entry is $\text{Ty} \rightarrow \text{Expr}$ (or $\text{Dim} \rightarrow \text{Expr}$); *_typed proves its scheme at every instance
D — System F $(\wedge, [\tau], \vee)$	rejected	the toy FTy with rank measures what it would add; its prenex fragment <i>is</i> instantiation of families

model	verdict	evidence
E — higher rank	rejected	every candidate use has rank ≥ 2 (<code>applyBoth_rank</code> , <code>existential_rank</code>) and a rank-1 replacement (<code>applyBoth_</code> <code>replacement</code>)

Why C needs no kernel support: a use site always has *closed* argument types. Every declaration’s expected type is frozen and closed since Phase 1, and `infer` is bottom-up, so finding the instance of a scheme is one-way *matching* of the scheme’s pattern against closed types — decidable, returning the unique substitution on the pattern’s variables (`matchTy_sound`, `matchTy_complete`, **formally proved** in `Surface/Poly.lean`). There is no unification of two open types, no let-generalization inside expressions, and no principal-type search: those problems arise when a definition’s type is inferred from its body, and BDL definitions carry their signature (mappings do; a helper `fn` would). Dimension polymorphism (`sum : list (q d) → q d, min at q d`) uses the same mechanism with dimension pattern variables; no kind system, no `Type + Dim` universe, because the dimension algebra already lives in the primitive table.

The scheme model: `PTy` is the kernel’s type formers over type variables and dimension patterns (`PDim.const d | PDim.dvar n`); `Subst` maps both kinds of variable; `Scheme = (pattern, caps)` where `caps : List (Nat × Cap)` names, per variable, what it needs; `Scheme.instantiate 0 0 τ` matches then checks the capabilities and is sound (`Scheme.instantiate_sound`). The two failure points have designer-level explanations — *no instance* (“expected a collection”) and *capability failed* (“cannot compare functions”; “Mode values can be compared for equality, but they have no default order”) — and a nominal mismatch is reported by the kernel’s unique typing as “Brightness and Opacity are different concepts”, never as a unification residue, because there is no unification.

Capabilities: the audit that reverted a decision

Phase 9b’s first form generalized both `eq` and `lt` to every data type through a structural order on values: booleans `false < true`, numbers, `none < some`, pairs and lists lexicographically, concepts by representation. It was formally consistent — `Red_prim` held, determinism held — and the closed capability vocabulary collapsed to `{Data}`. Phase 9c audited it and rejected it.

The audit asked, for each type, whether `==` and `<` have a domain-natural meaning for a behavior designer:

expression	meaning	verdict
<code>temperature1 == temperature2, <</code>	magnitude comparison of one quantity	<code>Eq, Ord (q d; q Length < q Time</code> still rejected)
<code>brightness1 < brightness2</code>	the concept is a magnitude the designer declared ordered	<code>Ord by declaration</code> , through the representation
<code>model == mode2</code>	same mode	<code>Eq</code>
<code>model < mode2</code>	none — any order would come from a code, a constructor tag or a <code>ConceptId</code>	rejected
<code>pair1 == pair2</code>	same reading	<code>Eq</code>
<code>pair1 < pair2</code>	lexicographic order is a mathematical convenience with no design meaning	rejected
<code>list1 == list2</code>	same collection in the same order	<code>Eq</code>
<code>list1 < list2</code>	none	rejected
<code>optional1 == optional2</code>	both absent, or both present and equal	<code>Eq</code>
<code>None < Some x</code>	a constructor-tag artifact	rejected

So `Data ⇒ Eq` holds (extensionally on this type grammar: `Cap.eq_iff_data`), but `Eq ≈ Ord`. An implementation may need a total order on values for maps, canonical forms, serialization and tests; that is toolchain-internal and is not `<` in BDL. The kernel’s structural order (`Value.blr`) was deleted and `lt` restored to `lt (d : Dim)` on quantities — the Phase-4 form — which made Phase 9b *smaller*. The surface vocabulary became `Cap = data | eq | ord (Poly.Cap)`: `data` is what `delay/sync` need and what the `eq` proof field checks; `eq` coincides with `data` today and is kept as a separate name because it answers a different question and is what diagnostics say; `ord` is a *surface* capability — a quantity, or a concept the designer declared ordered (`OrdDecl`) and represented by a quantity (`Ty.ordB`) — with no kernel counterpart, because an ordered concept compares as `lt d` on `rep`, a term the kernel already admits. Enumerations follow the same rule: equality is natural, declaration order is never silently behavioral order.

Ordering in the library is evidence-indexed: `Ordered` τ is $q \ d$, or $sem \ s \ d$ for a concept declared ordered with representation $q \ d$ (well-formed against Θ : `Ordered.WF`); `ltAt o a b` is `lt d a b` on a quantity and `lt d (rep a) (rep b)` on a concept, so `min`, `max`, `clamp`, `inRange`, `inInterval` return one of their arguments with identity intact. The comparator escape hatch `minBy`/`maxBy` needs no declaration and loses nothing: `minBy_recovers_min` (**formally proved**) shows the comparator $\lambda a \ b. \ a < b$ makes `minBy` compute exactly `min`. Negative examples on realistic concepts: `mode1 == mode2`, `pair == pair`, `list == list`, `opt == opt` accepted; `mode < mode`, `pair < pair`, `list < list`, `None < Some`, `bool < bool` rejected at the surface (`lt_rejected`, `min_mode_rejected`) and unwritable in the kernel (`lt_only_on_quantities`).

Production adopted this as ADR-0025 and refined it as ADR-0026: two values of one concept compare as that concept (`==` always, `<` and the ordered equations only while the concept is declared ordered); two values of different concepts never compare, whatever their representations; a concept beside a *plain value of its representation* — `tilt < 10 deg` — is observed and compared as that representation, because the designer wrote a number. The remaining asymmetry is visible and intended: `min(o1, o2)` is refused while `min(o1, 0.5)` is not, and `Explain` shows the observation.

The definitional library

`Surface/StdLib.lean` is the reference for production's `bdl`-equations. Each entry is a closed de Bruijn term indexed by types: `idF` τ , `constF` $\tau \ \sigma$, `swapF` $a \ b$, `minF` o , `maxF` o , `clampF` o , `inRangeF` o , `inIntervalF` o , `minByF` τ , `maxByF` τ , `foldrF` $\tau \ \sigma$, `anyF` τ , `allF` τ , `containsF` $\tau \ h$, `mapF` $\tau \ \sigma$, `filterF` τ , `appendF` τ , `sumF` d , `optElimF` $\tau \ \sigma$, `mapOptF` $\tau \ \sigma$, `getOrElseF` τ , `zipF` $a \ b$; `listLit`, `oneOfE` (a finite-set literal is `contains` over a list literal); records as right-nested pairs with positional projections (`recTy`, `recE`, `projE`). `zip` is derived through `fold`, `take 1`, `drop 1` and `reverse`, folding over the reversed first list with the remaining second list and the accumulated pairs in a pair — not pretty, and the point is that it is derivable.

Every entry is a **combinator**: variables, literals, lambdas, applications, registered operators, the recursor and, since 9c, `rep` — no reference, no state, no transport, no `mk`. For combinators, proved once: typing is independent of the design and the grant and reads Θ only through write-once representation bindings (`HasType.comb_irrelevant`); the value is the same in every design at every tick under every input (`lib_eval_context_free`, from the kernel lemma `Ev.pure`: a pure term in a pure environment is context-free); the term is clocked in every domain (`lib_clocked`); nothing is constructed (`Comb.noConstruct`); and the four together as the inlining statement `lib_expansion`. This is what lets production inline an equation at each use without creating a declaration — a library entry as a declaration would be monomorphic and would enter the dependency graph (FVD-0094).

The evaluation specifications connect each entry to the mathematical function: `any_spec` (`List.any`), `all_spec`, `contains_spec` (`List.any (beq x)`), `map_spec`, `filter_spec`, `min_spec`, `max_spec`, `clamp_spec`, `inRange_spec` (by the compared magnitudes `Ordered.key`), each **formally proved** through a general `fold_spec`: the recursor computes `List.foldr g` whenever the step closure implements `g` on the reachable accumulators. From these, finite quantification is a fold: `forall_in_list` ($\forall x \in xs, \ P \times \text{iff all } xs \ P \text{ evaluates to true}$) and `exists_in_list`; a finite-set literal means membership (`oneOf_mem`, via `beq_iff`), and duplicates do not change the answer (`oneOf_dup_irrelevant`), so no uniqueness convention and no `Set` type. An interval is a pair with a convention and membership is a function of the pair. A predicate is an ordinary $\alpha \rightarrow \text{bool}$.

Nominality survives all of it: `generic_preserves_identity` — *any* family typed at $\alpha \rightarrow \alpha \rightarrow \alpha$, instantiated at concept s , rejects an argument of concept $s' \neq s$, the representations never consulted; `generic_preserves_dimension` for $q \ d$ vs $q \ d'$; `pair_projections_keep_concepts`; `map_keeps_concepts`; `eq_across_concepts_rejected`. Executed: `min` at `Brightness` typed and `Opacity` rejected though both are $q \ 0$ underneath (`exJ`), and the same for lengths and times (`exI`).

Enumerations, records, sets: what was encoded and what was deferred

Production's textual syntax declares enumerations with payloads and matches on them. Phase 9b encoded `LampMode` as a tag paired with an optional payload, with `match` as conditionals on the tag (`exM`, `executed`), and deferred a kernel `sum` type: nothing tested needs more than the encoding gives, and the cost of a kernel `sum` would be one more eliminator term former, exactly like `fold`. Production keeps user enums open (ISS-0005). Records are positional nested pairs; labels resolve to positions at elaboration; row polymorphism was rejected because no case needs a function generic over record shapes. Sets are list literals with `contains`; set algebra is list definitions when needed; nothing tested observes a canonical form.

The expressiveness ceiling

Total, first-order-data computation over booleans, quantities, nominal concepts, options, lists and pairs, with higher-order functions and one list recursor; generic definitions instantiated at closed types; no general recursion, no type abstraction in terms, no sums yet (encoded), no unbounded quantification. This is a design recommendation backed by the executed cases and the proved library, minimal among the tested candidates; it is not a minimality theorem. The seven questions the phase was asked have the answers: rank-1 by families constrained by {Data, Eq, Ord} is the weakest useful polymorphism; products belong in the data core; no Set type; finite $\forall/\exists$ reduce to folds, proved; no existentials (Phase 8a components hide by fresh instantiation); no user typeclasses; and the ceiling above.

The natural expression surface

Designers should be able to write

```
all reading in readings:
  reading in 10 deg .. 45 deg
```

while the semantics stays exactly the verified Phase-9 machinery. Phase 11 tested that binder syntax and closed ranges are *conservative* surface abstractions and proved it for a six-form fragment (Surface/Natural.lean). The surface model NatExpr has an embedded closed core term, a named binder local, application, the four binders all/any/map/filter x in xs : body with the element type chosen by local inference, a range x in $lo \dots hi$ with its Ordered evidence, and a coalesce $x \ ?? \ d$. Elaboration *desugar* is one-way, under a binder stack:

```
all x in xs: p      ↦ app2 (allF  $\tau$ ) (lam  $\tau$  p') xs'      (any, map, filter alike)
x in lo .. hi       ↦ app3 (inRangeF o) x' lo' hi'
x ?? d              ↦ app2 (getOrElseF  $\tau$ ) x' d'
```

A binder local is the kernel's lambda parameter: *desugar* resolves a name to its de Bruijn index — nearest binder first — and adds nothing. **Formally proved:** nearest-binder scoping, shadowing, an unbound name is an error rather than a free variable, ambient references are untouched (*desugar_local_nearest*, *desugar_shadow*, *desugar_unbound*, *desugar_core*); alpha-equivalence — renaming a binder and its occurrences to a name fresh in the expression and the stack leaves the elaboration unchanged (*desugar_rename*, *alpha*), which is what a Formula Composer generating fresh locals needs; no *mk* is introduced (*desugar_constructs*); typing of each form with the local *forced* to the collection's element type by inversion (*binder_local_type*) and range bounds forced to the value's nominal type (*range_bounds_forced*); evaluation exactly the library's (*binder_all_eval*, ... from *all_spec*, *any_spec*, *map_spec*, *filter_spec*; *natural_forall*, *natural_exists*; *range_eval* = $lo \leq x \wedge x \leq hi$); and clocks of the elaborated term are the operands' (*binder_clock*, *range_clock*).

Executed: the call form and the natural form elaborate to literally the same term; *angle in 10 deg .. 45 deg* after unit elaboration, with the boundary values checked; nested *all row in rows*: any v in row: $v > 0$; shadowing where the inner x ranges over ys ; alpha renaming under shadowing; *Tilt* ranges with *Tilt* bounds accepted, q *Angle* bounds rejected, and accepted through *rep*; a *MotorAngle* predicate rejected on a *Tilt* local; the required negatives (*all x in 5*: true, *filter x in xs*: 5, *angle in 2 s .. 3 s*, an unordered concept in a range, a binder variable outside its body). Verdicts: binders, *BinderKind*, the range node and *??* are surface *desugaring*; an interval type, general comprehension (generators, *yield*, *where*) and a general quantifier are removed — nested binders cover every required case, and the forms are finite list equations.

Production implements the forms (P11; ADR-0028's second amendment; **production implemented and tested**): *bdl-elab::formula::binder* lowers the three forms once to the equation library, the parser keeps them as their own nodes and the formatter keeps the spelling authored, binder locals are lexically scoped to the body, and the tests *natural_forms_lower_to_the_same_core_as_the_call_forms*, *binder_locals_are_elements_scoped_to_the_body* and *natural_form_mistakes_are_named_in_their_own_words* discharge the Phase-11 claims by differential elaboration; the Composer draws binders and ranges with local chips and *??* as a comparison node (*ComposeAction.binder*, *range*). The production recommendations recorded with the phase, which the implementation followed: *all*, *any*, *map*, *filter* and *in* as contextual keywords only in the binder head, so existing names keep parsing; the local visible in the body only; the natural form kept as authored and never reconstructed from *Core*; *..* binding tighter than *in* and looser than arithmetic, legal only as the right operand of *in*; local inference exactly *binder_local_type*; the Composer representing a binder as a node with a fresh local backed by *desugar_rename*; diagnostics in concept language ("5 is not a collection", "Mode values have no order, so *in lo .. hi* does not apply").

IV.4 Units, coordinates and charts

The literal `90 deg` in the lamp’s formula owns a unit; the reference `tilt` does not; and a designer who prefers to read tilt in radians changes a number on screen, not the design. Keeping those three things apart — the physical quantity, its coordinate in a unit, and a display preference — is the whole of what units need, and the chapter shows that none of it is a kernel construct: a unit is elaboration data, a coordinate is arithmetic against a scale, and an affine chart such as Celsius is a coordinate system whose conversions compose exactly. It also records the one conclusion the development revised, and why (Phases 10 and 10b), and the formal basis on which the Formula Composer of Part VI infers what a slot expects.

Three notions kept apart

notion	formal object	layer
physical quantity	a kernel value of type <code>q d</code> : a canonical magnitude and a dimension	kernel
unit coordinate	<code>inUnit q u = q / scale(u), a</code> dimensionless number	surface elaboration
display / authoring unit	<code>Presentation.preferred : ConceptId →</code> <code>Option Unit</code>	presentation

Collapsing any two of these is the mistake the phase was designed to avoid: a quantity with a unit inside it splits one quantity into many types; a coordinate that remembers its unit is a runtime tag nobody needs; a display preference that enters elaboration changes semantics when a designer changes what they like to see.

Linear units: no kernel construct

Six models were compared. Literal-only elaboration (A) is insufficient alone — a coordinate is needed for display, for normalization (`tilt` in degrees $\div 90$), and for export. `inUnit(q, u) : Scalar` by elaboration (B) and `withUnit(x, u) : Q[d]` by elaboration (C) were adopted: `inUnit` is `div q (lit d scale(u))`, typed `q (d - d) = q 0 (inUnitE_typed)` and rejecting a quantity of another dimension by inversion (`inUnitE_safe`); `withUnit` is `mul x (lit d scale(u))`, typed `q (0 + d) (withUnitE_typed)` and never a concept (`withUnitE_is_quantity`); a literal is `withUnit`. First-class runtime units (D) were rejected — no tested case delays, syncs, stores or compares a unit, and a UI dropdown is not a runtime value. Units in quantity types (E) were rejected — `1 m` and `100 cm` would differ in type while being one quantity (exA: same type, equal values). A kernel conversion primitive (F) was rejected — `convert x u v = inUnit (withUnit x u) v = x.scale(u)/scale(v)` (`convert_eq`, `ev_convertE`), transitive (`convert_trans`), identity on the same unit (`convert_self`), mismatch rejected never coerced (`inUnit_mismatch`). Unit operations construct nothing (`unitOps_no_construction`), and on a concept they go through `rep` — so `inUnit(TiltValue, deg)` obeys the representation discipline, and rebuilding a concept from a coordinate requires `mk` under the concept’s grant. Tilt and `MotorAngle` are both readable in degrees and stay distinct under every unit round trip (`nominal_distinct`); Brightness and Opacity likewise. Unit compatibility is not identity.

The numeric domain, in the open

The formal development says exactly where exactness holds. Unit laws are proved over an abstract scalar domain `Scalars K` — a commutative monoid with a division that cancels multiplication by a non-zero scale — and instantiated by `Sym`, the free abelian group on the generators `2, 3, 5, 127,  $\pi$` . Every registered scale is an element, exactly: `deg =  $\pi/180$  rad =  $\pi \cdot 2^{-2} \cdot 3^{-2} \cdot 5^{-1}$` , `inch = 127/5000 m`. Round trips (`inUnit_withUnit`, `withUnit_inUnit`), conversion and mismatch hold with π symbolic: `90 deg` in radians is $\pi/2$ (exD), not 1.5707.... The executable kernel’s magnitudes are naturals, so its registry uses integer scales relative to canonical sub-units — `0.1 mm, ms, arc-second, g, K/180` — under which `mm, cm, m, km, inch, deg, turn, ms, s, min, g, kg` and `K` are exact; round trip 1 is exact for every positive scale (`inUnit_withUnit_nat`) and round trip 2 exact when the scale divides the magnitude (`withUnit_inUnit_nat`), the strongest law integer division admits; radians are *not* in that registry because their scale is not an integer in any degree-compatible basis. Production uses IEEE doubles with radians canonical (ADR-0011; `bdl-elab::units`): neither round trip holds exactly there, and the document says so rather than hiding it (§VII.1).

A unit itself is `(id, dim, scale)` in a registry; its symbol is not part of it — two units with equal identity are the same unit whatever they are spelled — and `unitsFor reg d` is sound and complete relative to the registry (`unitsFor_sound`, `unitsFor_complete`).

Presentation, and the realistic formula

Presentation is a separate object; a design with a presentation is a pair, and every kernel judgment of the pair — typing, evaluation, dependencies, clocks — is literally the judgment of the design (`presentation_irrelevant_*`, by construction, `rfl`); what changes is the displayed number (`presentation_changes_display`). Ordering compares canonical magnitudes through `ltAt`, which mentions no unit; comparing *displayed* coordinates would not be safe, since integer display can identify distinct magnitudes (`display_may_identify_distinct`).

The recurring formula `Tilt → Brightness: tilt / (90 deg)` and `inUnit(tilt, deg) / 90` are both dimensionless and evaluate equal for every `tilt` (`normalizations_agree`, `exF`). The representation-aware form is the second — it names the unit the designer thinks in and keeps the divisor a plain number — and neither changes the concept: both consume `rep tilt`, and the result is wrapped as `Brightness` by the block’s own signature (ADR-0013).

Design guidance, not a theorem, for persistence: the *source literal unit* (`90 deg`) is semantic source — its scale enters elaboration — and lives in the formula text; the *concept preferred display unit* (`Tilt` shown as `deg`) is authoring metadata beside the display name and never enters elaboration; a *simulation-input display unit* is session state.

Unit ownership in authoring

The production Formula Composer made this a rule (ADR-0028): in `tilt / 90 deg`, `tilt` is a semantic reference and `90 deg` is a designer-authored quantity literal. **A literal owns an editable unit; a semantic reference does not; a display boundary may carry a presentation unit.** The reason is semantic ownership. A literal’s unit is part of what the designer *wrote* — its scale enters elaboration, so switching it with the quantity kept (`180 deg → 3.141592653589793 rad`) is a change of spelling that preserves meaning, and editing the coordinate is a change of meaning; both are the designer’s to make on the literal. A reference’s kind is its declaration’s: `tilt` is a `Tilt`, represented as an angle, and no unit is part of that; offering a unit pop-up on a reference would invite a rewrite of the representation binding of a concept from inside one formula, which is exactly the representation escape hatch the grant discipline forbids. The display of a reference’s *value* — in simulation, in the inspector — may use the concept’s preferred unit, which is presentation. This decision is traceable to `unitOps_no_construction`, `withUnitE_is_quantity` and `presentation_irrelevant_*`, not to taste: the Composer draws a literal as two fields (coordinate and unit pop-up) and a reference as a chip with its concept’s socket glyph, because those are the two semantic situations.

Affine units: the failure of scale-only units, and the revision

Temperature scales are affine: `canonical = scale · x + offset`. Phase 10 showed, executably, that the linear model cannot represent °C or °F — `0 °C = 273.15 K`, which no scale produces from 0 (`celsius_not_linear`) — and that an affine *literal* and an affine *coordinate* nevertheless elaborate exactly with no kernel change: `n °C ↦ lit Temp (n · s + off)`, `inUnit °C q ↦ (q - off) / s`, in the `K/180` basis where °C and °F have integer scales and offsets (`affLitE_typed`, `affInUnitE_typed`, `affine_roundtrip_ev`; `0 °C = 32 °F`, `100 °C = 212 °F` in `scales_agree`). What breaks is arithmetic on *absolute* values: `20 °C - 10 °C` is a difference of 10 K (`delta_is_linear`), while `10 °C + 10 °C` is well typed at `q Temp` and reads 293 °C (`sum_of_points_is_not_a_point`, `sum_well_typed`). The kernel’s dimension cannot separate a point on the scale from a difference.

Phase 10’s conclusion, as first recorded: the missing information is a *sort* — point or difference — on top of the dimension (`AffSort`, with `affAdd/affSub` as the affine-space rules: `point + point` refused, `point - point = delta`), and a Formula Composer cannot offer only difference units for a delta slot from the dimension alone (`delta_candidates_need_sort`). Affine conversion was called “safe now”; affine *arithmetic safety* was said to need the sort and was deferred, with the sort treated as potentially necessary for future production.

Phase 10b tested the smaller hypothesis that for *conversion* chart-specific information is intentionally erased at coordinatization and only the affine transformation structure between coordinate systems must be preserved. It holds, and the conclusion was revised. This revision is recorded as research evidence, not smoothed over.

Charts: coordinate erasure and conversion functoriality

Over one physical dimension, a chart is `(scale, offset)`, valid when `scale ≠ 0`, with

$$\text{reconstruct}_u(x) = s_u x + o_u, \quad \text{coord}_u(q) = (q - o_u) / s_u.$$

The scalar domain is an abstract field (Field K: commutative, with inverses of non-zero elements, negatives and fractions) instantiated by `Q`, exact rationals built in the development as a quotient of `Int` fractions with every law proved from `Int`’s ring identities — because core Lean’s `Rat` proves its algebra with `Classical.choice`, which this development has kept out of every proof. Concrete equalities are decided by cross-multiplication.

Formally proved, over any field:

law	theorem
chart left inverse coord u (<code>reconstruct u x</code>) = x	<code>chart_left_inverse</code>
chart right inverse reconstruct u (<code>coord u q</code>) = q	<code>chart_right_inverse</code>
conversion is affine: $C(u,v)(x) = (s_u/s_v) \cdot x + (o_u - o_v)/s_v$	<code>convert_is_affine</code> , <code>convertMap</code>
identity $C(u,u) = \text{id}$	<code>convert_identity</code>
composition $C(v,w) \circ C(u,v) = C(u,w)$ — from the chart laws alone, no algebra	<code>convert_compose</code> , <code>convertMap_compose</code>
inverse $C(v,u) \circ C(u,v) = \text{id}$, both ways	<code>convert_inverse</code>
difference map $f(y) - f(x) = a \cdot (y - x)$	<code>difference_map</code>
offset cancels $f(x + \delta) - f(x) = a \cdot \delta$	<code>difference_offset_cancels</code> , <code>difference_converts_linearly</code>
linear-part functoriality $L(u,u) = 1$, $L(v,w) \cdot L(u,v) = L(u,w)$	<code>linear_part_identity</code> , <code>linear_part_compose</code>
not additive when the offset is non-zero	<code>not_additive_of_offset</code>
erasure keeps conversion structure coord $v \ q = C(u,v)(\text{coord } u \ q)$	<code>unit_erasure_preserves_conversion_structure</code>
display switch preserves the quantity reconstruct $v \ (C(u,v) (\text{coord } u \ q)) = q$	<code>display_switch_preserves_quantity</code>
coordinate edit changes the quantity	<code>coordinate_edit_changes_quantity</code>

Compatible charts form a groupoid of affine isomorphisms (theorem-level; no category-theory framework was built). Unit coordinates erase chart identity while preserving affine coordinate change; differences inherit the linear part of that transformation. A conversion with a non-zero offset is *not* an additive homomorphism, and this document never calls Celsius-to-Fahrenheit one.

Celsius and Fahrenheit, exactly. Kelvin canonical, `celsius` = $\langle 1, 27315/100 \rangle$, `fahrenheit` = $\langle 5/9, 45967/180 \rangle$: $C(^{\circ}\text{C}, ^{\circ}\text{F})(x) = 9/5 \cdot x + 32$ and $C(^{\circ}\text{F}, ^{\circ}\text{C})(x) = 5/9 \cdot (x - 32)$ for every x (`CtoF_closed`, `FtoC_closed`); $0^{\circ}\text{C} = 32^{\circ}\text{F}$, $100^{\circ}\text{C} = 212^{\circ}\text{F}$, $-40^{\circ}\text{C} = -40^{\circ}\text{F}$, the $^{\circ}\text{C} \rightarrow ^{\circ}\text{F} \rightarrow ^{\circ}\text{C}$ round trip, $^{\circ}\text{C} \rightarrow ^{\circ}\text{K} \rightarrow ^{\circ}\text{F}$ equal to $^{\circ}\text{C} \rightarrow ^{\circ}\text{F}$, $\Delta 10^{\circ}\text{C} = \Delta 18^{\circ}\text{F}$ from two base points, $\Delta^{\circ}\text{C} = 5/9 \cdot \Delta^{\circ}\text{F}$, and the non-additivity counterexample $f(0+0) = 32 \neq 64$ (executed, `Experiments/AffineExamples.lean`).

Not a Celsius special case. A 10-bit ADC over 5 V, millivolts, and a calibrated reading with a 0.5 V zero instantiate the same theorems — composition through millivolts, inverse back to raw, a difference independent of base (`exH`); so does an encoder with home offset, `angle` = $45/512 \cdot \text{count} + 30^{\circ}$ (`exI`). Unit conversion and sensor calibration are one affine-map abstraction.

The scalar coordinate is chartless. `32` is a Fahrenheit coordinate of 0°C and a Celsius coordinate of 32°C (`coordinate_needs_chart`); the destination chart supplied to `reconstruct/convert` interprets it; no runtime tag is needed, and no theorem requires a unit to be data, delayed, synced, stored or compared — Phase 10’s rejection of runtime units re-audited and confirmed.

AffSort, revised. Conversion never takes a sort and the sort checker never takes a chart (`sort_orthogonal_to_conversion`, `conversion_orthogonal_to_sort`, by construction). Two claims coexist: (A) unit conversion is complete and correct without `point/delta`; (B) a domain-specific checker may still use `point/delta` to reject `point + point`, as *optional physical-arithmetic validation*, orthogonal to A. The design decision FVD-0106 was rewritten from “affine conversion works, but `point/delta` is missing information” to “affine conversion is complete as coordinate-change semantics; `point/delta` is additional validation information for restricting physical arithmetic.” `Ty.q d` is unchanged; no conversion theorem needed a sort in the type.

The production chart model

Production (`crates/bdl-elab/src/units.rs`, ADR-0028 and its amendment) realizes the registry as `UnitDef { id, symbol, dim, chart }` with

```

pub enum Chart {
  Linear { factor: f64 }, // canonical = coordinate × factor
  Affine { scale: f64, offset: f64 }, // canonical = coordinate × scale + offset, scale ≠ 0
}

```

and coord, reconstruct, linear_part on the chart; $\text{convert}(x, \text{from}, \text{to}) = \text{coord_to}(\text{reconstruct_from}(x))$ is the one conversion operation and hides the chart's shape. The Phase-10b laws are property-tested on f64 — identity, composition, inverse, display switch preserving the quantity, the difference map linear — within a few ulps, including the Celsius/Fahrenheit charts; the formal laws are exact. The affine charts exist as tested infrastructure and are **not** in the registry offered to formulas or the Composer's pickers (ISS-0004): conversion and display of an absolute temperature are safe, arithmetic on absolute temperatures is a separate validation concern, and the product does not offer 20 °C in a formula until the point/difference validation exists or a decision is taken not to need it. The symbol is presentation; identity is the id; inch is spelled out because in is the membership keyword.

The production contract for floating point, as recommended by the formal phase: for registered charts u, v, w and sampled x , $|C(u, u)(x) - x| \leq \epsilon|x|$, $|C(v, w)(C(u, v)(x)) - C(u, w)(x)| \leq \epsilon|x|$, $|C(v, u)(C(u, v)(x)) - x| \leq \epsilon|x|$, $|C(u, v)(x + \delta) - C(u, v)(x) - L(u, v)\delta| \leq \epsilon|\delta|$, and $|\text{reconstruct}_v(C(u, v)(\text{coord}_u(q))) - q| \leq \epsilon|q|$, with ϵ a few ulps scaled by the largest coefficient, plus exact-rational oracles — never a claim of exact f64 identity or composition.

The Formula Composer's formal basis

A structured composer needs, from the formal side, exactly three things: typed holes, expected-dimension propagation, and candidate-unit inference. Surface/Composer.lean gives them and no more.

PExpr is holes, known operands by dimension (a literal, a reference or a known equation result), and add/sub/mul/div; it is never executed. check is bottom-up on the known parts. solve pushes an expected dimension down by the local rules of each operator in the Dim group — add/sub: both sides get the result; mul: the other side gets $r - d$; div: numerator $r + d$, denominator $d - r$ — deterministically and in one pass; a product of two holes is reported *unsolved*, never searched. **Formally proved** for the whole fragment: solve_sound (filling the holes as solved makes the expression check at the expected dimension) and solve_complete (any filling that checks agrees with the solution — a one-hole equation in an abelian group has one solution); candidates_sound, candidates_complete and slot_sound (the units offered for a hole are exactly the registered units of its solved dimension); refCandidates_sound for declaration references. Executed: ? / (1 s) : Speed $\Rightarrow$? : Length with candidates mm, cm, m, km, inch (exG); Force $\times$? : Torque $\Rightarrow$? : Length (exH); the explanation Speed = [Length] / [Time] as slots (speed_slot). No unification, no symbolic algebra beyond the group laws.

The Composer does not need to make a unit part of an expression's type: an affine quantity literal needs the physical dimension, the chosen chart and the scalar coordinate, and changing the displayed chart is $\text{coord}_u(q) \rightarrow \text{coord}_v(q)$ with q preserved. What it *would* need beyond the dimension to offer only difference units for a delta slot is the point/difference sort — which is why that sort is recorded as optional validation information the editor could carry, not as anything in the kernel.

IV.5 Time: memory, causality and clock domains

The lamp remembers its last brightness when set down, and its temperature updates once a second while its tilt updates fifty times a second. The first needs memory with an explicit first value and a rule for which self-references are legal; the second needs a notion of *which values update together* that is independent of how often, and a way for one rhythm to read another that does not make the scheduler observable. This chapter derives all of temporal semantics from one primitive — read a clock domain at its previous activation — and shows why signal and event types, task-style execution, same-tick visibility and hidden resampling were rejected or reduced, and how the lossless cross-domain buffer falls out of the same two primitives plus ordinary list data (Phases 4, 5, 9a).

A Minimal Reactive Semantics

One primitive

The reactive kernel adds one expression form and no types:

delay *init e*.

Its meaning is given by a tick-indexed big-step evaluation relation $\text{Ev } \Delta \ I \ t \ \rho \ e \ v$: the value of e at tick t under a local environment ρ , where $I : \text{DeclId} \rightarrow \mathbb{N} \rightarrow \text{Value}$ supplies the inputs. Every declaration is a stream by interpretation. An unresolved declaration is an input and takes its value from I ; a realized declaration is evaluated from its body at the current tick; delay $\text{init } e$ evaluates e at tick t when read at tick $t + 1$, and init at tick 0:

$$\frac{d \notin \text{def}(\Delta)}{\text{Ev } \Delta \ I \ t \ \rho \ (d) \ (I \ d \ t)}$$

$$\frac{\text{Ev } \Delta \ I \ 0 \ \rho \ \text{init } v}{\text{Ev } \Delta \ I \ 0 \ \rho \ (\text{delay } \text{init } e) \ v} \quad \frac{\text{Ev } \Delta \ I \ t \ \rho \ e \ v}{\text{Ev } \Delta \ I \ (t + 1) \ \rho \ (\text{delay } \text{init } e) \ v}.$$

There is no signal type in Ty . Under this semantics a signal type would be inhabited by exactly the terms of the underlying type and would reject nothing. The information a reactive type would carry in a multi-domain setting is *which clock*, and the next section places that information in a judgment rather than a type. There is likewise no event type. Within one domain an input delivers at most one value per tick by construction, so an occurrence is a stream of optional type, and the streams of type $\text{Option } \tau$ are exactly the streams of multiplicity at most one. What separates an occurrence from an optional value — two occurrences falling in one observation interval — can only be seen when a source ticks faster than its observer. That is a cross-domain question and is treated at the end of this part.

Two restrictions on delay were forced by the totality proof rather than chosen. The delayed type must be a **data** type, one with no function type inside, because a delayed closure would have to persist across ticks and the logical relation for closures is tick-indexed and cannot be transported. And delay may occur only at **top level**, under no binder, because a delay under a lambda would re-evaluate its operand at the previous tick in an environment created at the current tick. Temporal state therefore belongs to declarations, and mappings are pointwise. This is the arrangement of pre in Lustre, where it lives in nodes rather than in functions [5], and its practical consequence is that a reusable stateful component is instantiated into fresh declarations rather than abstracted over.

Causality

The dependency graph on declarations comes in three variants. Structural dependency, DependsOn , records every reference in a body. Instantaneous dependency, InstDependsOn , excludes references under the delayed operand of a delay (the initial value is read at tick 0 and counts as instantaneous). A design is **causal** when its instantaneous graph is acyclic, witnessed by a bounded rank:

$$\text{Causal } \Delta := \exists \text{rank}, R. (\forall d. \text{rank } d < R) \wedge$$

$$\forall a \ b. \text{InstDependsOn } \Delta \ a \ b \rightarrow \text{rank } b < \text{rank } a.$$

On the delay-free fragment this coincides with structural acyclicity, so the earlier acyclicity condition is the timeless special case rather than a replaced requirement. A structural cycle every one of whose paths passes through a delayed operand — $A := \text{delay } \theta \ B$; $B := A$, or a self-delayed accumulator — is causal and runs at every tick. A cycle that is partly delayed is not causal. A strict cycle, one through neither a delay nor a lambda, has no value at any tick.

Evaluation is a partial function: one tick, one environment, one term, at most one value, with no hidden evaluation order, and this holds unconditionally. It is total on causal designs: if Δ is causal and globally well formed and the inputs are well typed, every declaration has a value at every tick, related to its type by a logical relation, by an induction on tick, rank, and derivation. An executable interpreter is proved sound for the relation, and every trace reported in Part IV was obtained by running it.

One gap should be recorded. A cycle guarded by a lambda, $A := \lambda x. \ A \ x$, is rejected by Causal , yet $\text{declRef } A$ does evaluate — to a closure; only applying it diverges. Causal is conservative for lambda-guarded cycles, and the negative theorem covers strict cycles only.

Derived operators

Every temporal operator of the surface language reduces to delay and the primitive operators. There is no independent kernel definition of count for the reduction to be proved equivalent to; what was done instead was to elaborate

each operator, check its typing and causality, and run it on a concrete input trace. In each row the declaration refers to itself. Each is a self-delayed cycle, the class that structural acyclicity forbade and causality licenses.

Surface	Declaration shape
previous x	delay init x
previous x without init	delay none (some x); the absence is pushed to consumers
hold init e	getD e (delay init self)
count e	ite (isSome e) (1 + delay 0 self) (delay 0 self)
since e	ite (isSome e) 0 (1 + delay 0 self)
once e	delay false self v isSome e
every n	modulo- n counter over delay
rise b	$b \wedge \neg \text{delay false } b$, as an optional Boolean

State has no identity of its own. A cell is a **delay** in a declaration body, and nothing refers to it because consumers refer to the declaration. There is consequently no notion of two writers to one cell in this kernel.

State preserves concept identity and dimension. The typing rule is $\text{delay} : \tau \rightarrow \tau \rightarrow \tau$ for data τ , so a delayed tilt is a tilt and a backward difference over a time step has dimension **Length** – **Time** with no derivative primitive. Provenance holds in the reactive setting too: if no signature announces a concept and no input carries it, no value at any tick carries it. Temporal state carries tags; it never creates them.

Initialization is semantic, not validation. Every **delay** carries an explicit initial value. Two toy relations without one show why: the first tick is either undefined or nondeterministic. Adding or removing a delay, or changing an initial value, is an edit.

Clock Domains

Identity, not rate

“Contact and orientation move with the interaction; temperature moves with the environment.” A designer can say this before any rate is known, and it is a statement about which quantities are updated together, not about how often. BDL records it as a nominal **clock domain**, **ClockId**, and treats rate as validation data that never enters the kernel.

The time model is one global base tick and a schedule $S : \text{ClockId} \rightarrow \mathbb{N} \rightarrow \text{Bool}$ saying at which global ticks each domain activates. A period n induces the schedule $t \bmod n = 0$; the schedule lives outside the design. Domain-local time is not a separate counter but the sequence of a domain’s activations.

Each non-agnostic declaration is assigned a domain by a **clock environment** $K : \text{DeclId} \rightarrow \text{Option ClockId}$; a declaration with no domain is a pure mapping that may serve any domain. The clock is interface data in every sense that matters — clients’ validity depends on it, it is frozen under refinement, and changing it is an edit — and it is stored as a projection beside the interface, exactly as a concept’s representation is stored in Θ rather than in the type. Whether to fold it into the interface record is churn rather than semantics.

Rate and identity are distinct. A clone of a domain with the identical schedule is a different domain, and a direct wire between them is rejected; a domain at the same rate but shifted in phase reads different values through a transport. Rate changes are validation-only. They change the induced schedule and hence the observed values, but no client’s well-formedness. This is where BDL departs from synchronous languages that recover clocks by inference [6]. The domain is authored, because the information needed to infer it does not arrive until realization binding, which in this workflow is the point at which the designer is least able to make the decision.

One transport primitive

Cross-domain reading is the second and last temporal form:

$$\text{sync}_\kappa \text{ init } e,$$

the value of e , evaluated in domain κ , at the last activation of κ strictly before the current tick, and *init* if there has been none. The multi-domain evaluation relation $\text{MEv } S \Delta I \kappa t \rho e v$ indexes evaluation by the domain in which it takes place, and its two transport rules are

$$\frac{\text{prevAct } S \ \kappa' \ t = \text{none} \quad \text{MEv } S \ \Delta \ I \ \kappa \ t \ \rho \ \text{init } v}{\text{MEv } S \ \Delta \ I \ \kappa \ t \ \rho \ (\text{sync}_{\kappa'} \ \text{init } e) \ v}$$

$$\frac{\text{prevAct } S \ \kappa' \ t = \text{some } t' \quad \text{MEv } S \ \Delta \ I \ \kappa' \ t' \ \rho \ e \ v}{\text{MEv } S \ \Delta \ I \ \kappa \ t \ \rho \ (\text{sync}_{\kappa'} \ \text{init } e) \ v}.$$

delay is sync at the expression's own domain: $\text{delay } \text{init } e \equiv \text{sync}_{\kappa} \ \text{init } e$ in domain κ , as an equivalence of the two relations. The kernel therefore has one temporal primitive — read a domain at its previous activation — and the single-domain semantics of the previous section is its diagonal. Under the always-active schedule, MEv coincides with Ev in every domain, so the earlier results are the one-domain special case rather than a replaced machine. A delay in a slow domain reads three global ticks back where a delay in a fast one reads one, with the same syntax.

A domain judgment Clocked $K \ \kappa \ e$ rejects every other cross-domain reference: a reference stays in its domain or is agnostic, a delay needs a domain, and $\text{sync}_{\kappa'}$ switches the domain of its operand. Typing is unchanged and is blind to domains; the direct wire between two domains at the same value type is well typed and rejected only by the domain judgment. Placing the domain in the type instead was tried and set aside. Every pure mapping would need clock polymorphism, and nothing the type rejects is missed by the judgment.

Strictly before

A transport sees only source activations strictly before the destination tick. That is a choice with an observable alternative, and the alternative was built. A transport that lets simultaneously active domains see each other's current values makes the scheduler order observable: two priorities between the domains give two outputs. The strictly-before rule has no such parameter, and multi-domain evaluation is deterministic with no order between simultaneously active domains appearing in the semantics. It also makes cross-domain causality free. A transport's operand is never instantaneous, so Causal Δ is unchanged and no cross-domain cycle can be instantaneous. Every crossing costs one destination-visible step; “synchronous sub-domains evaluated in one instant” are, in this model, the same domain.

Every sync carries an explicit initial value, used at a destination activation with no earlier source activation. Totality extends to the multi-domain case — a causal, globally well formed design with well-typed inputs has a value in every domain at every tick — so the first activation is deterministic with the stated initial values. Concept identity and dimension pass through transport untouched, by the typing rule; a crossing from `Tilt@fast` to `Tilt@slow` authorizes neither `Tilt -> MotorAngle` nor `q Length -> q Time`.

Cross-domain occurrences and lossless buffering

Phase 5 left one problem open, and it is worth recording in full because its resolution shaped the data language.

The problem. Within one domain an occurrence is a stream of optional type: an input delivers at most one value per tick, so `opt τ` streams are exactly the streams of multiplicity at most one. Across domains this fails. `sync` is a zero-order hold: a slow consumer of a fast event source sees the last value only. Phase 5's Counterexample C (`opt_loses_multiplicity_under_sync`, **executable**) has two fast events at ticks 1 and 2 and one event at tick 2 indistinguishable at the slow activation at tick 3, and a single event at tick 1 followed by a quiet fast tick *dropped* outright. The counterexample is against `sync` as an *event transport*, not against optional types.

The window model, on tick sets. What the destination should see is the source's occurrences at the source activations since the destination's own previous activation — the *window* `windowTicks S src dst t = srcTicks S src (prevAct S dst t) t`. Phase 5 proved (`buffer_from_log_and_cursor`) that this window equals the source's accumulated log read at the current tick minus the log length read at the previous destination activation: two single-instant reads, a sync of a source-side accumulator and a delay of a cursor. Policies are functions of the window; `latest`, `count`, and `count with latest` each identify distinct windows (`policies Lose information`), and only the list is injective. So multiplicity and order are observable, buffering is required to keep them, and buffering is a structured use of the existing state basis *plus list data* — which the kernel then lacked.

The failed abstractions, generalized. Phase 9a tested the models the brief named. Model A, `latest`; B, `count`; C, a fold such as a sum; D, a fixed tuple of the newest entries; E, the list. A–D are lossy, with witnesses (`latest_not_lossless: [1,2] vs [2]`; `count_not_lossless`; `sum_not_lossless: [1,4] vs [2,3]`; `modelD_not_lossless`). The general

statement behind A and D is `bounded_summary_not_lossless` (**formally proved**): any summary depending only on the newest k entries, for any fixed k , identifies a k -entry window with a $(k+1)$ -entry window. A lossless summary is injective and therefore unbounded (`lossless_iff_injective`). The claim discipline is precise: not that the list is the only lossless representation, but that any lossless one must be injective on windows and hence unbounded in size, and that an ordinary list is the smallest general sequence representation tested.

The object-language buffer. With list τ (the `Ty.list` constructor, `Value.list`, and six registered operators `nil`, `cons`, `length`, `take`, `reverse`, `head`, with `length` yielding `q 0`), the Phase-5 construction is five ordinary declarations (`Surface/Buffer.lean`):

```
log    @src : cons src (delay nil log)      -- source-side accumulator
logD   @dst : sync src nil log              -- the log, transported
seen   @dst : length logD                  -- log length now
cursor @dst : delay 0 seen                  -- log length at the previous activation
window @dst : reverse (take (seen - cursor) logD) -- the new entries, oldest first
```

Nothing new is evaluated: `delay`, `sync` and registered operators. `logD` is the one cross-domain read, an explicit `sync` with the explicit initial value `nil`. **Theorem M** (`buffer_window_correspondence`, **formally proved**): for every schedule, every input, every destination domain and every tick,

$$\text{MEv } S \Delta I \text{ dst } t \llbracket \text{window} \Downarrow \text{list } (\text{map } (I \text{ src}) (\text{windowTicks } S \text{ src } \text{dst } t)) \rrbracket,$$

whenever the six declarations are realized as above and `src` is an input. The strictly-before rule is untouched — the window at t contains source activations $< t$ only. The elaboration is well typed in any environment with the declared types (`buffer_elaboration_well_typed`) and well clocked with `src`, `log` in the source domain and the rest in the destination (`buffer_elaboration_well_clocked`). The list summary is lossless by injectivity (`buffer_lossless`); its entries are indexed by the window ticks in strictly increasing order (`window_to_list_preserves_order`), and for every predicate on values the number of entries satisfying it equals the number of window activations whose value does (`window_to_list_preserves_multiplicity`).

Executed (`Experiments/BufferAlternatives.lean`, on Phase 5’s multi-rate schedule with fast events at ticks 1 and 2): at the slow activation 3 the window is `[none, some 1, some 2]`, at 6 it is `[none, none, none]`, while `latest` sees `some 2` at 3 (`buffer_trace`); the two histories Phase 5 could not tell apart under `latest` are told apart by the window (`buffer_distinguishes_what_latest_identifies`); the elaborated design passes the typing, clocking and causality checkers. The buffer is also a Phase-8a system with no new binding kind: a sensor component provides its event *and* its log, a consumer requires the log through an ordinary transported binding with initial value `nil` and computes the window (`transport_trace`), whereas binding the event itself through `sync` yields `latest` and loses the tick-1 event (`latest_transport_loses`). “Buffer or not” is a choice of *which provided port to bind*, visible in the interface.

Capacity is a validation obligation. The kernel model is the unbounded log. A deployment has finite memory and must show that no window exceeds its capacity. `CapacitySufficient S src dst cap T` states that every window up to horizon T fits; it is decidable for a finite horizon; `requiredCapacity` computes the least sufficient value with a proof (`requiredCapacity_sufficient`); and for periodic schedules the horizon is unnecessary — a window never holds more than one destination period of source activations (`periodic_window_bound`), so one period is a sufficient capacity at every horizon (`periodic_capacity_sufficient`). Overflow policies are explicit functions of the unbounded window — `dropOldest cap`, `dropNewest cap` — and under a sufficient capacity both are the identity (`sufficient_capacity_preserves`, `bounded_buffer_agrees`); under an insufficient capacity they change the trace (`negE`, `executed`: the three-entry window at tick 3 becomes two entries under either policy, capacity 2 fails the check at horizon 6, the required capacity at horizon 30 is 3). Consequently the only overflow policy that preserves the kernel semantics is to *reject the deployment*; dropping is a semantic change and must be written by the designer as ordinary computation over the window if wanted.

What is still not a kernel type. An event stream is a data-typed declaration in a domain; an occurrence is its value at an activation; a lossless cross-domain view of it is the five declarations above; `latest`, `count`, `coalesce`, `drop`, `sample` are ordinary computations over `window`. None of this needs an `Event` type, a buffer primitive, a scheduler order, same-tick visibility, an implicit overflow rule or an effect system, and none was added. The negative examples of the phase also settle a subtler point: a clone of `fast` with the identical schedule is a different domain, and listing the value does not make the direct read legal (`negF`) — equal rate is not the same domain, as in Phase 5.

Production. Production follows this exactly, with two additions the formal model deliberately left to engineering (ADR-0024, ADR-0027; docs/spec/deployment-capacity.md). A collection in the generated core is `alloc::vec::Vec<T>` stored last element first, behind a `collections` feature of the runtime crate turned on exactly when the plan carries a list; a program without lists stays allocation-free and `Copy`. No capacity is imposed by the core and no core ever drops a value. *The design writes its bounds:* a remembered collection is bounded by the operators that bound it — `take cap ...` — and by nothing else, so the deployable window is the Phase-9a construction with `take cap` on the log and the count kept apart:

```
count @fast := 1 + delay 0 count
log    @fast := take cap (cons x (delay [] log))
logD   @slow := sync fast [] log
seen   @slow := sync fast 0 count
cursor @slow := delay 0 seen
window @slow := reverse (take (seen - cursor) logD)
```

With `cap ≥ required(fast → slow)` its window equals the unbounded construction’s at every tick — the production reading of `bounded_buffer_agrees`, differentially tested across the reference evaluator, the executable-IR interpreter and the generated core (corpus `bounded_buffer` against `buffer`, 7 and 3 000 ticks); with `cap` smaller the oldest values of a window are the ones missing (corpus `overflowing_buffer`). *The toolchain computes the bound and the requirement:* a sound static bound per declaration and cell (`bdl-exec-ir::bounds`), the required window capacity per crossing under the deployment schedule (`bdl-reactive::capacity`, the production `Capacity.lean`), a `CollectionsReport` with readiness, byte estimates and per-crossing requirements, and `deployment.*` diagnostics; an unbounded state is reported on a host and **refuses the artefact on a bounded-memory target** (`bdld compile --bounded-memory --period`). The static bound is an engineering analysis with no theorem behind it — recorded as such in §VII.1. A ring buffer specialized for transport buffers was considered and deferred: the window is not a construct the compiler can recognize without a surface form (ISS-0001), and a bounded refinement must be visible in the design to be equivalent — `take cap` is exactly that.

IV.6 Behavior composition: components and groups

A second lamp should reuse the first’s behavior without copying it, and a large design should be organized before anyone decides its interfaces. The first needs a component with a promised interface, instantiated with fresh identity and bound to other behaviors by ordinary realization steps; the second needs a grouping gesture that changes nothing the kernel checks. This chapter derives both from what §IV.1 already provides — realization plus renaming — proves that the flattened system is checked by the same judgments as any design, and keeps the modular-semantics theorem’s restrictions separate from what is proved (Phases 8a, 8b).

Behavior components

The requirements report. The phase began with a requirements report (docs/notes/behavior-system-requirements.md) rather than a formalization: six capability questions — can a behavior be reused, instantiated twice, bound to other behaviors, nested, opened for later completion, substituted — and, for each, whether the answer is `REQUIRED` of the kernel, `OPTIONAL`, or `DERIVABLE` from what exists. The outcome was that everything is derivable from Phase-1 realization plus renaming, and the formal development proves it.

Renaming. `Ren` bundles four renamings — declaration identities, concept identities, clock identities, output identities — and the kernel’s judgments are equivariant under it: `HasType.rename` (typing, given agreement of the three environments on the image; no injectivity needed), `Satisfies.rename`, `Clocked.rename` (the domain judgment, for a clock environment that agrees on the declared identities), and an abstract condition on evidence, `Evidence.Equivariant`. This is the machinery of Phase 8; nothing else is new.

Interfaces and components. A `Port` is a template declaration by local identity with the public part of its interface and its parameter `clock`. A `BehaviorInterface` has required ports (unresolved declarations a composer binds), provided ports (declarations offered to others), elaboration-time parameters (unresolved data-typed declarations bound to closed constants at instantiation), and clock parameters (the domains the template is written against). A `BehaviorComponent` is an interface, a template `Design` over local identities below a width, and a partition of its concepts and sinks into private (freshened per instance) and shared. `Realizes ev C` is a predicate over the *existing* judgments:

the template is a well-formed design; every required port is an unresolved declaration of the stated interface; every provided port is declared with it; parameters are unresolved, data-typed and clock-free.

Fresh instantiation. `fresh W k n = W · (k+1) + n` maps local identity `n` of instance `k` into the global space, with `decode` its inverse; the identities of two distinct instances never coincide (`inst_decl_disjoint`, Theorem A). The encoding is a device, not a semantics: any injective allocator would do, and production uses the same allocator its own ids come from. An instance is a component with a clock-parameter assignment; `Ren.inst` renames the template’s private identities into the instance’s range and its clock parameters to the assignment.

Bindings and flattening. A `Binding` realizes a destination port of one instance from a source — a port of another instance, or a closed constant — with an optional *transport*: `none` for a direct reference in the same or an agnostic domain, `some init` for sync from the source’s domain with an explicit initial value. A `BehaviorSystem` is a width, a list of instances, a list of bindings, the shared concept environment and the external sinks. `flatten` is the union of the renamed instances followed by the bindings applied as Phase-1 realization steps: the destination port is realized as `declRef src` or `sync c init (declRef src)`. The result is an ordinary `Design`. Every existing pass consumes it unchanged.

Theorems B–I (`Behavior/Preservation.lean`, **formally proved**): under `ComposeWF` — every instance realizes its interface, every binding is well formed (`BindingWF`: types agree, a direct binding’s source is in the destination’s domain or agnostic, a transported binding’s source has a domain), external sinks are driven by at most one instance — the flattening is globally well formed (`flatten_WF`), causal under an acyclic inter-instance graph (`flatten_causal`), well clocked (`flatten_wellClocked`), single-driver (`flatten_singleDriver`); open ports stay open (`open_port_stays_open`). Evidence must be `Monotone`, `Equivariant` and `PortSound` (a discharged commitment survives when a port copy is realized by a reference to a declaration of the same interface); a concrete compositional evidence model would discharge these once and is still open.

Semantics (Theorem J, `Behavior/Semantics.lean`): for wiring designs with closure-free inputs and *direct* bindings, the value of a declaration in an instance evaluated alone with a consistent modular input equals its value in the flattened system (`eval_flat_to_inst`, `eval_inst_to_flat`, `modular_iff_flat`). The restriction is exact and recorded: transported bindings under `MEv` need a domain-indexed input for the transported port, and higher-order bodies are not covered — the same obstacle in both directions.

Substitutability. `IFaceRefines A B`: `B` may replace `A` when every port `A` provides, `B` provides with the same type and clock, and every port `B` requires, `A` required. `substitute_composeWF`: replacing an instance’s component by a refining one preserves composition well-formedness. `toComponent` packages a whole system as a component of a chosen interface (the decidable side condition that the chosen interface is realized is not proved).

Counterexamples (`Experiments/BehaviorAlternatives.lean`): a name-based identity collides on double instantiation; a shared clock captured inside a template cannot be re-bound; a binding across domains without transport is rejected by the domain judgment; a binding of a required port to a wrong type is rejected; two instances driving one external sink violate single-driver. The lamp system — a source instance and two dimmer instances at different domains — is executed through the flattening.

Production (ADR-0021, ADR-0022; `bdL-system`; `docs/architecture/behavior-systems.md`): components, stored port contracts (“a component interface is a stored promise”: the contract is persisted and checked against the body, so an edit to the body that breaks the promise is reported at the contract), instances with fresh identity from the project’s allocator, direct and transported bindings, flattening with an *origin map* back to instances and ports so diagnostics land on the instance, composition validation, packaging, versions and substitution. There is one BDL: no system type checker, evaluator, clock judgment or code generator exists. Every project is a behavior system — a design with no components is the degenerate one — and the flat design is derived and never persisted. What production supports beyond the theorem fragment — transported bindings under the multi-domain evaluator, higher-order bodies — is *production-tested*, not proved (§VII.1).

Multi-clock limitations, exactly. Theorem J is proved for `Ev` (one domain) with direct bindings. The clock parameter list must cover the clocks of declarations and sinks; a sync clock appearing only inside a body is renamed to a fresh domain, harmless for `Ev` and `Clocked` but something an elaborator should collect. `InstAcyclic`, the inter-

instance graph condition for causality, is too coarse for extraction (below); a port-level graph would subsume both results.

Behavior groups

Three things that must not be confused. A designer who selects several mapping blocks and chooses *Group as Behavior* has made an authoring decision, not a semantic one. A designer who later chooses *Package as Component* has made a semantic decision: from that point the unit has an interface, can be instantiated with fresh identity, and can be bound.

	mapping / declaration	BehaviorGroup	BehaviorComponent
identity	DeclId	GroupId and a member list	template identities, fresh per instance
public interface	its own DeclInterface	none — only projections	required / provided ports, clock parameters
instantiation	—	none	fresh identities
flattening	—	none needed	flatten
semantics	kernel	none: authoring metadata	elaborates to declarations

The hierarchy mapping → group → component → system has exactly one semantic step, packaging, and that step is an elaboration into the Phase-8a system model.

Transparency. A BehaviorGroup is a GroupId and a list of member DeclIds; a GroupedDesign is a Design with a list of groups beside it; eraseGroups forgets the list. Every group operation — group, ungroup, addMember, removeMember, move, merge, split — acts on the list and leaves the design untouched, so every kernel judgment of the erased design is the *same proposition* before and after (Theorems A–G, each proved by `rfl` or `Iff.rfl`). That is not a weakness of the theorems; it is the point. Dependency edges are untouched; nesting is a relation on the flat group list with no recursive structure. In the Phase-1 vocabulary this is a fourth class below “validation-only”: nothing is rechecked because nothing changed.

Boundaries as projections. Over a finite enumeration of the design’s declarations: `crossIn` lists the non-members some member depends on (the required ports), `crossOut` the members some non-member depends on (the provided ports), `openMembers` the unresolved members, `drivenMembers` the members driving a sink, `privateMembers` the rest (Theorems H–L: `mem_crossIn`, `mem_crossOut`, `internal_not_crossIn`, `internal_only_not_crossOut`). The aggregate input socket a collapsed group shows is `crossIn ++ openMembers`; the output socket is `crossOut`. None is a declaration. Theorem H, `socket_no_fanout`, is what makes the socket safe: $a \in \text{crossIn}$ says *some* member depends on a and nothing about the others. Counterexample 6 shows what goes wrong if the socket is realized as a declaration every member reads (a member acquires an instantaneous dependency it never had); Counterexample 5 shows the tuple-returning-declaration encoding (the consumer of one mapping comes to depend on the inputs of all of them). The kernel has no tuples for boundaries, and this is a reason not to add them for this purpose — the products of Phase 9b are values, not sockets.

Inference rules and their counterexamples. Required = `crossIn`; provided = `crossOut`; private = members neither provided nor driving a sink; clock parameters = every clock the design uses (a group owns no domain, so nothing is captured — Counterexample 3 shows the reconnecting binding failing its clock condition when a clock is captured); physical sinks are not ports — a member that drives a sink keeps its drive edge inside the component and the sink stays external (Counterexample 4: converting a sink into a provided port yields a behavior with no physical effect). Two naive rules are rejected: every reference of a member as required exposes internal producers and changes behavior (Counterexample 1); only the members’ own open declarations hides external dependencies and leaves the template ill-typed (Counterexample 2).

Extraction. Extract builds two *restrictions* of the design: the component keeps the members and holds an unresolved copy of each crossing-in declaration as a required port; the residual keeps the non-members and holds an unresolved copy of each crossing-out member. It forms a Phase-8a system — instance 0 the residual, instance 1 the component, clock parameters mapped identically — with one direct binding per crossing declaration. Each binding is a Phase-1 realization step: the port copy is realized as a reference to the home copy. No body is translated, copied across the boundary, or rewritten. **Formally proved:** both templates realize their inferred interfaces (`restrict_realizes`,

needing Evidence. `InterfaceLocal`: a discharged commitment depends only on the interfaces of the referenced declarations); the system is a well-formed composition (`system_composeWF`) and its flattening a well-formed design (`flat_WF`, Theorems M–Q); causality needed its own argument — a group with both inputs and outputs is never `InstAcyclic`, yet the flattened instantaneous graph is the original graph with every crossing edge subdivided through a port copy, and doubling the original rank witnesses it (`flat_causal`: rank $2 \cdot r$ on home copies, $2 \cdot r + 1$ on port copies); open members stay open; private members are not provided ports, source no binding, and are never referenced by the residual (`private_unobservable`); each crossing-out member is its own provided port (`provided_iff`) — several independent mappings yield several independent ports; and behavior is preserved (Theorem R, `orig_iff_flat`): for a wiring design with closure-free inputs, an original declaration and its home copy in the flattened extraction evaluate to the same value at every tick.

Identity. Before packaging nothing is renamed. After packaging the templates keep the original identities, and the flattened system carries Phase-8a fresh identities: an original n lives at $W + n$ on the residual side and $2W + n$ on the component side, with a port copy on the other side where the boundary is crossed; a later second instance receives $3W + n$. The group’s identity never becomes a component identity.

Production (ADR-0019; `.bdl/authoring.json`; `ui/layout.json`). A group persists as an identity and its member list by identity, in the authoring sidecar — no types, clocks, formulas, outputs or sockets; anything derived is recomputed and not authoritative if stored. Collapsed/expanded state, position, size, colour, member ordering and socket placement are layout and affect no check. Group operations never trigger semantic invalidation: the workspace does not re-run typing, causality, clock, output or hardware checks and does not clear cached verdicts, because these are the same propositions (`group_is_identity_on_design`). *Package as Component* computes interface candidates from the member set and lets the designer widen them — promote a private member, add an open required port — but not narrow below the inferred required set or turn a sink into a semantic port; the flattened result is guaranteed well formed under the stated hypotheses, and the causality check passes even though the instance graph shows edges both ways, so packaging is not gated on `InstAcyclic`. What cannot be inferred and is asked: whether an unresolved member is the component’s input or an open internal; whether a private member should nevertheless be exported; whether a member’s physical sink is private to the component; display name and documentation. Packaging a group *inside* a component body is open (ISS-0007).

Hardware interaction

Components partition sinks into private and external. Private sinks are freshened per instance, so two instances of a component with a private status LED are two LEDs and two requirements; external sinks are shared and may be driven by at most one instance in a system (`ExternalSingleDriver`). Requirements generation (§IV.8) runs over the flattened design’s sinks, so the number of instances is exactly what the board must accommodate — a fact the hardware validator sees only through the flat design, which is the intended separation.

IV.7 The physical boundary: Sources, logical Outputs, provision and realization

Everything in §IV.1–IV.6 happens *between* two boundaries. The lamp’s tilt is a value the environment provides at the concept’s type; its `light` is the design’s intent that the product carry a brightness. This chapter is the account of both boundaries as one story: on the input side, the canonical type $() \rightarrow A$, the derived Source role, and provision by a raw reading and a pure transducer; on the output side, the logical Output with its single explicit driver, why the dual form $A \rightarrow ()$ is not a consumer, realization by a pure encoder and a machine sink, the raw command as the machine boundary, the adapter’s operation one step below it, and the explicit device clock. It ends with the chain drawn end to end and the strength of every arrow stated (Phases 6, 12, 13, 14, 15).

The canonical type of a relationship that reads nothing

A relationship declared with no inputs — mapping `TempSensor : RoomTemp`, mapping `boost : Brightness` — is typed by the kernel at its output: `expectedType = B`. Production (ADR-0029) gives the same relationship a *canonical type* $() \rightarrow B$, with $()$ the empty product, so that every relationship has one type shape `domain(inputs) → B`: `domain([]) = ()`, `domain([A]) = A`, `domain([A, B, ...]) = A × (B × ...)`. The surface spells mapping $f : B$ as shorthand for mapping $f : () \rightarrow B$; f , $f()$ and $f()$ are one reference; $()$ is refused as a value anywhere else, as an output, and as a value form; hover and Explain show the canonical type and the sentence “its canonical domain is $()$, the empty product; the kernel encodes $() \rightarrow B$ as B ”. Production’s `bdl_ir::Ty` has a `Unit` constructor for this — a real IR type with the

ordinary predicates — that never types a Core term, a representation or a runtime value. ISS-0014 asked whether the kernel and the canonical type agree by theorem.

Phase 12 (Surface/UnitDomain.lean) answers without adding a unit to the kernel. The canonical types live in an *interface layer* above it — CTy is the kernel's Ty plus unit, interface arrows and domain products — and the kernel interface type is the value of a normalization function on them (**formally proved**):

$$\text{elim}(\text{canonical}(\sigma)) = \text{encode}(\sigma), \quad \text{encode}(\langle [], B \rangle) = B, \quad \text{encode}(\langle A_1, \dots, A_n; B \rangle) = A_1 \rightarrow \dots \rightarrow A_n \rightarrow B,$$

where elim performs unit elimination ($() \rightarrow B \mapsto B$) and currying ($(A \times B) \rightarrow C \mapsto A \rightarrow (B \rightarrow C)$), is total by a weight both steps decrease, and has no value on the bare unit — $\text{elim} () = \text{none}$: the unit is never the type of anything. The encodings are inverse over every signature whose output is not an arrow, hence over every concept signature (decode_encode , $\text{canonicalOfKernel_encode}$, encode_injective , $\text{canonical_injective}$), which is production's inverse test as a theorem for all signatures. The categorical slogan $\text{Hom}(1, B) \cong B$ is thus, in this model, definitional at the kernel (the realization obligation of $() \rightarrow B$ is $\vdash e : B$ in the empty context, $\text{zero_input_obligation}$, by Iff.rfl), a theorem at the interface (the inverse pair), and a bijection denotationally ($\text{homUnit} : (1 \rightarrow \beta) \cong \beta$).

Why the unit is eliminated *before* the kernel is a theorem, not a preference. The literal alternative — $\text{realize} () \rightarrow B$ as a lambda over the unique argument and read it by application — is refused by typing: no term $\lambda x. \text{delay } i \ e$ has any type, for any binder domain, because delay and sync are typed only in the empty context ($\text{delay_not_under_binder}$, $\text{sync_not_under_binder}$), while $\text{delay init } e : B$ is a legal realization of $() \rightarrow B$ under the kernel encoding (zero_input_memory). A kernel unit binder would forbid memory in every zero-input declaration; a counter $\text{boost} : = \text{delay } 0 \ (\text{boost} + 1)$ would become untypable. lams_typed states the general case: a formula body checked in the context of its inputs is a realization of the encoded type under n binders, and for no inputs there is no binder.

The three spellings are one term (refForms_agree), so they share typing, evaluation and the clock judgment by reflexivity (HasType.refForms , Clocked.refForms). Two theorems say what the unique argument could have carried, and that it carries nothing: a reference's value at a tick is independent of the local environment — a realized declaration evaluates in the empty environment, an unresolved one is read from the input stream ($\text{Ev.declRef_env_irrelevant}$, $\text{MEv.declRef_env_irrelevant}$) — and two readings in one tick agree under any two environments ($\text{same_tick_same_value}$). There is no per-reference call to repeat; a zero-input declaration is evaluated as a declaration, once per activation, and the reader observes that value. No unit clock, no extra activation, no evaluation step.

The source role is a realization state, not a type shape. $() \rightarrow A$ describes the shape of a signature; a *source* is a declaration with no realization, whose value at every tick is provided by the environment: $\text{Source } \Delta d := \text{realizationOf } d = \text{none}$, and source_value gives that value as $I(d, \tau)$ in every domain and under every environment — indexed by the declaration and the tick alone, which is the formal sense in which the unit argument carries no temporal or environmental information. A resolved $() \rightarrow A$ never consults the input stream ($\text{resolved_not_source}$); it is not a sensor, and every mathematical $() \rightarrow A$ is not one either. Production's *simulation input* — Studio offers a declaration as an input when it is unresolved and unit-domain — is $\text{SimulationInput} := \text{Source } \wedge \text{UnitDomain}$, a narrowing the kernel does not need (its input stream provides a value for every unresolved declaration, arrow-typed ones included) and that is recorded as surface policy over the same semantics: on a unit-domain source the two agree ($\text{SimulationInput.value}$). No source kind exists in either model; production's $\text{Signature}::\text{is_unit_domain}$ is the one predicate, and the formal $\text{UnitDomain } \Delta d$ is the same predicate on the kernel type.

The two production consequences of the unit domain are corollaries of existing rules (**formally proved**). Only a data value — never a function — can be remembered or transported: a well-typed sync or delay of a reference forces the referenced type to be data, hence not an arrow, hence unit-domain ($\text{transport_needs_unit_domain}$, $\text{delay_needs_unit_domain}$) — production's diagnostic *f has inputs, so its value cannot be carried across timing domains* is this theorem's message. And a driver of a sink is a value of the accepted type, so when the sink accepts a concept the driver is unit-domain ($\text{driver_is_unit_domain}$, from DriveWF). The dual form $A \rightarrow ()$ is taken up below.

Executed (**executable example**, `Experiments/UnitDomainExamples.lean`): the lamp's signatures round-trip through both encodings and elim computes them; $\text{boost} := \text{delay } 0 \ (\text{boost} + 1)$ read as boost , $\text{boost}()$ and $\text{boost}()$ gives 0, 1, 3 at ticks 0, 1, 3, and the same under a non-empty local environment; the unresolved TempSensor reads the input while boost ignores it entirely; infer accepts $\text{delay } \dots : 00$ in the empty context and refuses a binder around it; sync of boost and of TempSensor is typed and sync of the one-input dimByTilt is refused.

Verdicts (**informed by FV**): unit in the canonical interface notation — keep, above the kernel; a kernel unit type, a unit runtime value, a unit term — remove; a source semantic kind — remove, derive from the realization state; the source surface role — keep in the surface; A -> () as a physical sink — remove (below); the output/drive boundary — keep. Production’s correspondence row for ADR-0029 reads *formally proved (model)* at the interface and *transcribed* at the kernel encoding, and ISS-0014 is resolved.

Source, Rule, Value: production’s reading of the realization state

Production reads the two kernel facts a declaration carries — does its type have inputs, does it have a realization — as one derived **relationship role** (ADR-0032, amended 2026-09-20; `bdl_model::RelationshipRole`):

```
role(x) = Rule    if x reads something           (its canonical type is an arrow)
        = Value   if x reads nothing and has a realization
        = Source  if x reads nothing and has none
```

The role is derived, never persisted, never authored, and distinct from every *state*: a declared rule, an invalid formula, a rule nothing applies, a driven value, a bound port, a clocked declaration are states that vary within a role and never move it. *Has a realization* means an attached definition — a formula whether or not it checks, the reference a binding made, memory, a constant, a parameter’s argument — and never a draft; the role follows the committed revision. The rule has one home, `MappingBlock::role`; the compiler states it in `MappingAnalysis.role`, the daemon on every `MapView.role` (protocol 0.20), and Studio maps the enum and re-derives nothing. The earlier arrangement, in which Studio derived a Source from the shape of the projection, was replaced by this one authority on 2026-09-20 and is not what the product does.

Only the Source case has a formal object: `Source Δ d := realizationOf d = none` on the unit domain (FVD-0118, `source_value`, `resolved_not_source`). *Rule* and *Value* are production’s names for the arrow type and the realized unit domain; the kernel has no need to distinguish a Value from a Rule beyond typing, and none was added. The role crosses the component boundary as a fact of *one design*: the relationship behind a required port is a Source of the body, provided through the port; a provided port realized inside is a Value; an unbound required port of an instance is a Source of the system (Theorem H, §IV.6); a base relationship a binding realizes is a Value at the top level because its flattened copy carries the binding. Production’s `relationship-roles.md` is the normative matrix.

Two consequences the kernel already fixed, restated in the role vocabulary. **A Source may drive an output.** `DriveWF` (*Physical outputs without arbitration*, below; FVD-0051) requires the driver’s type to equal the sink’s accepted type and its domain to be the sink’s; it mentions realization nowhere. So a Source — a value the environment supplies — may be passed straight to a light, and production’s output pass accepts it; a Rule is refused because its type is an arrow (`output.type_mismatch`). Any statement that *only a Value can drive an output* is stale. **A provisioned Source is a Value.** After Phase 13’s provision (below) the target has a realization, so the same rule gives the same answer: no `ProvisionedSource` role exists or is needed (ADR-0032’s amendment; FVD-0127, `provision_source_role`).

Applied rules. A Rule has no value at a tick — it is a function — so the simulator shows it no column, and it does work only when a Value applies it. Production states, per analysis, references (the declarations a realization names by identity — the kernel’s `DependsOn`) and `applied_by`, its **direct** inverse, never a transitive closure; a rule nothing applies is reported as `reactive.rule_unapplied` (informational) and the value that would apply it is offered as the action `rule.apply`. These are IDE and product facts over the dependency graph of §IV.1, not kernel kinds, and the monograph treats them as such.

Commitments, exactly. The kernel gives every declaration a commitment list and makes evidence for it part of well-formedness (§IV.1); Phase 13’s `provision_wf` needs evidence for each target’s commitments on its new realization (FVD-0128). Production authors no commitments at the snapshot: `require` is a reserved word without a production, every declaration’s commitment list is empty, and no commitment solver exists. The formal hypothesis is therefore *vacuous* in production today, and this document does not present a production commitment mechanism as implemented.

Physical outputs without arbitration

The model

A declaration computes a value; it does not move hardware. Physical effect happens only through an explicit **drive edge** from a declaration to a **logical output** — Phase 6 called it a *physical sink*, meaning *outside the behavior, terminal*, and Phase 14 (below) fixes the reading: it is the behavior’s semantic intent at the boundary, physically realized by a deployment-chosen machine sink. The Lean names are unchanged, and the relation `PhysicalOutput` — the value a logical output carries at a tick — keeps its historical name; the machine sink is Phase 14’s `p`:

- `OutputId` — the nominal identity of a sink, a logical actuator channel;
- $\Omega : \text{OutputId} \rightarrow \text{Option OutputSpec}$, a sink’s accepted type and clock, declared by the deployment;
- $\beta : \text{DeclId} \rightarrow \text{Option OutputId}$ — the drive edges, a write-once per-declaration projection of the same shape as `K`;
- $\text{DriveWF } \Omega \ K \ \Delta \ \beta$ — each edge is well formed when the driver’s expected type *equals* the sink’s accepted type and the driver’s clock is the sink’s clock;
- $\text{SingleDriver } \beta$ — at most one driver per sink: if $\beta \ \delta_1$ and $\beta \ \delta_2$ are both some o then $\delta_1 = \delta_2$;
- $\text{CompleteOutputs } \beta \ req$ — every required sink is driven.

Nothing was added to types, typing, the domain judgment, the evaluation relation, or the grant. The binding neither coerces nor converts nor synchronizes. A declaration typed `Tilt`, or bare `q Angle`, cannot drive a `MotorAngle` sink. A sink that accepts a representation type needs an explicit `rep`-typed declaration in front of it, so the distinction between semantic target and hardware representation stays visible. A slow driver reading a fast value must sync it upstream, since a fast driver cannot drive a slow sink and the edge never synchronizes.

Sink identity is nominal for the same reason concept identity is. Keying the binding by type collides when two servos accept the same type; keying by concept conflates a concept with a device, since one concept may feed several; keying by declaration makes two declarations that both mean the steering motor into two sinks, and the multiple-driver counterexample cannot even be stated. “Desired steering angle” is a value; “the steering motor” is a resource; the kernel keeps them in different sorts.

One final driver

The principle is *many contributors, one explicit final driver*. Take two declarations driving one sink, each globally well typed, well clocked, causal, and individually well formed. Only `SingleDriver` fails, and it fails globally rather than at either edge. What has gone wrong is semantic: with two drivers the physical output is not a function of the tick, and there is a tick at which the sink receives two values. With one driver, the physical output is a partial function of the tick, and with `SingleDriver` it is unique wherever it exists:

$$\begin{aligned} \text{PhysicalOutput } S \ \Delta \ I \ \Omega \ \beta \ o \ t \ v &:= \exists d \ spec . \\ &\beta \ d = \text{some } o \wedge \Omega \ o = \text{some } spec \\ &\wedge \text{MEv } S \ \Delta \ I \ spec . \text{clock } t \ [] \ (d) \ v. \end{aligned}$$

Contributors are dependencies, not drivers. `base + corr -> final -> motor` passes every check; priority is an ordinary conditional in the single driver; `blend`, `maximum`, and `clamp` are ordinary declarations of the target type. Why arbitration must be explicit is best shown rather than argued: first-wins, last-wins, and maximum over the same value graph give three different physical outputs. A hidden policy is a design decision made on the designer’s behalf.

Binding an unbound declaration to an undriven sink is a refinement and preserves `SingleDriver`. Binding to an already-driven sink is invalid. Retargeting a sink’s accepted type, renaming a sink, or detaching an edge invalidates an unchanged design. Partial designs may leave sinks undriven; executable designs may not.

What was removed

The earlier draft of BDL had an effect row on the behavior judgment, action requests as values, and per-context policies that allowed, suppressed, transformed, and arbitrated requests, resolved in a dedicated phase of each tick. Each was tried in toy form against the single-driver model, and the outcome is narrow. Direct effect rows — the set of sinks a declaration drives — are exactly the drive edges, and single-driver is exactly their pairwise disjointness. Propagated rows, which also include the sinks of everything a declaration reads, flag a valid design in which a

display reads the driver. Action values move the conflict into the collector that consumes them, which must then be a policy, which is the single driver by another name. None of this bears on richer effect systems; it says that these formulations add no rejection the single-driver rule lacks. Per-context action policies no longer exist as a mechanism, and the StateHandler cases that involved them — event-latched activation with exit-wins, state-local output choice, nested choice with an output — are ordinary declarations with one driver and were run as such.

Why $A \rightarrow ()$ is not a sink

Once zero-input relationships have the canonical type $() \rightarrow A$ (above), the dual form suggests itself: could a physical consumer be a relationship $A \rightarrow ()$, a function that takes a value and returns nothing? The kernel has no unit type, so the form cannot be written; Phase 12 records why it should not become writable (**formally proved**, denotationally). In a pure total language every function into the one-point type is the same function — `unit_codomain_collapse` : $\forall f g : A \rightarrow 1, f = g$, by function extensionality alone — so two “consumers” $\text{sink}_1 \text{ sink}_2 : A \rightarrow ()$ are indistinguishable (`consumers_indistinguishable`): nothing in a value of that type says *which* physical output receives A, or that anything receives it at all. The kernel’s evaluation relation has no effect component (`eval_independent_of_drives`); a derivation relates a tick, an environment, a term and a value, and the physical consequence of a value lives outside it. Naming a receiver needs an effect or output semantics, and BDL already has exactly one: the sink identity `OutputId`, the drive edge with `DriveWF` (the driver’s type equals the accepted type, in the sink’s domain), `SingleDriver` and `CompleteOutputs`. The receiver is named by the edge, the value delivered is the driver’s, of type A, and the driver is a unit-domain declaration whenever the sink accepts a concept (`driver_is_unit_domain`). This is a design result, not a preference: physical consumption stays on the drive boundary, and $A \rightarrow ()$ is removed from consideration.

Provision: how a Source gets its value on a product

A Source $s : () \rightarrow C$ is environment provision at the concept’s type: the kernel input gives a Temperature, and nothing in the design or in the kernel says how a value of Temperature comes to exist. On a product it does not. An ADC yields counts, a GPIO a level, an I²C sensor a register image. The gap between the abstract Source and the physical reading is the last boundary the formal development had to settle, and the shape of the answer is fixed by everything before it: no new kernel construct, no Source kind, no effect, and no way for the design to tell the difference.

The picture, before the theorems:

abstract Source	$s : () \rightarrow C$	unresolved; its value is the environment’s, at type C
raw provider	$r : () \rightarrow R$	a fresh unresolved declaration at the device’s raw type
device channel	$\text{tr} : R \rightarrow \text{Rep}(C)$	a pure term with its transfer function, from a device profile
provisioned Source	$s := \text{mk } C (\text{tr } r)$	an ordinary realization; s is now a Value

The design file keeps mapping `TempSensor : () → RoomTemp`; the deployment supplies `r`, `tr` and the realization; the consumer `tooHot` computes the same truth values from counts as it did from the induced temperature. The key result is a containment, and the direction matters: every trace of the provisioned design under a raw input is a trace of the abstract design under the induced input — **Trace(device) $\subseteq$ Trace(abstract)** — and this strict refinement is the *normal, intended* case. A saturating ADC never yields 451 K; the abstract design admits that observation and the provisioned one does not; nothing is wrong. Trace *equality* needs a **joint section** — one raw trace every channel transfers to what the abstract input gives its target — and the absence of a joint section is not a deployment error and is reported by no surface. Production’s proposal PRP-0001 asked for this construction in seven claims; Phase 13 built it, kept three of the claims, and corrected four. What follows is the construction and the corrections (**formally proved** unless marked; `Surface/Provision.lean`; FVD-0121 ... FVD-0130; PRP-0001 remains a draft, revised, and nothing in production implements it).

The construction. A *channel* is a representation type `rep`, a term `tr`, its transfer function `transfer` on values, and the coherence computes (the term computes the function on every raw-typed value); the term is **pure** — no `declRef`, `delay` or `sync`. A *device profile* is a raw type (sem-free data) and its channels; it mentions no concept. A *provision* is one fresh raw declaration `r` with its clock and an assignment of channels to target Sources — several targets may share one raw reading (an IMU image feeding pitch, roll and acceleration), and one target is the singleton case. Then

$$\text{provision}(\Delta, P) = \Delta [r \mapsto \langle \text{raw}, [] \rangle \text{ unresolved}] [\delta \mapsto \text{mk}_C (\text{tr } (r)) \text{ for each target } \delta : C],$$

with $\text{tr}(\text{declRef } r)$ alone at a representation-typed Source; the raw declaration’s kernel type is raw and its canonical type $() \rightarrow \text{raw}$ (above). Fitting is decidable: at $\text{sem } c$ the concept’s representation is the channel’s; at a representation type the types coincide. Nothing enters the kernel: provision is a function on environments built from DesignDecl , declRef , app and mk .

Why purity, and why a transfer function. Typing the channel term in the *empty* design already forbids reading a declaration ($\text{Channel.WF_refFree}$), but not memory: $(\lambda k. \lambda n. k) (\text{delay } 0 \ 1)$ is typed at $q_0 \rightarrow q_0$ and maps the raw value 7 to 0 at tick 0 and to 1 at tick 1 (exD, **executable example**). A transducer must be a function of the raw reading, so the profile condition is purity — equivalently, typed in the empty design and delay-free ($\text{pure_iff_delayFree_of_wf}$). The channel carries transfer beside tr because the *induced* abstract input must be a function: extracting it from per-tick existence of a value would be a choice principle, which this development does not use. The grant argument the proposal relied on is a theorem of the existing rules: a declaration typed $\text{sem } c$ is realized under $\text{Grant.of } (\text{sem } c) = \{c\}$ ($\text{realization_checked_under_own_grant}$, grant_of_sem), and a channel term typed under Grant.none constructs nothing ($\text{channel_constructs_nothing}$); the same $\lambda x. \text{mk RoomTemp } x$ is refused under the empty grant and accepted under the Source’s own (exC).

Structure. Provision is an environment refinement ($\text{provision_envRefines}$): each target keeps its identity and interface and goes from unresolved to realized, r is new, nothing else moves. The provisioned design is globally well formed (provision_wf) from the abstract design’s well-formedness, monotone evidence, the provision’s preconditions — and evidence for each target’s *commitments* on its new realization, a hypothesis the proposal did not state: a Source’s commitments are obligations on the profile. Causality is preserved with the rank shifted by one and r at the bottom; purity keeps the channel term edge-free (provision_causal). The domain judgment is preserved with $K \ r = K \ s$ ($\text{provision_wellClocked}$); no device clock is introduced, and a device with its own rate is a later sync.

Transparency. For every schedule, domain, tick, term that does not mention r , and local environment whose closures avoid r ,

$$\text{MEv } S \ \Delta \ I \ \kappa \ t \ \rho \ e \ v \Leftrightarrow \text{MEv } S \ (\text{provision } \Delta \ P) \ I' \ \kappa \ t \ \rho \ e \ v, \quad I = \text{induced}(\Delta, P, I'),$$

where the induced input gives each target the wrapped transfer of the raw reading and leaves every other identity as I' gives it ($\text{provision_transparent}$). The hypotheses the proposal lacked are stated: the raw input is typed at r and closure-free, and the abstract design mentions no r — true of every globally well-typed design ($\text{NoMention.of_globalWF}$). The observation boundary has two equivalent forms: syntactic ($r \notin e.\text{refs}$) and by typing (a term typed in the abstract design cannot name r , $\text{provision_transparent_typed}$). The proof is one simulation lemma over MEv with the closure invariant “no closure body mentions r ”, instantiated in both directions and once more for input congruence; the provisioned target’s value comes from the pure term’s canonical evaluation transported to any design, input, domain and tick (Transduces.mev , MEv.of_ev_pure) — at top level, where a realization is evaluated, which is what avoids a closure-equivalence theorem. Physical outputs are unchanged ($\text{provision_physicalOutput}$).

Trace abstraction, exactness, strictness. Every behaviour of the provisioned design under a raw input is a behaviour of the abstract design under the induced input ($\text{provision_abstracts}$): deployment *restricts* the abstract environment; it does not give the abstract design its meaning. The proposal’s converse — “when the transducer is surjective the trace sets are equal” — is wrong as stated: pointwise surjectivity is an existence per tick, and with a shared raw reading it is insufficient even in principle (id and succ from one reading are each onto, and the abstract pair (5, 9) has no witness, no_joint_witness). Equality needs a *joint section*, a raw trace every channel transfers to what the abstract input gives its target (provision_exact); for one channel a pointwise right inverse on typed values is one (JointSection.one). The strict case is executed: a saturating ADC never yields 451 K, the abstract design observes $\text{TempSensor} = 451 \text{ K}$, and no provisioned deployment does (exE).

Re-application and commutation. The proposal called provision “idempotent per Source”. After provision a target is realized, is no longer a Source, and the operation’s precondition fails because r is no longer fresh ($\text{provision_not_reapplicable}$; the Source role moves to r , $\text{provision_source_role}$). The totalized function does satisfy $P(P(\Delta)) = P(\Delta)$ ($\text{provision_idem_total}$) — only because a second pass overwrites every target with the same body; a second pass with a different term is not a refinement ($\text{provision_reprovision_not_refinement}$). Independent provisions commute *exactly*, as environment equality (provision_comm), and the channel assignment is a set (provision_perm).

Executed (`Experiments/ProvisionExamples.lean`): the identity GPIO channel on a `bool Source`; the thermistor $\tau = 2n + 250 \text{ K}$ on `TempSensor`, with the consumer `tooHot` computing the same truth values from counts as from the induced temperature and the provisioned design proved refining, causal and well clocked; rejected profiles; the impure typed term; the saturating ADC; one IMU image provisioning pitch and roll with level reading both.

What changed in the proposal (design recommendation): the purity condition; the transfer function beside the term; the commitment hypothesis; the transparency hypotheses; the joint-section exactness; “not re-applicable” for “idempotent”; shared raw as the primitive with the singleton as its case; the terminology — *abstract Source*, *provisioned Source*, *raw declaration*, never “monomorphised”, which §IV.3’s rank-1 polymorphism owns; and the removal of the dependency on designer-facing °C/°F (ISS-0004): the thermistor is a linear chart on counts and the language keeps kelvin. The output side, which the proposal left as a duality note, is Phase 14’s subject below: it is not a provision at all but a lowering, and the asymmetry is the result. Verdict: provision is a deployment/surface construction over existing kernel terms, not a kernel construct (FVD-0121). The proposal stays a draft, revised, for human review; nothing in production implements it.

A shared raw reading. One inertial-measurement image feeds pitch, roll and acceleration; one thermistor feeds one temperature. Phase 13 made the shared reading the primitive — a provision is *one* fresh raw declaration with its clock and an assignment of channels to target Sources — and the singleton the special case (`Provision.one`; FVD-0125), because the joint-section finding is invisible in the singleton: with one raw reading and two channels each pointwise onto its target, `id` and `succ`, the abstract pair (5, 9) has no raw witness (`no_joint_witness`, **counterexample**). The shared case is the realistic one, and stating every theorem for it and instantiating the singleton (`WF.one`, `provisionOne_transparent`, `induced_one`) costs nothing.

Where the pieces will live. Production has decided, in advance of building any of it (ADR-0032’s amendment; `relationship_roles.md` § Phase 13), what owns what when PRP-0001 is implemented: the device profile and raw type in a **device catalogue** (`hardware/`, never the Standard Library, which is an authoring catalogue of fragments); the transducer as a checked BDL formula; *Fits* in deployment analysis; the provision transformation in the compiler’s deployment pass; the raw input in the platform adapter; the induced input as a testing oracle. None of it exists at the snapshot (ISS-0016; **proposal / not implemented**): the first platform adapter *refuses* a design that has a Source (`adapter.inputs_unbound`) rather than supply a value for it, and the Source sheet of the authoring surface — which creates a Source over a concept the designer chooses — is authoring, not a device binding. The output side of the same boundary, below, has been built; the asymmetry is recorded in §VII.1.

What stays open (FVI-0020): memory in a transducer (a debouncing or filtering channel is not a function of the raw reading, and the transparency and exactness theorems are stated over functions — a stream-level theorem would be needed); a device clock, $\kappa \neq \kappa$ s with a sync at deployment (the transport exists; the construction would add one declaration per target); how a profile’s declared range discharges a Source’s commitments, once production authors any; output provision, the dual; whether `computes` is checked by the compiler or trusted from the catalogue with differential tests; and out-of-type raw readings, a validation question.

Behavior semantics and realization, once more. The separation this chapter makes explicit — on both sides now, provision on the way in and realization on the way out — is the one §IV.8 and Part V depend on. *Behavior semantics* is environment-provided inputs (`I(d, t)`), pure internal computation (`Ev/MEv`) and output obligations (`DriveWF`, `CompleteOutputs`, `PhysicalOutput`). *Realization* is sensor reads, ADCs and buses on the input side, and GPIO, PWM and device I/O on the output side — the platform adapter’s business, named nowhere in the kernel. A Source is not a read; a drive edge is not a write. Provision is the theorem that the input side of that separation can be crossed by a construction the design cannot observe; realization (below) is the theorem that the output side can be crossed by one the behavior cannot observe either — and that, unlike provision, it changes nothing at all in the behavior.

Realization: how an output reaches the world

The drive edge names a receiver and delivers a value of the accepted type; it says nothing about how the receiver is built. A behavior that drives `Light` accepts `Brightness` has not decided whether the product dims a lamp by a PWM duty cycle, a GPIO level through a relay, an I²C brightness register or a UART packet to a lighting controller. Phase 14 (`Surface/OutputRealization.lean`) is the account of why it never has to: the *logical output* of Phase 6 is exactly the platform-independence boundary, and the mechanism is a deployment construction downstream of it.

What an output is, after this phase. `OutputId, OutputSpec = (accepts, clock)`, the drive edge and `SingleDriver` are unchanged, and they are read as the **logical output**: the behavior’s intent that the product carry `C` at the output’s clock, terminal (no feedback), physically *realized* by deployment. Phase 6 called it a physical sink; the word meant *outside the behavior*, not *a device*, and FVD-0050 carries the dated amendment. Nothing about a device is in the spec — no `deviceKind`, no protocol — and no concept implies one: `Brightness` does not mean PWM, `SwitchState` does not mean GPIO. This is a design rule (FVD-0131) with an executed witness (`exH`): one `oLight` realized by PWM and by an I²C register has one behavior trace.

Three stages, and where each decision lives.

stage	what is decided	by whom
behavior design	$d : C \text{ drives } o \text{ accepts } C$	the designer
deployment / validation	the raw command type R_{out} , the encoder $\text{rep}(C) \rightarrow R_{\text{out}}$, its fit, the device’s requirements against the board	the deployment, the validator
lowering / machine	the encoder becomes an ordinary declaration driving a machine sink; the backend consumes the command	the compiler, the platform adapter

The encoder. An `Encoder` is $(\text{rep}, \text{raw}, \text{encode}, \text{transfer}, \dots, \text{computes})$: a **pure** term $\text{encode} : \text{rep} \rightarrow \text{raw}$ typed in the empty design under `Grant.none`, both `rep` and `raw` sem-free data, the transfer function on values, and the coherence `computes`. Purity is necessary for the same reason as on the Source side: $(\lambda k. \lambda n. k) (\text{delay } 0 \ 1)$ is typed at $q_0 \rightarrow q_0$ and encodes the same value as 0 at tick 0 and 1 at tick 1 (`exEFG`); a device transfer is a function of the current value or it is something else (a stateful adapter, §VII.4). The nominality boundary runs the other way from provision: the encoder *observes* the driver’s representation through `rep`, which needs no grant, and the encoder declaration, typed at a sem-free data type, has no grant at all (`encoder_constructs_nothing`, `encoder_decl_no_grant`, **formally proved**); $\lambda x. \text{mk } \text{Other } x$ is refused under `Grant.none` and accepted only under `Other`’s own grant, which no encoder holds. Fitting is decidable — `EFits`: at sem `c` the concept’s representation is the encoder’s; at a data type the types coincide.

Two models, one specification. The tempting construction — retarget `o.accepts` to the raw type and drive an encoder into it — is refuted: the existing edge $d \rightarrow o$ fails `DriveWF` because the driver is typed at the concept (`retarget_breaks_driveWF`, **formally proved**; executed in `exI`), so β would have to be rewritten and the abstract output’s meaning lost; Phase 6 already calls retargeting an edit. What survives is a pair. The *specification* is a relation on the unchanged design, $\text{RawCommand } S \Delta I \Omega \beta R \text{ t w}$: the command *specified* for realization R at tick t is transfer of what `o` carries (`PhysicalOutput`). The machine sink p does not occur in it — it says what the command is, not who receives it; that the lowered design’s p carries exactly this command is the theorem below, not the definition. This is the machine boundary — a relation on commands, not a term; the backend that consumes it is outside the semantics, and no effectful $R \rightarrow ()$, `Expr.write`, effect row or `I0` exists (FVD-0134; the reason is the $A \rightarrow ()$ result above). The *lowering* implements it in the existing kernel:

$$\Delta' = \Delta[e \mapsto \langle \text{raw}, [] \rangle := \text{encode}(\text{rep}(d))], \quad \Omega' = \Omega[p \mapsto \langle \text{raw}, \text{clock}_o \rangle], \quad \beta' = \beta[e \mapsto p], \quad K' = K[e \mapsto \text{clock}_o],$$

with `o`, `d`, the edge $d \rightarrow o$ and every other declaration untouched. The new edge is Phase 6’s *first binding* — a refinement of the drive edges that preserves single-driver (`lower_singleDriver`); the lowered design is well typed, globally well formed, causal (one new edge $e \rightarrow d$, `e` on top) and well clocked with `e` in the output’s clock (`lower_driveWF`, `lower_wf`, `lower_causal`, `lower_wellClocked`, **formally proved**).

Behavior preservation is literal, not observational. Off the fresh identity `e` the behavior environment *is* the original (`behavior_unchanged`, `definitional`); every pre-existing term evaluates identically under the *same* input — there is no induced input on the output side, because outputs do not feed evaluation (`lower_transparent`); every logical output, `o` included, carries the same value (`lower_physicalOutput_unchanged`). This is the asymmetry with provision: Source provision realizes an unresolved declaration and moves the source role to the raw reading; output realization adds downstream structure only, and no declaration changes state. `EnvRefines  $\Delta \Delta'$` holds, but trivially, and it is not the interesting relation; the interesting fact is $\Delta' \times = \Delta \times$ for every x the behavior knows.

Correspondence, directional. The central theorem is `lower_correspondence`: the machine sink’s value in the lowered design is exactly the specified command,

$\text{PhysicalOutput}(\Delta', p, t, w) \Leftrightarrow \text{RawCommand}(\Delta, R, t, w), \quad \text{i.e.} \quad \text{raw trace} = \text{transfer} \circ \text{abstract trace},$

tick by tick in the output's clock, under: a well-formed single-driver drive environment, a design that mentions no e (every globally well-typed design), inputs whose closures avoid e , and a logical output carrying typed representation values — which Phase 5's totality supplies in a causal, well-formed design (`output_value_typed`). The theorem is an equivalence of *observations at p* , but the construction is directional by nature: p is downstream of o , nothing flows back, and no “induced output” exists. Nor is any exactness needed: two abstract values may encode to one command — a 4-bit PWM sends duty 6 for 40 % and for 41 % (`exB_quantized`) — and a round-trip `decode(encode x) = x` is neither assumed nor provable; the profile-declared transfer *is* the intended realization (FVD-0135). Where provision needed a joint section because abstract inputs had to be *reached* from raw ones, realization only has to *deliver* what the behavior produced.

Platform independence, as a theorem — and its exact scope. Two realizations of one design — PWM and I²C for the light — evaluate every pre-existing term alike (`two_realizations_same_behavior`): the *same formal evaluation* (`MEv`) of every term that mentions neither fresh encoder, in either lowered design, under the theorem's hypotheses (both encoder identities fresh, inputs whose closures avoid them). That is the formal content of “the behavior is independent of the mechanism”; it says nothing about a compiler, a backend, a board or a physical device, whose correctness is the open codegen half (Part V, FVI-0022). Only the command types and traces differ (`exH`: duty 102 versus (`register 42, 40`) for the same 40 %). Independent realizations commute exactly, as environment equality (`lower_comm`). Both device profiles are admissible for the light on the Nano, and the PWM profile is not admissible for the relay — the fit fails, not the pins (`exH_admissible`). Admissibility is a conjunction of three judgments that never see each other: the encoder's typing `rep -> raw` under no grant, its fit, and Phase 7's solver on the device's requirements (`Admissible`, `admissible_satisfiable`, FVD-0139 superseding FVD-0137, which had omitted the typing); the encoder knows nothing about timers and the requirements nothing about brightness. The gap is executed: an Encoder value whose term is $\lambda n. \text{true}$ fits the light and allocates a PWM line and is not admissible (`exJ`, `admissible_needs_wf`). A solvable board is not an electrically correct device — no voltage, current, thermal or timing property is proved (§VII.4).

Single driver stays what it was. `SingleDriver` is about logical outputs: at most one behavior declaration drives o . A device that needs several pins — PWM plus direction, SDA plus SCL, TX plus enable — is a requirement list at validation, not several drivers; and a raw command may be structured — the H-bridge encodes `MotorSpeed = (forward?, magnitude) to (duty, direction)` with the kernel's pairs (`exD_hbridge`), no record type needed. A device that consumes several logical outputs at once is either several machine sinks committed in the same tick — the runtime already commits all outputs of a domain together — or one concept combined upstream, as Phase 6 places every combination of behaviors; the singleton realization is primitive by *decision* (FVD-0136) — a decision supported by the executed examples and the minimality argument, not a theorem that every atomic multi-output protocol reduces to upstream combination or per-tick batching (that question is FVI-0022) — and the analogy with the shared raw reading of Phase 13 fails for a stated reason: a reading physically arrives as one image and its split is real, a frame is assembled by the machine from values the behavior already produces separately.

Clocks. The machine sink is in the output's clock and the encoder with it; an encoder in another domain fails both `DriveWF` and the domain judgment (`exI`). A device that consumes at another rate needs an explicit sync in the lowered design — Phase 5's transport, never a resampling hidden in the binding — and a peripheral's carrier frequency (the PWM period, the bus clock) is device configuration for validation, not a `ClockId` (FVD-0138).

Executed (`Experiments/OutputRealizationExamples.lean`): the GPIO identity on a relay; 8-bit PWM on the light with the driver unchanged; the quantizing 4-bit PWM; a servo pulse $1000 + a \cdot 1000/180 \mu\text{s}$; the H-bridge pair; the misfits, the constructing encoder and the impure encoder refused; PWM versus I²C for one light; the structural theorems instantiated; both profiles admissible on the Nano and an ill-typed encoder that fits and allocates but is not admissible; Model A and the implicit clock crossing refuted.

What stays open (FVI-0022): stateful output adapters — slew-rate limiting, dithering, batching, hysteresis — and whether each is the behavior's, a stateful lowering's with a stream-level theorem, or the backend's; the device-clock variant; whether an *atomic* multi-value frame ever forces a many-to-one construction; the codegen half — abstract trace $\rightarrow$ raw command trace is proved, raw command trace $\rightarrow$ generated backend call is not (Part V); and commitments on outputs, which production does not author.

Production, at the snapshot. Phase 14 was consumed the day it closed (ADR-0036 (FVD-0131 ... FVD-0139); **production implemented and tested**). Realization is deployment data: a device binding names its profile in the device body — device `pwmLight : pwm_channel for light { realization pwm_duty8 }` — never on the output and never in the design graph. A profile is an encoder plus a requirement template: `OutputProfile { id, encoder, kind }` with `Encoder { rep, raw, encode }` a closed, pure, typed Core term whose purity check refuses `declRef`, `delay`, `sync` and `mk` structurally; quantization is admitted and no round-trip is asked. Admissibility is the three judgments of `Admissible`, reported separately with one diagnostic each, composed in `analyze_deployment`; the requirements come from the profile’s kind through the unchanged solver, and the hardening’s gap is kept refused by a test. The lowering emits one machine sink per valid chosen profile, due when the driver is, in the driver’s context under no grant — at the *plan* level, without the fresh `declIds` the model’s `lowerΔ` literally introduces (an alternative ADR-0036 records: plan-level sinks are not designer-visible entities with ids and canvas positions, and the theorems do not distinguish the two, since the sink is observably the same). The generated core’s `Tick` gains commands beside values and outputs; `TickTrace.commands` is the downstream view; simulation stays behavior-level; choosing a profile changes no analysis, no sample and no `Values/Outputs`, and that is a test (`changing_the_realization_changes_no_behavior_and_only_the_raw_trace`). Five profiles exist as witnesses — `pwm_duty8`, `pwm_duty4`, `i2c_level8`, `gpio_level`, `hbridge_signed` — versioned with the compiler, not a device catalogue (ISS-0017). Studio chooses on the Deploy page only; the Design page, the inspector and the canvas know nothing of profiles. Below the raw command a first platform adapter now exists (ADR-0037, Part V): it consumes `Tick.commands` on the Raspberry Pi Pico and is production-tested through recording sinks and a cross-build, and nothing formal is claimed for it — that arrow is the boundary FVI-0022 leaves open.

The adapter boundary and the device clock

Phase 14 stops at the raw command; production’s first embedded adapter (ADR-0037) starts there — `Tick.commands` → `adapter::apply` → an operation on a peripheral — with an explicit *boundary policy* (a finite duty in $0 \dots 255$ rounds half up and is applied, anything else is refused and the line holds its last value, a driver that is not due leaves the line) and a host-recorded operation trace. Phase 15 asked the narrow question the adapter made answerable: what is the smallest honest formal boundary beyond `RawCommand`? The answer is one step and no more (**formally proved**, `Surface/Adapter.lean`).

Policy, operation, line. A `Policy` is the adapter’s reading of a command — `accept : Value → Option Value`, partial by design; `duty8` is production’s policy on the kernel’s naturals, `clamp8` the alternative production refused, `total` applies as is. An `Op` is `set m`, refused or held — production’s `AdapterOp` as data; no term, no evaluation rule, nothing in Core. `AdapterOp S Δ I Ω β R spec P t op` is the operation on the realization’s sink at global tick `t`, over the *unchanged* design and gated by the output’s clock: held when the clock is not active, `set (P w)` when the specified command `w` is accepted, refused otherwise. The line is a fold over the operations: the start value until a command is accepted, the last accepted value thereafter; a refusal or an inactive tick changes nothing. Reject-and-hold is therefore a property of the fold — the only state at the boundary, and it lives in the adapter, below the realization, outside the behavior (FVD-0140).

What is proved. The operation is a function of the tick (`AdapterOp.det`); it is determined by what the lowered design’s machine sink carries — `lower_correspondence` carried one step further, in both directions (`adapter_of_sink`, `sink_of_adapter`); the raw command does not mention the policy, so no policy and no refusal reaches the behavior or the command (`adapter_downstream`, `two_policies_same_commands`); the line is a function of the tick, holds on refusal and on an inactive tick, carries an accepted command, and never carries a refused one (`Line.det`, `line_holds_on_refusal`, `line_holds_when_inactive`, `line_last_accepted`, `line_value_accepted`). Executed: duty 102 set and the line at 102; a dial at 120 % encodes to duty 306, refused under `duty8` with the line holding 0, set to 255 under `clamp8`, 306 under `total` — one command trace, three operation traces (`exA`, `exB`); an inactive tick holds (`exC`). Below the operation nothing is modelled — production’s `f64` rounding, the HAL, the register, the electrical world (FVI-0023); the generated Rust’s agreement with `AdapterOp` is **production-tested** (`host_adapter_operations_correspond_to_the_commands`), never proved.

The explicit device clock. Phase 14 refused an implicit crossing; Phase 15 builds the explicit one (`Surface/DeviceClock.lean`): the encoder declaration lives in a device domain `dc` and reads the driver’s representation through Phase 5’s transport,

$$e := \text{encode}(\text{sync}_{\kappa} \text{initRep}(\text{rep } d)), \quad e, p \text{ in } dc,$$

with c the output's clock and initRep a pure closed representation value for the ticks before c 's first activation. What must be explicit in the lowering is exactly dc and initRep ; both are deployment's choices, and the design still says nothing about a device; the carrier frequency stays configuration (FVD-0141). Every preservation theorem survives — the behavior is literally unchanged off e , the lowering is a refinement, well formed, well clocked with the transported operand in the output's clock, single-driver — and causality survives with *no new instantaneous edge at all*, because the transport is never instantaneous (lowerSync_causal). The correspondence samples strictly before: at a device tick t the sink carries transfer of the command the output specified at the last activation of c before t , or of initRep if there was none ($\text{lowerSync_correspondence}$). Executed: the device domain on odd ticks and the output's clock on even, the sink carrying 102 at tick 1 and 107 at tick 3 with prevAct naming 0 and 2; the initial representation before the first activation (exD_).

Stateful adapters, classified before generalized. No `StatefulEncoder` was added, because no case needed one. Slew-rate limiting is product-observable and is an ordinary declaration upstream of the logical output — `limited := if delay 0 limited + 10 < dial then delay 0 limited + 10 else dial ramps 10, 20, 30, 40` in the design, visible to the designer and the simulation, and the same pure encoder produces the ramped raw trace 25, 51, 76, 102 (exE_slew); servo smoothing and hysteresis are of the same kind; debouncing is Source-side (FVI-0020); PWM dithering modulates below the tick, where the model has no time, and is backend implementation; protocol batching is the backend's per-tick commit. A stateful primitive between the logical output and the raw command needs a value none of these can represent; none is known (FVD-0142, FVI-0025). Atomic multi-value frames likewise remain a stated criterion without a witness (FVI-0026).

The audit behind the phase. Phase 15 began by auditing every active open item against Phases 8a–14 and production 6be778b, because an open item is not a research backlog merely by having a number. Eleven were resolved, deferred or merged — the device-component library (FVI-0012) had its premise rejected by Phases 13/14, folding the clock into the interface (FVI-0013) was already decided against by FVD-0046, grant delegation (FVI-0018) is unnecessary under `constructs_granted`, the display-name table (FVI-0019) duplicates FVD-0014, the general `elim` (FVI-0021) is rejected by the canonical-interface design, interface-level references and evidence invalidation (FVI-0001, FVI-0015) are one blocked question until production authors a commitment; FVI-0022 was split into five questions with different answers (FVI-0023 ... FVI-0027), and FVI-0011 was narrowed to the now-concrete shared-configuration feasibility (RP2040 slices sharing a carrier) with its explanation half split off (FVI-0028). The triage table is in the formal repository's `docs/issues/README.md`.

Communication as state, catalogue profiles and assign

A realistic product — a host submitting motion work over TCP, motors reporting position, velocity and faults over CAN, jobs queued and cancelled, telemetry returned — looks, from the outside, like a system of messages, and the natural draft of a language for it has a `Message` or `Event` type, a queue, a transaction. Phase 16 tested the opposite hypothesis: that the behavior language needs *semantic state and its evolution* and nothing else, communication appearing only at the two boundaries of this section — provision on the way in, realization and the adapter on the way out. The hypothesis was attacked rather than assumed: every case was first encoded with the existing kernel and executed, and a primitive would have been admitted only against a concrete case the encoding could not represent without changing observable behavior. The kernel had already refused an event type three times on proofs about multiplicity (FVD-0036, FVD-0048, FVD-0085); this was the product-scale test.

The stress case, as state. Two designs of ordinary declarations, composed into one controller (`Experiments/CommunicationExamples.lean`): a *work queue* — `submit : list Job` (the jobs that arrived this tick, oldest first; a job is the pair $(\text{JobId}, \text{delta})$), `cancel : opt JobId`, `done : bool`; state queue : `list Job` (FIFO, bounded by `take cap`), `desired : Position` advanced by a job's `delta` when the head's identity changes, `overflow : bool` — and a *motion state* — `fb : opt (pos × vel × fault)`, a sample or none; `age` (ticks since the last sample), `fresh := age < 3`, `timedOut := age ≥ 6`, `held lastPos/lastVel`, `settled` (consecutive fresh in-tolerance samples), `doneM := settled ≥ 2`, a fault latched until `reset`, an `acked` held until the target changes. The composed design is proved globally well formed, causal (`done := delay false doneM` reads the previous tick, so there is no cycle) and well clocked. Then, by decide on the proved-sound interpreter: `Move(+10)`; `Move(+10)` and `Move(+10)` are queues of two and one and `desired` positions of 20 and 10 after the first completes (**multiplicity**, exA_multiplicity); `A`; `B` and `B`; `A` reach the same final target through different `desired` trajectories (**ordering**, exB_ordering); `[A, B]`, `[B, A]` and `[A]` arriving in one tick are three

queues — the batch is a list and nothing is lost (**same-tick multiplicity**, `exC_same_tick`); a queued job, the active job and an unknown id cancel as they should (`exD_cancellation`); the third job into a queue of two is refused and the refusal is a boolean the product observes (`exE_bounded`); a retried submission with the same `JobId` is one job under a design that deduplicates by id and two under one that does not — the design decides, because the identity is in the state (`exF_identity`); freshness, timeout, settling over repeated samples, the fault latch and the acknowledgement are counters and booleans under delay (`exG_freshness`, `exH_settling`, `exI_fault_latch`, `exJ_ack`); the composed controller completes job 1 after two settled samples and starts job 2 (`exQM_composed`); host submissions in a fast domain reach the controller’s slow domain through the Phase-9a window, flattened, in order (`exK_cross_clock`); and a paired axis — one `YPosition` on two motors with a prepare/prepare/commit protocol — is one logical output whose raw command is the pair (p, p) (`exN_pair`), or two realizations of the one output whose commands are the two transfers of one value (`paired_commands_of_one_value`), the frame order being the backend’s per-tick commit below the adapter operation (FVD-0148). No case needed a message, a queue primitive or a transaction (FVD-0143).

Three kinds of identity. A *transport* identity — a TCP sequence number, a CAN arbitration id used for routing, a retry token — is discarded by the provider’s channel, and a theorem says that nothing downstream can tell: two providers of one design whose raw inputs induce the same semantic Source trace evaluate every term that mentions neither raw declaration alike (`two_providers_same_behavior`, and `two_providers_same_outputs` at every logical output). “Same trace” is `SameTrace`: agreement of the induced inputs at every declaration other than the two raw ones, at every global tick; the raw types need not be related, and the executed instance frames the same submissions as (sequence number, batch) and as per-job (frame id, job) (`exL_theorem`, `exL_executed`). A *device* identity — a motor’s node id, a pin — is the machine sink and the raw declaration: deployment data. A *semantic* identity — the `JobId` a cancellation names — is in the state, because product behavior depends on it. The **replacement-invariance criterion** — a carrier replaced under the same semantic trace changes no behavior: TCP to USB CDC, CAN to another transport, polling to interrupts — is therefore three theorems, one per boundary: `two_providers_same_behavior` on the way in, `two_realizations_same_behavior` (Phase 14) at the output, `two_policies_same_commands` (Phase 15) at the adapter (FVD-0144).

Catalogue profiles and assign. If the profiles of §IV.7 are supplied by packages, what does a package contribute? In the model (Surface/Assignment.lean) a catalogue entry is a profile — Phase 13’s `DeviceProfile` for an input, Phase 14’s `DeviceOutputProfile` for an output — with an `Origin` (`builtin | package`), and the origin is a field no judgment reads: two entries with one profile give one lowered design, one raw command relation and one adapter operation relation (`assign_indistinguishable`, `assignSource_origin_irrelevant`; `executed`, `exP_origin`). The contract an entry must satisfy is the one the phases already state — $\text{InputContract } \theta \vdash \text{ch} := \text{ch.WF } \theta \wedge \text{Fits } \theta \vdash \text{ch}$, exactly what `Provision.WF` asks per target, and $\text{OutputContract } \theta \text{ accepts } P := P.E.WF \theta \wedge \text{EFits } \theta \text{ accepts } P.E$ — and a larger catalogue realizes and provisions more and nothing else (`realizable_mono`, `provisionable_mono`): a package extends the realizable world, not the language (FVD-0145). A deployment `assign MotorOutput` using `emm_v5.position_control` is Phase 14’s lowering with the entry’s encoder (`assignOutput  $\Delta$  en ... := lower  $\Delta$   $\Delta$  (en.realization ...)`), and `assign MotorPosition` using `emm_v5.position_feedback` is Phase 13’s provision with the entry’s channel: the design is literally unchanged off the fresh identities, every pre-existing term evaluates alike, and the assigned design is accepted by the unchanged judgments (`assignOutput_behavior_unchanged`, `assignOutput_transparent`, `assignOutput_checked`, `assignSource_transparent`; FVD-0146). Semantic admissibility and deployment feasibility stay separate: `Admissible  $\leftrightarrow$  OutputContract  $\wedge$  Feasible` (`admissible_iff`), and the executed pair profile satisfies the contract for `DesiredPos` on the Nano and on a one-pin board alike, being feasible on one only (`exQ_contract_not_feasible`; FVD-0147). Package resolution, versions, registries and signatures are engineering and are not modelled.

Where the attack found work, not a primitive. The output-side lowering into a device domain (`lowerSync`, Phase 15) samples the last command, so two commands specified between two activations of a *slower* device reach it as one; the occurrence-preserving crossing is the Phase-9a window mirrored into the lowering — a construction over existing primitives, not built (FVI-0024, amended). And every encoding rests on what a raw reading *means*: that a `list raw` reading delivers the occurrences since the previous tick in arrival order, that a transport retry is delivered once, that two raw sources are merged into one order, that the batch is bounded — the provider’s occurrence contract, which is what a Source device profile has to promise and which no record states (FVI-0029). Both are deployment work; neither is a language question.

The provider's contract and the occurrence-preserving window

Two boundaries were left below the state encodings of the previous section. The first is what a raw reading *means* when the transport carried several things since the previous tick — several commands, a retransmission, deliveries from two devices. The second is how an output whose every value matters reaches a device that activates more slowly than the output's clock: the device-clock lowering of §IV.7 samples, so two commands specified between two device activations reach the device as one. Phase 17 closed both with the existing kernel (`Surface/Provider.lean`, `Surface/OutputWindow.lean`); neither needed a message, a queue or a transaction.

The provider's occurrence contract. Below the raw reading sits the *provider*, the adapter's function from deliveries to the one value the Source reads. A delivery is a payload with a *transport identity* — a sequence number, a frame id, a retry token — that the transport owns and the design never sees. The provider keeps the identities it has delivered (adapter state, like the line of the adapter boundary) and delivers, per tick, the first cap payloads whose identity is fresh, in arrival order, together with a flag saying whether more were fresh (`provide`). That is the contract, in five clauses: one semantic occurrence per fresh identity — never per payload, so `Move(+10)`; `Move(+10)` with two identities is two commands (`dedup_of_fresh`) and a retransmission is nothing; arrival order within a raw source (`dedup_sublist`); retransmissions erased by identity below the boundary — a retransmission inserted anywhere in the delivery stream changes no batch at any tick and no behavior (`run_retry`, `retry_invisible`); a per-tick bound whose cut is a value the design reads (`overflow`), never a silent drop, the bound itself being a deployment assumption on the physical arrival rate decided as capacity is decided; and several raw sources as several raw readings, one provision each, with any merge an explicit deterministic policy (`mergeBySource`) that a design reading per source cannot tell from another (`perSource_of_interleaving`) — no physical total order is assumed. The reading is (`list raw, bool`), Phase 13's shared raw reading with two channels (`batchProvision`), so every provision theorem applies and equal batch streams are the same semantic trace (`sameBatches_sameTrace`); a scalar Source — a level, a sample — is the batch sampled (`Batch.latest`). Executed on the `auto_typer` designs: the host's submissions through the provider give the queue design the same queue and desired position under a stream with a retransmission as without (`exE_ingress`, `exE_theorem`); three jobs in one tick under `cap = 2` queue two and raise the flag in the design (`exE_overflow`); two motors' feedback from two raw readings, sampled, gives the motion state its `fb : opt Sample` with a retransmitted sample as `none` (`exF_feedback`); and the motion state, unchanged, is a behavior component instantiated on two axes in one system (`exG_two_axes`) — deployment adds no product behavior, the behavior layer already holds the controller.

The occurrence-preserving window. For a device that must receive every command, the crossing is Phase 9a's five declarations with Phase 14's encoder declaration as their source:

```
e      @c := encode (rep d)
log    @c := cons e (delay nil log)
logD   @dc := sync c nil log
seen   @dc := length logD
cursor @dc := delay 0 seen
window @dc := reverse (take (seen - cursor) logD) → p : list raw @dc
```

The behavior is literally unchanged off the six fresh identities (`lowerWindow_transparent`, six applications of one generic step, `update_transparent`); at every tick, in the device domain, the sink carries exactly the raw commands specified at the output-clock activations since the device's previous activation, in order and with multiplicity (`lowerWindow_correspondence` — Theorem M read through Phase 14's correspondence, for which the Phase-9a lemmas were generalized in place over any source value function); the batch is bounded by the Phase-9a capacity obligation on the crossing (`lowerWindow_bounded`); and the lowered design refines the abstract one and is well formed, causal, well clocked, drive-well-formed and single-driver. With the device on even ticks, the sampled lowering carries 104 at tick 2 where the window carries [102, 104] (`exA_window_vs_sample`); a repeated value appears twice (`exB`); a period-3 device with `cap = 2` is refused as a deployment infeasibility (`exC_capacity`). **State versus occurrence is the sink's type** — the device's consumption contract: a device that consumes `raw` is lowered by `lowerSync`, one that consumes `list raw` by `lowerWindow`; the logical output is one value stream in both and carries no flag (FVD-0152). The adapter's batch is a list of the adapter boundary's operations and the line after it the last accepted item (`batchOps`, `lineAfterBatch`; FVD-0153); two window realizations of one output carry batches that are pointwise the two transfers of one value (`paired_batches_of_one_window`), so the prepare/prepare/commit of a paired axis is the backend's order within one batch.

The Source side completed: provider state, the device clock, initialization, commitments, readings

What the input boundary still left open after the provider's contract — a transducer that needs memory, a device that samples in its own domain, a Source read before its first value, a range a profile is asked to guarantee, the computes obligation, a reading that is not a value of the raw type — Phase 18 closed with one lemma and six constructions (`Surface/SourceBoundary.lean`). The lemma is `MEv.congr_at`: two designs that keep every realization but one declaration's, read every other input alike, and give that declaration the same value at every tick in every domain, evaluate every term alike. The behavior sees a Source through its value trace and through nothing else; every answer below is that lemma with two traces computed.

Provider state is movable, so placement is visibility. A stateful transducer — a debouncer, a quadrature decoder, a low-pass filter — is a Mealy machine whose step is a pure BDL term over data (`Machine`). Run below the raw reading, its output is what the provider delivers; placed above it, it is one declaration with delay over the physical reading, and the Source is realized from its output. The Source's trace is the same either way, and every evaluation of the provisioned design holds in the upstream one (`provider_state_movable`, executed on the filter: `exA_filter`). The placement is therefore not a semantic question but one of *visibility*, and the criterion is stated with two formal facts behind it: a provider cannot read the design — its state is a function of the raw stream alone (`Machine.run_congr`), and a channel term mentions no declaration, so the fault latch that reads reset or the settling count that reads target cannot be a provider's (`exA_latch_reads_design`); and a parameter that changes the semantic trace on one physical stream — the debounce threshold, a hysteresis band, a filter constant — is the product's whenever the product's specification fixes it (`exA_debounce_param`, `exA_hysteresis_param`). Decoding an edge from two consecutive phase readings is the provider's; accumulating and homing a position, which reads reset, is the design's (`exA_quadrature`). Two providers running different machines with one output stream give one behavior (`stateful_providers_same_trace`). Nothing can be hidden *by* the placement, because the trace is what the behavior reads (FVD-0154).

The Source-side device clock and the first value. A device that samples in its own domain `pc` is provisioned by an explicit sync of the raw reading into the Source's domain, with an explicit initial value — the input dual of the output side's device clock (`provisionSync`): the Source carries the transfer of the reading at the last activation of `pc` strictly before, or of the initial value (`provisionSync_target`); the abstract design is unchanged (`provisionSync_transparent`); no new instantaneous edge, well clocked with the reading in `pc`, a refinement, globally well formed (`provisionSync_wf`). An occurrence-like Source — encoder edges, queued commands — crosses by Phase 9a's window over the raw reading (`provisionWindow_target`), the same five declarations and the same theorem, nothing lost (`exB_window_edges`). The initial value is an explicit `InitRep` under one of two policies a profile declares: a *supplied* raw value the Source reads until the first sample, or *unavailable* — none at an optional raw type, which the design reads as the sensor not having spoken (`exB_sampled`, `exB_unavailable`). A reading is never fabricated; activation gated on the first sample is not a construction, because a schedule does not read an input, and is the optional form (FVD-0155). The explicit initial value is the one principle both boundaries share.


Commitments at three evidence levels. A Source's commitment — a temperature at most 450 — is discharged by the provisioned realization when every value it takes, under readings satisfying an assumption `A`, has the property (`RangeSoundUnder`, `discharge_under`). The three levels are three assumptions and three ways of establishing them: *static*, `A` is typing and the transducer alone guarantees the property (a saturating ADC, `discharge_static`); *checked*, `A` is a validation the provider applies to every delivery, established by construction (`discharge_checked`); *trusted*, `A` is a range the profile asserts of the device, which nobody establishes here and the theorem carries as a visible hypothesis (`discharge_trusted`). A profile declaration is not evidence by itself: at the trusted level it is exactly the assumption named (FVD-0156). The computes obligation — that a profile's transfer function is what its term computes — stays a proof: derived when the function is the term's evaluation (`Channel.ofTerm`), decided at a finite raw type (`computesBool`, `computes_of_bool`); a separately supplied function at an infinite type is a claim to test, and production compiles the term (FVD-0157).

Readings that are not values. A malformed frame, a NaN, an invalid code, a count outside the declared range never become a Sem value: the checking provider refuses the delivery before the contract's deduplication and bound (`checkedProvide`), every delivered item is typed by construction (`checked_typed`), and the refusal crosses as none at an optional Source or as a flag beside it — semantic state if the product must react (`available`), a backend diagnostic if it must not (`exE_checked`). No exception semantics (FVD-0158). The richer profile a package may supply — transducer, machine, initial policy, delivery contract, assumed range, requirements — leaves the assignment reading Phase 13's

profile and nothing else (`assignSource_profile_only`), so `assign` stays deployment-only; and the motion controller of the previous section, fed by a batch provider sampled through `chLatest` and by a scalar provider through the identity channel, has one `fb` trace and one `doneM` trace (`exF_two_providers`). With this, FVI-0020 — the Source side’s open item since Phase 13 — is closed remainder by remainder.

The physical boundary as one whole

Read end to end, one value’s path from the world back to the world is the following chain, and each arrow has its own strength. The point of drawing it is that the strengths differ; flattening them into one “verified pipeline” would be the most misleading sentence this document could contain.



arrow	what it is	formally modelled	formally proved	production	still open
physical world $\rightarrow$ raw reading r	the provider delivers what the transport carried since the previous tick: the fresh occurrences in arrival order, bounded, with an overflow flag	Phase 17 Provider: <code>dedup</code> , <code>provide</code> , the reading (<code>list raw</code> , <code>bool</code>) as a two-channel provision (<code>batchProvision</code>); a scalar reading is the batch sampled	<code>dedup_of_fresh</code> , <code>run_retry</code> , <code>retry_invisible</code> , <code>provide_overflow_iff</code> , <code>sameBatches_sameTrace</code>	not built at the snapshot: no device provides a Source’s value (ISS-0016; a slice was in flight, uncommitted)	which bound a physical arrival rate needs (a deployment assumption); the device catalogue for inputs
raw reading $\rightarrow$ logical Source, with a device that keeps state, samples in its own domain, promises a range, or delivers a malformed value	provider state as a machine; the Source-side device clock; commitment discharge at three evidence levels; the checking provider	Phase 18 Machine, <code>provisionSync</code> , <code>provisionWindow</code> , <code>RangeSoundUnder</code> , <code>checkedProvide</code>	<code>provider_state_movable</code> , <code>provisionSync_transparent</code> , <code>provisionSync_wf</code> , <code>discharge_static / _checked / _trusted</code> , <code>checked_typed</code>	not built	the device’s own bound or range (a deployment assumption)
raw reading $\rightarrow$ logical Source ($s := mk\ c\ (tr\ r)$)	provision by a pure transducer	Phase 13 Provision	<code>provision_envRefines</code> , <code>provision_wf</code> , <code>provision_transparent</code> , <code>provision_abstracts</code> ; equality under <code>JointSection</code>	not built (PRP-0001 draft); the adapter refuses a design with a Source	stateful transducers, a device clock, commitment discharge (FVI-0020)
logical values $\leftrightarrow$ behavior	delay, sync, the library, components	§IV.1–IV.6	determinism, totality on causal designs, the correspondences of §IV.3–IV.6	reference evaluator, exact transcription; generated core tested against it	closure equivalence, Theorem J beyond its fragment
driver $d \rightarrow$ logical Output o	the drive edge	Phase 6 DriveWF, <code>SingleDriver</code> , <code>CompleteOutputs</code>	<code>single_driver_output_deterministic</code>	<code>bd1-output</code> ; drive light by brightness	commitments on outputs

arrow	what it is	formally modelled	formally proved	production	still open
logical Output → raw command w	realization by a pure encoder and a machine sink	Phase 14 lowerΔ, RawCommand	behavior_unchanged, lower_transparent, lower_correspondence, two_realizations_same_behavior, admissible_needs_wf	ADR-0036: profiles, three-judgment admissibility, SinkPlan, Tick. commands; tested	the plan-level lowering is not the model's fresh declarations (observably the same, unproved as such)
logical Output → raw command batch (a slower device)	the occurrence-preserving crossing: Phase 9a's window over the encoder into the device domain, a list raw sink	Phase 17 lowerWindow	lowerWindow_transparent, lowerWindow_correspondence, lowerWindow_bounded, the structural theorems	not built	a device that acknowledges; the initial representation (FVI-0024)
abstract sink operation → peripheral operation	the HAL call, the register	not modelled: the operation is where the semantics stops (FVD-0140)	not proved (FVI-0023)	ADR-0037 (amended): apply(tick, sinks...) on the RP2040 and on the Arduino Nano, recording sinks on the host, both cross-built in CI; tested	stateful adapters, a device clock, atomic frames, the numeric policy at the boundary
raw command → abstract sink operation	the adapter's policy reads the command; the line holds the last accepted value	Phase 15 Policy, AdapterOp, Line (below the raw command, over the unchanged design)	AdapterOp.det, adapter_of_sink, line_value_accepted, line_holds_on_refusal	ADR-0037 duty8, AdapterOp, TickTrace.adapter; production-tested (host_adapter_operations_correspond_to_the_commands)	the f64 rounding before the range check; anything below the operation (FVI-0023)
peripheral operation → physical world	a register write becomes light or motion	not modelled	not proved, and not testable by this project's means	the firmware runs; no bench measurement is recorded	an empirical matter of electronics, outside BDL

The symmetry of the two halves is real and limited. Both are constructions over designs that add a pure term the behavior cannot observe; both keep the logical object — the Source, the logical Output — as the boundary the designer sees; both refuse effects, kinds and device knowledge in the kernel. They are not the same abstraction. Provision *realizes* an unresolved declaration, moves the source role to the raw reading, and needs an induced input and a joint section because abstract values must be *reached* from raw ones; realization adds structure *below* a value that already exists, changes no declaration's state, needs no induced anything, admits quantization, and only has to *deliver* what the behavior produced. The input side is proved and not built; the output side is proved and built through the raw command; the last arrow of the output side is built and not proved. That asymmetry is the current state, and every chapter that touches the boundary says which arrow it is on.

Two remarks across §IV.1–IV.7

Why the signature is a design object

The most consequential decision in BDL is treating a declared relationship as visible product intent. In programming, a signature is often documentation and a static contract around code. In BDL it can precede any code-like definition and remain useful on its own: $?f : \text{Tilt} \rightarrow \text{Brightness}$ says that the designer has committed to a causal design relationship and to its semantic boundary, and does not say how the mapping is computed. Clients are typed against the type view and depend on the interface only through it and through monotone evidence, so the partial commitment is not a weaker form of a complete one. It is the form on which everything downstream already rests.

Nominal identity in three places

BDL makes the same choice three times. On the axis of quantity, `sem` makes concept identity nominal: two concepts of equal representation are distinct, and moving between them is a declared relationship. On the axis of time, `ClockId` makes temporal identity nominal: two domains of equal rate are distinct, and moving between them is a `sync` with an initial value. On the axis of effect, `OutputId` makes sink identity nominal: two sinks of equal accepted type are distinct, and driving one is an explicit edge. In each case the identity is independent of representation, of rate, and of type respectively; in each case the crossing is a visible artifact rather than a compiler action; and in each case the alternative — identity by representation, by rate, or by type — collided or was ambiguous for a mechanized reason. The three are not all placed alike. Concept identity is in the type, while clock and sink identity are projections beside the interface checked by separate global judgments, because putting the clock in the type forces polymorphism on every pure mapping. Symmetry was not a design goal; it is what remained.

IV.8 Validation and deployment

A lamp that is typed, causal, clock-consistent and output-complete may still not fit the microcontroller it is to run on, and may carry a collection its board's memory cannot hold. Neither question is a question about the design alone, and neither is a typing rule. This chapter concludes the derivation with the decidable validation layers that answer them without touching the kernel, keeps the three verdicts — behavior-valid, executable, deployable — apart, and states the two boundaries below deployment that are held closed by tests rather than proofs (Phases 7, 9a, 14).

Three verdicts that must not be confused

verdict	what it says	depends on	re-established when	who states it
behavior-valid	every realization satisfies its interface; the instantaneous graph is acyclic; every reference respects domains	the design	a refinement preserves it; an edit reopens dependents	<code>analyze(snapshot)</code> ; the kernel judgments of §IV.1–IV.5
executable	behavior-valid, and every required output has exactly one driver	the design and its required outputs	any drive or output change	<code>CompleteOutputs</code> , <code>SingleDriver</code> ; <code>bdl-output</code>
deployable	executable, and the target carries every derived requirement, every chosen realization is admissible (its encoder well typed, fitting the output's concept, its kind's requirements placed), and every collection has a sufficient bound under the schedule	the design <i>and</i> a target, the chosen profiles and a schedule	every design change, every board or profile change	<code>analyze_deployment(snapshot, target)</code> ; <code>bdl-hardware</code> , <code>bdl-output::realization</code> , <code>bdl-reactive::capacity</code> , <code>bdl-exec-ir::bounds</code>

The first two are monotone in the sense of §IV.1 — a refinement preserves them — and the third is not: six actuators fit a Nano and a seventh, a monotone extension of the design, does not; a fast source added to a slow consumer can exceed a window bound that was sufficient. Deployability is evidence about a *pair*, re-solved from scratch, and never merged with the evidence that survives refinement. Production reports the two in different places — the status line and the inspector for the first two, the Deploy page for the third — and words them differently, so that a valid design is never mistaken for a deployable one or the reverse (ADR-0006, ADR-0015 (FVD-0057, FVD-0062)).

Target-specific hardware validation

Everything to this point is board-independent. A design that is typed, causal, clock-consistent, and output-complete may still not fit the microcontroller it is to run on, and that question is answered by a validation layer that never touches the kernel.

Resources, capabilities, requirements

A target is a finite table. A **resource** — a pin — has an identity, a list of **capabilities** from a shared vocabulary (digital in/out, PWM, analog in, interrupt, the I2C, SPI, and UART lines), and, per capability, the **unit** backing it, such as the timer behind a PWM pin or the peripheral behind a bus line. A **hardware description** is a list of resources with a per-capability sharing policy: bus lines are shareable, everything else is exclusive.

A design’s needs are **requirements**. Each has an identity, one capability, an optional fixed resource for a manual pin choice, and an optional membership in a group with a unit relation, *same* or *distinct*:

$$\text{Requirement} = \langle id, cap, \\ fixed : \text{Option ResourceId}, \\ group : \text{Option } (\mathbb{N} \times \text{UnitRel}) \rangle.$$

Requirements are generated from the logical outputs of the previous section by a **device kind** (the requirements half of a device profile; the encoder half is the realization section above): an H-bridge channel needs a PWM line and a digital output; an I2C sensor needs SDA and SCL on the same unit; a quadrature encoder needs two interrupt lines. The pipeline is

$$\text{OutputId} \rightarrow \text{DeviceKind} \rightarrow \text{Requirements} \rightarrow \text{solve} \rightarrow \text{Assignment},$$

and the sink never enumerates pins, so swapping the board is re-solving the same requirements with the design untouched. Requirement identity is independent of sink and declaration identity: one sink generates several requirements, and two identical PWM needs must be two distinct variables.

An **assignment** is a list of requirement–resource pairs. It is **partially valid** when every entry is supported — the resource has the capability and any fixed choice is respected — and every pair is **compatible**: two entries on the same resource must have the same capability and it must be shareable, and two entries in the same group must have equal units under *same* and different units under *distinct*. It is **valid for** a requirement list when partially valid and covering exactly that list; the target is **satisfiable** when such an assignment exists.

A sound and complete solver

Every constraint is unary or binary, so validity is prefix-closed, and an exhaustive depth-first search that prunes on unary support and pairwise compatibility with the current prefix is complete as well as sound. Both are proved: if the solver returns an assignment it is valid for the requirements, and if a valid assignment exists the solver returns one. Feasibility of a finite instance is therefore decidable, and the solver runs inside the proof checker by *decide*. The instance is small enough that a general constraint solver [7] is unnecessary; the claim is about this scope, not about embedded allocation in general.

Extension of a target — adding capabilities, units, or sharing to existing resources — preserves every valid assignment. Removing a resource, adding or strengthening a requirement, or fixing a pin may not.

Case study: an Arduino Nano

The Arduino Nano’s digital and analog pins are encoded as a table: PWM on D3, D5, D6, D9, D10, and D11 backed by timers 2, 0, 0, 1, 1, and 2; external interrupts on D2 and D3; I2C on A4 and A5. A design with four H-bridge motor channels and one I2C inertial sensor is satisfiable, and the solver’s output is the assignment a tool would present:

```
M1 -> D3 / D0      M2 -> D5 / D1
M3 -> D6 / D2      M4 -> D9 / D4
IMU -> A4 / A5
```

Two counterexamples fix the boundary between kernel and validation. Seven independent PWM actuators pass every kernel condition — globally well formed, well clocked, causal, every drive edge well formed, single-driver, output-complete — and are unsatisfiable on the Nano, which has six PWM pins. The same requirements are satisfiable on a larger mock board with six more PWM pins on three more timers. Two interrupt lines and six PWM lines meet every capability *count* exactly, two and six, and are unsatisfiable, because D3 is both the only second interrupt pin and one of the six PWM pins. Capability counting is not feasibility. Further examples establish that two PWM requirements pinned to the same pin are rejected, that two I2C sensors on A4/A5 are accepted since allocation is not

all-different, that a manual pin choice can turn a satisfiable design unsatisfiable while a consistent one is honoured, that four PWM lines required on independent timers are unsatisfiable on the Nano’s three timers though six PWM pins exist, and that TX and RX pinned to one UART unit stay together.

The explanation facility reports the first dead end under greedy placement. For the seven-actuator design it names the seventh actuator and, for each PWM pin, the actuator blocking it. This is *a* conflict under one placement order, not a minimal unsatisfiable core, and it is meaningful only when the solver has already returned no assignment.

Two kinds of evidence

Hardware feasibility is evidence about the pair (design, target), and it is not monotone in the sense the preservation theorems require. Six actuators are satisfiable; adding a seventh — a monotone extension of the design by a declaration and its sink — is not. The distinction drawn earlier between evidence that survives refinement and evidence that is rechecked after every change is here concrete. Commitments discharged compositionally survive *EnvRefines*; deployability on a target is re-solved; and the two are never merged. This is why the workspace reports them as different states.

What the layer does not model should be stated as plainly. Voltage, current, thermal budgets, memory, processor load, deadlines, bus bandwidth, torque, travel, and PWM frequency values are outside it. Some of these need summation constraints that are not binary, and while the architecture leaves room for a compatibility predicate over sets, nothing here establishes it. Units are one integer per capability per resource, which models which timer or which UART but not timer modes.

Deployment capacity

The third kind of validation evidence, added in Phase 9a, is deployment capacity: whether every bounded collection in the design has a bound sufficient for the target’s memory and the schedule. Like feasibility and unlike a commitment, it is not monotone — adding a fast source to a slow consumer can exceed a bound that was sufficient — and it is re-established after every change. §IV.5 gives the formal result (*sufficient_capacity_preserves*, *periodic_capacity_sufficient*; **formally proved**) and production’s realization (*bd1-exec-ir::bounds*, *refusal at deploy*; ADR-0027). The workspace state *hardware-feasible* of Part III is, in production, the conjunction of allocation and capacity.

Provision and realization as deployment architecture

§IV.7’s two constructions are deployment constructions, and they sit in the same place in the picture — between *executable* and *deployable* — with opposite implementation status. Provision takes the executable design and a device profile per Source and yields another design, provably an environment refinement of the first, in which every abstract Source is a Value computed from a raw declaration; *Fits* — the concept’s representation is the channel’s — is a deployment check beside allocation and capacity. Nothing of it is built at the snapshot (PRP-0001, ISS-0016); the placement is recorded so that when it is built the transducer is a checked term in the deployment pass and not host code in the adapter, and the first adapter refuses a design with a Source rather than guess. Realization takes the same executable design and a chosen profile per driven output and yields the lowered design with one encoder and one machine sink per output; *admissibility* — the encoder’s typing, *EFits*, and the kind’s requirements placed by the solver — is a deployment check beside allocation and capacity, composed in *analyze_deployment* and reported per judgment on the Deploy page. It is built (ADR-0036), and the artefact is refused when a chosen profile is not admissible (*backend.realization_invalid*). The third verdict of the table above is therefore, in production, allocation *and* admissibility *and* capacity; the first two are the pair (design, target, profiles) and the last is (design, target, schedule), and none of them is monotone.

The generated-code boundary

Below deployability is one more boundary, and it is the weakest link in the trust chain. The formal semantics is a tick-indexed relation; production’s reference evaluator is its executable transcription; the generated *no_std* core (Part V) is an implementation of that evaluator which lowers representation — dense slots, static structs, *¶64* fields, inlined lambdas — and may not change meaning. The core is held to the evaluator by a 22-case differential corpus, golden files, a determinism check, property-generated designs and an error-tick agreement check (**production implemented and tested**), and since output realization the corpus compares the core’s raw commands (*Tick.commands*) against the encoder applied to the reference evaluator’s logical outputs. No theorem relates the generated

code to Ev/MEv, none relates the generated commands to RawCommand, and none relates the static list-bound analysis to the capacity theorems: a generated-code refinement proof and a proved bound analysis are the open items that would close the gap (§VII.4). Differential tests are evidence for the corpus; they are not a proof, and this document never infers formal correctness of generated code from them.

Below the generated core there is now one more layer, and the trust chain gets one link weaker still. The platform adapter (Part V; ADR-0037) applies each tick's raw commands to peripherals on the Raspberry Pi Pico or the Arduino Nano. What holds it to the raw command trace is a recording host counterpart — the same generated apply over mock sinks, so a host trace carries the operation sequence the firmware would perform — and a cross-build of each family's firmware in CI. That is **production implemented and tested**; it is not a theorem, the formal model stops at the command (FVD-0134), and no bench measurement of a physical effect is recorded anywhere in either repository. The numeric policy at that boundary is likewise a tested production decision, not a formal one: production computes in f64, the encoders emit f64 commands (pwm_duty8 carries 127.5 for 50 %), and the adapter's explicit, deterministic duty8 policy rounds a finite in-range duty to the nearest whole value and *refuses* an out-of-range or non-finite one, holding the line's last value rather than clamping — a third numeric domain beside the kernel's Nat and production's f64, recorded in §VII.1.

Open hardware questions

The validation layer covers discrete pin and peripheral allocation with unary and binary constraints, and the explanation facility reports a first dead end under one placement order. What it does not cover is stated once, in §VII.4: minimal unsatisfiable cores; voltage, current, thermal and timing budgets, which need summation constraints that are not binary; timer modes and PWM frequency values; the device catalogue that provision presupposes on the input side and that the five witness profiles stand in for on the output side; and a device that provides a Source's value on a board. Board files exist for the Nano, a larger mock board and the Raspberry Pi Pico (rp2040_pico), the last with a generated adapter behind it; runtime loading of board files, a second microcontroller, flashing and telemetry are roadmap items, not records of anything that exists.

Part V — From Semantics to an Executable Toolchain

Part IV is a specification: it says what a design means and what may be decided from what. This Part is how that specification becomes a toolchain — KCN-judu/BDL, a Rust workspace with a Flutter authoring environment in front of it — without the toolchain ever redefining the language. It is organized by *responsibility* rather than by crate: for each layer of the tooling path of Part II, what truth it owns, what it consumes, what it emits, and what it must never decide. Crate names appear as implementation references, and the map of crates is the last chapter. The production described is the pinned snapshot (081296d, Appendix F); it follows the formally developed semantics and is not itself formally verified, and the one property the whole system is built to keep obvious is that *BDL semantics flows downward; implementation mechanisms never flow upward and redefine the language*.

layer	owns	consumes	emits	never decides
authoring surfaces (Studio, Code view, the language server)	presentation, interaction, layout, drafts	the compiler's projections and verdicts	edits against a revision; text as typed	a type, an identity, a dimension, a clock, a driver, a placement
the semantic model and persistence (bdl-model, bdl-text, bdl-layout)	stable identities, the surface model, revisioned edits, the derived role, the project on disk	edits, source files, the identity sidecar	a new revision, the textual projection, the placement of unpositioned entities	a semantic verdict; a layout that means anything
elaboration and checking (bdl-syntax, bdl-elab, bdl-equations, bdl-check)	the kernel's terms: Θ , interfaces, Core formulas; every type, grant and dimension judgment	the model	the Design IR with diagnostics	anything about a board or a device
analysis (bdl-reactive, bdl-output, bdl-hardware, bdl-compiler)	dependency, causality, clocks, drive edges, the single driver; the target-relative deployment analysis with its three verdicts	the Design IR; a board file; chosen profiles	ProjectAnalysis, DeploymentAnalysis	a design's meaning from a board
the reference evaluator (bdl-reactive::eval)	the executable definition of runtime behavior	the Design IR and an input trace	values, outputs and commands per tick	a second semantics
lowering and code generation (bdl-exec-ir, bdl-lower, bdl-codegen-rust)	the executable IR, the generated no_std core, the adapter plan and glue	the checked design and the solved deployment	a crate whose step yields Tick { values, outputs, commands }, a manifest, firmware for a target	a change of meaning
the platform adapter (bdl-runtime-adapter, a HAL binding per family)	the tick loop on a board, the numeric policy at the boundary, faults and start-up levels	Tick.commands and the adapter plan	peripheral operations	anything upstream of a command; a pin; a value the core did not produce
the IDE service and daemon (bdl-ide-db, bdl-ide, bdl-lsp, bdl-daemon)	one semantic project with overlays, one revision stream, the semantic queries	the model and the analysis	projections, verdicts, edit plans, tokens, the protocol	a fact the compiler did not state

Four trust layers

layer	owns	never does
Flutter Studio	presentation, interaction, layout, ephemeral render state	compute type validity, concept identity, dimensions, causality, clocks, output ownership, hardware feasibility, simulation

layer	owns	never does
Rust compiler (bdl and crates)	the canonical project model, every semantic judgment, diagnostics, simulation, allocation, code generation, the placement of entities that have no position yet	render, decide where a placed node goes
generated Rust core	deterministic executable behavior: domain step functions, state, output values, and — after realization — the raw commands of the chosen profiles	touch hardware, know about tasks or executors
platform adapter	activating the compiled schedule, applying each tick's raw commands to peripherals, the numeric policy at that boundary, faults and start-up levels	interpret BDL semantics, read a declaration, choose a pin, invent a value

Studio may render an edit optimistically, but the truth comes back from the compiler as a *projection*; Studio never holds a second copy of the language (ADR-0001). bdl is a process boundary (ADR-0002), speaking protobuf over framed stdio (ADR-0007), with one canonical revision stream per opened project.

V.1 Authoring surfaces and the semantic project

A design is authored as a graph or as text, and both are views of one project that the daemon owns. What makes that possible is a persistence model in which every semantic fact has exactly one home and identity is independent of both spelling and position.

A project is one directory (ADR-0023, superseding ADR-0020's third project *kind*): bdl.toml (name, schema, compiler version — no kind), src/**/*.*.bdl (the authored design and system — the semantic source), .bdl/identities.json (source key → stable id, allocators, the flat-id freshening table — tool-owned), .bdl/authoring.json (behavior groups by identity), ui/layout.json (positions, viewports, group boxes — presentation). “Source of truth” is a per-fact statement: the semantic model never depends on geometry, the graph never stores a semantic fact, the text never stores a position; erasing both sidecars and the layout changes no semantic fact; erasing the sources loses the design. A legacy JSON project is migrated in place the first time it is opened, identities and layout kept.

bdl's coordinator is the single imperative loop: it owns the Session, applies edits serially through the pure apply_edit, and hands immutable snapshots to analyses. Every semantic edit is sent against the revision Studio holds and refused if the project has moved on; every response and event carries its revision; revisions are strictly monotone (undo produces a new revision), so staleness is a < comparison. Overlays — a definition draft per mapping, a text document per file — never create a revision; AnalysisSnapshot composes committed + overlays, runs analyze, and is stamped (revision, generation); a query holding an old snapshot keeps seeing the world it started in.

V.2 Elaboration and checking

The elaborator is where the surface of Part III becomes the kernel of Part IV, and the checker is the authority every later pass trusts. The pipeline is a sequence of explicit passes, each with a stated input and output, and diagnostics are first-class outputs rather than failures: an incomplete project is the normal case.

```
load/parse → identity resolution → signature resolution → surface elaboration
→ type checking → semantic-construction (grant) checking → dimension checking
→ dependency analysis → causality → clock domains → physical outputs
→ hardware requirement generation → hardware allocation → realization admissibility
→ reactive lowering (with one machine sink per chosen realization) → Rust code generation
→ [target] adapter plan → generated adapter glue and firmware
```

Each pass has an explicit input and output type and is pure where practical (docs/architecture/compiler-pipeline.md). Diagnostics are first-class outputs: an incomplete project is the normal case, never a fail-fast, which is the engineering face of the signature-first position. apply_edit classifies every operation as a *refinement* (dependents' established facts remain valid) or an *edit* (dependents must be re-validated) and reports an explicit Invalidation set — Interface | Realization | Semantic | Reactive | Clock | Output | Deployment — with the originating declarations (ADR-0009); incremental analysis subscribes to these categories, so the refinement-versus-edit distinction of §IV.1 is an engineering asset, not folklore.

Numerics are IEEE f64 (ADR-0011), a recorded deviation from the kernel's Nat: division by an exact zero and any non-finite result fail the tick with a structured error, equality is exact, and the production numerics are not what the formal theorems are about (§VII.1).

Formula elaboration (ADR-0013, extended by ADR-0025): a designer-facing formula language $- + - * / < = > >= == != \&\& || !, \text{if}, \text{match}$ over Bool/Option/numbers/counts, let, calls, list and pair literals, rules $x \Rightarrow \dots$ as equation arguments, membership $x \text{ in } [\dots]$, the slot $?$ — parsed by a hand-written Pratt parser into a surface AST that is never reused as Core. Input names are the concepts' display names, re-resolved on every analysis; inputs appear as their representations ($\text{rep } (\text{var } i)$) and the whole formula is wrapped in $\text{mk } B$ under the declaration's own grant — the elaborator never emits mk of any other concept, and the checker refuses one anyway. Units elaborate through the registry; equations are matched and inlined at the use; bdl-check re-derives the Core term's type and is the authority.

The surface editor and kernel are connected by an elaboration function from surface designs to kernel environments with diagnostics. These passes were stated in the formal development before any elaborator existed; the production pipeline above realizes them — resolution and formula elaboration in bdl-elab , temporal lowering and domain assignment in bdl-lower and bdl-reactive , output binding in bdl-output , hardware validation in bdl-hardware , normalization and erasure in bdl-exec-ir and bdl-codegen-rust . The list is kept because it is the specification those crates were written to, and because it says what each pass may and may not consult.

Name and signature resolution. Resolve concepts to identities, Mapping signatures to declaration interfaces, contexts, device kinds, units, and imported components. At this stage an unresolved $\text{Tilt} \rightarrow \text{Brightness}$ mapping already has a stable identity and an expected type, and every reference to it is typed.

Formula elaboration. For each Mapping with a definition, check the definition against the declared codomain under the grant of the declaration's own signature. A scalar formula attached to a Brightness output elaborates to $\text{mk}_{\text{Brightness}}(\dots)$ because the surrounding signature announces the concept; the constructor is inserted by the elaborator and is legal only there. Units elaborate to scaled dimensioned literals. Curve, example, and component definitions elaborate to the same realization form.

Temporal lowering. Surface temporal phrases lower to the declaration shapes of the derived-operator table. A context lowers to an activation declaration, an entry declaration, gated and reset local state, and a conditional in each output's driver. A reusable stateful component is instantiated into fresh declarations at each use, because temporal state cannot occur under a binder.

Domain assignment. Domains are declared, so this pass records the clock environment and checks the domain judgment. A reference across domains without a transport is reported at the reference, in the designer's names for the two domains, with the two legitimate resolutions — transport the value with a stated initial value, or move the reader to the source's domain — stated in terms of their behavioral consequence.

Output binding. Record the drive edges and check type and clock equality, single-driver, and, for executable designs, completeness. Where several contexts would drive one output, the elaborator emits one driver whose body selects among them; the selection is visible in the elaborated design.

Hardware validation. Generate requirements from device kinds, run the solver against the selected target, and attach the assignment or the explanation to the bindings.

Normalization and erasure. After checking is complete, concept identities, dimensions, and domains carry no computational content and may be erased; erasure is proved sound for identities and for dimensions. Nominal wrappers introduced by elaboration cancel, $\text{rep}(\text{mk}_C e) \rightsquigarrow e$, and the flattened program is checked under the universal grant because every construction was authorized at its declaration. Unfolding all realizations into a closed term preserves typing on the delay-free fragment and agrees with tick-by-tick evaluation on first-order designs; the higher-order case holds up to closure equivalence and was not formalized. Erasure is applied selectively at boundaries the checker cannot see through — supplied blocks, device bindings, the public interface of generated code — where wrappers are retained so that the host compiler continues to check what BDL cannot.

The kernel never depends on normalization to decide type equality; it is not dependently typed, and normalization is an analysis and code-generation instrument. Floating-point arithmetic is not associative, so any symbolic normal-

ization over the reals must record the resulting numerical deviation as an obligation rather than silently altering the property being checked.

V.3 The reference evaluator

The formal semantics of §IV.5 is a tick-indexed relation, not a program. Production needs a program, and it needs one whose meaning is *defined* to be that relation rather than approximated from it — otherwise the generated code would be held to nothing. The reference evaluator (`bd1-reactive::eval`) is that program: the executable definition of runtime behavior, transcribed from `Ev/MEv` with the two-phase tick, the previous/next state and the strictly-before rule for `sync`, and used in three roles. It is what the Simulate page runs, so a designer sees the semantics and not an approximation of it; it is the *differential oracle* against which every generated core is compared value by value and tick by tick (ADR-0016); and since output realization it also carries the encoder over each realized output’s value, so that the generated raw commands are compared to it as well. It is not a second semantics: where the evaluator and the model could disagree — numerics are `f64`, the model’s are naturals — the disagreement is recorded as a deviation in Part VII, not resolved in the evaluator’s favour.

For the same design, schedule and input trace, the reference evaluator and the generated program must agree value by value, tick by tick (`crates/bd1-compiler/tests/backend/*.rs`). The corpus: `lamp` and `lamp-with-output`, pure arithmetic with dimensions and a count, collections (`any`, `sum`, `map`, `filter`, `zip`, `contains`, `clamp`, `getOrElse/head`, structural equality, a list in a state cell), `buffer` (the Phase-9a window as five declarations), `bounded_buffer` and `overflowing_buffer`, semantic `rep/mk` with a Boolean concept, `booleans/comparisons/strict if/options`, `delay`, a cycle broken by `delay`, `sync` across two domains in both directions with the slow domain on period 2, a domain-agnostic declaration shared by two domains, division by zero, non-finite result, missing input, a design with no domain — 22 cases, each generated, `cargo` checked as a `no_std` library, built with its bridge, run and compared, with golden files pinning the generated bytes and a determinism check that generating twice gives identical output; a 3 000-tick run; property-based generation of designs (250 per run in process, a seeded batch of 8 through the full toolchain); and the storage-order audit of the reversed `Vec`. Comparison is exact (`f64` bit-for-bit in effect; the generated code performs the same IEEE operations in the same order — no reassociation, fusion or CSE — and the JSON transport is exact); errors must fail at the same tick with a failure the reference could report, computed by starting the reference’s traversal at every due declaration.

V.4 Lowering and the generated core

The evaluator interprets a design; a product cannot afford to. Lowering turns the checked relationships into a first-order executable IR — dense slots for clocks, state, inputs, outputs and machine sinks, lambdas inlined, the recursor a single runtime loop — and code generation prints a `no_std` Rust crate that is statically laid out and, without collections, allocation-free. What the core must preserve is stated once (`docs/spec/runtime-semantics.md`) and repeated here because it is the contract every platform adapter inherits.

`bd1-exec-ir` is a first-order IR with dense slots: clock slots, state cells addressed by `StateCellId { decl, path }` (the declaration and the expression path of the `delay/sync` inside its realization, so identity survives unrelated edits), input slots for unresolved declarations, output plans with their driver. Lambdas and applications are inlined as `let` bindings; the recursor becomes `Fold { elem, acc, step, init, list }` with one closure per recursor applied by the runtime — never a closure *value*; `rep/mk` become `Unwrap/Wrap` on a per-concept newtype. Lowering fixes an evaluation order from the plan; the reference evaluator’s own tests show its result is independent of the host’s processing order and the generated program’s traces equal the reference’s.

The generated core

```
pub const DESIGN: &str; pub const HAS_DOMAINS: bool; pub const CLOCK_COUNT: u16;
pub const CLOCK_k: ClockSlot;
pub struct SemN(pub Repr); // one per concept carried
pub struct Cells { pub cell_k: Option<T>, ... } // temporal state, None until first written
pub struct State { pub cells: Cells }
pub struct Inputs { pub decl_n: Option<T>, ... } // unresolved declarations
pub struct Values { pub decl_n: Option<T>, ... } // every declaration, None when not due
pub struct Outputs { pub output_n: Option<T>, ... }
pub struct Commands { pub command_<device>: Option<R>, ... } // the raw command per realized output
pub struct Tick { pub values: Values, pub outputs: Outputs, pub commands: Commands }
```

```
pub fn init() -> State;
pub fn step(state: &mut State, active: ActiveDomains, inputs: &Inputs) -> Result<Tick, RuntimeError>;
```

A program without lists is Copy, statically sized, laid out by the compiler: no graph, no map, no allocation, no traversal at runtime. Types: $q\ d \rightarrow f64$ (the dimension is static and in the manifest), $bool, nat \rightarrow u64$, $sem\ s \rightarrow SemN$, $opt\ \tau \rightarrow Option<T>$, $list\ \tau \rightarrow Vec<T>$ (last element first), $\tau \times \sigma \rightarrow (T, S)$; a function type has no runtime representation. Symbols derive from stable ids (`decl_17`, `Sem3`, `cell_0`), never from display names. The core is `#![no_std] #![forbid(unsafe_code)]` and mentions no HAL, pin, peripheral or board; cargo check --lib of every corpus crate is a test.

Embedded execution model

The formal semantics is a tick-indexed relation; production gives it a deterministic step function per domain and states what the generated core must preserve (`docs/spec/runtime-semantics.md`):

- **Declarations are not tasks** (ADR-0004). Never one async task per declaration. Per clock domain, one deterministic $step(prev_state, inputs) \rightarrow (next_state, outputs)$: hardware timer or interrupt $\rightarrow$ activate domain $c \rightarrow step \rightarrow$ publish c 's snapshot $\rightarrow$ commit the physical outputs owned by c . Meaning never depends on executor task order.
- **Two phases**. *Read*: every declaration due this tick is evaluated (lazily, memoized) with temporal forms yielding their cell's committed value, or the initial value if the cell was never written. *Write*: every temporal site whose writing domain is active — the owner's domain for `delay`, `src` for `sync` — evaluates its operand in the same read mode into the *next* state. Nothing is updated in place; reads never see writes of the same tick.
- **Previous/next state**. `step` reads only `prev`, writes only `next`, and assigns at the end — not at all on error. Every `delay` carries its explicit initial value; a cell is `Option<T>`, `None` until first written; there is no implicit zero.
- **The sync snapshot rule**. `sync src init e` reads the last committed snapshot of `src` from an activation strictly before the current tick, `init` if there was none, and never invokes `step_src` recursively. When two domains are ready at the same instant, each observes only the other's previously committed activation — the scheduler's order is unobservable, exactly as in the formal model (`scheduling_order_observable` is the counterexample for the alternative). In generated code a `sync` cell's writer is the *source* domain's slot.
- **Output commit timing**. Outputs are built after the write phase from the driving declarations' values (`output_values`), and committed by the domain that owns the output; a realized output's raw command is the profile's encoder applied to that value, in the same tick, `None` when the driver was not due.
- **Allocator requirements**. A core that carries a list needs a global allocator (ADR-0024); the manifest records `requires_allocator`, per-cell bounds, `state_bytes_max` and `tick_bytes_max`; the platform adapter declares an arena sized from the manifest (`state_bytes_max + tick_bytes_max`, rounded up to KiB) for a bounded design, none for a scalar-only design, and refuses an input-bounded or unbounded design for a board.
- **Collection bounds**. Static, sound per declaration and cell (`bd1-exec-ir::bounds`); the window capacity model per crossing; refusal of an unbounded state on a bounded-memory target (ADR-0027, §IV.5).
- **Cost discipline**. Moves, not clones: a use analysis per expression tree moves a local referenced once in its scope, clones any other reference, and always clones a local captured by a fold's closure, so a fold step that conses onto its accumulator moves it — `map`, `filter`, `append`, `sum`, `any`, `all`, `contains` clone nothing per element (clone-counting tests; measured $\times 15$ for 2 000 $\rightarrow$ 32 000 elements). A total conditional — neither branch can fail — is emitted as a Rust `if`, observationally the strict `ite`; a branch that can fail keeps `prim::ite` and its strictness. Borrowed reads for `length`, `head`, `take`, `==`, `<`. Commit-only writes: a tick clones no cell it does not read. What remains: `zip` clones its accumulator pair per element (ISS-0013); no fusion of `map` $\rightarrow$ `filter` chains, because each is linear and measured and a fused emission was not justified by the numbers.
- **The realization boundary**. A Source (an unresolved unit-domain declaration) is an *input slot* of `step`, filled from a simulation trace today and, when a device provides it, by the adapter (ISS-0016); a driven output is an *output field* of the step's result, and a realized output additionally a *command field* — the encoder's raw value — which the adapter applies to a peripheral. Neither is a function call inside the core: a zero-input relationship compiles to a zero-argument accessor of the committed value (its unit argument erased, ADR-0029), and a sink is a field, not an `A -> ()` callback (§IV.7). The core therefore has no device vocabulary at all; the encoder is a pure expression in the plan, and the peripheral is the adapter's.

- **Targets.** The core is target-independent and unchanged by a target except for one feature-gated `mod_adapter`; line; cargo check --lib of the core still compiles with no feature and no HAL. macOS and Windows are first-class hosts for the toolchain; the first embedded target is the Raspberry Pi Pico (next section).

Current cost figures are recorded, not optimized: a lamp core is ~95 lines and ~7 KB of source with a zero-byte State; the collections and buffer corpus cases allocate and their debug-build timings are two orders of magnitude above the allocation-free cases — a baseline for later evaluation, nothing more.

V.5 Platform adapters: one command plan, several host mechanisms

The adapter is the layer that turns an already-lowered raw command into a physical effect without the board reaching back into the design (ADR-0037; docs/architecture/embedded-adapter.md). Its architectural significance is negative: it *interprets nothing*. It reads `Tick.commands` and nothing else — no declaration, no output, no concept — and each command becomes one peripheral operation. Everything a design means was decided above it.

Board and generation. The first target is the Raspberry Pi Pico (`rp2040_pico`, family `rp2040`, thumbv6m-none-eabi), a board file like the others; the Arduino Nano (`arduino_nano`, family `arduino`, avr-none) is the second, described at the end of this chapter. `bdld compile --target rp2040_pico [--tick-micros N]` generates, beside the unchanged core, the adapter glue — `apply(tick, sink1, ...)`, one `&mut dyn PwmDuty8` or `&mut dyn Level` parameter per machine sink, in sink order, with the bindings as data and no map, name or string dispatch — and the firmware that constructs the sinks on the assigned pads, ticks the core, activates the schedule and applies each tick’s commands. Building the firmware is a cargo build for the embedded target; orchestrating that build inside `bdld`, flashing through `probe-rs` and telemetry back into Studio are roadmap priorities 2–4 and do not exist.

Solved deployment, consumed. Command sink identity is one stable route with no display name on it: the device binding’s id names the `Commands` field, the manifest sink and the requirement; the solver’s assignment names the pad (GP15); the target entry derives the peripheral from the pad number ($GP_n \rightarrow PIN_n$; PWM slice $(n/2) \% 8$, channel by parity) and checks the board file’s slice against it. Every step is checked and none has a fallback: an infeasible placement, an inadmissible realization, a sink without an assigned resource, a resource without the capability, a profile with no sink on this target (`i2c_level8`, `hbridge_signed`), a design with a Source (ISS-0016), an unbounded collection and a zero base tick are each refused by the adapter plan with a named code. A device that is placed but realizes nothing gets no sink and no peripheral; no value is invented for it.

Clock activation. The firmware invents no clock semantics. One Embassy Ticker fires at the base tick; at global tick t the slots with $t \% PERIODS[slot] == 0$ are active — the simulator’s rule, verbatim, in the runtime crate and tested against it; then one global step (ADR-0004) and `apply` in sink order, the interpreter’s order. A domain that is not due leaves its sinks untouched. No cross-domain synchronization happens in the adapter; the compiler resolved or refused it upstream.

The numeric policy at the boundary. Production computes in f64 and the encoders emit f64 commands; the peripheral takes integers. The conversion is one explicit, deterministic function of the value alone, `duty8`: a finite raw duty in $0 \dots 255$ rounds to the nearest whole duty, halves up; anything else — out of range, NaN, $\pm\infty$ — is refused and the line holds its last applied value, with the refusal recorded. No clamping and no `as cast` decides semantics: a design that commands 105 % has said something the profile did not promise to carry, and the peripheral holds rather than guesses. A `bool` command is applied as written. The PWM slice is configured with `top = 254`, so a duty d is high for exactly $d / 255$ of the carrier period; the carrier (≈ 30.6 kHz) is peripheral configuration, never a `ClockId` (FVD-0138). Every line starts low before the first tick; a failed tick latches — state unchanged, lines hold, the firmware waits for interrupts until reset — and a refused command is not a fault. These are adapter policies, recorded so they can be revisited with evidence; none of them is a truth of the language.

What is tested, and what is not. The host path records the same operations: `HostProgram::adapter_ops` applies the generated `apply` to mock sinks, so a host trace carries the operation sequence the firmware would perform, and `crates/bdl-compiler/tests/embedded_rp2040.rs` and the runtime’s own tests hold the glue, the policy and the schedule to that; the firmware cross-compiles in CI. That is the whole of the evidence: **production implemented and tested**. The correspondence from the raw command trace to the adapter’s operations is not a theorem, the correspondence from a register write to a physical effect is not measured, and the formal model stops one layer above at `RawCommand` (FVD-0134; FVI-0022). The adapter closes the *engineering* gap that PRP-0001 named for the input side and ADR-0036

for the output side — a transducer or an encoder as unchecked host code — by leaving the encoder in the checked plan and keeping the adapter free of semantics; it does not close the *formal* gap, and this document does not say it does.

Two families from one plan. The snapshot has two target families and one entry per family (`bdl-codegen-rust::targets::Entry`): the RP2040 over `embassy-rp`, asynchronous, an Embassy Ticker for the base tick and a collection arena sized from the manifest; and the Arduino family over `avr-hal`, synchronous, a blocking `tick_wait` after each step so that a tick lasts at least the base period, 8-bit PWM natively, no arena — a design with a bounded collection is refused for the AVR (`adapter.collections_unsupported`) — and a nightly-only cross-build (`--target avr-none -Zbuild-std=core`). The Nano’s board table (D3/D11 on timer 2, D5/D6 on timer 0, D9/D10 on timer 1) is checked against the board file’s units; an Uno would be a second table, a Mega a table of its own. What the two share is everything upstream: the plan that binds each machine sink to the solver’s resource, the generated glue `apply(tick, sink, ...)`, the numeric policy and the sink traits of the vocabulary crate (`bdl-runtime-adapter`, renamed from `bdl-runtime-embassy` when the second family arrived, because the crate says what it is rather than who first used it), and the host’s recorded operations. The architectural point is the one the whole Part makes: the executor mechanism — an async runtime or a blocking loop, a timer or a busy wait — is not BDL semantics, and a second family adds an entry and touches nothing above it.

V.6 The IDE service, the daemon and the protocol

The compiler’s verdicts reach a designer through one more layer, and it is the bridge into Part VI: a service that answers, in the compiler’s own terms, what is at a position, what may be written here, what a name refers to, what would fix a finding — for a text editor, a canvas and an external language server alike.

Studio (visual) and text editors (textual) are two projections of one semantic model and consume one service. `bdl-ide-db` holds the ground state — the committed snapshot plus overlays for unsaved drafts and text buffers — and produces immutable, stamped `AnalysisSnapshots`; `bdl-ide` answers diagnostics, hover and Explain, scope-aware completion, definition/references/rename across files, semantic actions and edit plans, invalidation previews, symbols, semantic tokens and the Formula Composer’s three queries, all in BDL-owned types keyed by `EntityRef` and role — never by text position or node id; `bdl-lsp` and `bdld` are thin adapters at the edge (ADR-0017). The daemon speaks a typed, versioned protobuf protocol over framed `stdio` (ADR-0002, ADR-0007). Every semantic edit is sent against the revision the client holds and refused if the project has moved on; every response carries its revision; revisions are strictly monotone, so staleness is a `<` comparison (ADR-0009). Protocol versions are additive within a minor series and volatile: the version at the snapshot and what each recent minor added are in Appendix F, not here.

Incremental invalidation. `apply_edit` classifies every operation as a refinement or an edit and reports an Invalidation set — Interface | Realization | Semantic | Reactive | Clock | Output | Deployment — with the originating declarations. That is the refinement-versus-edit distinction of §IV.1 as an engineering asset: what is re-analysed after an edit is decided by a model of which established facts an edit can disturb, not by folklore. A proved per-category preservation theorem does not exist (§VII.4).

One classifier for text. Highlighting looks lexical and is not: whether `tilt` is a Source or a Value, whether `deg` is a unit or a parameter, whether `all` is a binder word here, whether `clamp` is the library’s — each is a fact of the snapshot, the same fact that drives hover, Explain, completion and the canvas. So there is one classifier, `bdl-ide::tokens`: a lexical layer from the lossless syntax tree (total on any text, so a file mid-edit keeps its reading) merged with a semantic layer from the snapshot’s projections and the derived role, into one sorted, disjoint stream of byte-range tokens. Its vocabulary is the LSP semantic-token vocabulary — the standard types and modifiers, themed by every LSP editor without configuration — plus two BDL additions, `unit` and `slot`, with BDL’s distinctions expressed as *modifiers* on standard types (a Source is `variable + source`), under a versioned legend. The language server encodes the stream as `textDocument/semanticTokens` in the client’s position encoding through a pure encoder for UTF-8, UTF-16 and UTF-32; the daemon sends the same stream structured to Studio; conversion between the two is lossless and lives in `bdl-ide`. No Dart, TypeScript or grammar file decides a class; no colour crosses a boundary (ADR-0035). The rule is the one ADR-0001 sets for semantics generally — Rust classifies, clients render — applied to the one place a client would most naturally have kept its own copy of the language.

V.7 The implementation map

For reference, the crates and the direction of dependency among them.

The dependency direction is strict and acyclic: `model` → `ir` → {`syntax` → `elab`, `check` → `equations` → `elab`, `check` → `reactive` → `output`} → `compiler` → `ide-db` → `ide` → {`lsp`, `daemon`}; `protocol` sits between `compiler` and `daemon`; `text` and `layout` hang off `system` and are joined by `daemon`, `lsp` and the CLI; `hardware` depends on `model` only — it never sees Δ — and `compiler` joins the two; `lower` → `codegen` hang off `exec-ir`; `runtime-core` depends on nothing and is what generated code links against. A crate exists only where a real boundary exists.

crate	owns
<code>bdl-model</code>	stable ids · surface model · revisioned edits (<code>apply_edit</code> , <code>pure</code>) · persistence · quantity vocabulary · the derived relationship role (<code>RelationshipRole</code> , <code>MappingBlock::role</code>)
<code>bdl-ir</code>	the Design IR and the Reactive Core IR: the kernel's <code>Ty</code> , <code>Expr</code> , <code>Prim</code> , <code>environments</code> , <code>transcribed</code>
<code>bdl-diagnostics</code>	<code>Diagnostic</code> , <code>Span</code> , <code>stable codes</code> , <code>deterministic order</code>
<code>bdl-syntax</code>	<code>Logos lexer</code> · <code>event parser</code> (recursive descent + Pratt) · <code>Rowan lossless CST</code> · <code>typed AST</code> · <code>lowering</code>
<code>bdl-equations</code>	the equation library: rank-1 schemes with type and dimension variables and { <code>Data</code> , <code>Eq</code> , <code>Ord</code> }, first-order matching, one closed Core builder per equation
<code>bdl-elab</code>	<code>concepts</code> → Θ · <code>signatures</code> → <code>interfaces</code> · <code>formulas</code> → Core with equations inlined at their instance · the unit registry and charts
<code>bdl-check</code>	Core typing · <code>Grant</code> · <code>realization vs interface</code> · <code>pretty-printing</code> ; the authority for every type judgment
<code>bdl-reactive</code>	<code>dependency graph</code> · <code>causality</code> · <code>Clocked</code> · the reference evaluator · <code>simulation</code> · <code>window capacity</code>
<code>bdl-output</code>	<code>DriveWF</code> · <code>SingleDriver</code> · <code>CompleteOutputs</code> · <code>output_values</code> · <code>realization</code> : the profile registry, <code>Encoder { rep, raw, encode }</code> with its well-formedness and purity checks, the three judgments of admissibility
<code>bdl-hardware</code>	<code>capabilities</code> · <code>resources</code> · <code>hardware</code> · <code>device</code> → <code>requirements</code> · <code>boards</code> · <code>solve/diagnose</code>
<code>bdl-exec-ir</code>	the executable IR: <code>slots</code> , first-order expressions, <code>evaluation plan</code> ; <code>interpreter</code> ; <code>static list bounds</code>
<code>bdl-lower</code>	<code>reactive lowering</code> : <code>Design IR</code> → <code>Exec IR</code> (<code>clock</code> , <code>state</code> , <code>input and output slots</code> ; <code>inlining</code> ; <code>order</code>)
<code>bdl-codegen-rust</code>	<code>Exec IR</code> → <code>owned Rust AST</code> → <code>printed crate</code> + <code>host bridge</code> + <code>bdl-manifest.json</code> ; <code>adapter</code> and <code>targets:: { rp2040, arduino }</code> ; the generated adapter glue and one firmware entry per target family
<code>bdl-compiler</code>	<code>analyze(snapshot)</code> , <code>analyze_deployment(snapshot, target)</code> (<code>allocation</code> , <code>admissibility</code> , <code>capacity</code>), <code>compile(snapshot, options)</code> with the realizations lowered to sinks, <code>target:: adapter_plan</code> binding each sink to its assigned resource, the <code>collections report</code>
<code>bdl-system</code>	<code>components</code> · <code>instances</code> · <code>bindings</code> · <code>freshening</code> · <code>flatten</code> → <code>ProjectSnapshot</code> + <code>origins</code> · <code>packaging</code>
<code>bdl-text</code>	<code>persistence</code> : <code>source discovery</code> · <code>identity sidecar</code> and <code>reconciliation</code> · <code>load_workspace</code> · <code>item-level write-back</code> · <code>legacy migration</code>
<code>bdl-layout</code>	<code>deterministic, incremental placement of entities without a position</code> ; <code>never semantics</code>
<code>bdl-library</code>	the Standard Library: items whose fragments instantiate ordinary concepts and Sources through the ordinary edits; <code>search</code> ; a multi-library set
<code>bdl-ide-db</code>	IDE ground state: <code>IdeHost</code> · <code>overlays</code> · <code>EntityRef/EntityRole</code> · <code>projections</code> · <code>immutable stamped AnalysisSnapshot</code> · <code>cancellation</code>

crate	owns
<code>bdl-ide</code>	semantic IDE queries over a snapshot: diagnostics, hover/explain, completion, references, rename, actions (incl. <code>rule.apply</code>), invalidation preview, symbols, the one token classifier (tokens), the Formula Composer queries
<code>bdl-lsp</code>	an LSP adapter only (ADR-0017)
<code>bdl-protocol</code>	protobuf schema · framing · conversions
<code>bdl-daemon</code>	<code>bdld</code> : session, coordinator, transport, analysis push, layout on open and commit; <code>bdld check compile simulate</code>
<code>bdl-runtime-core</code>	<code>no_std</code> vocabulary of every generated core: <code>ActiveDomains</code> , <code>ClockSlot</code> , <code>RuntimeError</code> , checked numerics; feature collections: list operators and the recursor over <code>alloc::Vec</code>
<code>bdl-runtime-host</code>	std harness: <code>DynValue</code> , JSON run request/trace over stdio, cargo driver; mock sinks that record what the firmware's apply would do (<code>TickTrace.adapter</code>)
<code>bdl-runtime-adapter</code>	the platform adapter's target-independent vocabulary (<code>no_std</code> , depends on <code>bdl-runtime-core</code> only): the numeric policy <code>duty8</code> , the sink traits <code>PwmDuty8 / Level</code> , <code>apply_*</code> , <code>CommandFault</code> , <code>schedule::active</code>
<code>bdl-runtime-embassy-rp</code>	the RP2040 binding over <code>embassy-rp</code> — PWM slices, GPIO lines, the arena, <code>halt</code> — kept outside the workspace so the HAL's dependency tree never enters the host lockfile; built only into generated firmware
<code>bdl-runtime-arduino</code>	the AVR binding over <code>avr-hal</code> — the ATmega timers as 8-bit PWM, digital pins, a blocking <code>tick_wait</code> , no arena — outside the workspace for the same reason; nightly-only

Planned and designed but not implemented: `bdl-component` (supplied Rust component contracts; “supplied Rust cannot drive outputs”, ADR-0005). A second target, build orchestration in `bdld`, flashing and telemetry are roadmap items behind the adapter.

Part VI — Designing Studio and the IDE

Parts IV and V established what a BDL design means and how a toolchain computes it. This Part asks the question that neither answers: how should a professional environment for industrial designers *expose* that meaning, so that a person who owns a product’s form can own its behavior in the same sitting? The answer is a design argument, not a feature list. It starts from the user and the workflow, states the principles a tool for that user must satisfy, says which professional software Studio borrows its conventions from and what it refuses to borrow, and then develops the one architectural idea that organizes everything else — the compiler’s semantics projected into a canvas and a text editor alike — before turning to the surfaces themselves: the canvas and its roles, the gestures and context menus, the Code view, the Formula Composer, and the Deploy page. Every claim about what designers *find* is a hypothesis; the last chapter says so, and Part VII lists the studies that would test them.

VI.1 The user, the workflow and the constraints

The user of Part I is a product or interaction designer with a mature medium for form and none for behavior, working across sensing, logic and actuation, typically on a small board, often alone or with one engineer, and rarely with the vocabulary of embedded programming. The workflow is the lifecycle of Part III: name the product’s concepts; draw and declare the relationships between them, leaving each undefined until there is something to say; refine relationships locally by formula; say where the environment enters and where the design leaves; give time its domains only where the product has more than one rhythm; organize a growing design into behaviors and components; simulate; then, and only then, choose a board, realize the boundary, and generate.

Three constraints follow from that user and that workflow, and they are the ones the rest of this Part answers. The design must be *legible at every stage*, including the stages at which most of it is undefined, because the designer’s work is largely done in those stages. The tool must speak the *product’s vocabulary* — brightness, tilt, a rule nothing applies yet, an output with two drivers — and hold the implementation vocabulary in reserve for the engineer who inherits the design. And the tool must meet *professional expectations*: dense, precise, keyboard-and-pointer interaction, stable object identity, undo, persistent project context, and no modal wizardry, because the people it is for already use professional tools all day.

Those constraints rule out the obvious starting points. Studio cannot be a **node editor** of the kind visual-programming environments provide, because in those editors a node is a computation and an edge an execution step, and the design of Part III is neither: its nodes are concepts and relationships, its edges are signatures and dependencies, and drawing order means nothing. It cannot be a **code editor** with a canvas bolted on, because a designer’s first artifact is a relationship declared without a body, which a text editor shows as an error and a canvas shows as an object. It cannot be an **Arduino-style IDE**, whose whole vocabulary — `loop()`, `analogRead`, a pin number — is the operational structure Part I identified as the cost. And it cannot be a **generic visual-programming environment** with BDL as one more block palette, because the facts that matter on the canvas — which concept a socket carries, whether a relationship is a Source, whether an output is contested — are semantic verdicts the palette cannot compute. Studio is a *behavioral CAD* system: the canvas holds the product’s concepts, relationships and outputs the way a CAD canvas holds parts, and every semantic fact about them is computed elsewhere and rendered here.

VI.2 Principles

The principles below are the constraints of VI.1 made specific. Each names the chapter of Part IV whose result it rests on and the chapter of this Part where it becomes visible.

Preserve product intent over implementation detail. The canvas shows that brightness follows tilt; the formula, the units, the timing domain and the drive edge live in the inspector of the selected object, and the kernel’s vocabulary — identities, grants, judgments — appears only in Explain. Nothing that is machinery is drawn as if it were design (§IV.1; VI.5, VI.6).

Support incomplete designs as first-class states. A declared rule, a Source with no value yet, an open formula, a rule nothing applies, a slot in a formula, a file that does not build: each is a drawn, explained, saveable state, and none is an error. The tool checks the design as far as it can be checked and says what remains (§IV.1; VI.6, VI.9).

Identity is stable; names are labels. Every canvas object, every socket and every reference is keyed by the model's stable identity, never by a display name or a text position, so that renaming is a refactoring, a reference edge survives an edit, and a selection survives a re-analysis (§IV.1; VI.4).

The graph and the text are views of one project. Design, Code and Split are three presentations of one semantic project with one revision stream; an object authored in either is indistinguishable in the other once synchronised, and a file that does not build leaves the graph at its last good state rather than erasing it (Part V; VI.8).

Diagnostics speak the product's language. *This adds an angle and a time; Mode values have no default order; this output already has a driver — combine the two brightness values before connecting it.* A finding is attached to the object and the field it concerns, and a fix is offered where the compiler can plan one (VI.7).

Deployment stays outside Design. Boards, pins, device kinds and realization profiles belong to the Deploy page, because behavior semantics is separate from deployment (§IV.7, §IV.8); the canvas and the inspector know nothing of them (VI.11).

Semantic truth comes from the compiler and the IDE service. Every verdict on every surface — a socket's hue, a dashed outline, a hover card, a completion list, a token's colour, an offered fix — is a projection of one analysis of one model (Part V; VI.4).

The UI never reimplements the language. Studio holds no parsed formula tree, no unit rule, no type judgment and no tokenizer; building any of them in the presentation layer was considered, and each was refused because a second definition of the language drifts the day either changes (VI.4, VI.8, VI.9).

VI.3 Professional software references: what is borrowed, what is refused

Studio does not invent its interaction grammar. Its design record (production's ADR-0012 and docs/architecture/studio-ui.md) names three references, and this monograph adds two analogies for the audience it addresses. In each case what matters is the boundary: what Studio takes, and what it deliberately leaves.

DaVinci Resolve: pages and a persistent project. Resolve organizes a whole workflow — edit, colour, sound, deliver — as *pages* in workflow order on an always-visible page bar, with one project context that persists across them and a fixed library-centre-inspector arrangement per page. Studio borrows the structure whole: one window, one project, the pages Design, Simulate, Deploy and Library in the order the lifecycle of Part III takes them (and Monitor as a placeholder for telemetry that does not exist yet), a status line above the bar, a project manager on the left and settings on the right. What it borrows is the *progressive disclosure by page*: a page owns a stage of the work and exposes only that stage's decisions, which is why realization profiles are a Deploy-page concern and never appear on Design. What it does not borrow is Resolve's content model: nothing in Studio is a timeline, and the pages do not transform an artifact through stages — the design is one object seen from four sides.

Blender's node editor: the canvas grammar. Blender's node editor supplies the anatomy of a node — a header, collapse, coloured sockets on the left for inputs and the right for outputs — and the gestures: links dragged socket to socket as curves, drop-on-empty discards, box selection, shift-click to extend, a context menu on right-click, dense keyboard-and-pointer interaction. Studio borrows the anatomy and the gestures because designers who have used any node editor already know them. It deliberately does not borrow the *meaning*: in Blender a graph is a computation and a socket's colour is a data type; in Studio a graph is a diagram of product relationships, a socket's colour is a concept identity (*Tilt* and *MotorAngle* are two colours and never connect), an edge is a signature or a dependency and never an execution step, a node is never an arithmetic operator, and dragging a node over a link does not insert it — because a link is typed by concept and silent insertion would be a semantic edit. The interaction grammar is useful; the execution semantics that usually comes with it is exactly what BDL refuses.

CAD and EDA systems: identity and the intent/manufacturing line. For an industrial designer the natural reference is not a programming tool at all but the CAD system already on the desk, and the analogy is instructive in both directions. What Studio shares with CAD and EDA practice: objects with stable identity that survive renaming and re-analysis; precise selection with a selection grammar (click, shift-click, box, select all); contextual editing in an inspector attached to the selection rather than in modal dialogs; and, above all, the separation of *design intent* from *physical realization* — a part is designed before its manufacturing process is chosen, a schematic before its board layout, and a BDL behavior before its board, its devices and its realization profiles. That separation is the

CAD analogue of the physical boundary of §IV.7, and it is why the Deploy page exists. What Studio does not take from CAD is geometry as meaning: node position is layout only, placed by a deterministic layout service and never entering the model, whereas in CAD the position *is* the design. The analogy is this monograph’s framing for its audience; production’s own design record cites Resolve, Blender and the platform’s human-interface guidelines.

Professional IDEs and the language server protocol: the semantic toolset. From IDEs Studio takes the expectations a professional has of any editor of a formal language — hover, go to definition, find references, completion, diagnostics with quick fixes, semantic highlighting, format — and takes them in the standard form: the IDE service speaks the LSP vocabulary, and an external editor uses it through a language server. The difference, and the central idea of this Part, is that in Studio these capabilities are not the text editor’s: they serve the canvas equally, because both are projections of one service (VI.4).

The platform’s human-interface guidelines. The look is the host platform’s — system font and type ramp, flat hairline controls, a sidebar–content–inspector arrangement, accent colour reserved for selection and the default action — implemented without a third-party widget kit, so that Studio reads as a native professional application and not as a web app in a window. This is a matter of expectation, not of semantics, and the record treats it as such.

VI.4 Semantic projection: one service, two surfaces

The organizing idea of Studio is that the compiler’s semantics is *projected* into every surface, and that no surface computes any of it. The canvas is not a drawing that the compiler later interprets, and the Code view is not a text that the canvas later parses; both are renderings of one analysis of one model, obtained through one service:




Figure 5: Semantic projection. One compiler decides; one service answers; two surfaces render the same answers. Nothing on either surface is computed by the surface.

The alternative — a Code view that uses the compiler and a canvas that guesses — was the shape of the first Studio, and it was replaced for reasons that are architectural, not cosmetic. **Consistency:** a name means one thing; the role a canvas node wears, the hover card in the Code view and the completion list in the formula field all read MappingAnalysis.role and never re-derive it. **Rename** is one operation on one identity, and every surface follows because none of them holds the name as meaning. **References** are the kernel’s DependsOn, drawn as reference edges on the canvas and listed under a name in the Code view, never scanned from formula text by either. **Output actions** — *Drives*, *Connect*, the eligible drivers of an output — are the service’s answer to a question about the drive edge, so the inspector’s chooser and a future gesture on the canvas cannot disagree. **Diagnostics** are placed by the compiler on entities and fields, and a finding appears as a red mark on the canvas, a row in the inspector and an underline in the text because it is one finding with three renderings. **Context menus** offer what the service says is possible for the selected object, not what a hard-coded menu guesses. The practical consequence is that adding a semantic capability once — the derived role, the unapplied-rule finding, semantic tokens — makes it appear everywhere at the same moment, and the risk the architecture exists to remove is the one every dual-surface tool eventually meets: two definitions of the language, one of them in the presentation layer, diverging quietly.

Every semantic fact is designed for exactly one primary level and may echo at the next only as detail behind the first (ADR-0018):

level	answers	may use	may not use
1 Canvas	what does the product do; what depends on what; what is still open; where does behavior become physical	object silhouette, socket shape, socket hue, links, grouping, containment, line style, the object’s own state; one state word where unavoidable	type labels, ids, compiler words, badges, counts

level	answers	may use	may not use
2 Inspector	what does the selected object mean; what can I change; what will the change affect	designer vocabulary: <i>Meaning, Value, Unit, Reads, Produces, Relationship, Used by, affects, checked again</i> ; diagnostics in product language attached to the field they concern	SemanticId (production's name; ConceptId in the formal development since Phase 22), DeclId, Ty, Grant, realization, Clocked, SingleDriver, solver, protocol, revision, enum names
3 Explain	why was this accepted or refused; what did the surface form elaborate into; which rule applies	all of the above, kernel notation, Core IR, diagnostic codes, revision	—

Explain is one collapsed disclosure at the end of the inspector, never open by default; nothing in levels 1–2 depends on it. This is the progressive-disclosure claim of Part I made operational: the kernel is reachable, and it is never in the way.

VI.5 The canvas as a diagram of product relationships

Socket hue = concept identity. In a node editor a socket's colour is its data type; in BDL the type that matters is the nominal concept, so each concept gets a stable hue derived from its `SemanticId` — deterministic, never from the name, with lightness chosen per hue so every identity clears 3:1 contrast in both appearances. A link is accepted only between sockets of the same concept: while a link is dragged every compatible socket gains a faint halo, the one under the pointer a strong halo, and an incompatible socket shows the forbidden cursor. The canvas shows the nominal typing rule without a diagnostic. Tilt and MotorAngle, both angles, have two hues and never connect.

Socket shape = value form. Quantity $\circ$, on-off $\diamond$, count $\square$, collection $\boxed{}$ (a stack), grouped value $\square$ (a split square), optional value $\circ$ (a ring with a hole); a concept whose value form is not yet chosen is a hollow ring. The same glyph, drawn by the same code, appears in the library, in chips, toggles and pop-ups. No type words are written beside a socket anywhere.

Unresolved state. A declared rule — realization none with inputs — is drawn with a dashed outline, an empty definition region and the one word *declared*; dashed survives selection; never red. A Source — realization none on the unit domain — is *not* dashed: it is complete, and it is drawn as the next section says. A definition the compiler cannot accept gets a red mark at the definition line, where the problem lives, and no word in the header. A relationship waiting on a concept whose value form is not chosen is a solid node whose read socket is hollow, with the inspector saying *checked once Temperature's value is decided*. These are the states of Part III's table, drawn.

What the canvas never means. Edges are dependency, not execution order; drawing order does not set output priority; node position is layout only (ADR-0003) — an entity without a position is placed by the daemon's layout service on open and on commit, deterministically, without moving anything placed, and Studio arranges nothing at render time. No node type exists per arithmetic operator: formulas live in the inspector. Dragging a node over a link does not auto-insert it, because links are typed by concept and silent insertion would be a semantic edit.

Header colour = category: concept, mapping, context, output, transport — muted; identity hues are the only saturated marks on the canvas. Interaction follows Blender's node editor conventions; a concept in use cannot be deleted without the banner naming its dependents.

The three roles of §IV.7 are the canvas's first distinction, drawn without colour alone. A **Source** has no input socket and one output socket, never a $()$ port; its header carries the word *Source*, an entry glyph — an arrow crossing a boundary tick — and a solid bar on the node's **left** edge, the environment side, the mirror of a sink's bar on the right, with a green header strip as the redundant cue. It is not dashed: nothing is missing. Assistive technology hears *name, Source: a value entering the behavior model from the environment, provides C*. A **Rule** wears the word *rule*; when nothing applies it, its output socket is hollow and the header says *not applied*. A **Value** is the plain node. The inspector says what an object *is*, not what it lacks: a Source's *Meaning* section has *Role: Source* with one sentence, its output section is *Provides* (a relationship's is *Produces*), and its *Relationship* section opens with *Realization: provided by the environment; no device is bound yet* — never *calls, reads hardware* or *sensor*. Hover and Explain speak the formal

vocabulary: role: Source, type: () -> C, definition *none* — *provided by the environment, observed once per activation*; a resolved () -> A cites resolved_not_source.

Edges. A **signature edge** runs socket to socket and is the interface; a **reference edge** is a thin neutral line into a relationship's formula line for every relationship the formula names, drawn from `MappingAnalysis.references` — the kernel's `DependsOn` by identity — never from a reading of the formula text, not editable, not a drop target (ADR-0034 (FVD-0005)). *Produces* is the signature; *carried by*, in the Simulate probe, is a value per tick. Studio consumes references and `applied_by` and never inverts or scans anything.

A behavior system has three zoom levels on one canvas: the system of instances, an instance's body, and the flat design. A group is a box that can be collapsed to its aggregate sockets (`crossIn ++ openMembers in, crossOut out` — §IV.6), moved, merged and split with no revision and no re-check; *Package as Component* opens a sheet that shows the inferred required, provided and private members and lets the designer widen but not narrow them. An instance node shows its required and provided ports with the same socket glyphs; a component's body is edited in its own scope with its own drafts and completion. The packaging sheet teaches by showing: the boundary the designer sees on the collapsed group is exactly the interface the component receives.

VI.6 Roles, gestures and context menus

Each visual interaction on the canvas corresponds to one semantic design object, and the correspondence is the reason the gesture exists. A **Concept node** is a `SemanticId` with its value form, and its socket is the one place its values enter and leave a relationship. (Production at the snapshot draws the concept as a node; the formal development has since decided the projection that replaces it — Sem blocks and mapping blocks, no concept node, FVD-0163 and §IV.2 — and production's canvas follows in its next slice.) A **relationship node** is a declaration; its input sockets are its signature's inputs and its output socket the concept it produces; dragging a link from a concept's socket into an input socket edits the signature — *adds the concept to what the relationship reads* — and dragging into an empty output socket sets what it produces. The **three roles** are drawn without colour alone: a Source with no input socket, a boundary bar on the environment side and the word *Source*; a Rule with the word *rule* and, when nothing applies it, a hollow output socket and the words *not applied*; a Value as the plain node. A **logical Output** is a terminal node at the right with a boundary bar of its own and one input socket for its one driver. **Groups** are boxes that collapse to their boundary sockets, and **instances** show their promised ports with the same socket glyphs (§IV.6).

Selection follows the grammar a professional expects — click, shift-click to extend, box selection, select all — and is Studio's own state: nothing semantic depends on what is selected, and a re-analysis never loses it, because selection is keyed by identity. **Context menus** are built from two things and nothing else: the kind of the object under the pointer, and the semantic actions the service reports for it. A relationship's menu offers rename and delete and the fixes the compiler planned; a concept in use cannot be deleted without the banner naming its dependents; the empty canvas offers *Add Concept* ► from the library's categories and *Add Source* ► to the Source sheet. A menu entry never appears because a UI author guessed it might apply.

High-level gestures that resolve to precise semantics. The tool may offer a gesture that is higher-level than the kernel provided it never lies about the semantics. Connecting a value to an output is the clearest case: an output's inspector offers *Connect*, a pop-up of the relationships eligible to drive it — those producing its concept in its domain, Sources and Values, never Rules — so that the designer chooses among semantically eligible candidates and never among all nodes. The gesture reads as *this output shows the brightness*; what it commits is one drive edge, checked by type and domain equality and refused with the word *contested* if the output already has a driver. The same shape governs *Add a value that applies X* on a rule nothing applies: ready when each read concept has exactly one producing value, a chooser when one has several, blocked with the reason when one has none. In every case the designer's gesture is resolved by the service into the unique semantic object it denotes, or into a choice the designer makes, and never into a guess.

A finding may carry a **fix**, and a fix is a semantic action the IDE service plans and the daemon applies as an ordinary edit: *Make empty domain explicit* on the legacy shorthand, *Insert explicit sync* at a cross-domain reference, and *Add a value that applies X* (`rule.apply`) on a rule nothing applies — ready when each read concept has exactly one producing value, a choice when one has several, blocked with the reason when one has none or the arguments update in two timing domains; blocked, with the instance named, for a rule inside an instance, because the value belongs in the

component's source. Actions and entity hover are requested in the system context only, on the flat entities. A fix chosen on the Simulate page selects the object, asks for its actions and applies the one of that kind when it arrives ready. The tool never guesses.

VI.7 Diagnostics, fixes and the three information levels

The three information levels of VI.4 are the rule for where a fact is said; findings and fixes are the rule for how a problem is said. A finding is placed by the compiler on the entity and the field it concerns, worded in the product's language, and rendered once per surface; where the IDE service can plan a repair, the finding carries a fix, and a fix is a semantic action the daemon applies as an ordinary edit against the current revision — *Make empty domain explicit* on the legacy shorthand, *Insert explicit sync* at a cross-domain reference, *Add a value that applies X*. Two findings deserve mention because they show the vocabulary at work: an output with two drivers is *contested*, on the output, with the two claimants named; and a rule nothing applies is an *information* finding, never an error, because a design with an unapplied rule is legal and merely says less than it could.

VI.8 The Code view: the same project as text

The same open project is shown as a graphical behavior model, as its source files, or both side by side (ADR-0023). Switching is a view change — no conversion, import, export or project type. Graph → text: a canvas operation is a semantic edit on the model, and the textual projection is the minimal item-level splice of the sources, keeping comments and formatting outside the touched item; computed for the Code view on request and written on save. Text → graph: a Code-view edit is the whole text of one file against a revision; if it builds, declarations are bound to their identities by reconciliation (retained, allocated, dropped, renamed, ambiguous, duplicate) and the project moves to a new revision; if not, the committed project stays and the draft is held with the loader's faults — malformed text never erases the graph. An entity authored first in Code and one authored first in Design are indistinguishable once synchronised. The Code pane shows an out-of-step banner with the daemon's reasons and syncs selection in Split.

The Code view is an IDE surface, not a text box (protocol 0.22). Completion at the caret offers what the compiler service knows is allowed at that spot — a concept after `:`, a clock after `@`, the items at an item start, and inside a formula the inputs, the other relationships (a rule as a call, a Source or a Value by name), the locals, the equation library, units after a number; a hover card names the declaration under the pointer with what it produces, its state and its role; $\mathbb{X}$ -click or F12 goes to a declaration in this or another file; $\uparrow$ F12 lists every place that names it; *Format* lays a file out canonically as one edit and leaves a file that does not parse as it is. Every one of these is the same query the language server answers — `bdl_idc::navigation`: one answer to *what is at a position*, resolved through the projection map and the elaborator's input environment, never by spelling, with the authored entity winning over the flattened copies it backs — over the text as typed. The significance is architectural rather than a feature count: textual and graphical authoring share one semantic project *and* one compiler service, so a designer who moves between the canvas and the text meets one set of names, one set of verdicts and one set of fixes, and the text editor holds no second definition of the language.

Ownership boundaries, stated once and relied on everywhere: source semantics in `src/**/*.bdl`; identity in `.bdl/identities.json`; authoring metadata (groups) in `.bdl/authoring.json`; layout in `ui/layout.json`; viewport and session state in the running Studio only. Studio state is three things kept apart — semantic projection (owned by the compiler, only rendered), editor state (selection, open inspector, pending requests), rendering state (drag position, hover, zoom, never persisted or sent).

The Code view and the definition editor's text field colour BDL text by the IDE service's semantic tokens (Part V): one request per text shown and per pause in typing, the spans of the unchanged prefix and suffix shifted while a keystroke's answer is on its way, and never a stale span over changed text; only the latest generation's answer lands. Studio's `SyntaxTheme` maps the token *names* to styles in the canvas's own category encoding; Studio classifies no BDL syntax itself (ADR-0035). This is the broader architectural rule made visible at the character level: semantic authority stays in Rust, and clients render projections.

VI.9 The definition editor and the Formula Composer

The Formula Composer is not a second language and not a graphical replacement for formulas. It is structured authoring of the same equation language of §IV.3, drawn from the compiler's own projection of a formula, with

slot expectations computed by the local dimension inference of §IV.4 and the natural forms of §IV.3 drawn as what they are — a binder over an indented body, a range between two ends — and elaborated as what they desugar to. The design decision beneath it is that an incomplete formula is ordinary draft text with a slot ? in it, so that every structured action is a byte-range edit of that text and the Text view and the Formula view can never disagree.

The definition editor has two projections of one draft, chosen with a **Formula | Text** control (ADR-0028). The Text view is a field over the compiler’s verdict, with completion, hover and fixes from the IDE service. The Formula view draws the compiler’s `FormulaProjection` as the expression it is — `clamp( [Tilt] ÷ [90][deg ▼], [0], [1] )` — never as a graph inside the node.

Architecture. The draft text of a mapping’s definition is an overlay in the IDE host (`MappingDefinitionDraft`), never a revision; `bdl-ide::formula` builds the projection from the *effective* definition text on every request and never stores it:

```
FormulaProjection  formula_projection(&snapshot, mapping)
                  = trace_formula_in (bdl-elab: the surface tree and the type of every
                    sub-expression, on success and on failure)
                    → FormulaNode tree: tree-path ids, byte ranges, kinds, actual types,
                      expected types by local dimension inference (solve),
                      diagnostics placed on the innermost node, the slots in source order

SlotInfo           formula_slot(&snapshot, mapping, node)
                  = the projection's expectation at `node` + units_for(dim)
                    + reference candidates by type + equations by scheme match and capability

ComposeResult      compose(&snapshot, mapping, source, ComposeOp)
                  = the text a structured action makes: byte-range edits of the draft,
                    parenthesised by precedence, a unit switched with the value kept
```

`FormulaTrace` is the elaborator’s typing trace, kept on success and on failure, so a formula that does not check still has a projection for the parts that do. A `FormulaNode` carries its *actual* type (what the sub-expression is) and its *expected* type (what its position demands, by `solve`), and node identity is the tree path — ephemeral, stable within one draft generation, not an entity, never in `.bdl/identities.json`. `bdld` carries the projection with every `DefinitionDraftAnalysis` (one round trip per keystroke, the same generation) and answers `GetFormulaProjection` (the committed definition, no overlay), `GetFormulaSlot` and `ComposeFormula` (protocol 0.12).

Every structured action is a text edit. `compose` answers a `ComposeOp` — fill a slot, wrap in an operator, call an equation with the component as first argument, set a literal’s unit, set a literal’s coordinate, remove — with byte-range edits of the draft and the whole new source; Studio puts it into the draft through the ordinary `DefinitionDraftChanged`, and the ordinary analysis follows. One commit path, one conflict model, one set of diagnostics. Studio stores the mode, the selection and the open pop-up — never a parsed tree, an inferred dimension or a unit rule; building the tree in Dart was rejected because precedence, parenthesization and unit conversion would move into Studio (ADR-0001).

Encodings. A *slot* is a dashed hollow chip (dashed = not decided, as the canvas’s *declared* node); a *reference* is a chip with its concept’s socket glyph; a *literal* is two fields, the coordinate and a unit pop-up listing the compiler’s units for the literal’s own dimension, where the pop-up switches the unit and keeps the quantity and the coordinate field makes a new quantity in the same unit — only a literal has a unit pop-up (§IV.4); *operators* are their glyphs, a *call* its name and parentheses, and */ or / not* their words, a *choice if* with its condition over indented then */ else* branches (with true / false buttons and truth-valued references offered on a slot that expects true or false), a *binder* a head over an indented body with italic local chips, a *range* two ends around `..`, and an *opaque form* (`let`, `match`, a block, a rule as an argument, a collection or grouped literal, `delay/sync`) its text in monospace, selectable and edited as text. A finding is a red underline on the component *and* its row under the field — one diagnostic, two projections. Beneath the field, for the selected component: the compiler’s sentence — *Expected: an angle, because an angle ÷ an angle = a dimensionless quantity* — with the kernel’s notation behind `Explain`; for a slot, a number entry whose unit pop-up holds the units of the expected dimension (none for a dimensionless slot; none, with a sentence, when the position is not determined), then references by type and equations folded; for a component, `+` `-` `×` `÷`, `Compare`, `Function` (the equations whose result fits, wrapping the component as first argument) and `Remove`.

Stale projection policy. A projection is current only when it is of exactly the text on screen and that text parsed; otherwise the field is dimmed with a notice — *Waiting for the compiler to read the formula...* while the verdict for this text is on its way, *The text cannot be read as a formula.* when it never will be, and then no tree is shown because none is invented — no component answers a click, no slot panel opens, no key acts, the reducer refuses a structured

action (or a second one in flight) and discards an answer for text that has moved on; **Edit as text** is the way out. Generation safety: every answer carries the draft generation it was computed for and is discarded if the draft has moved. A formula with a slot may be saved and is *invalid* until filled; commit, revert, conflict and detach are the Text view's, unchanged.

A walkthrough. The designer selects `dimByTilt : Tilt → Brightness` and chooses `Function → clamp`; the draft becomes `clamp(?, ?, ?)`, three dashed slots. Selecting the first, the slot panel says *Expected: a dimensionless quantity* — because `clamp`'s scheme unifies all three arguments with the result and the result is Brightness's representation — and lists references of that kind; the designer chooses `tilt`, and the compiler answers *This is an angle; a dimensionless quantity is expected*, red on the chip. Wrapping it in `÷` yields `clamp(tilt ÷ ?, ?, ?)`; the new slot's panel now says *Expected: an angle, because an angle ÷ an angle = a dimensionless quantity* and offers `rad`, `deg`, `turn`; the designer types 90 and picks `deg`. Filling 0 and 1 completes `clamp(tilt / 90 deg, 0, 1)`, which checks, and the node's red mark goes. Switching the `deg` pop-up to `rad` rewrites the literal to `1.5707963267948966 rad` and nothing else. Every step was a byte-range edit of one text; the Text view shows the same characters throughout.

The slot deserves its own record because it is the clearest example in BDL of authoring semantics distinct from runtime semantics. `?` is a `SlotExpr` of the textual grammar: it parses wherever a value may stand and formats as itself; it is never a value, an operator or a unit (`1 ? 2` is a syntax error). It exists in the *surface* because an incomplete structured formula must be ordinary draft text — the same state as any unfinished draft, savable, visible in the Text view, subject to the same conflict model — rather than a second representation held in an overlay of partial expressions (the rejected alternative, which would have needed its own persistence and its own protocol of operations). Elaboration refuses it with `formula.slot.empty` — *This slot is empty; it expects an angle* when the context says what it expects — and it never reaches Core, the checker, the evaluator or generated code; a kernel hole term was rejected because the kernel has no notion of an unfinished expression and needs none. Expected information flows *into* the slot from the formal Composer's solve: the operator above it and the known sibling determine its dimension; a product of two slots is *insufficient information* — the formal model reports it unsolved rather than searching, and Studio shows the position as not determined with a sentence. That two-hole case is not a limitation to be engineered away: local, deterministic inference is what makes the expectation explainable in one sentence, and a global solver would trade that for guesses.

VI.10 The Library and the Source sheet

The Library tab lists the Standard Library's items in two sections, *Concepts* and *Sources*, with names and search localized by id through Studio's ordinary localization pipeline. Inserting a Concept item is one request (`InstantiateLibraryItem`) the daemon plans as a fragment of ordinary objects and applies in one transaction — every step or none, one revision, one Undo — and the canvas draws what was created exactly as it draws what was made by hand.

The Source sheet (protocol 0.23). A Source is never created without a concrete concept, so every way of adding one — *Add Source* ► on the canvas, a Source row's double-click or drag, *New Source...* — opens one sheet whose single decision is the concept the Source provides: an existing concept of the design, by identity, or a new concept created in the same transaction. The Library's Source items are *presets* for that sheet — *Temperature Input*, *Tilt Input*, *Button Input*, *External Input* and the rest — each suggesting a concept name, a value form, a unit and a Source name; the daemon lists the design's concepts with the suggested value form first and infers nothing from names, units or dimensions. What is committed is mapping `tiltInput : () -> Tilt`, a Source by the derived rule with no preset, no flag and nothing about how it was made. The design consequence is that a Source is authored as what it is at the boundary — an input for a concept the designer chose — and never as a chosen device; the device is deployment's, and at the snapshot no device provides a Source's value (ISS-0016). The Library is an authoring catalogue; it says nothing about devices.

VI.11 Simulate, Deploy and the boundary of Design

The Deploy page exists because of §IV.7 and §IV.8. Behavior semantics is separate from physical realization, and the tool's pages follow the separation: nothing a designer decides on Design depends on a board, and nothing decided on Deploy changes a behavior. Realization is the clearest case. A device card lists every profile whose encoder fits the output's concept, shows the three admissibility judgments — well typed, fits, placed — with the analysis's sentence when one fails, and choosing a profile changes the raw command the machine will receive and nothing on the

canvas, in the inspector or in the simulation (§IV.7; **production implemented and tested**). Source binding would live on the same page for the same reason, and the page is honest about its absence: no device provides a Source's value at the snapshot, and a design with a Source cannot yet run on a board (ISS-0016). Hardware feasibility is the third example: a board's capacity is a verdict about the pair (design, board), re-solved on every change, reported here and never merged with the design's own validity (§IV.8).

Simulate runs the reference evaluator over structured inputs per tick (including list and pair values, and concept values in the design's terms — `Brightness(0.5)`), showing every declaration's value and every output; the simulator does not get a separate interpretation. *Deploy* is the target-relative read model (ADR-0015): board selection reruns deployment analysis only, and the workspace reports validity and deployability in different places; the readiness matrix (bounded-memory refusal, unbounded state, window capacity) is shown here, and so is **realization**: each device card lists every profile whose encoder fits the output's concept, shows the three judgments — well typed, fits, placed — and the analysis's sentence when one fails, and a chosen profile changes nothing on the Design page, in the inspector or in the simulation (ADR-0036). The Deploy page is the only surface that knows a profile exists; the board list includes the Raspberry Pi Pico, the one target a generated adapter exists for. *Simulate* lists the Sources as its inputs and refuses to step until each has a value; a probe on a concept names what carries it and, on a rule, where it is applied. *Library* is the previous section's. *Monitor* — telemetry from a deployed core — is a placeholder page; the adapter it would listen to exists, the telemetry path does not (roadmap priority 4).

VI.12 What Studio does not decide, and what is not claimed

Studio computes no type validity, concept identity, dimension, causality, clock, output ownership, hardware feasibility or simulation; it holds no parsed formula tree and no unit rule; it arranges nothing at render time. Every one of those is a projection from the daemon, keyed by `EntityRef` and role, never by text position or node id. That is the discipline that lets the same facts appear as socket hues, inspector sentences and Explain notation without three implementations of the language.

Everything in this Part is design rationale, precedent and intended benefit. None of it is evidence that designers find Studio easier, understand a design faster, or make fewer mistakes with it; those are the empirical questions of Part VII, and no study has been run. The workspace states, the vocabulary, the levels and the gestures are hypotheses about a professional practice, made as precise as the semantics beneath them allows, and stated so that a study could refute them.

Part VII — Evidence, Rejected Designs, and the Open Agenda

The reader now knows what BDL is, why it is shaped as it is, how it is built and how it is exposed. This Part is the accounting. It states, construct by construct, how the formal model and the production implementation correspond and where they deliberately differ; it retells the designs that were tried and failed as the arguments they are; it bounds the minimality claims; and it separates what is still open into three classes that must not be confused — formal questions a theorem would settle, engineering work production has decided and not yet done, and empirical questions only a study would answer. The evidence vocabulary is the preface’s; Appendix D is the ledger row by row.

VII.1 Formal ↔ production correspondence

The formal development is a specification (ADR-0010): production transcribes it and discharges it by tests; a theorem never proves the Rust or Dart code. This chapter states, construct by construct, the formal object and its result, the production object and its status at the snapshot, how strong the correspondence is, every known deviation with its reason, and the proof gap that remains. The classification of each row is one of: **exact** (the production object is the transcription and the tests discharge the stated property); **faithful implementation** (production implements the semantics and is tested against it, without a theorem about the production code); **approximation** (production’s domain differs, with a stated contract); **deliberate product projection** (a production reading of a kernel fact, with no separate formal object); **production-only** (an engineering concern the model does not speak to); **formally proved, not implemented**; **intentionally deferred**; **unresolved divergence** (none at the snapshot). The snapshot is 081296d (Appendix F); the correspondence page of the formal repository (`docs/project/production-correspondence.md`) and production’s `formal-correspondence.md` carry the same rows with the FVD and ADR identifiers.

Construct by construct

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
declarations, interfaces, refinement	DeclEnv, DeclInterface, DeclLeq, EnvRefines; local_refinement_preserves_global_typing, _wf (FVD-0005 ... FVD-0016)	bdl-model snapshot, MappingBlock; apply_edit with refinement/edit classification and Invalidation sets (ADR-0009); implemented, tested	faithful implementation	production’s seven invalidation categories are an engineering refinement of the two-way split	a per-category preservation theorem
commitments and evidence	commitments : List PropertyId, Evidence, Evidence. Monotone, Satisfies (FVD-0001, FVD-0009, FVD-0010, FVD-0017)	Interface. commitments, always empty; require reserved	intentionally deferred	production authors no commitments; the evidence relation is abstract	a concrete compositional evidence model; interface-level references (ISS-0003)
typing through the type view	HasType, infer_sound/_complete/_unique (FVD-0006)	bdl-check over bdl-ir (Ty, Expr, Prim transcribed); implemented, tested	exact	—	—
nominal concepts, representation, grant	sem s, 0, rep/mk, Grant.of, constructs_granted (FVD-0019 ... FVD-0030)	elaborator inserts rep on inputs and one mk under the declaration’s grant; checker refuses others (ADR-0013); semantic corpus	exact	—	—

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
dimensions	q d, algebra in Prim.ty (FVD-0031)	bdl-model::Dim (eight bases), primitive typing; product-language diagnostics	exact	—	—
numerics	Nat magnitudes in the executable kernel (FVD-0102)	IEEE f64; division by an exact zero and non-finite results fail the tick (ADR-0011)	approximation	proofs need decidable saturating naturals; products measure reals; NaN/∞ have no kernel counterpart	not planned to close; f32 on device a further deviation (ISS-0006)
units and charts	Units.lean, Charts.lean, Rational.lean: exact symbolic scales, exact Q charts, the groupoid laws (FVD-0101 ... FVD-0110)	UnitDef { id, symbol, dim, chart }, Chart:: {Linear, Affine}, one convert; laws property-tested on f64 within ulps; affine charts as infrastructure, not offered (ISS-0004)	approximation	$\pi/180$ is not rational; exactness is what the formal laws are about	the point/difference validation, or a decision not to need it
Composer inference	Composer.lean solve, candidates; solve_sound, _complete (FVD-0104)	bdl-ide::formula projection, slot, compose (ADR-0028); formula_composer.rs	faithful implementation	two-hole operands unsolved by design in both	—
equation library, schemes	Stdlib, Poly, Generic; matchTy_sound/_complete, lib_expansion, *_spec (FVD-0088 ... FVD-0097)	bdl-equations: rank-1 schemes with {Data, Eq, Ord}, matching, one Core builder per instance inlined at the use (ADR-0025)	exact	—	—
equality and order	eq on data with the proof field; lt on quantities; Ty.ordB, Ordered (FVD-0098 ... FVD-0100)	Concept::ordered, ordered concept; ADR-0026's mixed-comparison rule	exact, plus a deliberate product projection	the concept-beside-plain-value case is a production decision the model does not make	recorded as intended
natural binder syntax	Natural.lean: desugar, desugar_rename, binder_local_type, binder*_eval (FVD-0111 ... FVD-0114)	bdl-elab::formula:: binder; natural_forms_lower_to_the_same_core_as_the_call_forms	exact	no parser is modelled; the elaborator's local inference is the output of binder_local_type, described not proved	—
reactive semantics	Ev, Ev.det, reactive_total, Causal (FVD-0034 ... FVD-0042)	bdl-reactive reference evaluator, two-phase tick, StateCellId (ADR-0004)	exact (the evaluator is the executable definition)	Causal is conservative for lambda-guarded cycles in both	causality precision
clock domains, sync	MEv, Clocked, strictly-before, delay_is_sync_own, single_domain_embedding (FVD-0043 ... FVD-0048)	the Clocked pass, the sync snapshot rule, step_in_order	exact	—	—

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
logical outputs, drive edge	OutputId, DriveWF, SingleDriver, CompleteOutputs; single_driver_output_deterministic (FVD-0050 ... FVD-0056, FVD-0131)	bdL-output, PhysicalOutput { accepts, clock, required } (the historical name of the logical output); drive light by brightness; a Source or a Value may drive, a Rule may not	exact	the production type name says <i>physical</i> ; the object is the logical output and no device fact is on it	commitments on outputs
relationship roles	only Source Δ d (FVD-0118) has a formal object	RelationshipRole { Source, Rule, Value }, one derived predicate stated on every projection (ADR-0032 amended, protocol 0.20)	deliberate product projection	Rule and Value are production's names for the arrow type and the realized unit domain	—
applied_by, unapplied rules	DependsOn (FVD-0005, FVD-0012)	MappingAnalysis. references and its direct inverse applied_by; reactive.rule_unapplied; rule.apply	production-only over an exact base	—	—
unit-domain canonical type	UnitDomain.lean: elim_canonical, decode_encode, zero_input_obligation, refForms_agree, source_value (FVD-0115 ... FVD-0120)	bdL_ir::Ty::Unit, Ty::of_signature, kernel_of_signature, Signature::is_unit_domain (ADR-0029); ..._the_encodings_ are_inverse, ..._argument_is_erased	exact at the interface; the kernel encoding transcribed	Ty::Unit is an IR type that never types a Core term, representation or runtime value — exactly elim unit = none	—
simulation inputs	the kernel input provides for every unresolved declaration	inputs are the Sources (Source $\wedge$ UnitDomain, SimulationInput.value)	deliberate product projection	an environment provides values, not functions	—
hardware validation	solve_sound, solve_complete, satisfiable_iff_solve, diagnose (FVD-0057 ... FVD-0063); admissibility's third judgment (FVD-0139)	bdL-hardware, board files (Nano, a larger mock board, the Raspberry Pi Pico), target-relative deployment analysis (ADR-0006, ADR-0015)	exact for the finite fragment	numeric constraints out of scope in both	MUS; numeric constraints; runtime loading of board files

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
list data, buffer, capacity	Ty.list, Theorem M buffer_window_correspondence, Capacity.lean (FVD-0083 ... FVD-0087)	Vec last-element-first behind the collections feature, take cap bounds, bdl-reactive::capacity, bdl-exec-ir::bounds, bounded-memory refusal (ADR-0024, ADR-0027); corpus buffer, bounded_buffer, overflowing_buffer	faithful implementation	the static bound analysis has no theorem behind it	a proved bound analysis, or a surface window form (ISS-0001)
products, fold, toList, drop	Phase 9b kernel (FVD-0088 ... FVD-0091)	bdl-ir gains exactly these; Fold in exec IR with one runtime closure	exact	—	—
behavior systems	Theorems A–I; Theorem J restricted (FVD-0064 ... FVD-0073)	bdl-system: freshening from the project allocator, bindings, flatten with origins, stored contracts (ADR-0021, ADR-0022)	faithful implementation	production supports transported bindings under MEV and higher-order bodies beyond the theorem fragment	Theorem J under MEV with a domain-indexed modular input
behavior groups	Theorems A–G (rfl), H–L, M–R (FVD-0074 ... FVD-0082)	.bdl/authoring.json, no invalidation on group edits, the packaging sheet (ADR-0019)	faithful implementation	Theorem R is proved for single-domain wiring designs without transported bindings; production packages beyond it	a port-level causality graph
Source provision	Provision.lean: provision_envRefines, _wf, _causal, _wellClocked, _transparent, _abstracts, _exact, _not_reapplicable, _comm (FVD-0121 ... FVD-0130)	not implemented (PRP-0001 draft; ISS-0016); ownership decided in relationship-roles.md § Phase 13; the adapter refuses a design with a Source; the Source sheet (0.23) is authoring over a chosen concept, not a device binding	formally proved, not implemented	—	stateful transducers, device clock, commitment discharge (FVI-0020); the input device catalogue

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
output realization: encoder, admissibility, machine sink	OutputRealization.lean: Encoder.WF, EFits, Admissible; behavior_unchanged, lower_transparent, lower_correspondence, two_realizations_same_behavior, admissible_needs_wf, retarget_breaks_driveWF (FVD-0131 ... FVD-0139, FVD-0137 superseded)	ADR-0036: DeviceBinding.realization, OutputProfile { id, encoder, kind }, Encoder { rep, raw, encode } with well_formed and purity, DeviceRealization { check, hardware_placed } composed in analyze_deployment, one SinkPlan per valid chosen profile, five witness profiles; output_realization.rs, the Deploy chooser	faithful implementation (the encoder and the judgments are transcriptions; the lowering is at the plan level)	production lowers to a plan-level sink without the fresh DeclIds of lowerΔ; observably the same, not proved as such; the profiles are witnesses, not a catalogue	a proof that the plan-level lowering is the model's; the device catalogue (ISS-0017)
generated raw commands	RawCommand — a relation on the unchanged design, the machine boundary (FVD-0134)	Tick.commands beside values and outputs; TickTrace.commands; the corpus compares them to the encoder over the reference evaluator's outputs (Exec IR v3)	faithful implementation, tested only	the generated command is compared to the encoder applied to the evaluator's output, not to RawCommand	codegen correspondence (FVI-0022)
the platform adapter	— (not modelled: the semantics stops at the raw command)	ADR-0037 (amended): apply(tick, sinks...) on the RP2040 over Embassy and on the Arduino Nano over avr-hal, the solver's pad per sink, duty8 reject-and-hold, the compiled schedule or a blocking tick loop, halt on fault, the arena from the manifest (RP2040 only); recording mock sinks on the host; both firmwares cross-built in CI	production-only, production implemented and tested ; never formally proved	a third numeric domain (f64 command → integer register) with a stated policy; no bench measurement of a physical effect	raw command → physical effect (FVI-0022; ISS-0017); a device for a Source (ISS-0016)
A -> () as a consumer	unit_codomain_collapse, consumers_indistinguishable (FVD-0119)	() refused in output position	exact	—	—

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
generated Rust	— (no theorem)	bdl-lower + bdl-codegen-rust, the no_std core, the host bridge; 22-case differential corpus, golden files, property tests (ADR-0016)	faithful implementation, tested only	the core lowers representation; a lowering bug would be a semantic bug	a generated-code refinement proof
project persistence, identities	—	src/**/* .bdl canonical, .bdl/identities.json, .bdl/authoring.json, ui/layout.json; migration in place; complete authoring state saved (ADR-0023, ADR-0030)	production-only	the model has no files	—
the protocol, the IDE service, highlighting, the Code view	—	protobuf over framed stdio; bdl-ide-db/bdl-ide/bdl-lsp; one token classifier on the LSP vocabulary; one answer to <i>what is at a position</i> serving the language server, the formula field and the Code view (ADR-0007, ADR-0017, ADR-0035; protocol 0.22)	production-only	—	—
Studio's semantic projection	—	canvas, inspector, Explain, Composer, Simulate, Deploy, Library, Code/Split; Studio decides no semantic fact (ADR-0001, ADR-0018)	production-only	—	usability: open empirical questions
Standard Library, the Source sheet	Source Δ d (FVD-0118)	LibraryItem → Fragment → ordinary objects, one-transaction instantiation; a Source created over a chosen concept, presets suggesting names and value forms (protocol 0.23)	production-only over an exact base	an authoring catalogue; not the device catalogue provision needs	the device catalogue
temporal modifiers, contexts	the elaboration cases of Phase 4, executed (FVI-0009)	not offered (ISS-0010)	design recommendation, not implemented	—	handler-scoped clocks; a preservation theorem for the elaboration

construct	formal object · result	production object · status	correspondence	deviation and its reason	remaining gap
supplied computation blocks	—	designed (ADR-0005, component-boundary.md); nothing built	proposal / not implemented	—	—

Deviations, stated once with their reasons

deviation	why	risk	evidence that bounds it	closing the gap
kernel <code>Nat</code> magnitudes vs production <code>f64</code>	proofs need decidable, saturating naturals; products measure reals	truncation in the kernel's executed examples is not production's rounding; NaN/∞ and division by zero have no kernel counterpart	production fails the tick on exact-zero division and non-finite results; exact comparison; the formal theorems are about structure, not rounding	not planned to close; <code>f32</code> on device is a further recorded deviation (ISS-0006)
exact symbolic unit scales (<code>Sym</code> , π as a generator; exact <code>Q</code> charts) vs approximate floats	$\pi/180$ is not rational; exactness is what the formal laws are about	round trips and compositions hold only within ulps in production	toleranced property tests of identity, composition, inverse, display switch and difference law; exact-rational oracles for registered charts	not planned; the contract is stated and tested
kernel <code>Nat</code> registry has no radian; production is radian-canonical	no integer scale for radians against degrees	none semantic; the registries differ in which units are exact	documented per registry	none needed
theorem-level semantics for direct bindings on wiring designs vs production's transported bindings and higher-order bodies in systems	Theorem J's obstacle: a domain-indexed input for a transported port	production behavior beyond the fragment is tested, not proved	flattening plus the reference evaluator; corpus systems	Theorem J under <code>MEv</code> with a domain-indexed modular input
formal buffer model (unbounded, capacity as a validation predicate) vs production static list bounds and bounded-memory refusal	the kernel has no capacity; production must decide deployability	the static bound analysis has no theorem behind it	differential tests of <code>bounded_buffer</code> vs <code>buffer</code> at 7 and 3 000 ticks; <code>overflowing_buffer</code>	a proved bound analysis, or a surface window form (ISS-0001)
formal <code>Ev/MEv</code> vs generated Rust	the core lowers representation (slots, structs, <code>f64</code> , inlined lambdas)	a lowering bug would be a semantic bug	22-case differential corpus, golden files, determinism, property-based designs, error-tick agreement	a generated-code refinement proof (§VII.4)
abstract evidence conditions (<code>Monotone</code> , <code>Equivariant</code> , <code>PortSound</code> , <code>InterfaceLocal</code>) vs production's empty commitments	the formal model is parametric in evidence; production authors none	a future production evidence model could violate a condition unnoticed	nothing to violate today; <code>require</code> is reserved	a concrete compositional evidence model, before commitments are authored
mixed comparison rule (ADR-0026)	the formal order policy is between concept values only; production had to decide the concept-beside-plain-value case	<code>min(o1, 0.5)</code> accepted while <code>min(o1, o2)</code> refused	a tested matrix; <code>Explain</code> shows the observation	recorded as intended

deviation	why	risk	evidence that bounds it	closing the gap
affine charts hidden from formulas	conversion is safe; point/difference validation does not exist	a designer cannot write $20\text{ }^{\circ}\text{C}$	ISS-0004 open	the point/difference validation, or a decision not to need it
<code>bdl_ir::Ty::Unit</code> exists; the formal kernel <code>Ty</code> has no unit	production wants one type language for the interface and the kernel encoding; the formal model keeps the canonical types in <code>CTy</code> above the kernel	none semantically: <code>Ty::Unit</code> never types a Core term, a representation or a runtime value, which is exactly <code>elim unit = none</code>	<code>elim_canonical</code> ; the production tests that the encodings are inverse and the argument is erased	not planned: the formal <code>CTy</code> is the record that <code>Unit</code> is an interface type
simulation inputs are unresolved <i>unit-domain</i> declarations; the kernel <code>Input</code> provides for every unresolved declaration	an environment provides values, not functions	none: on a unit-domain source the two agree (<code>SimulationInput.value</code>)	<code>source_value</code> , <code>resolved_not_source</code>	not planned; recorded as surface policy
three roles in production; one formal role	production needs a name for the arrow type and the realized unit domain in every surface; the kernel does not	none: the role is a function of two kernel facts	<code>MappingBlock::role</code> tests; <code>the_role_is_one_answer_across_the_component_boundary</code>	none needed
provision proved; nothing provisions	the construction was audited before production builds a device catalogue	none today; the risk is building the transducer as adapter host code — the adapter refuses a design with a <code>Source</code> instead	PRP-0001 revised; ownership decided in advance	implement PRP-0001 (ISS-0016)
the model lowers a realization to fresh declarations <code>e, p</code> ; production lowers to a plan-level <code>SinkPlan</code>	designer-visible sinks would need ids, canvas positions and diagnostics; the plan is the lowest layer that preserves <code>d, o, d → o</code>	the equivalence of the two lowerings is observed by tests, not proved	<code>changing_the_realization_changes_no_behavior_and_only_the_raw_trace</code> ; the corpus's commands	a proof that the plan-level sink is <code>lowerΔ</code> 's machine sink, or a decision that the observation suffices
the raw command is an <code>f64</code> ; the peripheral takes an integer	production computes in <code>f64</code> (ADR-0011); a register is a count	a third numeric domain at the adapter, with rounding and refusal that the model does not have	the explicit <code>duty8</code> policy and its tests; reject-and-hold recorded as a revisitable policy	ISS-0006 (numeric representation on device) stays deferred
raw command proved; adapter tested	the formal model stops at <code>RawCommand</code> by decision (FVD-0134)	a faithful adapter is a matter of tests and hardware, not of the theory	recording sinks, the cross-build; no bench measurement	FVI-0022, ISS-0017

The snapshot

What is implemented, partial and planned at `aa6e7f4`, with the protocol history and the milestones of the last week, is Appendix F, so that this chapter's classification survives the next milestone and the snapshot is updated in one place.

VII.2 Rejected designs, as arguments

The kernel was obtained by a method: each candidate construct — from the earlier draft of the language, from the production roadmap, or from the standard repertoire of typed functional languages — was formalized in its smallest plausible form together with its alternatives, and each was asked the same questions. What does it reject that the others accept? What does it accept that it should not? Is it a special case of another? Which operations on a design are refinements under it? A construct entered the kernel only when every tested alternative failed for a stated reason, recorded as a theorem or an executed counterexample. The failures are among the strongest results in the record, and they are best read as arguments — a design pressure, a candidate that seemed to meet it, the counterexample

that broke it, and what was kept instead — before they are read as rows. Ten of them follow; the tables after them are the complete ledger, and the construct-by-construct table with every row’s evidence is docs/kernel/minimality.md in the formal repository.

Concepts by representation. *Pressure:* a designer’s `Tilt` and a motor’s `MotorAngle` are both angles, and a language that types them both as `q Angle` is simpler. *Candidate:* structural typing by representation, then — when that admitted the wire `motorTarget := tiltSensor` — identity as interface metadata checked by a direct-wire rule. *Failure:* the metadata checker is evaded by η -expansion, $(\lambda x. x) \text{ tilt}$ having the same flow and no direct wire, and the compositional form of the checker turned out to have the rule shapes of nominal typing itself (`bweak_evaded_by_eta`, §IV.2). *Kept:* the nominal type `sem s`, an ordinary constructor over an internal identity, with erasure proved sound so that generated code carries no residue.

Unrestricted construction of concept values. *Pressure:* a formula must be able to observe a tilt’s number and produce a brightness, so `rep` and `mk` must exist. *Candidate:* both available everywhere. *Failure:* $\lambda x. \text{mk Motor (rep x)}$ is a well-typed `Tilt -> MotorAngle` with no declaration and no visible crossing (`hidden_crossing_inside_unrelated_body`). *Kept:* observation free, construction licensed by the *grant* of the realizing declaration’s own signature — a capability the designer already wrote, so no annotation was added (§IV.2).

Reactive and event types. *Pressure:* the earlier draft had `Signal  $\tau$` and `Event  $\tau$` because time-varying values and occurrences are different things to a designer. *Candidate:* both as kernel types. *Failure:* under a tick semantics a signal type is inhabited by exactly the terms of τ and rejects nothing; within one domain an occurrence is an optional value and the streams of type `opt  $\tau$` are exactly the streams of multiplicity at most one; the one thing an event type would have to carry — multiplicity across rates — is a window, expressible as five declarations over `delay`, `sync` and lists once list data exists, and proved to compute exactly the window (`buffer_window_correspondence`, §IV.5). *Kept:* ordinary typed declarations, one temporal primitive, and the buffer as a surface elaboration.

Same-tick visibility between domains. *Pressure:* two domains active at the same instant might reasonably see each other’s current values. *Candidate:* a transport that reads the current activation. *Failure:* the scheduler’s order becomes observable — two priorities between the domains give two outputs (`scheduling_order_observable`). *Kept:* strictly-before: a transport sees only source activations strictly before the destination tick, which also makes cross-domain causality free (§IV.5).

Output arbitration at runtime. *Pressure:* a safety override and an interaction both reach for one light. *Candidate:* action requests as values, per-context policies, a resolve phase in each tick — the draft’s whole request-and-policy model. *Failure:* first-wins, last-wins and maximum over the same value graph give three physical outputs (`hidden_arbitration_observable`); effect rows duplicate the drive edge or flag a valid design; action values move the conflict into a collector that must itself be a policy. *Kept:* one explicit driver per logical Output, with every combination of behaviors an ordinary relationship upstream, visible on the canvas (§IV.7).

Structural order on all data. *Pressure:* the equation library wanted `min`, `max`, `clamp` at every type. *Candidate:* Phase 9b’s first form generalized `lt` to every data type through a structural order — booleans, options, lexicographic pairs, concepts by representation — and it was formally consistent. *Failure:* the audit asked what `mode1 < mode2` and `None < Some x` mean to a behavior designer, and the answer was nothing: any such order would come from a code, a constructor tag or an identity (`lt_rejected`, `min_mode_rejected`). *Kept:* the order reverted to quantities, with ordered concepts by declaration at the surface — a decision reversed by its own evidence, and recorded as such (§IV.3).

Affine units as missing information. *Pressure:* 0°C is 273.15 K, which no scale produces from zero, so the linear unit model cannot represent Celsius. *Candidate (Phase 10’s own conclusion):* a point/difference sort on the dimension is the missing information, and conversion is unsafe without it. *Failure of the conclusion:* the smaller hypothesis — that coordinatization erases chart identity and only the affine transformation between charts must be preserved — holds: the chart laws, the groupoid laws and the difference law are proved without a sort, and conversion never takes one (`convert_compose`, `difference_map`, `sort_orthogonal_to_conversion`). *Kept:* conversion complete as coordinate change; the sort demoted to optional arithmetic validation (FVD-0106 rewritten; §IV.4).

Retargeting the output to the raw type. *Pressure:* a light is dimmed by a duty cycle, so the obvious realization is to make the output accept the duty and drive an encoder into it. *Candidate:* `o.accepts := raw`. *Failure:* the existing `edge d -> o` fails `DriveWF` because the driver is typed at the concept, so the drive environment would have to be

rewritten and the abstract output’s meaning lost (`retarget_breaks_driveWF`). *Kept*: the logical Output unchanged, and realization as a lowering that adds an encoder declaration and a machine sink *below* it, with the behavior literally unchanged (§IV.7).

Fit and allocation as admissibility. *Pressure*: a device profile is admissible when its encoder fits the concept and the board can carry its requirements — the first Phase 14 form. *Candidate*: `FitsAndAllocates`. *Failure*: an Encoder value whose term is $\lambda n. \text{true}$ fits the light, allocates a PWM line, and is not typed `rep -> raw` at all (`exJ`, `admissible_needs_wf`). *Kept*: three judgments that never see each other — the encoder’s typing, its fit, a solvable board — with the gap kept refused by a production test (FVD-0139 supersedes FVD-0137; §IV.7).

An implicit device clock. *Pressure*: a device consumes at its own rate. *Candidate*: let the encoder live in another domain and resample implicitly. *Failure*: the drive edge and the domain judgment both fail (`exI`), for the same reason the hidden transport was refused in §IV.5. *Kept*: the explicit variant — the encoder in a device domain reading the driver through Phase 5’s sync with a stated initial representation — proved to preserve every judgment with no new instantaneous edge (`lowerSync_causal`, `lowerSync_correspondence`; §IV.7), and the carrier frequency kept as configuration, never a `ClockId`.

The pattern across the ten is the same: the candidate was the obvious construct, the failure was mechanized, and what was kept was smaller. That is also why the record keeps the failures. A reader who meets only the retained design would not know why the obvious construct is absent, and would be tempted to add it.

Kernel constructs retained

construct	why it could not be derived	evidence
<code>DeclId</code> and frozen expected types	references resolve by identity; retyping invalidates clients	probes 1–6, <code>local_refinement_preserves_global_wf</code>
nominal <code>sem s</code> with <code>0</code> , <code>rep</code> , <code>mk</code> under a grant	the numeric baseline accepts the invalid wire; a metadata checker is η -evaded; concept-as-declaration has category errors	Phase 2/3 Models A–C, <code>constructs_granted</code> , <code>temporal_state_preserves_semantic_identity</code>
<code>q d</code> with the algebra in <code>Prim.ty</code>	erasure accepts <code>length + time</code>	<code>dimension_mismatch_rejected</code> , <code>counterexampleB_baseline_accepts_length_plus_time</code>
<code>delay</code> , <code>sync</code> (one temporal read)	totality forces data-typed, top-level state; <code>delay</code> is <code>sync</code> at the own domain	<code>reactive_total</code> , <code>delay_is_sync_own</code> , <code>scheduling_order_observable</code>
<code>ClockId</code> , <code>Clocked</code>	a direct wire is ambiguous without it; equal rate is not the same domain	<code>equal_rate_not_same_domain</code>
<code>OutputId</code> , drive edge, <code>SingleDriver</code> , <code>CompleteOutputs</code>	two drivers make an output non-functional; hidden policy is observable	<code>multiple_direct_drivers_rejected</code> , <code>hidden_arbitration_observable</code>
<code>list τ</code> and its operators	a lossless window is unbounded sequence data	<code>bounded_summary_not_lossless</code> , Theorem M
<code>prod</code> , <code>pair</code> , <code>fst</code> , <code>snd</code>	function encodings are not data and cannot be state	<code>arrow_not_delayable</code> , <code>church_fst_rank</code>
<code>fold</code>	one eliminator derives every collection operation; operators never apply closures	<code>fold_total</code> , the <code>*_spec</code> family
<code>eq</code> at every data type (proof field)	boolean/concept/pair/list equality was unwritable	<code>Value.beq_iff</code>
<code>lt</code> on quantities	ordering is a property of magnitudes	<code>lt_rejected</code> , <code>lt_only_on_quantities</code>

Rejected constructs

candidate	reason it was considered	counterexample or theorem	verdict
Signal τ in <code>Ty</code>	the earlier draft’s reactive types	inhabited by exactly the terms of τ ; rejects nothing	REMOVE
Event τ (single domain)	occurrences as a type	<code>opt τ</code> streams are the multiplicity- ≤ 1 streams	REMOVE
Event τ (across domains), a buffer primitive	multiplicity across rates	the window is five declarations over <code>delay/sync/</code> lists; Theorem M	REMOVE (surface desugar)

candidate	reason it was considered	counterexample or theorem	verdict
effect rows	output effects in types	duplicate the drive edge or false-positive	REMOVE
action requests / action values	outputs as values	relocate the conflict without resolving it	REMOVE
runtime output arbitration	several drivers	hidden_arbitration_observable: the hidden policy is observable	REMOVE (single driver)
clock-indexed types <code>Signal[c, τ]</code>	domain in the type	clocked_type_forces_ polymorphism on every pure mapping	REMOVE
same-tick cross-domain visibility	“synchronous sub-domains”	scheduling_order_observable: the scheduler order becomes observable	REMOVE (strictly-before)
scheduler / staging rule	ordering simultaneous domains	unnecessary under strictly-before	REMOVE
explicit machine state / <code>Step</code>	a state machine object	Phase 4/8: state is gated self-delayed declarations; components hold state by instantiation	REMOVE
StateHandler kernel	contexts as a construct	elaborates to activation, entry, gated and reset state, conditional drivers (tested behavior only)	KEEP IN SURFACE-DESUGAR; nesting/handler-scoped clocks untested
state identity as a kernel notion	naming state	StateCellId by declaration and expression path is an engineering device; the kernel has none	REMOVE
user-defined typeclasses / dictionaries	constrained generics	closed <code>{Data, Eq, Ord}</code> covers the library; comparators recover the rest (<code>minBy_recovers_min</code>)	REMOVE
structural ordering on all data	Phase 9b’s first form	<code>mode < mode, None < Some</code> , lexicographic pairs have no design meaning	REMOVE (9c)
Set type	finite sets	<code>oneOf_mem, oneOf_dup</code> irrelevant: membership is a fold	REMOVE
Interval type	ranges	a pair with a convention; <code>inRange_spec</code> ; no case stores or compares a range	REMOVE
record type / row polymorphism	records	nested pairs with positional projection; no case needs shape-generic functions	REMOVE
kernel polymorphism (type variables, $\forall, \lambda, [\tau]$)	generics	prenex fragment = family instantiation; <code>matchTy_sound/_complete</code>	REMOVE (surface families)
higher-rank types	polymorphic function arguments	every candidate rank ≥ 2 with a rank-1 replacement	REMOVE
existential types	hiding	Phase-8a instantiation hides; encoding is rank 2	REMOVE
general recursion	loops	total language; <code>fold</code> is the one eliminator	REMOVE
general quantification in expressions	$\forall x : \text{Real}$	finite $\forall/\exists$ are folds (<code>forall_in_list</code>); unbounded forms belong to a commitment layer	REMOVE (surface finite forms)
general comprehension (generators, <code>yield</code> , <code>where</code>)	list syntax	nested binders cover every required case	REMOVE
runtime unit values	units as data	no case delays, syncs, stores or compares a unit	REMOVE

candidate	reason it was considered	counterexample or theorem	verdict
units in quantity types	unit-indexed q	1 m and 100 cm would differ in type	REMOVE
a kernel conversion primitive	convert	it is the composition of <code>inUnit</code> and <code>withUnit</code>	KEEP IN STANDARD LIBRARY
affine quantity sort in <code>Ty.q</code> (<code>Ty.q d sort</code>)	°C arithmetic	no conversion theorem needs it; <code>sort_orthogonal_to_conversion</code>	REMOVE from types; optional validation
kernel hole term	the Composer's slot	the surface <code>?</code> is refused by elaboration; holes belong to the editor	REMOVE
semantic <code>BehaviorGroup</code>	groups as objects	every group operation is the identity on the design	MOVE TO AUTHORING/UI
component runtime object	components at runtime	flattening produces an ordinary design; no system evaluator exists	REMOVE
aggregate socket as a declaration / tuple boundary	collapsed groups	Counterexamples 5 and 6: spurious dependencies	REMOVE
kernel <code>sum</code> type	enums	encoded as tag × optional payload (<code>exM</code>)	DEFER
kernel unit type / unit value / unit term	the literal reading of <code>() -> B</code>	<code>elim unit = none</code> ; a unit binder around memory is untypable (<code>delay_not_under_binder</code>); the kernel encoding permits memory (<code>zero_input_memory</code>)	REMOVE — KEEP <code>()</code> IN THE INTERFACE LAYER
a source semantic kind	sensors as a kind of declaration	the source role is <code>realizationOf d = none</code> ; a resolved <code>() -> A</code> never reads the input (<code>resolved_not_source</code>)	REMOVE / DERIVE
<code>A -> ()</code> as a physical consumer	the dual of <code>() -> A</code>	every pure total <code>A -> 1</code> is one function (<code>consumers_indistinguishable</code>); the drive edge names the receiver	REMOVE
numeric rates in the kernel	periods	a rate induces a schedule; validation data	MOVE TO VALIDATION
buffer capacity, overflow policy	bounded memory	<code>sufficient_capacity_preserves</code> , <code>negE</code> : only refusal preserves semantics	MOVE TO VALIDATION; explicit take in the design
preferred display unit	UI preference	<code>presentation_irrelevant_*</code>	MOVE TO AUTHORING/UI
interface <code>semanticRole</code> field with a direct-wire checker	concept identity as metadata beside the type	evaded by η -expansion: $(\lambda x. x)$ tilt has the same flow and no direct wire; the compositional form duplicates nominal typing	REMOVE (kernel); the concept stays the surface object
a source kind / <code>MappingKind::Source</code> / a <code>Source</code> keyword	sensors as a kind of declaration	the source role is <code>realizationOf d = none</code> ; a resolved <code>() -> A</code> never reads the input (<code>resolved_not_source</code>); a persisted kind could disagree with the shape after a definition is added	REMOVE (kernel, text, protocol); KEEP the <code>Source</code> as a derived surface role
a <code>ProvisionedSource</code> role or kind	a <code>Source</code> realized by a device	after provision the target has a realization and is a <code>Value</code> by the same rule (<code>provision_source_role</code>)	REMOVE
singleton-only provision (one raw declaration per <code>Source</code>)	the simple case first	the IMU case is the realistic one; the joint-section finding is invisible in the singleton (<code>no_joint_witness</code>)	REMOVE; shared raw is primitive, the singleton its case

candidate	reason it was considered	counterexample or theorem	verdict
a stateful transducer as the profile condition	debouncing and filtering in the profile	a typed impure term maps one raw value to two results at two ticks (exD); the induced input must be a function	DEFER (FVI-0020); purity is the condition today
“idempotent” provision	re-applying a deployment	the precondition fails after provision (provision_not_reapplicable); a second pass with another term is not a refinement	REMOVE the word; not re-applicable
Sensor as the library category	the everyday word	not every Source is a sensor; the template must carry no semantics the text cannot	MOVE TO AUTHORING/UI: the category is <i>Sources</i>
the Standard Library as the device catalogue	one catalogue	authoring fragments and deployment profiles are different objects with different owners	REMOVE; a device catalogue is separate (not built)
a Dart tokenizer or TextMate grammar as Studio’s highlighter	instant colour	a second definition of the language in the presentation layer, blind to role, unit-by-position and contextual keywords	REMOVE; one classifier in bdl-ide (ADR-0035)
retargeting the logical Output to the raw type (o.accepts := raw, drive the encoder into it)	the obvious way to realize an output	retarget_breaks_driveWF: the existing edge d -> o fails DriveWF because the driver is typed at the concept; the abstract output’s meaning is lost; exI	REMOVE — the logical output is never retargeted (FVD-0132)
an effectful encoder or device write inside the behavior (R -> (), Expr.write, an effect row, IO)	the machine as a term	consumers_indistinguishable: a pure R -> () cannot name a receiver; the behavior environment would depend on the device; RawCommand needs none of it	REMOVE — the machine boundary is a relation (FVD-0134)
a device kind in OutputSpec	one record per output	exH: one oLight realized by PWM and by I ² C has one behavior trace; a kind in the spec would make that two designs	REMOVE (never add) — the mechanism is deployment data (FVD-0131)
fit + allocation as admissibility (FitsAndAllocates)	the first Phase-14 form	exJ, admissible_needs_wf: an encoder λn. true fits Brightness and allocates a PWM line and is not typed rep -> raw	REMOVE — admissibility is typing ∧ fit ∧ a solvable board (FVD-0139 supersedes FVD-0137)
“closed and well-typed” as the encoder condition	typing without purity	(λk. λn. k) (delay 0 1) is typed at q ₀ -> q ₀ and encodes 7 as 0 at tick 0 and 1 at tick 1 (exEFG)	REMOVE — purity is the profile condition (FVD-0133)
an encoder that constructs a concept	λx. mk Other x	refused under Grant.none; encoder_constructs_nothing, encoder_decl_no_grant	REMOVE (FVD-0133)
an implicit clock crossing at the encoder (the encoder in another domain, a hidden resample)	a device at its own rate	DriveWF and WellClocked both fail (exI); Phase 6’s Counterexample G already refuses a hidden resample in the binding	REMOVE — a device clock is an explicit sync; a carrier frequency is configuration, not a ClockId (FVD-0138)

candidate	reason it was considered	counterexample or theorem	verdict
a required injectivity or round-trip of the encoding	“lossless” realization	a 4-bit PWM sends duty 6 for 40 % and for 41 % (exB_quantized); $\text{decode}(\text{encode } x) = x$ is neither assumed nor provable	REMOVE — correspondence is directional, quantization admitted (FVD-0135)
a many-to-one output lowering as a primitive (lowerMany)	atomic device frames	the H-bridge is a structured command from one concept (exD_hbridge); RGB is three per-tick commands or one Color upstream	DEFER — the singleton is primitive by decision (FVD-0136); the atomic-frame case is FVI-0022
a runtime dispatch table or symbol strings in the adapter	binding sinks by name	a display name would become an identity and binding a runtime lookup (ADR-0037)	REMOVE (production) — generated static parameters in sink order
clamping an out-of-range raw duty at the adapter	keep the peripheral moving	hides that the design commanded what the profile did not promise (ADR-0037)	REMOVE (production policy, revisitable) — reject and hold
Message τ / Event τ / Packet τ / Stream τ / Channel τ for a product-scale communication case	the auto_typer system looks like messages	every case — multiplicity, ordering, same-tick batches, cancellation, a bounded queue, freshness, faults, acks, cross-clock delivery — is state over the existing kernel (CommunicationExamples.lean)	REMOVE (FVD-0143, the fourth rejection)
a queue primitive	pending work	take cap (filter ... (append (drop-on-done prev) arrivals)) — library applications over one delay	REMOVE (FVD-0143)
a transaction / atomic-frame primitive for prepare–prepare–commit	a paired axis on two motors	one output with a pair command (exN_pair); two realizations agree tick by tick (paired_commands_of_one_value); the frame order is the backend’s per-tick commit	REMOVE (FVD-0148)
transport identity in a Source’s type	correlation and retries	discarded by the channel; two_providers_same_behavior, exL_theorem: unobservable; identity the product observes is semantic and lives in the state (exF_identity)	REMOVE (FVD-0144)
a Package object with versions, names or registries inside a judgment	package-provided profiles	the judgments take a profile; the origin is unread (assign_indistinguishable)	REMOVE (FVD-0145)
one admissibility verdict mixing contract and board	package compatibility	admissible_iff; exQ_contract_not_feasible: the same contract on two boards, feasible on one	KEEP SEPARATE (FVD-0147)
deduplication by payload at the provider	“the same command twice”	Move(+10); Move(+10) with two transport identities is two commands (dedup_of_fresh, exA_occurrences); only the transport can say what is a retry	REMOVE (FVD-0150)
an implicit total order across raw sources	one merged batch	no physical fact supplies one; source order is one explicit policy and a per-source reading is insensitive to it (perSource_of_interleaving)	REMOVE (FVD-0151)

candidate	reason it was considered	counterexample or theorem	verdict
a silent drop at the provider's bound	bounded memory	the flag is a value the design reads (<code>provide_overflow_iff</code> , <code>exE_overflow</code>)	REMOVE (FVD-0149)
an "event mode" flag on the logical output; an <code>Event/Stream</code> type for the output crossing	occurrence outputs	the window over the encoder into the device domain (<code>lowerWindow_correspondence</code>); the sink's type selects the lowering (<code>exA_window_vs_sample</code>)	REMOVE (FVD-0152)
a batch or transaction primitive at the adapter	a slower device receiving several commands	<code>batchOps : List Op</code> , <code>lineAfterBatch a fold</code> ; <code>paired_batches_of_one_window</code>	REMOVE (FVD-0153)
memory in the transducer term; a stateful provider primitive	debouncing, decoding, filtering	a machine below the reading is movable upstream with one trace (<code>provider_state_movable</code>); placement is visibility; a state that reads the design is the design's	REMOVE (FVD-0154)
a hidden provider clock; a fabricated initial reading; activation gated on the first sample	a device that samples in its own domain	an explicit sync with an explicit <code>InitRep</code> (<code>provisionSync_target</code>); supplied or unavailable; a schedule does not read an input	REMOVE (FVD-0155)
a profile range as evidence by declaration	vendor ranges	three evidence levels; the trusted one is a named assumption (<code>discharge_trusted</code>)	REMOVE (FVD-0156)
a "trusted" flag on computes; a proof language for transfer functions	package profiles	the field is a proof, derived from the term (<code>Channel.ofTerm</code>), decided at <code>bool</code> (<code>computes_of_bool</code>)	REMOVE (FVD-0157)
exception semantics for malformed readings	NaN, bad frames	the checking provider refuses; none or a flag above the boundary (<code>checked_typed</code> , <code>exE_checked</code>)	REMOVE (FVD-0158)

Decisions that changed

Not rewritten to look inevitable:

original position	evidence that broke it	replacement
Signal and Event as kernel primitives (draft)	a signal type rejects nothing; occurrences are optional streams; cross-domain multiplicity is a window over lists	ordinary typed declarations; the window as five declarations
runtime output arbitration (draft)	the hidden policy is observable	one explicit driver per sink; combination is ordinary computation
<code>eq/lt</code> at every data type through a structural order (Phase 9b)	<code>mode < mode</code> , <code>None < Some</code> , lexicographic pairs have no design meaning (Phase 9c audit)	<code>eq</code> on data; <code>lt</code> on quantities; order by declaration at the surface
affine units need a point/difference sort — "missing information" (Phase 10)	chart laws, groupoid laws and the difference law hold without a sort; conversion never takes one (Phase 10b)	conversion complete as coordinate change; point/delta is optional validation
a third project kind, text, beside JSON kinds with conversion (ADR-0020)	two persistence forms for one meaning are two authorities that drift	one project; Design, Code and Split are views (ADR-0023)
an authoring overlay of partial expressions (considered for the Composer)	a second representation of the formula to keep in step with the text	the slot ? in the text; every action a byte-range edit (ADR-0028)

original position	evidence that broke it	replacement
a relationship without inputs as “a value declaration of type B”, tested by <code>inputs.is_empty()</code> at every layer (production before ADR-0029)	no record said what the type of such a relationship <i>is</i> ; six layers each carried their own special case	one canonical type <code>domain(inputs) -> B</code> with <code>()</code> for no inputs, one predicate <code>is_unit_domain</code> , the kernel encoding by unit elimination proved exact (Phase 12)
a fixed-capacity list type with an overflow error (considered)	a runtime failure the formal model does not have; a type depending on a deployment fact	<code>Vec</code> under a validated bound; refusal on bounded targets (ADR-0024/0027)
Studio deriving the Source role from the shape of the projection (ADR-0032, first form)	two derivations of one fact drift the day either changes; the IDE service and Studio had already answered differently for port-backed declarations	one derived role in the model, stated by the compiler and the daemon on every projection; Studio maps the enum (ADR-0032 amended, 2026-09-20)
a BDL-specific semantic-token vocabulary in the language server (concept, mapping, output, ...)	every LSP client needed a mapping per name	the LSP standard names with BDL distinctions as modifiers; <code>unit</code> and <code>slot</code> the only additions (ADR-0035)
admissibility = <code>fit</code> $\wedge$ allocation (FVD-0137, Phase 14’s first form)	an encoder that fits and allocates but is ill-typed (<code>!n. true</code>) passed; wrong on the same evidence	admissibility = the encoder’s typing $\wedge$ <code>fit</code> $\wedge$ a solvable board (FVD-0139; ADR-0036 keeps the gap refused by a test)
a Source item that creates a concept and its Source in one drag (the first Standard Library form)	choosing an existing concept created a redundant one; a Source over a concept nobody chose	the Source sheet: the concept is the designer’s decision, existing or new in one transaction; the items are presets (protocol 0.23)
<code>drive light = brightness</code>	the first spelling of the drive edge	the relation reads backwards in a design; <code>drive light by brightness</code> says who drives what; the old spelling stays as legacy syntax with a hint and an opt-in migration

VII.3 Minimality, bounded

Two questions are asked of every construct, and they have different answers more often than not.

Kernel minimality asks whether the construct is needed for some design to be representable or unambiguous, with a theorem or counterexample as the argument. By this test `Signal`, `Event`, `effect rows`, `action values`, a `StateHandler`, a `group`, a `unit`, a `binder`, a `polymorphic type`, a `Source kind` and a `provision primitive` are all *unnecessary*: each is either derived from what exists or rejected outright, and the kernel band of §IV.1’s architecture figure is what remains.

Designer-facing vocabulary minimality asks whether the construct captures a distinct design concept that cannot be expressed cleanly as a property of an existing one — the cognitive-budget test of Part I. By this test a **Source**, a **behavior** (`group`), a **natural binder**, a **unit** on a literal, a **choice** in the Formula Composer and a **context** are all *valuable*: each names something a designer means, and each is kept at the surface, in presentation, or as a design recommendation for a surface not yet built. That a construct desugars away is a fact about the kernel; it is never a reason to remove it from the product. The Source is the clearest case — nothing in the kernel distinguishes it, the product draws it first — and the provision result of §IV.7 is what makes the two facts consistent: the surface concept is exactly a realization state, and the state changes at deployment by a construction the kernel already admits.

There is a third distinction the output side makes unavoidable: **kernel minimality is not whole-system capability**. Phase 14 added, above the kernel, an encoder, a machine sink, a lowering, three admissibility judgments and a machine boundary, and production added a profile registry, a plan-level sink, generated commands and a platform adapter. None of it entered the kernel, and none of it *could* have entered the kernel on the evidence — retargeting the output is refuted, an effectful term cannot name a receiver, and the lowering is expressible in the existing kernel with the behavior literally unchanged. The encoder was not promoted into the kernel because production needs one; it was kept out because the theorems show it is deployment structure. The full production architecture is therefore much larger than the kernel band of §IV.1’s figure and is supposed to be: the claim of minimality is a claim about the kernel, and the whole system’s capability is measured by what the constructions above the kernel can do without adding to it.

The verdicts below say which kind of minimality each row is about. REMOVE is a kernel verdict and a surface verdict at once only where the rejected table says so.

VII.4 The open agenda

The agenda has two halves that must not be confused: questions of semantics and formal correctness, which a theorem or a counterexample would settle, and questions about designers, which only a study would. Items production or the formal development has already answered are gone from this list; each remaining item names what exists and what would resolve it, and none is hidden in a “future work” sentence.

Formal limits, stated once

Causality is conservative for lambda-guarded cycles: `Causal` rejects them and the negative theorem does not cover them. The agreement between unfolding and tick-by-tick evaluation is proved for first-order wiring designs; higher-order closure equivalence is not. Theorem J (modular semantics) is proved for direct bindings on wiring designs with closure-free inputs; transported bindings under `MEv` and higher-order bodies are not covered, in either direction, for the same reason — a domain-indexed input for a transported port. Theorem R (extraction preserves behavior) is likewise proved for single-domain wiring designs without transported bindings. The `StateHandler` reduction covers the tested behavior and not contexts with their own clocks or independently clocked nesting. `InstAcyclic` is too coarse for extraction; a port-level graph would subsume `flatten_causal` and `flat_causal`. Four conditions rest on the abstract evidence relation (`Monotone`, `Equivariant`, `PortSound`, `InterfaceLocal`), and production discharges none of them because it authors no commitments. Hardware validation covers discrete pin and peripheral allocation with unary and binary constraints, and no numeric electrical, thermal or timing property; the explanation facility reports a first dead end, not a minimal core. Interface-level references — commitments that mention other declarations — are not modelled, so the dependency graph is over realizations only. Whether a realization may delegate its grant to a higher-order argument is untested. The buffer correspondence assumes an input source. The equation library’s evaluation lemmas assume the predicate value implements a Boolean function; the executed cases discharge it. The exact unit model is related to production floating point by a stated contract, not a theorem. Sums are encoded, not added. Two-hole operands are outside `solve` by design. No parser is modelled for the natural surface; its type annotations are the output of local inference, described, not proved. Provision is proved for pure transducers at the target’s clock, and for shared raw readings with a joint section where equality is claimed. Realization is proved for a pure, stateless encoder in the output’s clock, one logical output per machine sink, and stops at the raw command relation: the correspondence from the raw command trace to a platform adapter’s operations, and from those to a physical effect, is not a theorem, and the model does not contain the adapter at all. No theorem covers the Rust or Dart code, and none covers a board.

Open formal questions

Each names what exists and what would resolve it.

1. **A general edit/invalidation relation** (FVI-0015, deferred and merged into FVI-0001 — blocked until production authors a commitment). Refinement is a preorder with proved client stability; an arbitrary edit is outside it and forces a recheck of dependents. Production classifies edits into seven invalidation categories (ADR-0009). Missing: a formal relation that says, per edit kind, exactly which established facts survive — the theory behind incremental re-analysis. Would resolve: a proved per-category preservation theorem.
2. **Commitments that depend on interface references** (FVI-0001, deferred; ISS-0003). A commitment is a property id discharged by evidence over the realization; a commitment mentioning *another declaration* has no model, so interface-level cycles are invisible. Would resolve: an interface-level dependency relation and its interaction with refinement — and, before it, any production authoring of commitments at all.
3. **Lambda-guarded causality precision** (FVI-0003, deferred: closures cannot outlive a tick, so the conservative count loses only data-dependent dead closures, which the language does not want). `Causal` counts a reference under a lambda as instantaneous. Would resolve: a causality judgment with application sites and its totality theorem.
4. **Higher-order closure equivalence** (FVI-0004, deferred as churn: unfolding is not the executable definition). `unfolds_preserves_eval` holds for wiring designs. Would resolve: an equivalence of closures across unfolding, or a first-order normalization of higher-order realizations before unfolding.

5. **Full multi-clock component semantics** (FVI-0006). Theorem J under MEV with transported bindings needs a domain-indexed consistent modular input and a proof of its existence. Production supports the case; the theorem does not.
6. **Nested and stateful component elaboration limits** (FVI-0007, ready with FVI-0006; FVI-0009 deferred until the surface has contexts; ISS-0007, ISS-0010). toComponent's side condition is decidable and unproved; packaging inside a component body is not modelled; contexts with handler-scoped clocks and independently clocked nesting were not examined, and production offers neither temporal modifiers nor contexts. Would resolve: an elaboration with a preservation theorem, or a counterexample forcing a kernel construct.
7. **Affine physical arithmetic validation** (FVI-0017, deferred: coordinate semantics is solved, the validation is optional; ISS-0004). The point/difference sort as an operation-level validation, its rules (affAdd, affSub) at the surface, and the Composer offering only difference units for a delta slot. Would resolve: the validation and a decision to offer °C in formulas, or a decision not to need it.
8. **Hardware: minimal unsatisfiable cores** (FVI-0028, deferred — explanation, not correctness). diagnose reports a first dead end. Would resolve: a minimal core with a proof of minimality, or a decision that the first dead end is the better explanation.
9. **Numeric and shared-configuration hardware feasibility** (FVI-0011, narrowed and now concrete: RP2040 PWM slices share top and divider, masked today by one fixed carrier). Out of the solver's scope; a design can be allocated and still exceed a current budget. Would resolve: a numeric constraint layer beside the finite solver, with its own soundness.
10. **A generated-code refinement proof**. The core is held to the reference evaluator by differential tests, and its raw commands to the encoder over the evaluator's outputs. Would resolve: a proof that lowering plus code generation refines Ev/MEV and that the generated Commands are RawCommand — or a verified evaluator — closing the largest deviation of §VII.1. A proved static bound analysis is the same gap on the capacity side.
11. **The output boundary beyond a pure encoder** (FVI-0022 split into FVI-0023 ... FVI-0027; ISS-0017). Phase 15 answered the first two up to the abstract sink operation and the explicit device clock; Phase 16 added that the paired axis is not an atomic-frame case (FVD-0148); Phase 17 built the occurrence-preserving crossing to a slower device (lowerWindow, FVD-0152) and the adapter's batch (FVD-0153); open: below the operation (FVI-0023), a device that acknowledges and the initial representation (FVI-0024, narrowed), a stateful witness (FVI-0025), the atomic-frame criterion (FVI-0026); deferred: output commitments (FVI-0027). Originally: What Phase 14 leaves open, and the first platform adapter does not close — it applies each command independently and untouched: *stateful output adapters* (slew-rate limiting, PWM dithering, protocol batching, servo smoothing, hysteresis), and whether each belongs to the behavior as an ordinary declaration with delay, to a stateful lowering with a stream-level correspondence theorem, or to the backend; *a device clock different from the output clock* — the explicit-sync variant of the lowering (FVD-0138) and the line between an activation clock and a carrier frequency, which is configuration; *atomic multi-value frames* — whether a device that must receive several logical outputs in one indivisible frame (a display controller taking a Mode and a Level the behavior drives as two outputs, refusing a frame with one) ever forces a lowerMany beyond per-tick batching or upstream combination, which FVD-0136 decides for the singleton without settling; *the codegen and adapter correspondence* — abstract trace → raw command trace is proved, raw command trace → generated command → adapter operation → physical effect is not, and the last arrow is not even a testable statement inside this project; and *commitments on outputs*, which production does not author and whose discharge by an encoder's declared transfer would be the output analogue of FVD-0128.
12. **The input boundary beyond a pure transducer** (FVI-0020; PRP-0001, ISS-0016) — resolved by Phase 18, remainder by remainder: stateful transducers are machines movable upstream with one trace (FVD-0154); the Source-side device clock is an explicit sync or window with an explicit initial value (FVD-0155); commitment discharge has three evidence levels (FVD-0156); computes is derived from the term and decided at a finite type (FVD-0157); out-of-type readings are refused and cross as none or a flag (FVD-0158). Freshness was behavior state since Phase 16 and the occurrence contract Phase 17's (FVI-0029). What stays outside the development: the bound or range a physical device needs — a deployment assumption — and the production slice (ISS-0016).
13. **Enums and sums** (ISS-0005). Encoded as tag × optional payload; production keeps user enums open. Would resolve: a case that needs match exhaustiveness beyond the encoding, and then one eliminator term former like fold; or a decision that the encoding is the language.

14. **A surface form for occurrence windows** (ISS-0001; FVI-0012 resolved — a device is a deployment profile, not a component), so a bounded refinement can be recognized and a ring representation offered without a hidden transformation.
15. **Several candidate definitions with one active** (FVI-0014, deferred; ISS-0002): whether this is a surface convenience over a write-once kernel realization.
16. **Grant delegation** (FVI-0018) — resolved: unnecessary under `constructs_granted` and `lib_expansion`; the `display_name_table` (FVI-0019) — resolved by FVD-0014; the general `elim` (FVI-0021) resolved by the canonical-interface design; folding the clock into the interface (FVI-0013) rejected by FVD-0046.
17. **A mathematical specification backend and a verification backend.** The kernel is a specification; a backend rendering a design as a mathematical document for engineering hand-off does not exist; commitments beyond typing have an evidence slot and no discharge mechanism; the exact chart model and the choice-free rational field are the first pieces of a verification story about quantities.

Production engineering work

Not formal questions; work production has decided and not yet done, listed so that *not yet built* is never mistaken for *not yet understood*. In production's priority order at the snapshot: a device that provides a Source's value on the board — the adapter's remaining half, and the implementation of PRP-0001 (ISS-0016; the first adapter refuses a design with a Source); I²C and H-bridge sinks on the RP2040 (`i2c_level8` and `hbridge_signed` are refused for the board today) and the device catalogue beyond the five witness profiles (ISS-0017); build orchestration in `bold` for the generated crate and its firmware; flashing through `probe-rs`; telemetry back into Studio and the Monitor page, today a placeholder; a timer-driven tick on the AVR and a third target family (an ESP32-S3) after the second proved HAL independence; the core's numeric representation on device (ISS-0006, deferred); supplied Rust components (ADR-0005, designed, not built); and the items with no physical-boundary content — projection deltas and a persisted edit history (ISS-0009), a structural diagnostic entity for outputs (ISS-0008), a linear `zip` in the core (ISS-0013), the compiler's diagnostic sentences in every locale (ISS-0015), runtime loading of board files, and the Studio gaps production's own status page lists.

Empirical questions, explicitly

The interaction model of Part III and Part VI is fully described, largely built, and no part of it has been evaluated with users. The following are hypotheses, to be tested, and nothing in this document — no theorem, no test, no differential corpus, no working screenshot, no example a reviewer happened to read correctly — is evidence for any of them:

- that industrial designers can *read behavior intent* from a BDL artifact — the canvas, the text, or both — without translating it back into device terms first;
- that they can *author* useful product behavior without prematurely reasoning in device APIs, and that the point at which a design meets a device (the Deploy page, a realization profile) is late enough and visible enough;
- where they need *escape hatches* — supplied code, a raw value, a device-specific command — and whether the boundary constructions of §IV.7 put those hatches where designers reach for them;
- that an unresolved typed relationship is a natural stopping point, that *signature-first partiality* helps actual design work rather than merely being permitted, and at what task complexity it is taken up;
- that a **third party can recover design intent** from the behavior graph and the textual relationships — the *reverse-readability* question: the present system suggests that a reviewer who did not author a design may infer what the product is meant to do from its concepts, relationships and drives, and one such reading has been observed anecdotally during this project; that observation is not evidence of general readability, and a study that hands a BDL artifact to reviewers and measures recovered intent against the author's is the only thing that would be;
- that Source, Rule and Value are learnable as one distinction, and that *not applied* with an offered value is understood;
- that socket hue for identity and socket shape for value form are read correctly without type labels;
- that the three information levels place each fact where a designer looks for it, and that Explain is opened rarely and usefully;
- that the Formula Composer's slot expectations and candidate lists reduce the semantic-translation cost the workshop exposed, and that designers read `all` reading in `readings`: better than `all(readings, reading => ...)`;

- that product-language diagnostics (“This adds an angle and a time”; “Mode values have no default order”) are understood and acted on;
- that Design, Code and Split are experienced as views of one thing, and that completion, navigation and colour in the Code view help rather than distract;
- that groups, packaging and instances match how designers organize behavior;
- that a designer distinguishes a valid design from a deployable one, and a chosen realization from a changed behavior, when the workspace reports them apart;
- learnability across the vocabulary; productivity against a Node-RED- or Arduino-style baseline; cognitive load in the sense of the cognitive-dimensions framework.

The studies that would answer them are restated below with their measures. Neither has been run.

Expected cognitive advantages, as hypotheses

Lower viscosity. Changing a transfer function edits one Mapping definition rather than a procedural chain of read/compute/store/write nodes.

Reduced hidden dependencies. Semantic signatures expose what a relationship consumes and produces, while the kernel rejects incompatible connections before device code exists.

Reduced premature commitment. An unresolved declaration allows a designer to commit to a relationship without committing to its implementation, sensor, or exact parameters.

This last item requires a qualification that the rest do not. Signature-first authoring is not claimed to be the spontaneous habit of every designer. The author, working on a hardware project, adopted it when the complexity of a subsystem exceeded what could be held in view at once, and did not adopt it on simpler tasks, where a signature and its definition were written in a single motion. The hypothesis is therefore narrower than a claim about how designers think. It is that conventional tools provide no legal position for an undefined relationship, so the strategy cannot be adopted even when it would help, and that BDL supplies that position at no cost. Whether inexperienced designers take up the strategy once it is available, and whether doing so improves their designs, is an empirical question and is included in the comparative study.

Better role expressiveness. Contexts, Mappings, and temporal modifiers correspond to product-design concepts rather than generic program-control constructs.

Improved error locality. A mismatch is attached to a product relationship or binding rather than surfacing later as an embedded runtime fault, and a hardware infeasibility is attached to the binding that causes it.

Comparative study

A first controlled study should compare BDL against at least two baselines: a statechart-based prototyping environment and a node-based or Arduino-style implementation workflow. Participants should be industrial-design students and practitioners with limited professional software-engineering experience.

Tasks should include specifying a sensor-to-actuator mapping; adding temporal qualification such as “for 300 ms”; adding an orthogonal safety override that competes with an interaction context for one output; replacing a sensor with a different sample rate; relating a slowly updated quantity to a fast interaction, which forces a cross-domain decision; choosing a board that cannot accommodate the design; choosing a realization for an output and saying what changed; modifying a mapping late in the task; and, for the reverse-readability question, reading a design one did not write and stating what the product does.

Primary outcomes should not be limited to task time or a usability scale. More important measures are semantic errors in the final behavior; the number of implementation-only concepts participants must manipulate; time to detect an impossible or conflicting behavior; fidelity between verbal design intent and the elaborated model; the number and diversity of behavior alternatives explored; the quality of handoff to an engineer who did not observe the authoring session; subjective confidence calibrated against actual correctness; and whether, and at what level of task complexity, participants declare a relationship before defining it when the tool permits both.

The last measure tests the hypothesis stated above rather than assuming it. Because signature-first authoring may be a strategy that appears only above a complexity threshold, the tasks should vary in scale, and the default state of

a newly created relationship is itself a manipulable factor: a block that opens onto an empty formula editor and one that opens onto a signature with an explicitly legal undefined body invite different first actions. A small comparison of these two defaults is considerably cheaper than the full study.

The interaction model of Part III adds hypotheses of its own. Whether the single-driver diagnostic leads participants to an explicit combination block they can later read, whether the cross-domain question is answered correctly for a safety condition, and whether participants distinguish a valid design from a deployable one when the workspace reports them separately, are each measurable in the tasks above.

Field study

A controlled study cannot establish whether the representation fits real design practice. A second phase should embed the tool in a semester-long product-design studio or an industry project. The study should observe where unresolved declarations persist, which concept types designers invent, where they request escape hatches, how often the single-driver condition is met by a combination rule the designer finds natural, and how often engineers reinterpret or replace BDL artifacts during implementation. This field evidence is necessary before claiming that the language is native to industrial design rather than merely pleasant to its authors.

Risks of the design

Semantic-type proliferation. If every semantic distinction creates a visible type, the editor may become bureaucratic. The system needs reusable type libraries and sensible defaults, and the study should record which types designers invent.

Formula anxiety. Not every designer wants to write equations. Formula authoring must coexist with curves, examples, and direct manipulation.

Hidden elaboration. A context that elaborates to an activation declaration, an entry declaration, gated state, and a conditional driver is, in the elaborated design, several objects the designer did not draw. Hiding this can make runtime behavior mysterious. The explanation view of the interaction model is the proposed answer; whether designers use it, and whether it explains what they came to ask, is untested.

False confidence. A formally typed diagram can look verified even when no physical property has been checked, and a typed, causal, clock-consistent, output-complete design can be unplaceable on the chosen board. The workspace states of the interaction model exist to keep these apart, and the risk is sharpest for obligations discharged by declaration rather than by analysis.

Complexity migration into tooling. Much of what was removed from the kernel — event policies, output selection, context semantics — reappears as elaboration. The kernel is smaller and better understood; the elaborator is larger, and it exists (Part V) and is held to the kernel by differential testing rather than by proof (§VII.1); below it the realization lowering and the platform adapter add two more layers that are tested and not proved. The claim that the surface is “only syntax” over the kernel is a claim about tested cases, and the untested cases are the ones most likely to demand a kernel extension.

Usability hypotheses unestablished. Every statement in this document about what designers find natural is a hypothesis, including every sentence of the interaction model. The one anecdote reported is a single author’s practice on a single project.

Intellectual Context and Related Work

The reader now has the whole system in view, and this chapter places it against the traditions it intersects: what BDL borrows from each, and where its abstraction boundary differs. It comes after the argument rather than before it so that each comparison is made against a system the reader already understands. Each subsection ends with that difference, stated as a difference in *artifact* or *author*, not as a claim of superiority; the chapter is context for the design, not a case for its novelty.

Industrial design tools and physical prototyping

Phidgets reduced the implementation cost of physical interaction by presenting hardware components through a uniform software abstraction [1]. d.tools integrated physical prototyping, statechart-based behavior, testing, and analysis for designers [2]. Exemplar addressed the same authoring cost from the opposite direction, letting designers demonstrate sensor behavior and having the system infer the recognizer [8]; in BDL that technique is one of several ways to supply a definition, attached to a signature that already fixes the relationship’s semantic boundary. The Arduino ecosystem [9] is the de facto medium of the workshop reported in Part I, and its `loop()/digitalRead/analogWrite` vocabulary is the operational structure whose missing “place for a fact” motivated the language. The difference from these systems is not that they could not represent behavior. It is that their dominant representation still asks designers to formulate substantial portions of behavior in an operational structure; BDL investigates whether the typed semantic relationship can be the primary artifact, with operational machinery elaborated underneath.

Visual programming and end-user programming

Node-RED [10] is the closest deployed representative of the flow-graph medium that Part I analyses: nodes are computations wired by message passing, and the graph is a dependency view onto an event loop. Scratch and its block-language descendants [11] showed that syntax can be removed as an obstacle without removing the program-counter model, and Part III’s argument is precisely that the program counter, not the syntax, is the cost for designers. End-user software engineering [12] documents the tension between low-threshold authoring and the errors that follow from the absence of static structure; BDL’s answer is to make the static structure — concept identity, dimension, domain, single driver — the medium, and to make it legible through the three information levels rather than through diagnostics after the fact. Cognitive dimensions [4] supplies the vocabulary — viscosity, hidden dependencies, premature commitment — in which §VII.4’s empirical questions are posed.

Synchronous languages and functional reactive programming

The kernel’s temporal basis is that of the synchronous dataflow tradition. Lustre’s *pre* with an initial value, in a declaration-per-stream setting, is the delay primitive here [5]; the rule that a value computed at an instant is visible at the next instant, applied across domains, is the strictly-before rule; and causality as a static property follows the same line [13]. Esterel established the synchronous hypothesis under which these semantics are deterministic [14], and SCADE industrialized the tradition with a qualified code generator [15] — the destination that §VII.4’s “generated-code refinement proof” would move BDL toward. Where these languages recover clocks by a clock calculus [6], BDL requires domain identity to be declared, for the reason given in Part VII. Zélus extends the lineage to hybrid systems [16], which is the direction in which the continuous-dynamics limitation would have to be addressed.

Functional reactive programming established behaviors and events as compositional abstractions for time-varying computation [17], and FrTime gave a dynamic dataflow embedding with formal semantics [18]. BDL’s reactive core is much more restricted, and the restriction is a result rather than a starting point: under a tick semantics with one domain, a signal type rejects nothing and an event type is an optional-valued stream, so neither appears in the kernel; across domains, the event buffer is five declarations over *sync* and lists (§IV.5).

Model-based design, statecharts and systems engineering

LabVIEW established the graphical dataflow instrument-control paradigm; Simulink with Stateflow combines dataflow blocks with hierarchical state machines; Modelica models physical systems through acausal equations with units and dimensions [19], [20], [21]. These systems are considerably more capable than BDL. The distinction is in the primary artifact and the intended author: in each of them the artifact is an executable model whose blocks denote computation and the author holds an engineering model; BDL’s artifact is a set of typed relationships that need not yet compute anything, and the author holds a product model. Statecharts [22] are the representation that BDL’s

contexts replace for the designer: a context is a named situation with an activation condition, elaborated to gated and reset state and a conditional driver, rather than a transition table the designer maintains. SysML and model-based systems engineering address precise system structure, requirements, and verification at a broader level [23]; BDL is intended as a front end for early design that could later export into such representations, not as a replacement.

The model-driven-architecture ladder is a useful comparison because BDL’s boundary sits at a different rung. Classic MDA moves a computation-independent model through a platform-independent and a platform-specific model to an implementation, and the platform-independent artifact is already a software model — classes, components, messages. BDL’s artifact is a behavior design plus a realization and deployment step, and the behavior design is itself executable, formal and simulatable *before* any software ontology appears: its objects are concepts, relationships, domains and outputs, and the elaborator, the generated core and the platform adapter are the rungs below it. BDL places the platform-independent boundary above the conventional software-implementation ontology; it does not replace MDA, and a BDL design could be the input to an MDA pipeline that starts where BDL’s deployment step ends.

Typed holes, live and structured editing, projectional editing

Hazelnut demonstrated that incomplete structured terms can remain statically meaningful under bidirectional typing [24], and Hazel extended this to live evaluation around holes [25]. BDL takes the idea that incompleteness is a first-class static state and relocates it: the kernel object is a named declaration whose realization is optional; holes are positional in the Hazelnut tradition and named here at the declaration level, while the Formula Composer’s slot ? (§IV.4, Part VI) is a positional hole *inside* a definition, elaborated as an expression with an unresolved type and never shipped. The stability result concerns clients of a declaration under refinement rather than the typing of the incomplete term itself. Projectional editors [26] edit an AST directly and render it; the Composer is deliberately not one — the text remains the source of truth and the Composer is a projection over it (ADR-0028), which is why byte-range edits, ephemeral node identity and a stale-projection policy exist at all. Cognitive dimensions [4] frames the trade-off these editors make and the one BDL makes.

Dimensional typing and units

Dimensional typing follows the units-of-measure line begun by Kennedy [27]. What is specific here is the placement — the algebra in primitive operator types with no dimension-specific rule, and the separation of dimension from nominal identity — together with the mechanized observation that the numeric baseline is the erasure. §IV.4 goes further than the units-of-measure literature usually does in two respects: units are *coordinates on a dimension* with a symbolic exact scale group, and affine units are charts whose conversions form a groupoid of affine maps with a proved point/difference decomposition; Modelica’s `displayUnit` [21] is the closest deployed analogue of §IV.4’s presentation layer.

Effects

An earlier draft of BDL borrowed the separation between operation and interpretation from algebraic effects [28] and anticipated scoped effects for context-sensitive interpretation [29]. The kernel has no effect system. The negative result is stated narrowly in §IV.7 and does not bear on effect systems in general; it bears on the formulations tried for this design problem, where a single explicit driver per output expressed everything the request-and-policy model expressed and made visible what it hid.

Resource allocation

The hardware validation layer is a finite constraint satisfaction problem with unary and binary constraints [7], and its solver is a plain backtracking search whose soundness and completeness are proved. No claim is made relative to the embedded co-design literature; what is specific to the layer is architectural — the design is never an input to the solver, and feasibility is kept as a separate, non-monotone kind of evidence — rather than algorithmic.

Formal verification, verified compilation and differential testing

The kernel is mechanized in Lean 4 [3] without external libraries, and every result in this document rests on propositional extensionality and quotient soundness alone. The relationship between the kernel and the production compiler is *specification*, not *extraction*: production reimplements the reference evaluator in Rust and holds itself to it by differential testing [30] over a corpus (Part V). This is the weakest link in the trust chain (§VII.1) and the point where a verified-compiler approach in the CompCert tradition [31] would apply; §VII.4 lists a refinement proof from

lowering to Ev/MEv as an open problem. The choice-free rational field of §IV.4 was built because importing a general-purpose library would have brought classical choice into the axiom base; that discipline is a methodological choice of this project and not a claim about the libraries.

Embedded DSLs and toolchains

Production BDL's shape — a daemon with a project model, a compiler producing a `no_std` core, and a thin platform adapter — is a conventional embedded-DSL toolchain, and no novelty is claimed for it. What is specific is the placement of the semantics: the generated core is held to a reference evaluator that is itself held to a mechanized kernel, and the capacity, bounds and allocation questions are answered at validation rather than by a runtime allocator (§IV.5, Part V).

Closing

Modern industrial products increasingly combine physical form with sensing, computation, and control, yet designers still lack a behavior medium with the immediacy that CAD provides for geometry. BDL proposes that the missing medium should not be a friendlier version of procedural programming. It should be a language in which typed product relationships are first-class design artifacts, and in which an unresolved relationship is a legal state of the design rather than a defect in a program.

What that looks like in use is a workspace in which a designer names what the product is about, draws the relationships between those things, and leaves each undefined until there is something to say; refines them locally, by formula or curve or example, without the diagram changing shape; states time as a qualifier rather than as a timer; gives situations a name and a boundary rather than a transition table; is asked for one final target where two behaviors reach for one output, and for one stated way of seeing across a boundary where two quantities move at different speeds; and chooses a board last, receiving either a pin allocation or a conflict, with the design itself untouched either way.

The kernel that supports this is small, and it is small for reasons that were checked. A declaration has a frozen type, growable public commitments, and a write-once body; clients are typed against the type view, and refinement preserves what they established while edits reopen it. Semantic concepts are nominal, represented through a write-once binding, and constructed only where a signature announces them; dimensions are carried by the types of primitive operators. One temporal primitive reads a clock domain at its previous activation, and single-domain delay is its diagonal; evaluation is deterministic and total exactly on causal designs; every designer-facing temporal operator is a shape over it. Clock domains are nominal and checked by a judgment rather than a type; crossings are explicit, initialized, and strictly earlier. Physical outputs are nominal sinks with one driver each, and every combination of behaviors is ordinary computation upstream of the drive edge. A separate validation layer decides whether the design fits a board, and its evidence is kept apart from the evidence that survives refinement.

Each of these is minimal among the designs that were tested, and this document has tried to say, for each, what was proved, what was rejected by counterexample, and what was preferred. The elaborator, the editor, the realization lowering and a first platform adapter have been built and are described in Parts V–VI; the input-side device binding and the device catalogues have not; the last arrow, from a raw command to a physical effect, is tested and not proved; the studies have not been run, and they are where the claims about designers would be tested. The research question is not whether designers can be taught a simpler programming language. It is whether product behavior can become a *design material* whose structure is intuitive at the surface and rigorous underneath, and the kernel presented here is the part of that question that can now be stated precisely.

Appendix A — Formal notation

The notation is the Lean development's, kept uniform across Part IV; where an earlier phase used a different spelling the current one is used throughout and the old one is mentioned only in Appendix G.

symbol	Lean object	meaning
Δ	<code>DeclEnv : DeclId → Option DesignDecl</code>	the design: an environment of declarations
δ, DeclId	<code>DeclId</code>	a declaration's stable identity; display names are not modelled
<code>DesignDecl</code>	<code>(id, interface, realization)</code>	a declaration: identity, interface, optional realization
<code>DeclInterface</code>	<code>(expectedType, commitments)</code>	the frozen expected type and the monotone public commitment list
Δ^{ty}	<code>DeclEnv.tyView</code>	the type view: the expected type of each declared identity
Θ	<code>ConceptEnv : ConceptId → Option Ty</code>	the write-once representation binding of concepts
$C, \text{ConceptId}$	<code>ConceptId</code>	a concept's identity; <code>sem</code> s its nominal type
$q \ d$	<code>Ty.q Dim</code>	a physical quantity of dimension d (an exponent vector)
<code>Ty</code>	<code>bool nat arr sem q opt list prod</code>	the kernel types
G, Grant	<code>Grant, Grant.of τ</code>	the construction grant: the concepts in result position of a signature
<code>rep e, mk s e</code>	<code>Expr.rep, Expr.mk</code>	observe a concept's representation; construct a concept value under a grant
<code>declRef d</code>	<code>Expr.declRef</code>	a reference to a declaration by identity
<code>delay init e</code>	<code>Expr.delay</code>	last activation's value of e in the own domain, <code>init</code> before
<code>sync c init e</code>	<code>Expr.sync</code>	the value of e in domain c at c 's last activation strictly before now, <code>init</code> if none
<code>fold f z l</code>	<code>Expr.fold</code>	the list recursor, the one term former that applies a function value
$\Theta; \Delta; \Gamma \vdash_G e : \tau$	<code>HasType $\Theta \ \Delta \ G \ \Gamma \ e \ \tau$</code>	typing; reads Δ only through the type view
<code>Satisfies ev $\Theta \ \Delta \ \Gamma \ e \ S$</code>	<code>Satisfies</code>	e has S 's type under S 's grant and evidence discharges every commitment
<code>GlobalWF ev $\Theta \ \Delta$</code>	<code>GlobalWF</code>	every stored declaration sits under its identity and its realization satisfies its interface
<code>$\sqsubseteq$, InterfaceRefines</code>	<code>InterfaceRefines</code>	same type, commitments grow
<code>DeclLeq, EnvRefines</code>	<code>DeclLeq, EnvRefines</code>	the structural refinement order on declarations and environments
<code>Ev $\Delta \ I \ t \ p \ e \ v$</code>	<code>Ev</code>	single-domain tick-indexed evaluation under input I and local environment ρ
<code>MEv $S \ \Delta \ I \ c \ t \ p \ e \ v$</code>	<code>MEv</code>	multi-domain evaluation in domain c under schedule S
I	<code>Input : DeclId → Nat → Value</code>	the input stream: the environment's value for every unresolved declaration
S	<code>Sched : ClockId → Nat → Bool</code>	which domains activate at which global ticks
<code>prevAct $S \ c \ t$</code>	<code>prevAct</code>	c 's last activation strictly before t
K	<code>ClockEnv : DeclId → Option ClockId</code>	each declaration's timing domain; none is domain-agnostic
<code>Clocked $K \ c \ e$</code>	<code>Clocked</code>	the domain judgment

symbol	Lean object	meaning
Causal Δ	Causal	the instantaneous dependency graph is acyclic, witnessed by a rank
Ω	OutputEnv : OutputId $\rightarrow$ Option OutputSpec	each logical output's accepted type and clock (the "physical sink" of Phase 6)
β	DriveEnv : DeclId $\rightarrow$ Option OutputId	the drive edges, write-once per declaration
DriveWF, SingleDriver, CompleteOutputs	—	the output well-formedness conditions
Source Δ d	Source	realizationOf d = none: the environment provides d
CTy, elim, encode, decode	Surface/UnitDomain.lean	canonical interface types above the kernel and the normalization to kernel types
Channel, DeviceProfile, Provision	Surface/Provision.lean	a transducer with its transfer function; a raw type with channels; a raw declaration with its channel assignment
provision Δ P, induced Δ P I'	Surface/Provision.lean	the provisioned design; the abstract input induced by a raw input
Encoder, EFits, Realization	Surface/OutputRealization.lean	a pure encoder rep $\rightarrow$ raw with its transfer function; its fit to the output's accepted type; a logical output, its driver, a fresh machine sink and a fresh encoder declaration
RawCommand S Δ I Ω β R t w	Surface/OutputRealization.lean	the command specified for realization R at tick t on the unchanged design — the machine boundary
lower Δ , lower Ω , lower β , lowerK	Surface/OutputRealization.lean	the lowering: the encoder declaration, the machine sink, the new drive edge, the sink's clock
Admissible $\emptyset$ accepts H P	Surface/OutputRealization.lean	the encoder's typing $\wedge$ EFits $\wedge$ a solvable board — three judgments that never see each other
Hardware, Requirement, Assignment, ValidFor, solve	Validation/Hardware.lean	the target table, the design's needs, an allocation, its validity, the solver
CapacitySufficient, requiredCapacity	Validation/Capacity.lean	window capacity under a schedule

On categorical language. Where the text says that the canonical types form an interface layer whose normalization is inverse to encoding, or that compatible charts form a *groupoid* of affine isomorphisms, the claims are theorem-level statements in the development (decode_encode, convert_compose, convert_inverse) and the categorical words are explanatory. No categorical framework is formalized; “realization as a functor preserving composition” and “natural transformations between elaborations” are intuitions for future work, not results.

Appendix B — Theorem and result index, by concept

Every name is a Lean declaration in KCN-judu/BDL_FV (the file is given per group; theorem names are unqualified where the namespace is BDL). *Kind* is **T** theorem, **C** counterexample (a theorem whose content is a rejection), **X** executed example (decide/#eval), **D** definition-level fact (rfl/Iff.rfl, reported as such). Hypotheses that restrict a result are in the *scope* column; a theorem with no scope note holds as stated in Part IV.

Refinement and client stability (Core/Interface, Core/Decl, Core/Satisfaction, Core/Env)

name	kind	states	scope
InterfaceRefines_iff_semantic	T	the syntactic refinement order is complete for abstract evidence	—
DeclRefinesStar_iff	T	the lifecycle closure is the structural order plus well-formedness of the target	—
naive_breaks_wellformedness	C	strengthening a realized interface without re-verification breaks well-formedness	—
local_refinement_preserves_global_typing	T	refining one declaration preserves every typing judgment	the structural order only
local_refinement_preserves_global_wf	T	refining one declaration preserves global well-formedness	Evidence.Monotone
badEv_not_mono	C	non-monotone evidence is destroyed by a valid realization step	—
EnvRefines_update	T	an update at an identity is an environment refinement	—
HasType.mono_env, HasType.mono_concept, HasType.mono_grant	T	typing is monotone in the three environments	—

Typing, unfolding and dependency (Core/Typing, Core/Dependency)

name	kind	states	scope
infer_sound, infer_complete, HasType.unique	T	inference is sound, complete and unique	—
constructs_granted	T	a well-typed term constructs a concept only where granted	—
Unfolds.exists_of_acyclic, Unfolds.not_of_cyclic	T	unfolding exists iff the reference graph is acyclic	delay-free fragment
Acyclic.not_cyclic	T	the rank witness excludes cycles	—
unfolds_preserves_eval	T	unfolding agrees with tick evaluation	first-order wiring designs

Concept identity, representation and the grant (Core/Base, Experiments/SemanticTypeAlternatives, Experiments/RepresentationBindingAlternatives)

name	kind	states	scope
no_semantic_value_without_declaration	T	a Sem value originates only in a declaration of concept type	—
temporal_state_preserves_semantic_identity	T	delay carries a concept's tag and never creates one	—

name	kind	states	scope
HasType.erase, erase_not_injective, baseline_is_erased_modelA	T	a well-typed semantic term is well typed at its representation; the numeric baseline is the erasure	representations are data
semantic_rename_preserves_identity, rename_under_name_identity_breaks_client	T/C	display names are not identity; name-as-identity makes renaming destructive	—
bweak_evaded_by_eta	C	the <i>semanticRole</i> metadata checker is evaded by η -expansion	—
unrestricted_representation_binding_bypasses_semantic_identity, hidden_crossing_inside_unrelated_body	C	unrestricted mk/rep admits a hidden concept crossing	—
hidden_crossing_rejected_under_grant, representation_binding_does_not_enable_hidden_semantic_mapping	T	the hidden crossing is refused under the grant of an unrelated declaration	—
representation_change_is_edit_not_refinement	C	rebinding a concept's representation breaks existing realizations	—

Quantities and dimensions (Core/Base, Experiments/DimensionAlternatives)

name	kind	states	scope
dimension_mismatch_rejected	T	length + time is ill-typed	—
counterexampleB_baseline_accepts_length_plus_time	C	erasing dimensions accepts it	—
same_dimension_does_not_imply_same_semantic_identity, explicit_semantic_mapping_uses_dimensioned_formula	T	Tilt and MotorAngle at q Angle stay distinct	—

Reactive semantics (Core/Reactive, Experiments/ReactiveAlternatives)

name	kind	states	scope
Ev.det	T	evaluation is a partial function	—
reactive_total, fundamental	T	every declaration has a value at every tick	causal, globally well formed, well-typed inputs
evalF_sound	T	the interpreter is sound for Ev	—
Ev.not_of_strictCyclic	C	a strict cycle has no value	strict cycles only
Ev.tag_provenance	T	no signature announces a concept and no input carries it $\Rightarrow$ no value carries it	—
delay_not_under_binder, sync_not_under_binder	T	memory and transport are typed only in the empty context	—
arrow_not_delayable	T	nothing of function type can be delayed or transported	—
first_tick_undefined_without_init, first_tick_nondeterministic_without_init	C	delay without an initial value	toy relations

Clock domains and transport (Core/Clock, Experiments/ClockAlternatives)

name	kind	states	scope
MEV.det, multi_domain_total, mfundamental	T	multi-domain evaluation is deterministic and total	causal, globally well formed
delay_is_sync_own	T	delay is sync at the own domain	—

name	kind	states	scope
single_domain_embedding	T	under the always-active schedule MEv is Ev	—
scheduling_order_observable	C	same-tick cross-domain visibility makes the scheduler order observable	—
equal_rate_not_same_domain	C	a clone of a domain with the same schedule is a different domain	—
clocked_type_forces_polymorphism	C	a clock-indexed type forces polymorphism on every pure mapping	—
opt_loses_multiplicity_under_sync	C	sync as an event transport loses multiplicity	—
buffer_from_log_and_cursor	T	the window equals log-now minus log-at-previous-activation	—
policies_lose_information	T	latest, count, count+latest identify distinct windows	—

Lists, the buffer and capacity (Core/ListData, Surface/Buffer, Validation/Capacity, Experiments/BufferAlternatives)

name	kind	states	scope
buffer_window_correspondence (Theorem M)	T	the five declarations compute the window at every tick	src an input
buffer_elaboration_well_typed, buffer_elaboration_well_clocked	T	the elaboration types and clocks	—
buffer_lossless, window_to_list_preserves_order, window_to_list_preserves_multiplicity	T	the list summary is injective, ordered, multiplicity-preserving	—
bounded_summary_not_lossless, lossless_iff_injective	T	any summary of the newest k entries is lossy; lossless $\Leftrightarrow$ injective	—
latest_not_lossless, count_not_lossless, sum_not_lossless, modelD_not_lossless	C	models A–D lose windows	—
requiredCapacity_sufficient, periodic_window_bound, periodic_capacity_sufficient	T	the least sufficient capacity; one period suffices for periodic schedules	—
sufficient_capacity_preserves, bounded_buffer_agrees	T	under sufficient capacity every overflow policy is the identity	—
negE, negF, buffer_trace, transport_trace, latest_transport_loses	X	executed windows, overflow, clone-domain refusal, the component transport	—

Products, the recursor, equality and order (Core/Base, Surface/Poly, Experiments/PolyAlternatives)

name	kind	states	scope
Ty.prod_data	T	a pair is data iff both parts are	—
church_fst_rank, church_pair_prenex_one_projection	C	Church pairs need rank 2	toy System F
fold_total, mfold_total	T	the recursor is total on finite lists	—
Value.beq_iff	T	structural equality is equality on first-order values	—

name	kind	states	scope
Cap.eq_iff_data, Cap.ord_data, Cap.ord_not_data_converse	T	Data $\Rightarrow$ Eq; Ord $\Rightarrow$ Data; not conversely	—
lt_rejected, min_mode_rejected, lt_only_on_quantities	C	order on modes, pairs, lists, options, booleans is refused	—
minBy_recovers_min	T	the comparator escape hatch loses nothing	—

The equation library and rank-1 polymorphism (Surface/Poly, Surface/Stdlib, Surface/Generic, Experiments/EquationExamples)

name	kind	states	scope
matchTy_sound, matchTy_ complete	T	one-way matching is sound and complete	—
Scheme.instantiate_sound	T	an instantiated scheme checks its capabilities	—
instances_are_monomorphic	T	three uses of min are three kernel terms	—
applyBoth_rank, existential_ rank, applyBoth_replacement	C	higher-rank uses have rank-1 replacements	toy System F
HasType.comb_irrelevant, lib_ eval_context_free, lib_clocked, Comb.noConstruct, lib_ expansion	T	a combinator's typing, value and clock are context-free; it constructs nothing; inlining is sound	—
fold_spec, any_spec, all_spec, contains_spec, map_spec, filter_spec, min_spec, max_spec, clamp_spec, inRange_spec	T	each entry computes the mathematical function	the predicate value implements a Boolean function
forall_in_list, exists_in_list, oneOf_mem, oneOf_dup_ irrelevant	T	finite quantification and set membership are folds	—
generic_preserves_identity, generic_preserves_dimension, pair_projections_keep_ concepts, map_keeps_concepts, eq_across_concepts_rejected	T	nominality and dimensions survive generics	—
exA ... exM	X	the thirteen required equation cases	—

Natural surface (Surface/Natural, Experiments/NaturalExamples)

name	kind	states	scope
elab_local_nearest, elab_ shadow, elab_unbound, elab_core	T	scoping of binder locals	—
desugar_rename, alpha	T	alpha-equivalence of the elaboration	—
desugar_constructs	T	no mk is introduced	—
binder_local_type, range_ bounds_forced	T	the local's type is the element type; range bounds at the nominal type	—
binder_all_eval, binder_any_ eval, binder_map_eval, binder_ filter_eval, natural_forall, natural_exists, range_eval	T	evaluation is the library's	—
binder_clock, range_clock	T	clocks are the operands'	—

Units, coordinates and charts (Surface/Units, Surface/Composer, Surface/Affine, Surface/Rational, Surface/Charts)

name	kind	states	scope
inUnitE_typed, inUnitE_safe, withUnitE_typed, withUnitE_is_quantity	T	coordinate extraction and quantity construction type as stated	—
convert_eq, convert_trans, convert_self, inUnit_mismatch	T	conversion is the composition; transitive; identity; never coerces	—
inUnit_withUnit, withUnit_inUnit, inUnit_withUnit_nat, withUnit_inUnit_nat	T	round trips, exact over the symbolic group; the Nat laws	the second Nat law needs the scale to divide the magnitude
unitOps_no_construction, nominal_distinct	T	unit operations construct nothing; concepts stay distinct under round trips	—
unitsFor_sound, unitsFor_complete	T	the registry's units for a dimension	—
presentation_irrelevant_typing, presentation_irrelevant_eval, presentation_changes_display, display_may_identify_distinct	D/T	presentation is semantically irrelevant	—
normalizations_agree	T	tilt / 90 deg and inUnit(tilt, deg) / 90 agree	—
solve_sound, solve_complete, candidates_sound, candidates_complete, slot_sound, refCandidates_sound	T	local dimension inference and candidate units	the + - × ÷ fragment with holes
celsius_not_linear, sum_of_points_is_not_a_point, sum_well_typed, delta_is_linear, delta_candidates_need_sort	C/T	the linear model cannot represent °C; absolute-value arithmetic is the problem	—
affLitE_typed, affInUnitE_typed, affine_roundtrip_ev, scales_agree	T	affine literals and coordinates elaborate exactly	K/180 basis
chart_left_inverse, chart_right_inverse, convert_is_affine, convert_identity, convert_compose, convert_inverse	T	the chart and groupoid laws	any field
difference_map, difference_offset_cancels, difference_converts_linearly, linear_part_identity, linear_part_compose, not_additive_of_offset	T	differences carry the linear part; offsets are not additive	any field
unit_erasure_preserves_conversion_structure, display_switch_preserves_quantity, coordinate_edit_changes_quantity, coordinate_needs_chart	T	erasure keeps conversion; the coordinate is chartless	—
sort_orthogonal_to_conversion, conversion_orthogonal_to_sort	D	the point/difference sort and conversion are independent	—
CtoF_closed, FtoC_closed, exH, exI	X	Celsius/Fahrenheit, the ADC calibration, the encoder offset, exactly	—

Physical outputs (Core/Output, Experiments/OutputAlternatives)

name	kind	states	scope
multiple_direct_drivers_rejected	C	two drivers of one sink violate <code>SingleDriver</code>	—
single_driver_output_deterministic	T	with one driver the physical output is unique where it exists	—
first_output_binding_is_monotone	T	binding an unbound declaration to an undriven sink is a refinement	—
hidden_arbitration_observable	C	first-wins, last-wins and maximum give three outputs	—
driver_is_unit_domain	T	a driver of a concept-accepting sink is unit-domain	—

Unit domain and the source role (Surface/UnitDomain, Experiments/UnitDomainExamples)

name	kind	states	scope
elim_canonical, encode_decode, decode_encode, canonicalOfKernel_encode, encode_injective, canonical_injective	T	the normalization is inverse to encoding	signatures whose output is not an arrow
zero_input_obligation	D	the realization obligation of $() \rightarrow B$ is $\vdash e : B$	—
zero_input_memory, lams_typed	T	memory is legal in a zero-input declaration; a formula body is a realization under n binders	—
refForms_agree, HasType.refForms, Clocked.refForms	D	$f, f(), f(())$ are one term	—
Ev.declRef_env_irrelevant, MEv.declRef_env_irrelevant, same_tick_same_value	T	a reference's value is independent of the local environment; two readings in one tick agree	—
source_value, source_reads, resolved_not_source, SimulationInput.value	T	a Source's value is the input's; a resolved $() \rightarrow A$ never reads it	—
transport_needs_unit_domain, delay_needs_unit_domain	T	only a unit-domain declaration can be transported or remembered	—
unit_codomain_collapse, consumers_indistinguishable, eval_independent_of_drives	T	$A \rightarrow ()$ cannot name a receiver	denotational
homUnit	D	$(1 \rightarrow \beta) \equiv \beta$	—

Source provision (Surface/Provision, Experiments/ProvisionExamples)

name	kind	states	scope
realization_checked_under_own_grant, grant_of_sem, channel_constructs_nothing	T	the grant boundary of the profile	—
Channel.WF_refFree, pure_iff_delayFree_of_wf	T	typing in the empty design forbids references; purity is delay-freedom	—
exD	X	a typed impure transducer is tick-dependent	—
provision_envRefines, provision_tyView_eq, realizeAt_typed	T	provision is an environment refinement	WF

name	kind	states	scope
provision_wf	T	the provisioned design is globally well formed	monotone evidence; evidence for the targets' commitments
provision_causal, provision_wellClocked	T	causality and clocks are preserved, $\kappa \ r = \kappa \ s$	NoMention $\Delta \ r$
simulate	T	the design-simulation lemma behind transparency	closures avoid r
provision_transparent, provision_transparent_typed, provision_decl_transparent, provision_physicalOutput	T	the design cannot tell the provisioned environment from the abstract one	RawInput, NoMention, $r \notin e.\text{refs}$
provision_abstracts	T	every provisioned trace is an abstract trace	as above
provision_exact, JointSection.one	T	trace equality under a joint section	JointSection
exE, sat_never_451	X	the strict case: a saturating ADC	—
no_joint_witness, exF	C/X	shared raw readings: pointwise surjectivity is not enough	—
provision_not_reapplicable, provision_idem_total, provision_reprovision_not_refinement, provision_source_role	T	not re-applicable; the totalized function is idempotent only trivially	—
provision_comm, provision_perm	T	independent provisions commute; the assignment is a set	independence; nodup keys
Transduces.mev, MEV.of_ev_pure, Ev.pure	T	a pure term's canonical evaluation transfers to any design, input, domain and tick	—

Output realization (Surface/OutputRealization, Experiments/OutputRealizationExamples)

name	kind	states	scope
encoder_constructs_nothing, encoder_decl_no_grant, grant_of_semFree_data, Encoder.WF_refFree	T	the nominality boundary of the encoder: observes through rep, constructs nothing, no grant	Encoder.WF
retarget_breaks_driveWF	T	Model A refuted: retargeting o.accepts breaks the existing edge	raw sem-free
RawCommand.det	T	the machine command is a function of the tick	SingleDriver
behavior_unchanged, lower_transparent, lower_decl_transparent, lower_physicalOutput_unchanged	T	the behavior is literally unchanged off e; every pre-existing term and every logical output evaluates alike under the same input	NoMention, inputs avoid e
lower_envRefines, lower_outputEnv_extends, lower_singleDriver, lower_driveWF, lower_completeOutputs, encoderBody_typed, lower_wf, lower_causal, lower_wellClocked	T	the lowered design is a refinement: well typed, well formed, causal, well clocked, single-driver, complete for $p : \text{req}$	WF, DriveWF, SingleDriver
lower_correspondence, encoder_value, output_value_typed	T	raw trace = transfer $\circ$ abstract trace, tick by tick in the output's clock; typed values from totality	OutputTyped; $\emptyset$.WF, Causal, GlobalWF

name	kind	states	scope
two_realizations_same_behavior, lower_comm	T	the same evaluation of every pre-existing term in two lowered designs (the formal side of platform independence); independent realizations commute exactly	both e fresh, inputs avoid them; distinct e, p
lowered_interfaces, admissible_satisfiable, admissible_needs_wf	T	the encoder's canonical type $() \rightarrow \text{raw}$; admissibility = typing + fit + a solvable board, and the typing cannot be dropped	—
exJ	C/X	an encoder that fits and allocates but is ill-typed is not admissible	—
exA_gpio ... exD_hbridge, exB_quantized	X	GPIO, PWM, quantizing PWM, servo, H-bridge	—
exEFG, exI	C/X	misfits; the constructing and the impure encoder; Model A; the implicit clock crossing	—
exH, exH_structure, exH_admissible	X	one light by PWM and by I ² C: one behavior trace, two commands; both admissible on the Nano	—

The adapter boundary and the device clock (Surface/Adapter, Surface/DeviceClock, Experiments/AdapterExamples)

name	kind	states	scope
AdapterOp.det, adapter_downstream, two_policies_same_commands	T	the operation is a function of the tick; policies do not reach the command	SingleDriver
adapter_of_sink, sink_of_adapter	T	the operation is the policy's reading of the lowered sink's value	Phase 14's hypotheses; the tick active
Line.det, line_holds_on_refusal, line_holds_when_inactive, line_last_accepted, line_value_accepted	T	reject-and-hold as a fold; the line never carries a refused command	SingleDriver
lowerSync_transparent, lowerSync_envRefines, lowerSync_singleDriver, lowerSync_driveWF, syncBody_typed, lowerSync_wf, lowerSync_causal, lowerSync_wellClocked	T	the device-clocked lowering preserves everything; no new instantaneous edge	WF, InitRep.WF, NoMention
lowerSync_correspondence	T	the sink carries the command sampled strictly before, or the initial representation's transfer	Phase 14's hypotheses + TyVal rep i.value
exA_policies, exA_set, exB_refusal, exC_held	X	one command, three policies; hold on refusal and on an inactive tick	—
exD_device_clock, exD_initial, exD_structure	X	sampling at prevAct; the initial representation; the structural theorems	—
exE_slew	X	slew-rate limiting as behaviour state upstream with the same pure encoder	—

Communication as state, catalogue profiles and assign (Surface/Assignment, Experiments/CommunicationExamples)

name	kind	states	scope
admissible_iff, WF.contract, WF.one_of_contract	T	admissibility is the contract and feasibility; Provision.WF asks the input contract per target	—
OutputEntry.realization_of_profile, assign_indistinguishable, assignSource_origin_irrelevant	T	the origin of a profile is unread: one profile, one lowered design, one command relation, one operation relation	$en_1.profile = en_2.profile$
assignOutput_behavior_unchanged, assignOutput_transparent, assignOutput_checked, assignSource_transparent	T	assign is the lowering / the provision and edits no behavior; the assigned design is accepted by the unchanged judgments	Phase 14's / Phase 13's hypotheses
realizable_mono, provisionable_mono	T	a larger catalogue realizes and provisions more	entry inclusion
induced_avoids, induced_congr, two_providers_same_behavior, two_providers_same_outputs	T	Source-side non-interference: same semantic trace (SameTrace), same evaluation and same outputs	$Provision.WF \times 2, NoMention \times 2, RawInput \times 2$
paired_commands_of_one_value	T	two realizations of one output specify two transfers of one value	$SingleDriver, R_1.o = R_2.o$
exA_multiplicity, exB_ordering, exC_same_tick, exD_cancellation, exE_bounded, exF_identity	X	the work queue: multiplicity, order, same-tick batches, cancellation by semantic id, the bounded queue, dedup by id	—
exG_freshness, exH_settling, exI_fault_latch, exJ_ack, exQM_composed	X	the motion state: age, freshness, timeout, settling, the fault latch, the acknowledgement; the composed controller	—
exK_cross_clock, exL_theorem, exL_executed	X	host submissions through the window, flattened; two providers, one Source trace	—
exN_pair, exN_two_motors, exP_origin, exQ_contract_not_feasible	X	the paired axis as one pair command; builtin vs packaged; the same contract on two boards	—

The provider's contract and the output window (Surface/Provider, Surface/OutputWindow, Experiments/ProviderExamples, Experiments/OutputWindowExamples)

name	kind	states	scope
dedup_sublist, provide_items_sublist, dedup_of_fresh	T	what is kept keeps the transport's order and multiplicity; distinct identities are all delivered	—
mem_dedup_seen, mem_seenAfter, dedup_retry, run_retry, batch_retry	T	the remembered identities; a retransmission changes nothing at any tick	Retry
provide_length_le, batch_length_le, provide_overflow_iff, provide_items_of_no_overflow	T	the bound, the exact flag, nothing dropped without it	—

name	kind	states	scope
mergeBySource_interleaving, perSource_of_interleaving	T	source order is an interleaving; per-source reading is merge-insensitive	sources tagged by index
batch_value_tyVal, rawInput_ rawInput, sameBatches_ sameTrace, retry_invisible	T	the provider's reading is a Phase-13 raw input; equal batches are the same trace; a retransmission is invisible	typed payloads
update_transparent, NoMention. update	T	one fresh realized declaration is invisible to terms avoiding it	NoMention
lowerWindow_transparent, lowerWindow_decl_transparent, lowerWindow_encoder_at	T	the window design evaluates the abstract design unchanged; the encoder carries the trace's transfer	Fresh, WF, NoMention ×6
lowerWindow_correspondence, lowerWindow_batch_rawCommands, lowerWindow_batch_unique, lowerWindow_order_ multiplicity, lowerWindow_ bounded	T	the sink carries the raw commands of the window, in order, with multiplicity, bounded by capacity	RepTrace, SingleDriver, CapacitySufficient
lowerWindow_envRefines, lowerWindow_singleDriver, lowerWindow_driveWF, lowerWindow_wellClocked, lowerWindow_causal, lowerWindow_wf	T	the window design is a well-formed, causal, well-clocked, single-driver refinement	Phase 14's hypotheses
batchOps_length, lineAfterBatch_accepted, lineAfterBatch_refused, paired_batches_of_one_window	T	the adapter's batch; the paired axis batched	SingleDriver, one output
exA_occurrences, exB_order, exB_all_ticks, exC_merge, exD_ bound	X	occurrences, retries, order, two raw sources, the bound	—
exE_ingress, exE_theorem, exE_ overflow, exF_feedback, exG_ two_axes	X	host ingress into the queue design through the provider; feedback from two raw readings; the motion state on two axes	—
exA_window_vs_sample, exB_ multiplicity_order, exC_ capacity, exD_batch, exE_paired, exF_correspondence, exF_ structure, exF_transparent	X	window vs sample; capacity decided; the adapter batch; the paired axis; the theorems instantiated	—

The Source-side boundary (Surface/SourceBoundary, Experiments/SourceBoundaryExamples)

name	kind	states	scope
MEV.congr_at, MEV.declRef_env, MEV.declRef_nil	T	designs agreeing off one declaration whose value trace agrees evaluate every term alike	—
machine_upstream, Machine.run_ typed, Machine.run_congr	T	the upstream declaration computes the machine's run; the run is a function of the stream alone	single domain
below_source_trace, upstream_ source_trace, provider_state_ movable, stateful_providers_ same_trace	T	the same Source trace below or above the reading; every evaluation transfers; equal output streams, one trace	Provision.WF, typed streams

name	kind	states	scope
provisionSync_target, sampled_mev, provisionSync_transparent, provisionSync_causal, provisionSync_wellClocked, syncRealizeAt_typed, provisionSync_envRefines, provisionSync_wf	T	the sampled Source-side device clock with an explicit initial value: correspondence, transparency, structure	the Source in a domain; typed reading and initial value
provisionWindow_target	T	the occurrence-like Source over Phase 9a's window of the raw reading	Realized, typed reading
realizeAt_only_transfers, discharge_under, discharge_static, discharge_checked, discharge_trusted	T	commitment discharge at three evidence levels	RangeSoundUnder
checked_items_ok, checked_typed, checked_refused_iff, filter_length_lt_iff	T	the checking provider: delivered items validated and typed; the flag exact	ok implying typing
Value.beq_sound, Value.beqList_sound, computes_of_bool	T	structural equality is sound; a passed check at bool is computes	—
assignSource_profile_only	T	the richer profile does not change the assignment	equal Phase-13 profiles
exA_filter, exA_filter_run, exA_filter_theorem, exA_debounce_param, exA_quadrature, exA_hysteresis_param, exA_latch_reads_design	X	provider state below or above; the parameters in the trace; the latch that reads the design	—
exB_sampled, exB_structure, exB_unavailable, exB_window_edges	X	the sampled and the unavailable initial value; edges windowed without loss	—
exC_static, exC_trusted, exC_checked, exC_provider_checks	X	one property at three evidence levels	—
exD_ofTerm, exD_bool, exD_bool_computes, exE_checked, exF_two_providers	X	computes from the term and decided; refused readings into an optional Source; the motion state under two providers	—

Hardware validation (Validation/Hardware, Experiments/HardwareAlternatives)

name	kind	states	scope
solve_sound, solve_complete, satisfiable_iff_solve	T	the solver is sound and complete; feasibility is decidable	finite instances, unary and binary constraints
Hardware.Extends preservation	T	extending a target preserves valid assignments	—
motor_control_sat_on_nano, seven_pwm_unsat_on_nano, seven_pwm_sat_on_big, multifunction_overlap_unsat, seven_pwm_design_semantically_valid	X	the Nano case study and its two counterexamples	—
diagnose	D	a first dead end under greedy placement	not a minimal core

Behavior systems and groups (Behavior/*, Experiments/BehaviorAlternatives, Experiments/GroupAlternatives)

name	kind	states	scope
HasType.rename, Satisfies.rename, Clocked.rename	T	the judgments are equivariant under renaming	—
inst_decl_disjoint (Theorem A)	T	instances never share identities	—
binding_satisfies (Theorem C)	T	a binding is an ordinary realization the checker verifies	—
flatten_WF, flatten_causal, flatten_wellClocked, flatten_singleDriver, open_port_stays_open (Theorems D–I)	T	the flattening is a well-formed design	ComposeWF; InstAcyclic for causality; Monotone, Equivariant, PortSound evidence
eval_flat_to_inst, eval_inst_to_flat, modular_iff_flat (Theorem J)	T	modular and flat evaluation agree	restricted: single domain, wiring designs, closure-free inputs, direct/constant bindings
substitute_composeWF	T	replacing an instance by a refining component preserves composition	—
group_is_identity_on_design and Theorems A–G	D	every group operation is the identity on the design	—
mem_crossIn, mem_crossOut, internal_not_crossIn, internal_only_not_crossOut, socket_no_fanout (Theorems H–L)	T	boundaries are projections; sockets add no dependency	finite enumeration
restrict_realizes, system_composeWF, flat_WF, flat_causal, private_unobservable, provided_iff (Theorems M–Q)	T	extraction yields a well-formed system	Evidence.InterfaceLocal
orig_iff_flat (Theorem R)	T	extraction preserves behavior	restricted: single-domain wiring designs, no transported bindings

Appendix C — Decision index

Every formal design decision by its stable identifier, with the Part of this document that discusses it, the Lean files that carry its evidence, and the production record it supports, audits or bears on. The records themselves — choice, alternatives rejected, formal reason — are docs/decisions/ in the formal repository; the map from the retired ledger numbers is Appendix E. Status is *accepted* for every record except FVD-0137, superseded by FVD-0139 in the Phase 14 hardening (fit-and-allocate is not admissibility); FVD-0106 carries a Phase-10b amendment, FVD-0050 a Phase-14 amendment (the *logical output* reading), FVD-0136 a hardening amendment (a decision, not a theorem), and FVD-0098 records the partial reversal of FVD-0090’s lt generalisation.

id	decision	status	main section	formal evidence	production correspondence
FVD-0001	Commitments are atomic labels; evidence is abstract	accepted	§IV.1	Phase 0: Core/Interface, Core/Satisfaction	ADR-0010 (supports)
FVD-0002	InterfaceRefines freezes the type and grows commitments; not logical implication	accepted	§IV.1	Phase 0: Core/Interface, Core/Satisfaction	ADR-0010 (supports)
FVD-0003	strengthen on a realized declaration carries a re-verification premise	accepted	§IV.1	Phase 0: Core/Interface, Core/Satisfaction	ADR-0010 (supports)
FVD-0004	DeclInterface. commitments is a List, not a Finset	accepted	§IV.1	Phase 0: Core/Interface, Core/Satisfaction	ADR-0010 (supports)
FVD-0005	References are by DeclId inside Expr (declRef), not a separate DesignExpr	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0006	The declRef typing rule reads Δ.tyView only	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0007	Realization is write-once; “detach” is not a refinement step	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0009 (supports)
FVD-0008	The structural order (DeclLeq, EnvRefines) is separated from the invariant (GlobalWF)	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0009	Evidence takes the environment as an argument	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0010	Evidence.Monotone is a kernel-imposed constraint on the validation layer	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0011	Semantics at Phase 1 is unfolding to a reference-free term	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0012	Acyclicity is witnessed by a rank function	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0010 (supports)
FVD-0013	Phase 0's Artifact (bare list of ids) removed	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	FV-only
FVD-0014	Display names are not in the kernel	accepted	§IV.1	Phase 1: Core/Base, Core/Decl, Core/Typing, Core/Env, Core/Dependency	ADR-0008 (supports)
FVD-0015	Kernel ontology migrated from holes to declarations	accepted	§IV.1	Phase M: Core/* (the rename)	ADR-0010 (supports)
FVD-0016	Monotone refinement is distinct from arbitrary editing	accepted	§IV.1	Phase M: Core/* (the rename)	ADR-0009 (supports)
FVD-0017	Commitments and validation obligations share PropertyId for now	accepted	§IV.1	Phase M: Core/* (the rename)	ISS-0003 (bears-on)
FVD-0018	Evidence.Monotone is a stability condition, not a definition of validity	accepted	§IV.1	Phase M: Core/* (the rename)	ISS-0003 (bears-on)
FVD-0019	Concept identity lives in the type: Ty.sem : ConceptId → Ty	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0020	ConceptId is independent of DeclId and of display names	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0021	No introduction/elimination forms for concept types in Phase 2	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0022	Explicit semantic mappings are ordinary declarations	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0023	Concept identity change is an edit	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0025	Future constraint for Phase 3: representation binding must not defeat concept identity	accepted	§IV.2	Phase 2: Core/Base, Experiments/SemanticTypeAlternatives	ADR-0013 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0024	Canonical closed inhabitants replaced by unresolved declarations	accepted	§IV.2	Phase 2: Core/Base, Experiments/ SemanticTypeAlternatives	ADR-0013 (supports)
FVD-0026	Unrestricted symmetric mk/rep rejected	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0027	Representation types are sem-free (ConceptEnv.WF)	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0028	Observation is unrestricted; construction is licensed by the realized declaration's signature	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0029	Representation binding is a separate, write-once concept environment $\emptyset$	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0030	Unfolding preserves typing under <code>Grant.all</code> , not under the client grant	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0031	Dimensions in <code>Ty</code> as <code>q d</code> ; algebra in <code>Prim.ty</code> ; no dimension rule	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0011 (supports), ADR-0013 (supports)
FVD-0032	Units are surface: elaborated to scaled dimensioned literals	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0028 (supports)
FVD-0033	Concept identity is not indexed by dimension	accepted	§IV.2	Phase 3: Core/Decl, Core/Typing, Experiments/ RepresentationBindingAlternatives, Experiments/ DimensionAlternatives	ADR-0013 (supports)
FVD-0034	One temporal primitive: <code>delay init e</code>	accepted	§IV.5	Phase 4: Core/Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0035	delay is data-typed and top-level (empty context)	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0036	Signal is not a type; Event is opt	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0037	Causality replaces blanket acyclicity: Causal on InstDependsOn	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0038	Explicit initial value on every delay	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0039	State has no identity; state is structural	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0040	Representation types are data (ConceptEnv.WF strengthened)	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0041	Temporal changes are realization edits	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0042	The reactive semantics is a relation, not yet a machine	accepted	§IV.5	Phase 4: Core/ Reactive, Experiments/ ReactiveAlternatives	ADR-0004 (supports), ADR-0016 (supports)
FVD-0043	Time is one global tick with a schedule; no rates, no timestamps in the kernel	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ ClockAlternatives	ADR-0004 (supports)
FVD-0044	Nominal ClockId, stored per declaration in ClockEnv K; none = domain-agnostic pure mapping	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ ClockAlternatives	ADR-0004 (supports)
FVD-0045	One transport primitive sync src init e, reading strictly before	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ ClockAlternatives	ADR-0004 (supports)
FVD-0046	The clock is interface data held in a projection, not a record field	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ ClockAlternatives	ADR-0004 (supports)
FVD-0047	Rates, drift, jitter, latency, buffer capacity, value age are validation	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ ClockAlternatives	ADR-0004 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0048	Event transport = window read; buffering derived, Event still not a primitive	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ClockAlternatives	ISS-0001 (bears-on)
FVD-0049	The logical relation is generic in the application relation	accepted	§IV.5	Phase 5: Core/Clock, Experiments/ClockAlternatives	FV-only
FVD-0050	Physical sinks have nominal identity (OutputId), separate from ConceptId and DeclId	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0051	A drive edge is a per-declaration write-once projection β , checked by type and clock equality	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0052	Single-driver is a global invariant; completeness is the executable condition	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0053	No runtime arbitration, no implicit priority, no merge policy	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0054	Effect rows and action values rejected for this kernel	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0055	First binding is a refinement; rebinding is an edit; a second driver is invalid	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0056	Sinks are terminal	accepted	§IV.7	Phase 6: Core/Output, Experiments/OutputAlternatives	ADR-0005 (supports)
FVD-0057	Hardware feasibility is a validation layer over Design $\times$ Target, not typing	accepted	§IV.8	Phase 7: Validation/Hardware, Experiments/HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)
FVD-0058	Resources carry capabilities and per-capability units; sharing is a per-capability policy	accepted	§IV.8	Phase 7: Validation/Hardware, Experiments/HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0059	Requirements are independent variables with nominal RequirementId, optional fixed resource, optional unit relation	accepted	§IV.8	Phase 7: Validation/ Hardware, Experiments/ HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)
FVD-0060	Validity is unary support plus pairwise compatibility; the solver is exhaustive DFS, proved sound and complete	accepted	§IV.8	Phase 7: Validation/ Hardware, Experiments/ HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)
FVD-0061	Hardware extension is monotone; requirement extension, strengthening, fixing, and resource removal are revalidation triggers	accepted	§IV.8	Phase 7: Validation/ Hardware, Experiments/ HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)
FVD-0062	Deployment feasibility is environment-sensitive evidence, not Evidence. Monotone	accepted	§IV.8	Phase 7: Validation/ Hardware, Experiments/ HardwareAlternatives	ADR-0006 (supports), ADR-0015 (supports)
FVD-0063	Numeric electrical/timing constraints deferred	accepted	§IV.8	Phase 7: Validation/ Hardware, Experiments/ HardwareAlternatives	FV-only
FVD-0064	Behaviour components are surface objects; the kernel is unchanged	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0065	A port is a template declaration by identity, with its public interface and clock	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0066	Instantiation renames every identity the template owns; concepts and sinks are partitioned into internal (fresh) and global (shared)	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0067	Binding is a Phase-1 realization step	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0068	Composition well-formedness is stated on interfaces, never on bodies	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0069	Causality across instances is a validation condition on the inter-instance direct-binding graph	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0070	Clock parameters are nominal variables substituted by κ at instantiation; rates never enter	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0071	Evidence must be equivariant and port-sound	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0072	Hierarchy is packaging, not a tree constructor	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports), ADR-0022 (supports)
FVD-0073	Theorem J is proved on the single-domain wiring fragment with direct/constant bindings	accepted	§IV.6	Phase 8a: Behavior/ *, Experiments/ BehaviorAlternatives	ADR-0021 (supports)
FVD-0074	A behaviour group is authoring metadata; eraseGroups is a projection	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0075	Group operations are semantic no-ops, not refinements or edits	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0076	Aggregate sockets are projections of DependsOn	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0077	Boundary inference: required = crossing-in, provided = crossing-out, private = the rest without a sink, clocks = all clocks	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0078	Extraction = two restrictions of the design reconnected by Phase-8a bindings	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0079	Extraction causality is proved by subdividing the original graph, not by <code>InstAcyclic</code>	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0080	Template realization needs interface-local evidence	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0081	Identity: templates keep original identities; instances are fresh; the group id is never a component id	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (supports)
FVD-0082	Nested groups are a relation on the flat group list	accepted	§IV.6	Phase 8b: Behavior/ Group, Boundary, Extract, ExtractPreservation	ADR-0019 (bears-on), ISS-0007 (bears-on)
FVD-0083	<code>Ty.list τ</code> is a kernel data type; the object-language buffer needs it and nothing else	accepted	§IV.5	Phase 9a: Core/ ListData, Surface/ Buffer, Validation/ Capacity	ADR-0024 (supports)
FVD-0084	Six list operators, registered through <code>Prim.ty: nil, cons, length, take, reverse, head</code>	accepted	§IV.5	Phase 9a: Core/ ListData, Surface/ Buffer, Validation/ Capacity	ADR-0024 (supports)
FVD-0085	The buffer is a surface elaboration into five declarations over <code>delay/sync</code>	accepted	§IV.5	Phase 9a: Core/ ListData, Surface/ Buffer, Validation/ Capacity	ISS-0001 (bears-on)
FVD-0086	Capacity is validation; overflow policies are explicit; only rejecting the deployment preserves semantics	accepted	§IV.5	Phase 9a: Core/ ListData, Surface/ Buffer, Validation/ Capacity	ADR-0027 (supports)
FVD-0087	Buffered transport is a Phase-8a binding choice, not a transport kind	accepted	§IV.5	Phase 9a: Core/ ListData, Surface/ Buffer, Validation/ Capacity	ADR-0021 (supports)
FVD-0088	Products enter the kernel as value composition: <code>prod, pair, fst, snd</code>	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0089	The list recursor fold is a term former, not a registered operator	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0090	eq at every data type, with the data proof in the syntax	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0091	toList : opt $\tau \rightarrow$ list τ and drop are registered operators	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0092	Rank-1 polymorphism is definitional: families instantiated by matching; no type variable in the kernel	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0093	Constraints: the closed vocabulary {Data}; no user-defined classes	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0094	The equation library is a set of combinators, inlined at use sites	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0095	Sets, intervals, records, predicates, finite quantifiers are surface	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0096	Sums are encoded; a kernel sum is deferred	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ISS-0005 (bears-on)
FVD-0097	Existentials are not needed: hiding is Phase-8a instantiation	accepted	§IV.3	Phase 9b: Surface/ Poly, Stdlib, Generic, Experiments/ EquationExamples	ADR-0025 (supports)
FVD-0098	Ordering is a quantity comparison; the kernel has no structural order	accepted	§IV.3	Phase 9c: Surface/ Poly (Cap)	ADR-0025 (supports), ADR-0026 (supports)
FVD-0099	The surface capability vocabulary is {Data, Eq, Ord}; Eq = Data today; Ord is by declaration	accepted	§IV.3	Phase 9c: Surface/ Poly (Cap)	ADR-0025 (supports), ADR-0026 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0100	Ordered library entries take Ordered evidence; comparators recover them	accepted	§IV.3	Phase 9c: Surface/ Poly (Cap)	ADR-0025 (supports), ADR-0026 (supports)
FVD-0101	Units remain entirely surface; coordinate extraction and quantity construction are elaborated quantity arithmetic	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	ADR-0028 (supports)
FVD-0102	Unit semantics are stated exactly over an abstract scalar domain; the kernel's Nat and production's floats are models	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	ADR-0028 (supports)
FVD-0103	A unit is an identity with a dimension and a scale; spelling is presentation	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	ADR-0028 (supports)
FVD-0104	The Formula Composer's formal basis is typed holes with local bidirectional dimension inference — no unification	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	ADR-0028 (supports)
FVD-0105	Preferred display units are presentation, not design	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	FV-only
FVD-0106	Affine conversion is complete as coordinate-change semantics; point/delta is optional physical-arithmetic validation	accepted	§IV.4	Phase 10: Surface/ Units, Composer, Affine	ISS-0004 (bears-on)
FVD-0107	Unit coordinates are an erasure that preserves the affine coordinate change	accepted	§IV.4	Phase 10b: Surface/ Rational, Charts	ISS-0004 (bears-on)
FVD-0108	Conversions form a groupoid of affine isomorphisms; differences carry the linear part	accepted	§IV.4	Phase 10b: Surface/ Rational, Charts	ISS-0004 (bears-on)
FVD-0109	The exact scalar domain is a choice-free rational field built in the development	accepted	§IV.4	Phase 10b: Surface/ Rational, Charts	ISS-0004 (bears-on)

id	decision	status	main section	formal evidence	production correspondence
FVD-0110	Unit conversion and sensor calibration are one affine-map abstraction	accepted	§IV.4	Phase 10b: Surface/Rational, Charts	ISS-0004 (bears-on)
FVD-0111	Binder syntax desugars to the Phase-9 library applied to a lambda; a binder local is the lambda parameter	accepted	§IV.3	Phase 11: Surface/Natural	ADR-0028 (supports)
FVD-0112	Ranges are surface nodes desugared to inRange; no interval type or value	accepted	§IV.3	Phase 11: Surface/Natural	ADR-0028 (supports)
FVD-0113	No general comprehension, no general quantifier	accepted	§IV.3	Phase 11: Surface/Natural	ADR-0028 (supports)
FVD-0114	$x \text{ ?? } d$ desugars to getD	accepted	§IV.3	Phase 11: Surface/Natural	ADR-0028 (supports)
FVD-0115	The canonical type domain(inputs) $\rightarrow$ B with domain([]) = () lives above the kernel; the kernel interface type is its normalization	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0029 (supports)
FVD-0116	() $\rightarrow$ B is realized at B in the empty context; no unit binder	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0029 (supports)
FVD-0117	$f, f(), f(())$ are one reference; a reading is not a call	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0029 (supports)
FVD-0118	The source role is a realization state, not a type shape and not a kind	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0032 (supports)
FVD-0119	A $\rightarrow$ () is not a physical sink; the drive edge is	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0029 (supports), ADR-0032 (supports)
FVD-0120	Transport and memory of a relationship require the unit domain — as corollaries	accepted	§IV.7	Phase 12: Surface/UnitDomain	ADR-0029 (supports)
FVD-0121	Provision is a construction over designs, not a kernel construct	accepted	§IV.7	Phase 13: Surface/Provision, Experiments/ProvisionExamples	PRP-0001 (audits)
FVD-0122	The profile condition is purity: tr.Pure, i.e. typed in the empty design and delay-free	accepted	§IV.7	Phase 13: Surface/Provision, Experiments/ProvisionExamples	PRP-0001 (audits)

id	decision	status	main section	formal evidence	production correspondence
FVD-0123	A channel carries the transfer function and the term that computes it	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0124	The profile is generic in the concept; the Source's signature grants construction	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0125	Shared raw reading is primitive; the singleton is its special case	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0126	Trace equality needs a joint section; deployment is in general a strict refinement	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0127	Provision is not re-applicable; "idempotent" is the wrong word	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0128	A Source's commitments are obligations on the profile	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0129	Independent provisions commute exactly; the assignment is a set	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0130	Terminology: abstract Source, provisioned Source, raw declaration; not "monomorphised"	accepted	§IV.7	Phase 13: Surface/ Provision, Experiments/ ProvisionExamples	PRP-0001 (audits)
FVD-0131	A logical output is semantic intent; its physical mechanism is deployment data, never part of OutputSpec	accepted	§IV.7	Phase 14: Surface/ OutputRealization, Experiments/ OutputRealizationExamples	ADR-0015 (supports), ADR-0036 (supports)
FVD-0132	Output realization is a lowering that adds an encoder declaration and a machine sink; the logical output is never retargeted	accepted	§IV.7	Phase 14	ADR-0015 (supports), ADR-0036 (supports)
FVD-0133	An encoder is pure, consumes the representation, and constructs nothing	accepted	§IV.7	Phase 14	ADR-0005 (supports), ADR-0036 (supports)

id	decision	status	main section	formal evidence	production correspondence
FVD-0134	The machine boundary is the RawCommand relation; no effectful R -> () term exists	accepted	§IV.7	Phase 14	ADR-0016 (supports), ADR-0036 (supports), ADR-0037 (supports)
FVD-0135	Output correspondence is directional; no injectivity, exactness or round-trip is required	accepted	§IV.7	Phase 14	ADR-0036 (supports)
FVD-0136	The singleton output realization is primitive; a shared device batches per tick or is combined upstream	accepted	§IV.7	Phase 14	ISS-0017 (bears-on)
FVD-0137	Hardware requirements are a validation judgment separate from the encoder	superseded by FVD-0139	§IV.7	Phase 14	ADR-0015 (supports), ADR-0036 (supports)
FVD-0138	The machine sink is in the output's clock; a device clock is an explicit sync; a carrier frequency is not a ClockId	accepted	§IV.7	Phase 14	ADR-0037 (supports), ISS-0017 (bears-on)
FVD-0139	Deployment admissibility is the encoder's typing, its fit and a solvable board; the narrow fit-and-allocate predicate is not admissibility	accepted (supersedes FVD-0137)	§IV.7	Phase 14 hardening	ADR-0015 (supports), ADR-0036 (supports)
the open-item audit and Phase 15	every active item re-classified against Phases 8a–14 and production 6be778b (eleven resolved, deferred or merged; FVI-0022 split); the adapter boundary as policy, operation and line; the explicit device clock through sync; no stateful realization primitive				

id	decision	status	main section	formal evidence	production correspondence
FVD-0140	The adapter boundary is a policy, an abstract sink operation and a line, below the raw command and outside the behaviour	accepted	§IV.7	Phase 15: Surface/Adapter	ADR-0037 (supports)
FVD-0141	A device clock is an explicit sync lowering into the device domain; the carrier stays configuration	accepted	§IV.7	Phase 15: Surface/DeviceClock	ADR-0037 (supports)
FVD-0142	No stateful realization primitive without a non-encodability witness	accepted	§IV.7	Phase 15: Experiments/AdapterExamples	ISS-0017 (bears-on)
Phase 16	communication as state: the auto_typer stress case encoded and executed on the unchanged kernel; catalogue profiles with an unread origin; assign as the lowering / the provision; contract separate from feasibility; the paired axis without a transaction primitive		§IV.7	Phase 16	
FVD-0143	Communication artifacts are outside BDL semantics; communication history is state over the existing kernel	accepted	§IV.7	Phase 16: Experiments/CommunicationExamples	ISS-0001, ISS-0016, ISS-0017 (bears-on)
FVD-0144	Replacement invariance is three non-interference theorems, one per boundary; the input side is two_providers_same_behavior	accepted	§IV.7	Phase 16: Surface/Assignment	ADR-0032, ADR-0036, ADR-0037 (supports)
FVD-0145	A catalogue profile is a value satisfying the existing contract; its origin is unread; packages extend realizability, not the kernel	accepted	§IV.7	Phase 16: Surface/Assignment	ISS-0016, ISS-0017 (bears-on)

id	decision	status	main section	formal evidence	production correspondence
FVD-0146	assign is deployment-only: it is the lowering or the provision with a catalogue entry	accepted	§IV.7	Phase 16: Surface/ Assignment	ISS-0016, ISS-0017 (bears-on)
FVD-0147	Semantic admissibility and deployment feasibility are separate judgments; Admissible is their conjunction	accepted	§IV.7	Phase 16: Surface/ Assignment	ADR-0036 (supports), ISS-0017 (bears-on)
FVD-0148	A paired axis is one logical output with a pair command, or two realizations that agree tick by tick; commit order is below the operation	accepted	§IV.7	Phase 16: Surface/ Assignment	ISS-0017 (bears-on)
Phase 17	the provider's occurrence contract stated and proved; the occurrence-preserving output window over the encoder; the adapter's batch as a list of operations; FVI-0029 resolved		§IV.7	Phase 17	
FVD-0149	The provider's occurrence contract: one occurrence per fresh transport identity, arrival order, a bounded batch with an observable overflow flag; the raw reading is (list raw, bool)	accepted	§IV.7	Phase 17: Surface/ Provider	ISS-0016, ISS-0001 (bears-on)
FVD-0150	Deduplication is by transport identity below the boundary, never by payload; a retransmission is erased and a repeated command is not	accepted	§IV.7	Phase 17: Surface/ Provider	ISS-0016 (bears-on)

id	decision	status	main section	formal evidence	production correspondence
FVD-0151	Several raw sources are several provisions; a merged batch is an explicit deterministic policy; no physical total order is assumed	accepted	§IV.7	Phase 17: Surface/Provider	ISS-0016 (bears-on)
FVD-0152	Occurrence-preserving realization is Phase 9a's window over the encoder declaration into the device domain with a list raw sink; the sink's type selects it	accepted	§IV.7	Phase 17: Surface/OutputWindow	ISS-0017 (bears-on)
FVD-0153	The adapter's batch is a list of Phase 15's operations and the line after it a fold; no batch or transaction primitive	accepted	§IV.7	Phase 17: Surface/OutputWindow	ISS-0017 (bears-on), ADR-0037 (supports)
Phase 18	the Source-side boundary: provider state as a machine below or above the reading; the Source-side device clock with an explicit initial value; commitment discharge at three evidence levels; computes derived and decided; out-of-type readings refused; FVI-0020 resolved		§IV.7	Phase 18	
FVD-0154	Provider state is a Mealy machine below the raw reading, movable upstream with the same Source trace; placement is decided by visibility	accepted	§IV.7	Phase 18: Surface/SourceBoundary	ISS-0016 (bears-on), PRP-0001 (audits)

id	decision	status	main section	formal evidence	production correspondence
FVD-0155	The Source-side device clock is an explicit <code>sync</code> (or window) of the raw reading with an explicit initial value — supplied or unavailable	accepted	§IV.7	Phase 18: Surface/ SourceBoundary	ISS-0016 (bears-on)
FVD-0156	A provider discharges a commitment as evidence under an assumption; static, checked and trusted levels differ in who establishes it	accepted	§IV.7	Phase 18: Surface/ SourceBoundary	ISS-0016 (bears-on), PRP-0001 (audits)
FVD-0157	computes is proof-carrying; derived from the term, decided at a finite raw type; a supplied transfer function is a claim	accepted	§IV.7	Phase 18: Surface/ SourceBoundary	PRP-0001 (audits)
FVD-0158	An out-of-type reading is refused by the checking provider and crosses as none or a flag; no exception semantics	accepted	§IV.7	Phase 18: Surface/ SourceBoundary	ISS-0016 (bears-on)
FVD-0159	A concept is a nominal type; several declarations may produce values of it; alternatives are resolved by an explicit declaration	superseded by FVD-0161	§IV.2	Phase 19: Experiments/ ProducerAlternatives	ADR-0034 (supports)
FVD-0160	Producer uniqueness per concept is not a kernel invariant; signature-uniqueness is refuted	superseded by FVD-0161	§IV.2	Phase 19: Experiments/ ProducerAlternatives	ADR-0034 (bears-on)
FVD-0161	Each concept has one producer: <code>ProducerUnique</code> as a global invariant beside <code>SingleDriver</code>	superseded by FVD-0163	§IV.2	Phase 20: Experiments/ ProducerUnique	ADR-0034 (bears-on)
FVD-0162	A concept reference elaborates to its producer	superseded by FVD-0163	§IV.2	Phase 20 (deleted in Phase 22)	ADR-0034 (bears-on)

id	decision	status	main section	formal evidence	production correspondence
FVD-0163	The design objects are Sem blocks and mapping blocks; a concept is their type template; several Sem blocks of one concept are ordinary	accepted	§IV.2	Phase 21: Surface/ Sem, Experiments/ SemExamples	ADR-0034, ADR-0032, ADR-0041 (bears-on)
FVD-0164	The concept ladder: a concept is a type named by ConceptId, a Sem block is its instance; the vocabulary fixed in the code	accepted	§IV.2	Phase 22: the rename across BDL/	ADR-0013, ADR-0034, ADR-0041 (bears-on)

Appendix D — Evidence-strength ledger

One row per claim of the conceptual chapters, with the formal evidence, the production implementation and its tests, the decisions and production records, and the known limitation. Every entry in the *formal* column is **formally proved** unless it is an executed example or a counterexample by its name; every entry in the *production* columns is **production implemented and tested** at the snapshot (081296d); *not implemented* means **formally proved, not implemented**. Theorem names are in KCN-judu/BDL_FV.

claim	formal theorem / counterexample	production implementation	production tests	decisions · production records	known limitation
refining one declaration preserves every client	local_refinement_preserves_global_typing, _wf	bdl-model::apply_edit refinement classification	edit classification tests	FVD-0011–FVD-0016; ADR-0009	arbitrary edits are outside the order
typing consults only the type view	HasType.mono_env, refFree_env_irrelevant	bdl-check	pipeline tests	FVD-0004	—
concept identity is nominal; construction needs a grant	constructs_granted, no_semantic_value_without_declaration, temporal_state_preserves_semantic_identity	elaborator rep/mk insertion; checker	semantic corpus case	FVD-0019–FVD-0028	—
dimensions reject length + time	dimension_mismatch_rejected; erasure counterexample	Prim.ty typing	diagnostics tests	FVD-0031	—
evaluation is deterministic and total on causal designs	Ev.det, reactive_total, Ev.not_of_strictCyclic	reference evaluator	delayed_cycle, property tests	FVD-0036–FVD-0042	lambda-guarded cycles conservative
one temporal primitive; delay is sync own	delay_is_sync_own, single_domain_embedding	StateCellId cells; sync snapshot rule	sync, agnostic corpus	FVD-0044–FVD-0048	—
scheduler order is unobservable	MEv.det, scheduling_order_observable (alternative)	step_in_order	order-independence tests	FVD-0045	—
outputs have one explicit driver	multiple_direct_drivers_rejected, single_driver_output_deterministic	bdl-output	lamp_output	FVD-0050–FVD-0054	—
hardware feasibility is decidable, sound and complete for the finite fragment	solve_sound, solve_complete, satisfiable_iff_solve	bdl-hardware::solve	Nano cases A–H	FVD-0055–FVD-0063	numeric constraints out of scope
behavior systems flatten into a well-formed design	flatten_wf, flatten_causal, flatten_wellClocked, flatten_singleDriver	bdl-system::flatten	system tests, behavior-systems-correspondence.md	FVD-0064–FVD-0073; ADR-0021/0022	Theorem J restricted
groups are transparent; extraction preserves behavior	Theorems A–G (rfl), orig_iff_flat	.bdl/authoring.json, packaging	group tests	FVD-0074–FVD-0082; ADR-0019	single-domain wiring fragment

claim	formal theorem / counterexample	production implementation	production tests	decisions · production records	known limitation
the cross-domain window is expressible with lists	Theorem M buffer_window_correspondence; bounded_summary_not_lossless	five-declaration idiom; Vec core	buffer, bounded_buffer, overflowing_buffer	FVD-0083–FVD-0087; ADR-0024/0027	input source assumed
capacity is validation; only refusal preserves semantics	sufficient_capacity_preserves, periodic_capacity_sufficient, negE	bdl-reactive::capacity, bdl-exec-ir::bounds, refusal	3 000-tick differential	FVD-0086; ADR-0027	static bound unproved
products and one recursor suffice for the tested equations	fold_total, arrow_not_delayable, church_fst_rank	bdl-ir prod, Fold	collections corpus	FVD-0088–FVD-0091; ADR-0025	no minimality theorem
rank-1 polymorphism by matching, no kernel change	matchTy_sound, matchTy_complete, instances_are_monomorphic	bdl-equations	equation tests	FVD-0092; ADR-0025	—
the library adds no privilege	lib_expansion, Comb.noConstruct, lib_clocked, lib_eval_context_free	inlining at elaboration	—	FVD-0094	—
nominality survives generics	generic_preserves_identity, generic_preserves_dimension, map_keeps_concepts	checker on inlined terms	exI, exJ executed; mixed-comparison matrix	FVD-0092; ADR-0026	—
finite $\forall/\exists$ are folds	forall_in_list, exists_in_list	all/any equations	corpus	FVD-0095	—
equality is data; order is by declaration	Cap.eq_iff_data, Cap.ord_data, lt_rejected, min_mode_rejected, minBy_recovers_min	Concept::ordered; semantic.no_order	mixed-comparison matrix	FVD-0098–FVD-0100; ADR-0025/0026	mixed case is a production rule
unit operations need no kernel construct	inUnitE_typed/_safe, withUnitE_typed, convert_eq, unitOps_no_construction	bdl-elab::units coord/reconstruct/convert	unit tests	FVD-0101–FVD-0103	float approximation
presentation is semantically irrelevant	presentation_irrelevant_*	preferred units in authoring metadata	—	FVD-0105	—
Composer inference is sound and complete	solve_sound, solve_complete, candidates_sound/_complete	bdl-ide::formula	formula_composer.rs	FVD-0104; ADR-0028	two holes unsolved by design
affine conversion is a groupoid of affine maps; point/delta orthogonal	convert_compose, convert_inverse, difference_map, linear_part_compose, sort_orthogonal_to_conversion, CtoF_closed	Chart::Affine, convert	charts::tests (ulps)	FVD-0106 rev., FVD-0107–FVD-0110	affine charts hidden (ISS-0004)
binder syntax is conservative desugaring	desugar_rename, binder_local_type, binder_*_eval, range_eval, binder_clock	bdl-elab::formula::binder (P11)	natural_forms_lower_to_the_same_core_as_the_call_forms, parser tests	FVD-0111–FVD-0114; ADR-0028 second amendment	no parser in the model

claim	formal theorem / counterexample	production implementation	production tests	decisions · production records	known limitation
a Source is provisioned by a raw reading and a pure transducer without the design noticing	provision_transparent, provision_abstracts, provision_exact (joint section), provision_wf, provision_causal, provision_wellClocked, provision_not_reapplicable, provision_comm; exD, exE, no_joint_witness	not implemented	—	FVD-0121–FVD-0130; PRP-0001 (draft)	stateful transducers, device clocks open; the output side is Phase 14
a logical output stays independent of its mechanism: realization adds a pure encoder and a machine sink and changes nothing the behavior observes; raw trace = transfer ◦ abstract trace	behavior_unchanged, lower_transparent, lower_physicalOutput_unchanged, lower_correspondence, two_realizations_same_behavior, retarget_breaks_driveWF, encoder_constructs_nothing; exH, exB_quantized, exI	ADR-0036: DeviceBinding.realization, OutputProfile, Encoder { rep, raw, encode } with well_formed/purity, three judgments in analyze_deployment, one SinkPlan per chosen profile, Tick.commands; the Deploy chooser	output_realization.rs; changing_the_realization_changes_no_behavior_and_only_the_raw_trace; the corpus's Commands; the daemon and Studio e2e	FVD-0131 ... FVD-0139 (FVD-0137 superseded); ADR-0036	stateful adapters, device clock, atomic frames, the codegen half open (FVI-0022); the plan-level lowering observably, not provably, the model's; platform independence is the formal-evaluation statement only
the platform adapter applies the raw commands to a board	— (not modelled; RawCommand is the boundary, FVD-0134)	ADR-0037 (amended): the RP2040/Embassy and Arduino/avr-hal adapters, apply(tick, sinks...), duty8, the compiled schedule or the blocking loop, the arena, halt on fault	embedded_rp2040.rs, the AVR cross-build, the runtime's policy and schedule tests, recording mock sinks, the cross-build in CI	FVD-0134, FVD-0138 (the boundary consumed); ADR-0037	production implemented and tested only; raw command → physical effect never proved (FVI-0022)
() -> B is a conservative interface normalization whose kernel value is B	elim_canonical, decode_encode, canonicalOfKernel_encode, zero_input_obligation, lams_typed, refForms_agree, same_tick_same_value, Clocked.refForms	bdL_ir::Ty::{Unit, of_signature, kernel_of_signature, canonical_mapping_ty}, Signature::is_unit_domain	...the_encodings_are_inverse, ...argument_is_erased, the_shorthand_and_the_explicit_unit_domain_are_one_declaration	FVD-0115–FVD-0117; ADR-0029	elim proved on canonical types; () in output position refused by production
a unit binder in the kernel would forbid memory	delay_not_under_binder, sync_not_under_binder, zero_input_memory	delay/sync typed outside binders (transcribed)	exD	FVD-0116; ADR-0029 alternatives	—

claim	formal theorem / counterexample	production implementation	production tests	decisions · production records	known limitation
the source role is a realization state; A -> () cannot name a consumer	source_value, resolved_not_ source, SimulationInput. value, unit_ codomain_collapse, consumers_ indistinguishable, driver_is_unit_ domain, transport_ needs_unit_domain	RelationshipRole:: Source derived in the model and stated by the daemon (ADR-0032); the drive edge	the_role_is_one_ answer_across_the_ component_boundary; Studio simulation tests; exC, exE	FVD-0118- FVD-0120; ADR-0032	Rule and Value have no separate formal object
axiom discipline	every theorem on propext/Quot.sound; no sorry; no Classical.choice (audited per phase)	—	—	—	—

Appendix E — Migration map from the old decision numbers

Until 2026-09-20 the formal decisions were D-01 ... D-130, numbered in the order they were written in one ledger; the open items were unnumbered bullets, and for a few hours on the same day 0I-MN. The stable identifiers keep the sequence number (D-07 $\rightarrow$ FVD-0007) so that a citation in a commit message, an archived draft or a conversation can be followed by hand; the number encodes neither a phase nor a production record. Archived documents (paper/monograph/archive/, the dated entries of paper/monograph/NOTES.md) keep the old numbers and are not rewritten. Production’s records at the snapshot still cite the old numbers (ADR-0024, ADR-0025, formal-correspondence.md, status.md); this table resolves each. The canonical copy is docs/project/decision-id-migration.md in the formal repository.

old	new	title	production record
D-01	FVD-0001	Commitments are atomic labels; evidence is abstract	ADR-0010
D-02	FVD-0002	InterfaceRefines freezes the type and grows commitments; not logical implication	ADR-0010
D-03	FVD-0003	strengthen on a realized declaration carries a re-verification premise	ADR-0010
D-04	FVD-0004	DeclInterface.commitments is a List, not a Finset	ADR-0010
D-05	FVD-0005	References are by DeclId inside Expr (declRef), not a separate DesignExpr	ADR-0010
D-06	FVD-0006	The declRef typing rule reads $\Delta.tyView$ only	ADR-0010
D-07	FVD-0007	Realization is write-once; “detach” is not a refinement step	ADR-0009
D-08	FVD-0008	The structural order (DeclLeq, EnvRefines) is separated from the invariant (GlobalWF)	ADR-0010
D-09	FVD-0009	Evidence takes the environment as an argument	ADR-0010
D-10	FVD-0010	Evidence.Monotone is a kernel-imposed constraint on the validation layer	ADR-0010
D-11	FVD-0011	Semantics at Phase 1 is unfolding to a reference-free term	ADR-0010
D-12	FVD-0012	Acyclicity is witnessed by a rank function	ADR-0010
D-13	FVD-0013	Phase 0’s Artifact (bare list of ids) removed	FV-only
D-14	FVD-0014	Display names are not in the kernel	ADR-0008
D-15	FVD-0015	Kernel ontology migrated from holes to declarations	ADR-0010
D-16	FVD-0016	Monotone refinement is distinct from arbitrary editing	ADR-0009
D-17	FVD-0017	Commitments and validation obligations share PropertyId for now	ISS-0003
D-18	FVD-0018	Evidence.Monotone is a stability condition, not a definition of validity	ISS-0003

old	new	title	production record
D-19	FVD-0019	Concept identity lives in the type: $Ty.sem : ConceptId \rightarrow Ty$	ADR-0013
D-20	FVD-0020	<code>ConceptId</code> is independent of <code>DeclId</code> and of display names	ADR-0013
D-21	FVD-0021	No introduction/elimination forms for concept types in Phase 2	ADR-0013
D-22	FVD-0022	Explicit semantic mappings are ordinary declarations	ADR-0013
D-23	FVD-0023	Concept identity change is an edit	ADR-0013
D-24	FVD-0024	Canonical closed inhabitants replaced by unresolved declarations	ADR-0013
D-25	FVD-0025	Future constraint for Phase 3: representation binding must not defeat concept identity	ADR-0013
D-26	FVD-0026	Unrestricted symmetric <code>mk/rep</code> rejected	ADR-0013
D-27	FVD-0027	Representation types are sem-free (<code>ConceptEnv.WF</code>)	ADR-0013
D-28	FVD-0028	Observation is unrestricted; construction is licensed by the realized declaration's signature	ADR-0013
D-29	FVD-0029	Representation binding is a separate, write-once concept environment θ	ADR-0013
D-30	FVD-0030	Unfolding preserves typing under <code>Grant.all</code> , not under the client grant	ADR-0013
D-31	FVD-0031	Dimensions in Ty as $q\ d$; algebra in <code>Prim.ty</code> ; no dimension rule	ADR-0011, ADR-0013
D-32	FVD-0032	Units are surface: elaborated to scaled dimensioned literals	ADR-0028
D-33	FVD-0033	Concept identity is not indexed by dimension	ADR-0013
D-34	FVD-0034	One temporal primitive: <code>delay init e</code>	ADR-0004, ADR-0016
D-35	FVD-0035	<code>delay</code> is data-typed and top-level (empty context)	ADR-0004, ADR-0016
D-36	FVD-0036	<code>Signal</code> is not a type; <code>Event</code> is <code>opt</code>	ADR-0004, ADR-0016
D-37	FVD-0037	Causality replaces blanket acyclicity: <code>Causal on InstDependsOn</code>	ADR-0004, ADR-0016
D-38	FVD-0038	Explicit initial value on every <code>delay</code>	ADR-0004, ADR-0016
D-39	FVD-0039	State has no identity; state is structural	ADR-0004, ADR-0016
D-40	FVD-0040	Representation types are data (<code>ConceptEnv.WF</code> strengthened)	ADR-0004, ADR-0016
D-41	FVD-0041	Temporal changes are realization edits	ADR-0004, ADR-0016
D-42	FVD-0042	The reactive semantics is a relation, not yet a machine	ADR-0004, ADR-0016

old	new	title	production record
D-43	FVD-0043	Time is one global tick with a schedule; no rates, no timestamps in the kernel	ADR-0004
D-44	FVD-0044	Nominal <code>ClockId</code> , stored per declaration in <code>ClockEnv K</code> ; none = domain-agnostic pure mapping	ADR-0004
D-45	FVD-0045	One transport primitive <code>sync src init e</code> , reading strictly before	ADR-0004
D-46	FVD-0046	The clock is interface data held in a projection, not a record field	ADR-0004
D-47	FVD-0047	Rates, drift, jitter, latency, buffer capacity, value age are validation	ADR-0004
D-48	FVD-0048	Event transport = window read; buffering derived, Event still not a primitive	ISS-0001
D-49	FVD-0049	The logical relation is generic in the application relation	FV-only
D-50	FVD-0050	Physical sinks have nominal identity (<code>OutputId</code>), separate from <code>ConceptId</code> and <code>DeclId</code>	ADR-0005
D-51	FVD-0051	A drive edge is a per-declaration write-once projection β , checked by type and clock equality	ADR-0005
D-52	FVD-0052	Single-driver is a global invariant; completeness is the executable condition	ADR-0005
D-53	FVD-0053	No runtime arbitration, no implicit priority, no merge policy	ADR-0005
D-54	FVD-0054	Effect rows and action values rejected for this kernel	ADR-0005
D-55	FVD-0055	First binding is a refinement; rebinding is an edit; a second driver is invalid	ADR-0005
D-56	FVD-0056	Sinks are terminal	ADR-0005
D-57	FVD-0057	Hardware feasibility is a validation layer over $\text{Design} \times \text{Target}$, not typing	ADR-0006, ADR-0015
D-58	FVD-0058	Resources carry capabilities and per-capability units; sharing is a per-capability policy	ADR-0006, ADR-0015
D-59	FVD-0059	Requirements are independent variables with nominal <code>RequirementId</code> , optional fixed resource, optional unit relation	ADR-0006, ADR-0015
D-60	FVD-0060	Validity is unary support plus pairwise compatibility; the solver is exhaustive DFS, proved sound and complete	ADR-0006, ADR-0015

old	new	title	production record
D-61	FVD-0061	Hardware extension is monotone; requirement extension, strengthening, fixing, and resource removal are revalidation triggers	ADR-0006, ADR-0015
D-62	FVD-0062	Deployment feasibility is environment-sensitive evidence, not Evidence. Monotone	ADR-0006, ADR-0015
D-63	FVD-0063	Numeric electrical/timing constraints deferred	FV-only
D-64	FVD-0064	Behaviour components are surface objects; the kernel is unchanged	ADR-0021, ADR-0022
D-65	FVD-0065	A port is a template declaration by identity, with its public interface and clock	ADR-0021, ADR-0022
D-66	FVD-0066	Instantiation renames every identity the template owns; concepts and sinks are partitioned into internal (fresh) and global (shared)	ADR-0021, ADR-0022
D-67	FVD-0067	Binding is a Phase-1 realization step	ADR-0021, ADR-0022
D-68	FVD-0068	Composition well-formedness is stated on interfaces, never on bodies	ADR-0021, ADR-0022
D-69	FVD-0069	Causality across instances is a validation condition on the inter-instance direct-binding graph	ADR-0021, ADR-0022
D-70	FVD-0070	Clock parameters are nominal variables substituted by κ at instantiation; rates never enter	ADR-0021, ADR-0022
D-71	FVD-0071	Evidence must be equivariant and port-sound	ADR-0021, ADR-0022
D-72	FVD-0072	Hierarchy is packaging, not a tree constructor	ADR-0021, ADR-0022
D-73	FVD-0073	Theorem J is proved on the single-domain wiring fragment with direct/constant bindings	ADR-0021
D-74	FVD-0074	A behaviour group is authoring metadata; eraseGroups is a projection	ADR-0019
D-75	FVD-0075	Group operations are semantic no-ops, not refinements or edits	ADR-0019
D-76	FVD-0076	Aggregate sockets are projections of DependsOn	ADR-0019
D-77	FVD-0077	Boundary inference: required = crossing-in, provided = crossing-out, private = the rest without a sink, clocks = all clocks	ADR-0019
D-78	FVD-0078	Extraction = two restrictions of the design reconnected by Phase-8a bindings	ADR-0019

old	new	title	production record
D-79	FVD-0079	Extraction causality is proved by subdividing the original graph, not by <code>InstAcyclic</code>	ADR-0019
D-80	FVD-0080	Template realization needs interface-local evidence	ADR-0019
D-81	FVD-0081	Identity: templates keep original identities; instances are fresh; the group id is never a component id	ADR-0019
D-82	FVD-0082	Nested groups are a relation on the flat group list	ADR-0019, ISS-0007
D-83	FVD-0083	<code>Ty.list τ</code> is a kernel data type; the object-language buffer needs it and nothing else	ADR-0024
D-84	FVD-0084	Six list operators, registered through <code>Prim.ty: nil, cons, length, take, reverse, head</code>	ADR-0024
D-85	FVD-0085	The buffer is a surface elaboration into five declarations over <code>delay/sync</code>	ISS-0001
D-86	FVD-0086	Capacity is validation; overflow policies are explicit; only rejecting the deployment preserves semantics	ADR-0027
D-87	FVD-0087	Buffered transport is a Phase-8a binding choice, not a transport kind	ADR-0021
D-88	FVD-0088	Products enter the kernel as value composition: <code>prod, pair, fst, snd</code>	ADR-0025
D-89	FVD-0089	The list recursor <code>fold</code> is a term former, not a registered operator	ADR-0025
D-90	FVD-0090	<code>eq</code> at every data type, with the data proof in the syntax	ADR-0025
D-91	FVD-0091	<code>toList : opt τ → list τ</code> and <code>drop</code> are registered operators	ADR-0025
D-92	FVD-0092	Rank-1 polymorphism is definitional: families instantiated by matching; no type variable in the kernel	ADR-0025
D-93	FVD-0093	Constraints: the closed vocabulary <code>{Data}</code> ; no user-defined classes	ADR-0025
D-94	FVD-0094	The equation library is a set of combinators, inlined at use sites	ADR-0025
D-95	FVD-0095	Sets, intervals, records, predicates, finite quantifiers are surface	ADR-0025
D-96	FVD-0096	Sums are encoded; a kernel <code>sum</code> is deferred	ISS-0005
D-97	FVD-0097	Existentials are not needed: hiding is Phase-8a instantiation	ADR-0025
D-98	FVD-0098	Ordering is a quantity comparison; the kernel has no structural order	ADR-0025, ADR-0026

old	new	title	production record
D-99	FVD-0099	The surface capability vocabulary is {Data, Eq, Ord}; Eq = Data today; Ord is by declaration	ADR-0025, ADR-0026
D-100	FVD-0100	Ordered library entries take ordered evidence; comparators recover them	ADR-0025, ADR-0026
D-101	FVD-0101	Units remain entirely surface; coordinate extraction and quantity construction are elaborated quantity arithmetic	ADR-0028
D-102	FVD-0102	Unit semantics are stated exactly over an abstract scalar domain; the kernel's Nat and production's floats are models	ADR-0028
D-103	FVD-0103	A unit is an identity with a dimension and a scale; spelling is presentation	ADR-0028
D-104	FVD-0104	The Formula Composer's formal basis is typed holes with local bidirectional dimension inference — no unification	ADR-0028
D-105	FVD-0105	Preferred display units are presentation, not design	FV-only
D-106	FVD-0106	Affine conversion is complete as coordinate-change semantics; point/delta is optional physical-arithmetic validation	ISS-0004
D-107	FVD-0107	Unit coordinates are an erasure that preserves the affine coordinate change	ISS-0004
D-108	FVD-0108	Conversions form a groupoid of affine isomorphisms; differences carry the linear part	ISS-0004
D-109	FVD-0109	The exact scalar domain is a choice-free rational field built in the development	ISS-0004
D-110	FVD-0110	Unit conversion and sensor calibration are one affine-map abstraction	ISS-0004
D-111	FVD-0111	Binder syntax desugars to the Phase-9 library applied to a lambda; a binder local is the lambda parameter	ADR-0028
D-112	FVD-0112	Ranges are surface nodes desugared to inRange; no interval type or value	ADR-0028
D-113	FVD-0113	No general comprehension, no general quantifier	ADR-0028
D-114	FVD-0114	$\times \text{ } ?? \text{ } d$ desugars to getD	ADR-0028
D-115	FVD-0115	The canonical type $\text{domain}(\text{inputs}) \rightarrow B$ with $\text{domain}([]) = ()$ lives above the kernel; the kernel interface type is its normalization	ADR-0029

old	new	title	production record
D-116	FVD-0116	() -> B is realized at B in the empty context; no unit binder	ADR-0029
D-117	FVD-0117	f, f(), f(()) are one reference; a reading is not a call	ADR-0029
D-118	FVD-0118	The source role is a realization state, not a type shape and not a kind	ADR-0032
D-119	FVD-0119	A -> () is not a physical sink; the drive edge is	ADR-0029, ADR-0032
D-120	FVD-0120	Transport and memory of a relationship require the unit domain — as corollaries	ADR-0029
D-121	FVD-0121	Provision is a construction over designs, not a kernel construct	PRP-0001
D-122	FVD-0122	The profile condition is purity: tr.Pure, i.e. typed in the empty design and delay-free	PRP-0001
D-123	FVD-0123	A channel carries the transfer function and the term that computes it	PRP-0001
D-124	FVD-0124	The profile is generic in the concept; the Source's signature grants construction	PRP-0001
D-125	FVD-0125	Shared raw reading is primitive; the singleton is its special case	PRP-0001
D-126	FVD-0126	Trace equality needs a joint section; deployment is in general a strict refinement	PRP-0001
D-127	FVD-0127	Provision is not re-applicable; "idempotent" is the wrong word	PRP-0001
D-128	FVD-0128	A Source's commitments are obligations on the profile	PRP-0001
D-129	FVD-0129	Independent provisions commute exactly; the assignment is a set	PRP-0001
D-130	FVD-0130	Terminology: abstract Source, provisioned Source, raw declaration; not "monomorphised"	PRP-0001

Open items

old	new	title	production record
OI-01	FVI-0001	Interface-level references	ISS-0003
OI-02	FVI-0002	Delay/temporal boundaries	FV-only
OI-03	FVI-0003	Lambda-guarded instantaneous cycles	FV-only
OI-04	FVI-0004	Higher-order <code>unfolds_</code> <code>preserves_eval</code> (closure equivalence)	FV-only
OI-05	FVI-0005	Reusable stateful components	FV-only
OI-06	FVI-0006	Theorem J for MEV with <code>sync</code> bindings; higher-order Theorem J; a consistent modular input	FV-only

old	new	title	production record
OI-07	FVI-0007	toComponent realizes its chosen interface; port-level inter-instance causality	FV-only
OI-08	FVI-0008	Event multiplicity	FV-only
OI-09	FVI-0009	StateHandler with handler-scoped clocks / independently clocked nesting	ISS-0010
OI-10	FVI-0010	Ty.list and list operators, to write the buffer in the object language	FV-only
OI-11	FVI-0011	Minimal unsat cores; numeric hardware constraints; timer and PWM compatibility	FV-only
OI-12	FVI-0012	A reusable device-component library (Phase 8 surface)	ISS-0016
OI-13	FVI-0013	Folding ClockEnv into the DeclInterface record (churn only)	FV-only
OI-14	FVI-0014	Several candidate definitions with one active, as a surface convenience	ISS-0002
OI-15	FVI-0015	Environment-sensitive evidence and invalidation tracking for edits	ISS-0003
OI-16	FVI-0016	Representation binding for concepts	FV-only
OI-17	FVI-0017	Affine units: a point/difference sort as optional validation; the display-name table	ISS-0004
OI-18	FVI-0018	Grant delegation to higher-order mappings	FV-only
OI-19	FVI-0019	Display-name table for concepts	FV-only
OI-20	FVI-0020	Source provision: stateful transducers, a device clock, commitment discharge, output provision (resolved by Phase 18)	PRP-0001, ISS-0016
—	FVI-0022	Output realization: stateful adapters, a device clock, atomic multi-value frames, codegen correspondence, output commitments	ISS-0016
—	FVI-0023	Raw command → adapter operation, and what lies below the register	ISS-0017, ADR-0037
—	FVI-0024	Explicit device clock for a realized output: what remains after the sync lowering	ISS-0017
—	FVI-0025	Stateful output adapters: the classification and the missing non-encodability witness	ISS-0017
—	FVI-0026	Atomic multi-value frames	ISS-0017
—	FVI-0027	Output commitments and what an encoder must discharge (deferred)	ISS-0017
—	FVI-0028	Minimal unsatisfiable cores for hardware diagnosis (deferred)	FV-only

old	new	title	production record
—	FVI-0029	The provider's occurrence contract: batch delivery, deduplication of transport retries, merged order across raw sources, a per-tick bound (resolved by Phase 17)	ISS-0016, ISS-0001
OI-21	FVI-0021	Unit-domain normalization: elim beyond canonical types; the Input narrowing	ADR-0029

Appendix F — Production snapshot

Production snapshot as of 2026-09-20, commit `aa6e7ff4ee249556ca7561c854b43607f45f229fe` of `KCN-judu/BDL` (“test(studio): toggle-click with the host’s primary modifier in the gestures test”, the head of main after six Studio interaction commits — one pointer state machine, CAD selection, the type scale — that touch no boundary this monograph records, after the second platform adapter family and the vocabulary crate’s rename, the first embedded platform adapter, output realization, the Source sheet, the Code view as an IDE surface and the drive ... by ... spelling; protocol 0.24). What was checked: docs/README.md’s current snapshot, docs/project/status.md, docs/project/roadmap.md, docs/project/formal-correspondence.md, docs/architecture/{overview, relationship-roles, output-realization, embedded-adapter, syntax-highlighting, ide-service}.md, docs/spec/{protocol, textual-syntax, concept-library, hardware-model}.md, ADR-0032 with its amendments, ADR-0034 ... ADR-0037, PRP-0001, ISS-0016, ISS-0017, the change records of 2026-09 (code-view-ide, source-creation, output-realization, embedded-adapter, drive-by), library/std/concepts.toml, and the fixture sources under docs/fixtures/. The formal repository is described as of the working tree of this revision; its last commit before it is 8e65c63 (the Phase 14 hardening). The canonical location of the production hash in the formal repository is the snapshot field of docs/project/production-correspondence.md; this appendix and the title page repeat it.

Milestones since the snapshot `de8154f`, newest first

milestone	what it changed	records
Studio interaction (six commits)	one pointer state machine for the canvas, CAD-style selection, Concept → Output creation from the context menu, the type scale in MacType, status-line clusters, recaptured user-guide screenshots; no compiler, boundary or protocol change	Studio interaction record (docs/decisions), docs/user-guide/studio
the Arduino Nano over avr-hal	bdld compile --target arduino_nano: the second target family, one targets::Entry per family; a blocking tick loop with tick_wait, PWM on the ATmega timers (D3/D11, D5/D6, D9/D10), no arena — a bounded collection is refused (adapter.collections_unsupported); cross-built to an AVR ELF in CI on avr-hal’s pinned nightly; the Uno and the Mega as further board tables of the same base	ADR-0037 (amended); embedded-adapter.md
the adapter vocabulary crate	bdld-runtime-embassy renamed bdld-runtime-adapter: the sink traits, the numeric policy and the recorded operations are shared by every family, so the crate says what it is rather than who first used it; the production repository also carries this monograph as a reference mirror under reference/paper/	ADR-0037 (amended)
drive ... by ...	output driving is spelled drive light by brightness; drive light = brightness is legacy syntax with a hint (text.legacy_drive) and an opt-in bdld migrate-drive-by; the fixtures and the guide moved	change record 2026-09-drive-by

milestone	what it changed	records
the first embedded platform adapter	bdld compile --target rp2040_pico [--tick-micros N]: generated adapter glue and Embassy firmware for the Raspberry Pi Pico; bdl-runtime-embassy (vocabulary, since renamed) and bdl-runtime-embassy-rp (HAL binding, outside the workspace); the board file rp2040_pico; the explicit reject-and-hold numeric policy; the arena from the manifest; halt on a failed tick; TickTrace.adapter on the host; the firmware cross-built in CI	ADR-0037; embedded-adapter.md
output realization	DeviceBinding.realization; OutputProfile { id, encoder, kind } with Encoder { rep, raw, encode }, well_formed and purity; three-judgment admissibility (DeviceRealization { check, hardware_placed }); one SinkPlan per valid chosen profile; Tick.commands; five witness profiles; the Deploy page's chooser; Exec IR v3; protocol 0.24	ADR-0036; output-realization.md; ISS-0017
the Source sheet	a Source is created over a concept the designer chooses — existing, or new in one transaction; the Standard Library's Source items are presets (<i>Temperature Input</i> ...); CreateSource, ListSourceCandidates; protocol 0.23	change record 2026-09-source-creation
the Code view as an IDE surface	completion at the caret, a hover card, definition and references across files, <i>Format</i> as one edit — bdl_id::navigation, one answer to <i>what is at a position</i> for the language server, the daemon and Studio; protocol 0.22	change record 2026-09-code-view-ide

The milestones before these — syntax highlighting (ADR-0035; 0.21), relationship roles (ADR-0032 amended; 0.20), the unapplied rule (0.19), reference edges (ADR-0034; 0.18), the generalized Standard Library (0.17), the Source role (ADR-0032; 0.16), internationalization (ADR-0031), complete persistence (ADR-0030), the unit-domain normalization (ADR-0029; 0.14), the natural expression surface (ADR-0028's second amendment; 0.13), PRP-0001 — are described in the chapters and were the previous snapshot's; nothing in them moved.

Protocol

Protobuf over framed stdio; additive minors. 0.11 structured value forms · 0.12 the Composer's three queries · 0.13 binder and range nodes · 0.14 the unit type kind · 0.15 definition drafts in the system view · 0.16 Source template fields (deprecated) · 0.17 library items · 0.18 MappingAnalysis.references · 0.19 CreateMapping.definition?/clock_id? · 0.20 RelationshipRole, MappingView.role, MappingAnalysis.role / applied_by · 0.21 SemanticTokens · 0.22 SourceCompletion, SourceHover, SourceDefinition, SourceReferences, FormatSource · 0.23 CreateSource, ListSourceCandidates, LibraryItemView.preset · **0.24** SetDeviceRealization, DeviceView.realization, DeploymentAnalysis.realizations, MissingKind.REALIZATION_INVALID. Older clients are unaffected by each addition.

Implementation status, by area (production's own words, abridged)

Implemented: the language core through outputs and completeness; formula language v0 with the slot; the natural forms; one canonical type per relationship; drive ... by ... with the legacy spelling as compatibility syntax; the unit registry with linear charts offered and affine charts as tested infrastructure (ISS-0004); the data core and the equation library; the compiler and simulator with product-language diagnostics and Explain; the Rust backend with differential, golden and property tests over 22 corpus cases, now comparing Commands; collections at deployment (static bounds, window capacity, bounded-memory refusal); hardware and deployment analysis with the Nano, a larger mock board and the Raspberry Pi Pico; output realization (ADR-0036); behavior systems, groups, packaging,

versions and substitution; the unified project format with migration and complete persistence; the layout service; text ↔ graph synchronisation; the IDE service and LSP with a VS Code extension, semantic tokens, navigation and the Composer queries; the protocol and daemon at 0.24; the Standard Library (36 Concept items; 8 Source presets); localization (partial: the compiler's diagnostic sentences are English, ISS-0015).

Partial — the embedded runtime: two platform adapter families (ADR-0037, amended) — PWM duties and digital levels on the RP2040 with the compiled schedule and the arena, and on the Arduino Nano with a blocking tick loop and no arena; halt on fault; the host's recording of the same operations; both firmwares cross-built in CI. Not built: a device that provides a Source's value (ISS-0016: a design with a Source cannot run on a board yet), I²C and H-bridge sinks, stateful adapters and a device clock (ISS-0017), a timer-driven tick on the AVR, build orchestration, flash, telemetry, a third family (roadmap priorities 2–5).

Partial — Studio: the Design page with Design, Code and Split views; the canvas with roles, reference edges and groups; the inspector with the Formula | Text definition editor; the Formula Composer with references, literals, slots, operators, calls, binders, ranges, boolean logic and choices structured and richer forms as text; the Code view with colour, completion, hover, definition, references and format; the Source sheet; Simulate with Sources as inputs; Deploy on the 0.5 read model with the realization chooser; the Library; the conflict banner. Not built: the Monitor page (a placeholder), a device binding for a Source (ISS-0016), Explain over a protocol request, domain regions and cycle emphasis on the canvas, entity hover and fixes inside a component's source, a native menu bar.

Planned, designed, not implemented: supplied Rust components (ADR-0005); the input device catalogue and PRP-0001; the output device catalogue beyond the five witness profiles.

Not implemented by decision: user enums (ISS-0005), temporal modifiers and contexts (ISS-0010), a surface form for occurrence windows (ISS-0001), record syntax, forall/exists sugar, Set/interval/sum types, an A -> () consumer form, an effect or IO type, retargeting an output to a raw type.

Facts this document relies on that could move

Protocol version numbers; the count of library items and presets; the five profile ids; the list of structured Composer forms (let, match, blocks and rules are text at the snapshot); the roadmap order; the runtime crate names (bdl-runtime-adapter, bdl-runtime-embassy-rp and bdl-runtime-arduino at the snapshot); the set of target families and their board tables; the two known code-side drifts production recorded while auditing its own docs. When any of these moves, this appendix and the correspondence page move; the conceptual chapters do not.

Appendix G — Development chronology

The order in which the hypotheses were tested. Phase numbers are provenance, not structure: the conceptual chapters cite them so that a reader can find the report that holds the experiment, and the reports (docs/reports/ in the formal repository) are the local records — question, models tried, theorems with their hypotheses, executed cases, claim audit, verdicts — from which this document’s chapters were synthesized. A report is written once and corrected by strike-through when a later phase overturns a claim (Phase 9c did this to 9b’s structural order; Phase 10b to 10’s sort). Dates are the commits that added each report.

phase	date	question	main section	report
0	2026-09-14	a single persistent declaration	§IV.1	docs/reports/phase-00-a-single-persistent-declaration.md
1	2026-09-14	cross-declaration references	§IV.1	docs/reports/phase-01-cross-declaration-references.md
1 → M	2026-09-14	Migration: holes → declarations	§IV.1	docs/reports/phase-01m-migration-holes-to-declarations.md
2	2026-09-15	where does concept identity live?	§IV.2	docs/reports/phase-02-where-does-semantic-identity-live.md
3	2026-09-15	representation binding and physical dimensions	§IV.2	docs/reports/phase-03-representation-binding-and-physical-dimensions.md
4	2026-09-15	Reactive Core	§IV.5	docs/reports/phase-04-reactive-core.md
5	2026-09-15	Clock domains and synchronization	§IV.5	docs/reports/phase-05-clock-domains-and-synchronization.md
6	2026-09-15	Physical outputs and the single-driver discipline	§IV.7	docs/reports/phase-06-physical-outputs-and-the-single-driver-discipline.md
7	2026-09-15	Hardware constraint validation and resource allocation	§IV.8	docs/reports/phase-07-hardware-constraint-validation-and-resource-allocation.md
8a	2026-09-15	Behaviour as a first-class design object	§IV.6	docs/reports/phase-08a-behaviour-as-a-first-class-design-object.md
8b	2026-09-16	Behaviour grouping and component extraction	§IV.6	docs/reports/phase-08b-behaviour-grouping-and-component-extraction.md
9a	2026-09-17	List data and lossless buffered cross-domain events	§IV.5	docs/reports/phase-09a-list-data-and-lossless-buffered-cross-domain-events.md
9b	2026-09-17	Minimal data abstraction and the polymorphic equation language	§IV.3	docs/reports/phase-09b-minimal-data-abstraction-and-the-polymorphic-equation-language.md
9c	2026-09-18	Capability boundary audit: Data vs Eq vs Ord	§IV.3	docs/reports/phase-09c-capability-boundary-audit-data-vs-eq-vs-ord.md
10	2026-09-18	Unit coordinates and formula-assembly semantics	§IV.4	docs/reports/phase-10-unit-coordinates-and-formula-assembly-semantics.md
10b	2026-09-18	Affine coordinate erasure and conversion functoriality	§IV.4	docs/reports/phase-10b-affine-coordinate-erasure-and-conversion-functoriality.md

phase	date	question	main section	report
11	2026-09-18	Natural expression surface as conservative desugaring	§IV.3	docs/reports/phase-11-natural-expression-surface-as-conservative-desugaring.md
12	2026-09-18	Unit-domain normalization and the source boundary	§IV.7	docs/reports/phase-12-unit-domain-normalization-and-the-source-boundary.md
13	2026-09-20	Source provision by device transducers (PRP-0001 audit)	§IV.7	docs/reports/phase-13-source-provision-by-device-transducers-prp-0001-audit.md
14	2026-09-20	Output realization by device encoders — and, the same day, the hardening pass (FVD-0139 supersedes FVD-0137)	§IV.7	docs/reports/phase-14-output-realization-by-device-encoders.md
15	2026-09-20	The adapter boundary and the explicit device clock — after the audit of every open item	§IV.7	docs/reports/phase-15-the-adapter-boundary-and-the-explicit-device-clock.md
16	2026-09-20	Communication as state, catalogue profiles and the deployment-only assign — the auto typer stress case	§IV.7	docs/reports/phase-16-communication-as-state-catalogue-profiles-and-the-deployment-only-assign.md
17	2026-09-20	The provider’s occurrence contract and the occurrence-preserving output window	§IV.7	docs/reports/phase-17-the-provider-occurrence-contract-and-the-output-window.md
18	2026-09-20	The Source-side boundary: provider state, the device clock, initialization, commitments, computes, out-of-type readings	§IV.7	docs/reports/phase-18-the-source-side-boundary-provider-state-device-clock-commitments-and-readings.md
19	2026-09-21	May several declarations produce values of one nominal concept? Models A, B, C	§IV.2	docs/reports/phase-19-may-several-declarations-produce-one-concept.md
20	2026-09-21	One producer per concept: the invariant, refinement, composition, the concept reference	§IV.2	docs/reports/phase-20-one-producer-per-concept.md
21	2026-09-21	Sem blocks and mapping blocks: the design objects, the concept as their type template	§IV.2	docs/reports/phase-21-sem-blocks-and-mapping-blocks.md
22	2026-09-21	The concept ladder: concept = type (ConceptId), Sem block = instance; the renaming	§IV.2	docs/reports/phase-22-the-concept-ladder-and-the-renaming.md

Notation changed once: Phases 0–1 spoke of *holes* (holeId, holeRef); the post-Phase-1 migration (FVD-0015) replaced the vocabulary with declarations, and every later phase uses the current notation of Appendix A. Phase 9b’s

structural order and Phase 10's point/difference sort were revised by 9c and 10b; the earlier positions are kept in the reports and in §VII.2's *decisions that changed* table.

Production consumed the phases in this order, all on 2026-09-20 unless dated: Phases 1–7 and 9a–9c through the kernel transcription and ADR-0010/0024–0027 (2026-09-15 ... 18); Phase 8a/8b as ADR-0021/0022/0019 (2026-09-16 ... 17); Phase 10 as ADR-0028 (2026-09-18); Phase 11 under ADR-0028's second amendment; Phase 12 as ADR-0029 and ADR-0032; Phase 13 as the revised PRP-0001 (not implemented); Phase 14 as ADR-0036 the day it closed, followed by the first platform adapter (ADR-0037), which consumes the boundary Phase 14 defined and nothing formal beyond it.

Appendix H — Revision log

revision	what the record gained
2026-09, conference manuscript (Phases 0–7)	motivation, designer vocabulary, interaction model, the kernel through hardware validation, claim-strength discipline; the manuscript is archived beside this document
Phase 8a/8b — behavior systems and groups	components, fresh instantiation, bindings, flattening, substitutability; groups as authoring metadata and their extraction
Phase 9a — list data and lossless buffering	list τ , the object-language window, capacity as validation
Phase 9b/9c — data and equation language	products, the list recursor, rank-1 definitional polymorphism, the capability audit that reverted structural ordering
Phase 10/10b — units, charts, Formula Composer basis	unit coordinates, exact symbolic scales, local dimension inference, affine charts and the AffSort revision
production Formula Composer (P10a/P10b)	the projection, slots, compose operations, unit ownership
Phase 11 — natural expression surface	binders and ranges as conservative desugaring
the monograph rewrite (69caa3a)	the monograph structure, production architecture, correspondence and deviations, ledgers, open agenda
production P11 and ADR-0029 (3c6c8be)	the natural forms implemented; one canonical type per relationship, the empty product $()$ as the unit domain, protocol 0.13/0.14
Phase 13 — Source provision (PRP-0001 audit)	device channels and profiles, shared-row provision as an EnvRefines step, transparency, trace abstraction, joint-section exactness, the corrected claims
Phase 14 — output realization by device encoders	the logical output as the platform-independence boundary; the encoder, the RawCommand specification and the lowering; behavior preservation, directional correspondence, platform independence; the refuted models
Phase 14 hardening	Admissible requires the encoder’s typing (FVD-0139 supersedes FVD-0137; exJ); the RawCommand specification and the lowered machine sink kept apart in every sentence; the exact scope of two_realizations_same_behavior; FVD-0136 as a decision with the atomic-frame case open; “logical output” applied to the current pages, PhysicalOutput as the historical name; the production pin de8154f stated with its reason
Phase 12 — unit-domain normalization	$() \rightarrow B$ as an interface normalization whose value is the kernel type B; the source role as a realization state; $A \rightarrow ()$ shown unable to name a consumer
2026-09-20 — the conceptual restructuring	the conceptual restructuring; Phase 13 as the environment boundary (§IV.7); production at de8154f; the FVD/FVI identifiers; the indexes (Appendices B–E); the two kinds of minimality; the residue of the conference form removed; paper/README.md and main.typ describing the monograph pipeline
2026-09-20 — production at 6be778b	the physical boundary as one whole with the strength of each arrow (§IV.7); output realization implemented (ADR-0036) and the first embedded platform adapter on the RP2040 (ADR-0037) at production-test strength; the Source sheet, the Code view as an IDE surface, drive ... by ..., protocol 0.22–0.24; the physical-product and industrial-design position stated (Part I); the third minimality distinction (§VII.2); the open agenda split into formal, production and empirical, with the reverse-readability question; counts derived from the repository; the mirror policy for reference/paper/ in production

revision	what the record gained
2026-09-20 — the narrative restructure	the monograph reorganized as one argument in seven Parts: the design problem (I), the architecture with one central figure (II), one interactive physical product walked through the whole lifecycle (III), the formal model derived in dependency order with the phases as provenance only (IV), the toolchain by responsibility with the reference evaluator and two adapter families (V), Studio and the IDE as a design argument with its software references and the UX claims as hypotheses (VI), the evidence, the rejected designs as arguments and the open agenda (VII); related work after the argument; the front matter reduced to a preface; Appendix I (record conventions) and Appendix J (this section map); production re-pinned at 081296d — the Arduino Nano over avr-hal as the second adapter family, bdl-runtime-adapter; the formal development at Phase 15
2026-09-20 — Phase 16	communication as state: the auto_typer stress case encoded and executed on the unchanged kernel, the fourth and product-scale rejection of a message/event type; the three kinds of identity and the replacement-invariance criterion as three theorems (two_providers_same_behavior new); catalogue profiles with an unread origin, assign as the lowering / the provision, contract separate from feasibility, the paired axis without a transaction primitive (§IV.7, the architectural principle in Part II, §VII.2, §VII.4); FVI-0029; production re-pinned at aa6e7f4 (six Studio interaction commits, no boundary change)
2026-09-20 — Phase 17	the provider's occurrence contract stated and proved (one occurrence per fresh transport identity, retransmissions erased, a bounded batch with an observable flag, several raw sources as several provisions), composed with Phase 13; the occurrence-preserving output window over the encoder into the device domain with its correspondence, bound and structure; the adapter's batch as a list of operations; the motion state as a component on two axes (§IV.7, §VII.2, §VII.4, Appendices B, C, F, G); FVI-0029 resolved; the production tree's in-flight uncommitted Source slice noted, not cited
2026-09-20 — Phase 18	the Source side completed: the value-congruence lemma; provider state as a Mealy machine below or above the reading with one trace and the visibility criterion; the Source-side device clock as an explicit sync or window with an explicit initial value, supplied or unavailable; commitment discharge at three evidence levels with the trusted assumption visible; computes derived from the term and decided at a finite type; out-of-type readings refused into none or a flag; the motion state under two providers (§IV.7, §VII.2, §VII.4, Appendices B, C, F, G); FVI-0020 resolved
2026-09-21 — Phases 19–22	the concept ladder: the audit of producers (Phase 19), one producer per concept tried and set aside (Phase 20), Sem blocks and mapping blocks as the design objects (Phase 21, FVD-0163), the vocabulary fixed in the code — SemanticId is ConceptId (Phase 22, FVD-0164); §IV.2 gains the ladder, §III.1 names it, §VI.6 notes the projection production's canvas will follow

Appendix I — Record conventions: claim strength and identifiers

The monograph is also the authoritative record of BDL, and the conventions that make it one are gathered here so that the preface can stay short.

Claim strength

Every substantive claim carries one of the following labels, used in the text, in the ledgers of the appendices, and — with production’s own words — in production’s records (`docs/project/formal-correspondence.md` there, and the `fv` field of each ADR). The labels are never flattened into “BDL guarantees”.

label	means
formally proved	a named Lean theorem proves the stated property of the formal model; it never proves the Rust or Dart code
formally characterized under restricted hypotheses	a named theorem proves the property for a stated fragment (a wiring design, direct bindings, one domain); the restriction is part of the claim
mechanically executed example	a concrete design run through the proved-sound interpreter or solver inside the proof checker (<code>decide</code> , <code>#eval</code>); evidence for that input, not a theorem
counterexample / rejected by model	a candidate construct or claim was formalized and a theorem or executed case shows it wrong
informed by FV	a formal result or counterexample bounded an engineering choice; the choice itself is production’s
production implemented and tested	present at the production snapshot and exercised by a named test (<code>differential</code> , <code>golden</code> , <code>property</code> , <code>end-to-end</code>)
production architecture decision	an accepted production ADR; may or may not have formal backing, and the row says which
proposal / not implemented	a production proposal (PRP) or a formal construction with no implementation
design recommendation	guidance to production with no theorem behind it
open empirical question	a claim about designers or usability; no study has been run, and nothing in this document is evidence for it

“Minimal” is never claimed globally. Where the formal development says minimal it means *minimal among the tested candidates*, or *the smallest design found that supports the required cases*, and the text says which.

Identifiers

Formal design decisions are FVD-NNNN and formal open items FVI-NNNN (`docs/decisions/`, `docs/issues/` in the formal repository); production’s records are ADR-NNNN, PRP-NNNN and ISS-NNNN (`docs/decisions/`, `docs/proposals/`, `docs/issues/` in the production repository). Where a formal decision and a production record concern the same architectural question the text cites the two together — ADR-0032 (FVD-0118) — and no competing number is minted for a fact production already names. Until 2026-09-20 the formal decisions were numbered D-01 ... D-130 in a single ledger; Appendix E is the permanent map, and archived documents keep the old numbers. Theorem names are stable identities and never change to track a document id: `provision_transparent` is that theorem whether its record is cited as PRP-0001, FVD-0126 or a section of this document.

Snapshots

Production is described as of one pinned commit, stated in the preface and in Appendix F; the canonical copy of that hash in the formal repository is the `snapshot` field of `docs/project/production-correspondence.md`, and the title page repeats it. The formal development is described as of the working tree that contains the revision. Volatile production details — protocol minors, profile counts, crate names — live in Appendix F so that the conceptual Parts do not go stale with the next milestone.

Canonical source and the production mirror

There is one BDL monograph. Its canonical source is `KCN-judu/BDL_FV/paper/monograph/`; `paper.md` is the only file whose prose is edited, and `body.typ` and the PDF are generated from it there. The production repository carries a

copy under `reference/paper/` as a reference mirror, refreshed after a revision lands in the formal repository and never edited in place; it is not a second authority.

Appendix J — Section migration map from the previous revision

The revision of 2026-09-20 that preceded this one was organized as a topic taxonomy: fifteen Parts, one per concept, in the order the phases had produced them. This revision is organized as the argument of the preface. Every technical claim of the previous revision was classified before it was moved, and the table records where each of its sections now lives. *Retained* means the prose is in place under a new heading; *moved* means it is in a different Part; *merged* means it was joined with new prose or with another section; *appendix* means it left the argument for the reference apparatus; *removed* means the text was dropped — and only navigational or duplicated text was.

previous section	now	disposition
About this document: purpose, current result, scope	Preface; Part II <i>The state of the record</i> (the current-result and scope paragraphs)	merged
About this document: how to use	Preface <i>How to read this monograph</i>	merged (rewritten around the seven Parts)
About this document: claim strength; identifiers	Preface (short legend); Appendix I (full conventions)	appendix
About this document: revision log (short form)	Appendix H	removed as a duplicate; the full log is Appendix H
Part I Introduction; the Representation Problem; scope and non-goals	Part I — <i>The asymmetry of design media, The question, The working hypothesis, The Representation Problem, Scope</i>	retained, with two new chapters (<i>What BDL is not, Who BDL is for</i>)
Part II Concepts; Relationships and the three roles; Incomplete designs are legal; Formulas	§III.1–III.3	moved into the running example
Part II Time; Physical outputs; Behaviors, components and systems; the Standard Library; supplied computation	§III.6, §III.5, §III.7, §III.8	moved into the running example
Part II A representative scenario	Part III as a whole; its board and realization paragraphs §III.10	merged (the scenario is now the spine of Part III)
Part II What the workspace says; Looking underneath	§III.11, §III.12	moved
Part III Method; Scale and trust base	Part IV introduction	merged
Part III From Surface to Kernel: Architecture (the layer figure)	Part II <i>Three layers</i>	moved (the figure is now beside the architecture map)
Part III Declarations, Interfaces, and Refinement	§IV.1	retained
Part III Where the rest of the kernel is described	—	removed (navigation; the dependency order of Part IV replaces it)
Part IV Concept identity and physical quantities	§IV.2	retained
Part V The data and equation language	§IV.3	retained
Part VI Units, coordinates and charts	§IV.4	retained
Part VII Reactive semantics	§IV.5	retained
Part VIII Behavior systems	§IV.6	retained
Part IX The environment and the physical boundaries	§IV.7	retained
Part X Validation and deployment	§IV.8	retained
Part XI Four trust layers	Part V introduction	merged with the ownership table
Part XI Crates and semantic ownership	§V.7	moved to the end of Part V as the implementation map
Part XI The compiler as a pipeline; the elaboration passes the kernel implies	§V.2	merged
Part XI Executable IR and lowering; the generated core; embedded execution model	§V.4	merged
Part XI The platform adapter	§V.5, with the second family added	retained
Part XI Differential testing	§V.3 (the reference evaluator, new)	merged

previous section	now	disposition
Part XI Project persistence and the daemon	§V.1	moved
Part XI The IDE service, the daemon and the protocol	§V.6	retained
Part XII Three information levels	§VI.4 (with the semantic-projection diagram)	merged
Part XII The canvas; Source, Rule and Value on the canvas; groups and components on one canvas	§VI.5	merged
Part XII Semantic actions	§VI.6	merged with the gesture and context-menu account
Part XII Colour in the Code view; Design, Code and Split	§VI.8	merged
Part XII The Library	§VI.10	retained
Part XII The definition editor and the Formula Composer; partial expressions	§VI.9	retained
Part XII Simulate, Deploy, Library, Monitor	§VI.11	retained
Part XII What Studio does not decide	§VI.12, with the hypotheses stated	merged
—	§VI.1–VI.3, §VI.7 (the user, the principles, the software references, diagnostics)	new
§VII.1 Formal ↔ production correspondence	§VII.1	retained
Part XIV Two kinds of minimality	§VII.3	moved
Part XIV Kernel constructs retained; rejected constructs; decisions that changed	§VII.2, preceded by the ten arguments	merged
Part XV The open agenda	§VII.4	retained
Intellectual context and related work	after Part VII	retained (moved after the argument)
Appendices A–H	Appendices A–H	retained; Appendix F re-pinned; Appendix H extended
Closing	Closing	retained

References

- [1] S. Greenberg and C. Fitchett, “Phidgets: Easy Development of Physical Interfaces through Physical Widgets,” in *Proceedings of the 14th Annual ACM Symposium on User Interface Software and Technology*, ACM, 2001, pp. 209–218. doi: 10.1145/502348.502388.
- [2] B. Hartmann *et al.*, “Reflective Physical Prototyping through Integrated Design, Test, and Analysis,” in *Proceedings of the 19th Annual ACM Symposium on User Interface Software and Technology*, ACM, 2006, pp. 299–308. doi: 10.1145/1166253.1166300.
- [3] L. de Moura and S. Ullrich, “The Lean 4 Theorem Prover and Programming Language,” in *Automated Deduction – CADE 28*, in Lecture Notes in Computer Science, vol. 12699. Springer, 2021, pp. 625–635. doi: 10.1007/978-3-030-79876-5_37.
- [4] T. R. G. Green and M. Petre, “Usability Analysis of Visual Programming Environments: A 'Cognitive Dimensions' Framework,” *Journal of Visual Languages & Computing*, vol. 7, no. 2, pp. 131–174, 1996, doi: 10.1006/jvlc.1996.0009.
- [5] N. Halbwachs, P. Caspi, P. Raymond, and D. Pilaud, “The Synchronous Data Flow Programming Language LUSTRE,” *Proceedings of the IEEE*, vol. 79, no. 9, pp. 1305–1320, 1991, doi: 10.1109/5.97300.
- [6] J.-L. Colaço and M. Pouzet, “Clocks as First Class Abstract Types,” in *Proceedings of the Third ACM International Conference on Embedded Software*, ACM, 2003.
- [7] R. Dechter, *Constraint Processing*. Morgan Kaufmann, 2003.
- [8] B. Hartmann, L. Abdulla, M. Mittal, and S. R. Klemmer, “Authoring Sensor-Based Interactions by Demonstration with Direct Manipulation and Pattern Recognition,” in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, ACM, 2007, pp. 145–154. doi: 10.1145/1240624.1240646.
- [9] Arduino, “Arduino Language Reference.” 2024.
- [10] OpenJS Foundation, “Node-RED: Low-code programming for event-driven applications.” 2024.
- [11] M. Resnick *et al.*, “Scratch: Programming for All,” *Communications of the ACM*, vol. 52, no. 11, pp. 60–67, 2009.
- [12] A. J. Ko *et al.*, “The State of the Art in End-User Software Engineering,” *ACM Computing Surveys*, vol. 43, no. 3, p. 21:1–21:44, 2011.
- [13] J.-L. Colaço, B. Pagano, and M. Pouzet, “A Conservative Extension of Synchronous Data-flow with State Machines,” in *Proceedings of the 5th ACM International Conference on Embedded Software*, ACM, 2005. doi: 10.1145/1086228.1086261.
- [14] G. Berry and G. Gonthier, “The Esterel Synchronous Programming Language: Design, Semantics, Implementation,” *Science of Computer Programming*, vol. 19, no. 2, pp. 87–152, 1992, doi: 10.1016/0167-6423(92)90005-V.
- [15] G. Berry, “SCADE: Synchronous Design and Validation of Embedded Control Software,” in *Next Generation Design and Verification Methodologies for Distributed Embedded Control Systems*, Springer, 2007, pp. 19–33.
- [16] T. Bourke and M. Pouzet, “Zélus: A Synchronous Language with ODEs,” in *Proceedings of the 16th International Conference on Hybrid Systems: Computation and Control*, ACM, 2013, pp. 113–118. doi: 10.1145/2461328.2461348.
- [17] C. Elliott and P. Hudak, “Functional Reactive Animation,” in *Proceedings of the Second ACM SIGPLAN International Conference on Functional Programming*, ACM, 1997, pp. 263–273. doi: 10.1145/258948.258973.
- [18] G. H. Cooper and S. Krishnamurthi, “Embedding Dynamic Dataflow in a Call-by-Value Language,” in *Programming Languages and Systems: 15th European Symposium on Programming*, Springer, 2006, pp. 294–308.
- [19] National Instruments, “LabVIEW.” 2024.
- [20] The MathWorks, “Simulink and Stateflow.” 2024.

- [21] Modelica Association, “Modelica Language Specification, Version 3.6.” 2023.
- [22] D. Harel, “Statecharts: A Visual Formalism for Complex Systems,” *Science of Computer Programming*, vol. 8, no. 3, pp. 231–274, 1987, doi: 10.1016/0167-6423(87)90035-9.
- [23] Object Management Group, “OMG Systems Modeling Language (SysML) Version 2.0.” [Online]. Available: <https://www.omg.org/spec/SysML/2.0>
- [24] C. Omar, I. Voysey, M. Hilton, J. Aldrich, and M. A. Hammer, “Hazelnut: A Bidirectionally Typed Structure Editor Calculus,” in *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*, ACM, 2017, pp. 86–99. doi: 10.1145/3009837.3009900.
- [25] C. Omar, I. Voysey, R. Chugh, and M. A. Hammer, “Live Functional Programming with Typed Holes,” *Proceedings of the ACM on Programming Languages*, vol. 3, no. POPL, p. 14:1–14:32, 2019, doi: 10.1145/3290327.
- [26] M. Voelter, J. Siegmund, T. Berger, and B. Kolb, “Towards User-Friendly Projectional Editors,” in *Software Language Engineering (SLE)*, in LNCS, vol. 8706. 2014, pp. 41–61.
- [27] A. Kennedy, “Relational Parametricity and Units of Measure,” in *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, ACM, 1997, pp. 442–455. doi: 10.1145/263699.263761.
- [28] G. D. Plotkin and M. Pretnar, “Handling Algebraic Effects,” *Logical Methods in Computer Science*, vol. 9, no. 4, 2013, doi: 10.2168/LMCS-9(4:23)2013.
- [29] Z. Yang, M. Paviotti, N. Wu, B. van den Berg, and T. Schrijvers, “Structured Handling of Scoped Effects,” *Proceedings of the ACM on Programming Languages*, vol. 6, no. ICFP, 2022, doi: 10.1145/3547654.
- [30] W. M. McKeeman, “Differential Testing for Software,” *Digital Technical Journal*, vol. 10, no. 1, pp. 100–107, 1998.
- [31] X. Leroy, “Formal Verification of a Realistic Compiler,” *Communications of the ACM*, vol. 52, no. 7, pp. 107–115, 2009.